



Department of Computer Science

Lab Manual

of

MACHINE LEARNING

MCA521-4

Class Name: IV MCA

Master of Computer Application

2026

Prepared by:
Kuheli Begum
2547230

Verified by:

Dr. Rupali Sunil Wagh Dr. Helen K Joy Dr. Shivangi Singh
--



Department Overview

Department of Computer Science of CHRIST (Deemed to be University) strives to shape outstanding computer professionals with ethical and human values to reshape nation's destiny. The training imparted aims to prepare young minds for the challenging opportunities in the IT industry with a global awareness rooted in the Indian soil, nourished and supported by experts in the field.

Vision

The Department of Computer Science endeavours to imbibe the vision of the University “**Excellence and Service**”. The department is committed to this philosophy which pervades every aspect and functioning of the department.

Mission

“To develop IT professionals with ethical and human values”. To accomplish our mission, the department encourages students to apply their acquired knowledge and skills towards professional achievements in their career. The department also moulds the students to be socially responsible and ethically sound.

Introduction to the Programme

Master of Computer Applications (MCA) is a post-graduate programme of CHRIST (Deemed to be University) and approved by the All India Council of Technical Education (AICTE). It is a two-year programme spread over six trimesters to bridge the chasm between computer studies and applications, bringing within reach of enthusiastic youngsters an excellent group of experienced and dedicated staff and experts, and an enviable infrastructure. MCA at CHRIST (Deemed to be University) strives to shape outstanding computer professionals with ethical and human values to reshape the nation's destiny. The training programme aims to prepare young minds for challenging opportunities in the IT industry with a global awareness rooted in the Indian soil, nourished and supported by experts in the field. The Department is committed to the motto “Excellence and Service,” and this philosophy pervades every aspect of its functioning.

Programme Objective

This programme is conceptualised and designed based on the strong commitment of the department to provide better quality education to the students. The principal objectives of this course are to:

- Heighten technological awareness
- Train future industry specialists
- Encourage effective software development
- Provide research training
- Involve students in innovative research



- Produce graduates with a two-year professional education in Computer Science with technical, professional, and communications skills.

Ethics and Human Values

1. Only proprietary or open-source software would be used for academic teaching and learning purposes.
2. Copying of programs from the internet, friends or from other sources is strictly discouraged since it impairs development of programming skills.
3. Unique Practical (Domain based) exercises ensure that the students don't get involved in code plagiarism.
4. Projects undertaken by students during the course are done in teams to improve collaborative work and synergy between team members.
5. Projects involve modularization which initiates students to take individual responsibility for common goals.
6. Passion for excellence is promoted among the students, be it in software development or project documentation.
7. Giving due credit to sources during the seminar and research assignment is promoted among the students
8. The course and its design enforce the practice of good referencing technique to improve the sense of integrity.
9. Courses involving group discussions and debates on ethical practices and human values are designed to sensitize the students in dealing with customers and members within the Organization.

Programme Outcomes:

- P01: Foundation Knowledge: Apply knowledge of mathematics, programming logic, and coding fundamentals for solution architecture and problem-solving.
- P02: Problem Analysis: Identify, review, formulate, and analyse problems primarily focussing on customer requirements using critical thinking frameworks.
- P03: Development of Solutions: Design, develop and investigate problems with as an innovative approach for solutions incorporating ESG/SDG goals.
- P04: Modern Tool Usage: Select, adapt, and apply modern computational tools such as the development of algorithms with an understanding of the limitations including human biases.
- P05: Individual and Teamwork: Function and communicate effectively as an individual or a team leader in diverse and multidisciplinary groups. Use methodologies such as agile.
- P06: Project Management and Finance: Use the principles of project management such as scheduling, and work breakdown structure and be conversant with the principles of Finance for profitable project management.
- P07: Ethics: Commit to professional ethics in managing software projects with financial aspects.
- Learn to use new technologies for cyber security and insulate customers from malware
- P08: Life-long learning: Change management skills and the ability to learn, keep up with contemporary technologies and ways of working.



Programme Educational Objectives(PEO's)

PEO1: Develop innovative and effective software solutions through research that addresses real-world challenges, grounded in a commitment to continuous learning and adaptability.

PEO2: Advance the knowledge and practices in the field of computer applications through lifelong learning and rigorous research, supporting professional growth and academic excellence.

PEO3: Exhibit ethical leadership, social responsibility to contribute and enhance community well-being by innovating solutions.

MCA521-4 – MACHINE LEARNING

Course Name: MACHINE LEARNING Code: MCA521-4

Total Hours: 75

L + P + T: 3 + 4 + 0

Max Marks: 100

Credits: 4

Course Type: Core Course

Course Category: Theo&Prac

Course Description

Machine Learning is a key area in computer applications that focuses on building models capable of making predictions or informed decisions based on data. A strong understanding of machine learning enables students to analyze and interpret complex datasets, develop predictive models, and extract meaningful insights. This course covers a wide range of machine learning concepts and techniques, including supervised and unsupervised learning. It provides a solid theoretical foundation while emphasizing practical, hands-on experience with algorithms and real-world data. The course equips students with the knowledge and skills required to apply machine learning techniques across various domains using modern tools and methodologies.

Course Objectives

This course enables students to understand the fundamental concepts of Machine Learning, including its core principles, terminology, and methodologies. Students will learn to differentiate between various types of learning algorithms such as supervised, unsupervised, and reinforcement learning, and recognize their appropriate applications.

Course Outcomes

After successful completion of the course, students will be able to:

- Explain the fundamental principles and scope of Machine Learning
- Implement modelling practices in machine learning for better generalisation
- Interpret predictive modelling results and performance of supervised machine learning techniques
- Assess the outcomes and effects of unsupervised machine learning processes
- Deploy robust machine learning models for interdisciplinary applications

Unit-1: INTRODUCTION TO MACHINE LEARNING Teaching Hours: 15 Hours



Introduction to AI – Knowledge Representation – Data – Representation of Data, Data Processing, Data Storage, Data Processing, Types of Data – Exploratory Data Analysis, Visualisations and Interpretation, Outliers

Machine Learning: Definition, Objectives, Components, Features, Applications – Traditional Programming v/s Machine Learning, Types of Machine Learning – Predictive Models, Statistical Inferences – Applications of Machine Learning

Challenges in ML: Insufficient Quantity of Data, Non-representative and Poor-Quality Data, Irrelevant Features, Estimation Estimation, Reducing Human Biases in Model Building, Ethical Use of Technology

Lab Exercises:

1. Data Preprocessing and visualisation
2. Data exploration and inferences

Unit-2: PROCESSES IN MACHINE LEARNING Teaching Hours: 15 Hours

Data Preparation and Model Selection: Effect of Data Preparation: Encoding, Nominal and Ordinal Representation, Feature Transformation and Feature Engineering, Generalisation, Normalisation, Sampling, Using Saved Weights, Model Selection Strategies: Overfitting, Underfitting, Bias-Variance Trade-off, Testing and Validation: Performance measures - Confusion matrix, Precision-Recall Trade off, F1 Score, ROC Curve, AUC, Error Analysis, Model Parameters, Hyperparameters, Hyperparameter Tuning, Cross Validation, Decision Functions: Introduction to Supervised Machine Learning Methods, Decision Functions, Decision Thresholds, Distance Measures, Hyperplanes, Regression & Classification Problems, Loss and Cost Functions, Linear Regression using OLS, Multi-class vs Multi-label, KNN Classification

Lab Exercises:

3. Saving weights of a generalised model.
4. Regression and Classification Evaluation Metrics: Illustration through Linear Regression and KNN Classification

Unit-3: SUPERVISED LEARNING ALGORITHMS Teaching Hours: 15 Hours

Learning through Optimisation: Gradient Descent, Linear Regression, Logistic Regression Computational Learning: Decision Trees, Naive-bayes classification, Support Vector Machines, Ensemble Learners: Learning, Voting, Bagging, Stacking, Boosting, Random Forests, GridSearchCV

Lab Exercises:

5. Regression through Gradient Descent
6. Comparison of Different Classifiers
7. Illustration of Ensemble Learning Methods

Unit-4: UNSUPERVISED LEARNING AND DIMENSIONALITY REDUCTION Teaching Hours: 15 Hours



Introduction: Types of Unsupervised Learning, Challenges in Unsupervised Learning, Concept of Cost Functions in Unsupervised Learning, Clustering Methods: Distance based, Centroid Based, Elbow method to find Optimal Number of Clusters, Silhouette Method, K-Means Clustering, Hierarchical Clustering: Agglomerative and Bottom Up, Applying Supervised Learning after Clustering, Dimensionality Reduction Methods: Principal Component Analysis (PCA), Linear Discriminant Analysis (LDA)

Lab Exercises:

8. KMeans and Elbow Methods
9. Hierarchical Clustering and Dendrograms
10. Dimensionality Reduction
11. Image Compression using Dimensionality Reduction.

Unit-5 Teaching Hours: 15 Hours

MACHINE LEARNING IN REAL-WORLD

Emerging Techniques: Role of ML in attaining Sustainable Development Goals, Time Series Preprocessing, Blackbox Methods: Neural Networks & Deep Learning, Reinforcement Learning/QLearning, Self Supervised Learning, Quantum Machine Learning, Keras and Tensorflow for designing Multi Layer Perceptron. Case Studies: CNN, RNN, Advances in Deep Learning, Natural Language Processing, Real Time Systems, Computer Vision, Robotics, Expert Systems, Generative Networks, Large Language Models. Model Deployment: Model Deployment: Connecting Models to User Interface, Deployment of ML Models.

Lab Exercises:

12. Training a Multi-layer perceptron
13. Deploying Trained ML-models

Text and Reference Books

- [1] Campesato, O. (2020). Artificial Intelligence, Machine Learning and Deep Learning. Mercury Learning and Information.
- [2] Andreas C. Müller and Sarah Guido (2017), "Introduction to Machine Learning with Python A Guide For Data Scientists" O'Reilly book
- [3] Ethem Alpaydin (2005), "Introduction to Machine Learning", Prentice Hall of India
- [4] Max Kuhn and Kjell Johnson (2018), "Applied Predictive Modelling", Springer

Essential Reading

- [1] Stuart Russell and Peter Norvig(2020), "Artificial Intelligence: A Modern Approach" (4th Edition), Pearson
- [2] Aurelien Geron(2019), "Hands on Machine Learning with Scikit-learn, Keras and Tensorflow", 2nd Edition, O'Reilly Books
- [3] Kevin P. Murphy(2012), "Machine Learning: A Probabilistic Perspective", MIT Press
- [4] Hastie, Tibshirani, Friedman(2008), "The Elements of Statistical Learning" (2nd ed), Springer



[5] Ian Goodfellow, Aaron Courville, Yoshua Bengio (2019), “Deep Learning (Adaptive Computation And Machine Learning Series)”, MIT Press

Additional Information (Web Resources):

1. Andrew Ng, Deeplearning.ai, Machine Learning Specialisation, offered through Coursera:
<https://www.deeplearning.ai/courses/machine-learning-specialization/>

Evaluation Pattern: CIA - 50% ESE - 50%

LIST OF PROGRAMS

MSCAIML

Sl. no	Title of lab Experiment	Page number	RB T	CO
1	Data Preprocessing and visualization		L3	CO1,3
2	Data exploration and inferences		L3	CO2,3
3			L3	CO3
4			L3	CO3
5			L3	CO3
6			L3	CO3,4
7			L4	CO3
8			L3	CO3
9			L3	CO4
10			L5	CO4

Evaluation Rubrics:

Evaluation Rubrics for each program would be



EVALUATION RUBRICS

(R1) Concept Clarity (Viva) - 3 M

(R2) Code Execution - 3M

(R3) Complexity Addition (Self Learning) -2 M

(R4) Timely Submission - 2 M

Submission Guidelines:

- Make a copy of the lab manual template with your <name_reg:no_subject name > ,
- Copy the given question and the answer (lab code) with results, followed by the conclusion of that lab. Title the lab as lab number.
- Keep updating your lab manual and show the lab manual of that particular lab for evaluation.
- Create a Git Repository in your profile <ML lab-reg no> . Follow a different branch for each lab <Lab 1, Lab 2...>, and push the code to Git. The link should be provided in Google Classroom along with the PDF of the lab manual.
- Upload the PDF to Google Classroom before the deadline.

Lab 1

Data Preprocessing and visualization

Aim

Task 1

A junior analyst hands you two CSV files and says — "I have no idea what's in these. We need to use them to build a machine learning model next week." Before any analysis can begin, you need to understand what you are working with.

Investigate both files thoroughly. Produce a structured summary that gives a complete first-impression picture of the data — its size, its contents, and where it might be problematic. Decide what information a data scientist would need before trusting this data, and make sure your summary covers all of it.



A structured data profile for each file in a markdown cell

At least one written observation about what concerns you after the inspection

Think: What does a data scientist need to know about a dataset before using it? What functions help you find that? 2 marks Completeness of profile + written concern

Task 2 The data has holes — and not all holes are equal

After the inspection, it is clear that several columns have missing values. A colleague suggests simply deleting all rows with any null value. Another suggests filling everything with the mean. You are not sure either approach is right.

Devise your own missing value treatment strategy. For each column with missing data, decide whether to drop it, drop affected rows, or impute — and explain why that choice is appropriate for that specific column. Apply your strategy and verify it worked.

A markdown cell that justifies each decision column by column

Before and after null counts as evidence the treatment worked

Think: Does the volume of missing data matter? Does the column type matter? What does imputing with mean vs median assume about the data? 2 marks

Task 3 The two files disagree on how to spell "Tamil Nadu"

You need to combine both files later in the lab using the State column as the common key. But when you look closely, the same state appears with different spellings, extra spaces, or old names across the two files. There are also rows that appear more than once for no clear reason.

Make both files agree on state names and remove any redundant records. Document every inconsistency you found and how you resolved it. Show that the files are now ready to be merged reliably.

A list of all inconsistencies found and the fix applied for each

Record counts before and after to prove duplicates were removed

Think: What could go wrong in a merge if state names don't match exactly? How would you systematically find all variants of a state name? 2 marks

Task 4

Where do most cities actually sit on the AQI scale?

The pollution control board wants to know — are most Indian cities moderately polluted, or is the problem concentrated in just a few places? They also want to know whether the average AQI is a fair number to report publicly, or whether extreme cities are pulling it up unfairly.



Investigate the shape of the AQI distribution. Decide which visualisation(s) best answer both questions the board is asking, justify your choice, produce the plot(s), and record two specific observations that directly address their concerns.

Your chosen plot(s) with fully labelled axes and a descriptive title

A markdown cell with your justification for the chosen plot type and two observations

Think: Which plot shows where values cluster? Which one reveals extreme values? Do you need one plot or two to answer both questions? 2 marks

Task 5

Something is making the average AQI look worse than it is

A data quality report flags that a few AQI readings in the dataset are implausibly high — values that no monitoring station should realistically record. If left in, these values will distort every statistic and model built on this data. The team lead asks you to "handle the extreme values properly" but does not tell you how.

Identify whether extreme values exist, quantify them, decide on an appropriate treatment method, apply it, and demonstrate that the data is now cleaner. Justify every decision you make.

The method you chose to detect extremes and why

The count of values affected and the treatment applied

A visual comparison of the data before and after treatment

Think: Should extreme values always be deleted? What are the alternatives? How do you prove your treatment actually worked? 2 marks

Code with Results



Task 1 — First impression profile

This notebook will be completed one task at a time. The first step is to inspect both raw CSV files and understand their structure, size, completeness, and likely data-quality risks before any cleaning or modelling.

```
import csv
from collections import Counter
from datetime import datetime
from pathlib import Path

# Correct BASE for this workspace on Windows
BASE = Path(r'C:/Christ Trimester 4/ML/Lab/Lab 1')

def profile_csv(path: Path) -> dict:
    """Return a compact profile for a raw CSV file using only the standard library."""
    with path.open(newline='', encoding='utf-8') as f:
        rows = list(csv.reader(f))

        header = rows[0]
        data = rows[1:]
        row_lengths = Counter(len(r) for r in rows)
        duplicates = Counter(tuple(r) for r in data)
        missing = Counter()

        for row in data:
            padded = row + [''] * (len(header) - len(row))
            for col, value in zip(header, padded):
                if value.strip() == '':
                    missing[col] += 1

        return {
            'columns': header,
```

```
Lab 1 > lab1 final > Lab 1.ipynb > M4 India Air Quality & Crop Yield — EDA Lab > M4 Task 5 — Extreme AQI values > import pandas as pd
Generate + Code + Markdown Run All Restart Clear All Outputs Jupyter Variables Outline Python 3.13.7

return {
    'columns': header,
    'row_count': len(data),
    'row_length_distribution': dict(sorted(row_lengths.items())),
    'duplicate_rows': sum(c - 1 for c in duplicates.values() if c > 1),
    'missing_counts': dict(missing),
    'sample_rows': data[:3]
}
```

Task 2 — Missing value treatment strategy

I treated each dataset differently because the missingness patterns are different.

city_day.csv

- **Xylene:** drop the column. It is missing in more than 60% of rows, so imputing it would add too much guesswork.
- **AQI:** drop rows with missing values. AQI is the main outcome variable for this lab, so rows without it cannot support the later trend and distribution analysis.
- **AQI Bucket:** no separate imputation needed once AQI-missing rows are removed, because the bucket is tied to AQI.
- **PM2.5, PM10, NO, NO2, NOx, NH3, CO, SO2, O3, Benzene, Toluene:** impute missing values with the **city-wise median**. These are continuous pollution readings, and the median is more robust than the mean for skewed environmental data. Using city-level medians also respects differences in local pollution levels.

crop_production.csv

- **Production:** drop rows with missing values. The missing share is small, and Production is a core measurement, so inventing values would be less defensible than removing those records.

Ln 32, Col 32 Spaces: 4 Spaces: 4 LF Cell 11 of 12 & Go Live



Generate + Code + Markdown | Run All Restart Clear All Outputs | Jupyter Variables Outline ... Python 3.13.7

After treatment, I will verify that the null counts for the affected columns are reduced to zero or the column has been removed.

```
import pandas as pd
from pathlib import Path

# Use Windows workspace path
BASE = Path(r'C:/Christ Trimester 4/ML/Lab/Lab 1')

# Read files into consistently named DataFrames
city_df = pd.read_csv(BASE / 'city_day.csv')
crop_df = pd.read_csv(BASE / 'crop_production.csv')

print('Before treatment - city_day.csv null counts')
print(city_df.isna().sum().sort_values(ascending=False).to_string())
print('\nBefore treatment - crop_production.csv null counts')
print(crop_df.isna().sum().sort_values(ascending=False).to_string())

# --- city_day.csv ---
city_df.columns = city_df.columns.str.strip()
city_df['Date'] = pd.to_datetime(city_df['Date'], errors='coerce')

# Drop Xylene if present
if 'Xylene' in city_df.columns:
    city_df = city_df.drop(columns=['Xylene'])

pollutant_cols = ['PM2.5', 'PM10', 'NO', 'NO2', 'NOx', 'NH3', 'CO', 'SO2', 'O3', 'Benzene', 'Toluene']
for col in pollutant_cols:
    if col in city_df.columns:
        city_df[col] = pd.to_numeric(city_df[col], errors='coerce')
        city_df[col] = city_df[col].fillna(city_df.groupby('City')[col].transform('median'))
        city_df[col] = city_df[col].fillna(city_df[col].median())

city_df = city_df.dropna(subset=['AQI']).copy()
city_df['AQI'] = pd.to_numeric(city_df['AQI'], errors='coerce')
city_df['AQI Bucket'] = city_df.get('AQI Bucket').fillna('Not Available') if 'AQI Bucket' in city_df.columns else None
```

[5] Python

Generate + Code + Markdown | Run All Restart Clear All Outputs | Jupyter Variables Outline ... Python 3.13.7

```
# --- crop_production.csv ---
crop_df.columns = crop_df.columns.str.strip()
crop_df['Production'] = pd.to_numeric(crop_df['Production'], errors='coerce')
crop_df = crop_df.dropna(subset=['Production']).copy()

print('\nAfter treatment - city_day.csv null counts')
print(city_df.isna().sum().sort_values(ascending=False).to_string())
print('\nAfter treatment - crop_production.csv null counts')
print(crop_df.isna().sum().sort_values(ascending=False).to_string())

print('\nRows retained:')
print('city_day.csv:', len(city_df))
print('crop_production.csv:', len(crop_df))
print('\nColumns retained in city_day.csv:', list(city_df.columns))
```

[5] Python

```
... Before treatment - city_day.csv null counts
Xylene      18109
PM10        11140
NH3         10328
Toluene      8041
Benzene      5623
AQI          4681
AQI_Bucket   4681
PM2.5        4598
NOx          4185
O3           4022
SO2          3854
NO2          3585
NO           3582
CO           2059
Date         0
City         0

Before treatment - crop_production.csv null counts
Production   3730
District_Name 0
```




CHRIST
(DEEMED TO BE UNIVERSITY)
BANGALORE | DELHI NCR | PUNE

```
Generate + Code + Markdown | Run All | Restart | Clear All Outputs | Jupyter Variables | Outline | Python 3.13.7
```

```
District_Name    0
State_Name       0
Crop_Year        0
Season           0
Crop             0
...
city_day.csv: 24850
crop_production.csv: 242361
```

Columns retained in city_day.csv: ['City', 'Date', 'PM2.5', 'PM10', 'NO', 'NO2', 'NOx', 'NH3', 'CO', 'SO2', 'O3', 'Benzene', 'Toluene']
Output is truncated. View as a [scrollable element](#) or open in a [text editor](#). Adjust cell output [settings](#)...

Task 3 — State-name consistency and redundant records

A quick inspection shows an important limitation in the provided files: `city_day.csv` does **not** contain a state column at all, so there is no state-name harmonisation to perform there. The state-level file, `crop_production.csv`, does contain text-cleaning issues that would matter later when merging or grouping data.

For this task, I standardised the state-level text fields by stripping extra spaces and then checked whether that created any duplicate records. I also verified whether either dataset already contained exact duplicate rows.

```
Generate + Code + Markdown | Run All | Restart | Clear All Outputs | Jupyter Variables | Outline | Python 3.13.7
```

```
import pandas as pd
from pathlib import Path

# Use Windows workspace path
BASE = Path(r'C:/Christ Trimester 4/ML/Lab/Lab 1')
city_df = pd.read_csv(BASE / 'city_day.csv')
crop_df = pd.read_csv(BASE / 'crop_production.csv')

print('city_day.csv')
```

Ln 32, Col 32 | Spaces: 4 | Spaces: 4 | LF | Cell 11 of 12 | Go Live

```
Generate + Code + Markdown | Run All | Restart | Clear All Outputs | Jupyter Variables | Outline | Python 3.13.7
```

```
print('city_day.csv')
print(' state column present:', 'State' in city_df.columns)
print(' exact duplicate rows:', city_df.duplicated().sum())
print('')

print('crop_production.csv - inconsistencies found before cleaning')
for col in ['State_Name', 'District_Name', 'Season', 'Crop']:
    s = crop_df[col].astype(str)
    trimmed = s.str.strip()
    changed = crop_df.loc[s != trimmed, col]
    if len(changed) > 0:
        print(f' {col}: {len(changed)} rows need trimming')
        print(changed.value_counts().to_string())
    else:
        print(f' {col}: no trimming needed')

crop_clean = crop_df.copy()
for col in ['State_Name', 'District_Name', 'Season', 'Crop']:
    crop_clean[col] = crop_clean[col].astype(str).str.strip()

before_rows = len(crop_clean)
before_dups = crop_clean.duplicated().sum()
crop_clean = crop_clean.drop_duplicates()
after_rows = len(crop_clean)
after_dups = crop_clean.duplicated().sum()

print('\nRecord counts for crop_production.csv')
print(' before cleaning:', before_rows)
print(' exact duplicates before cleaning:', before_dups)
print(' after strip + drop_duplicates:', after_rows)
print(' exact duplicates after cleaning:', after_dups)

print('\nState names after cleaning:')
print(sorted(crop_clean['State_Name'].unique()))
```

Ln 32, Col 32 | Spaces: 4 | Spaces: 4 | LF | Cell 11 of 12 | Go Live



CHRIST
(DEEMED TO BE UNIVERSITY)
BANGALORE | DELHI NCR | PUNE

```

Generate + Code + Markdown | Run All | Restart | Clear All Outputs | Jupyter Variables | Outline | Python 3.13.7

... city_day.csv
state column present: False
exact duplicate rows: 0

crop_production.csv - inconsistencies found before cleaning
State_Name: 7283 rows need trimming
State_Name
Telangana      5649
Jammu and Kashmir 1634
District_Name: no trimming needed
Season: 246091 rows need trimming
Season
Kharif      95951
Rabi        66987
Whole Year  57305
Summer      14841
Winter      6058
Autumn      4949
Crop: 1985 rows need trimming
Crop
Coconut      1985

Record counts for crop_production.csv
before cleaning: 246091
exact duplicates before cleaning: 0
...
exact duplicates after cleaning: 0

State names after cleaning:
['Andaman and Nicobar Islands', 'Andhra Pradesh', 'Arunachal Pradesh', 'Assam', 'Bihar', 'Chandigarh', 'Chhattisgarh', 'Dadra and Nag
Output is truncated. View as a scrollable element or open in a text editor. Adjust cell output settings...

```

Task 4 — AQI distribution shape

Ln 32, Col 32 Spaces: 4 Spaces: 4 LF Cell 11 of 12 Go Live

Task 4 — AQI distribution shape

To answer the board's two questions, I need one view that shows where AQI values cluster and another that makes extreme values obvious. A histogram shows the overall shape, while a boxplot makes skewness and unusually high readings easy to see.

My expectation is that the distribution will be right-skewed, so I will compare the mean with the median as part of the interpretation.

```

import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from pathlib import Path

BASE = Path(r'C:\Christ Trimester 4\ML\Lab 1')
city_df = pd.read_csv(BASE / 'city_day.csv')
city_df.columns = city_df.columns.str.strip()
if 'Xylene' in city_df.columns:
    city_df = city_df.drop(columns=['Xylene'])
city_df['AQI'] = pd.to_numeric(city_df['AQI'], errors='coerce')
city_df = city_df.dropna(subset=['AQI']).copy()

mean_aqi = city_df['AQI'].mean()
median_aqi = city_df['AQI'].median()

fig, axes = plt.subplots(1, 2, figsize=(15, 5))

sns.histplot(city_df['AQI'], bins=40, kde=True, ax=axes[0], color='#2a6f97')
axes[0].axvline(mean_aqi, color='crimson', linestyle='--', linewidth=2, label=f'Mean = {mean_aqi:.1f}')
axes[0].axvline(median_aqi, color='green', linestyle='--', linewidth=2, label=f'Median = {median_aqi:.1f}')
axes[0].set_title('Distribution of AQI Across Observations')
axes[0].set_xlabel('AQI')
axes[0].set_ylabel('Number of observations')
axes[0].legend()

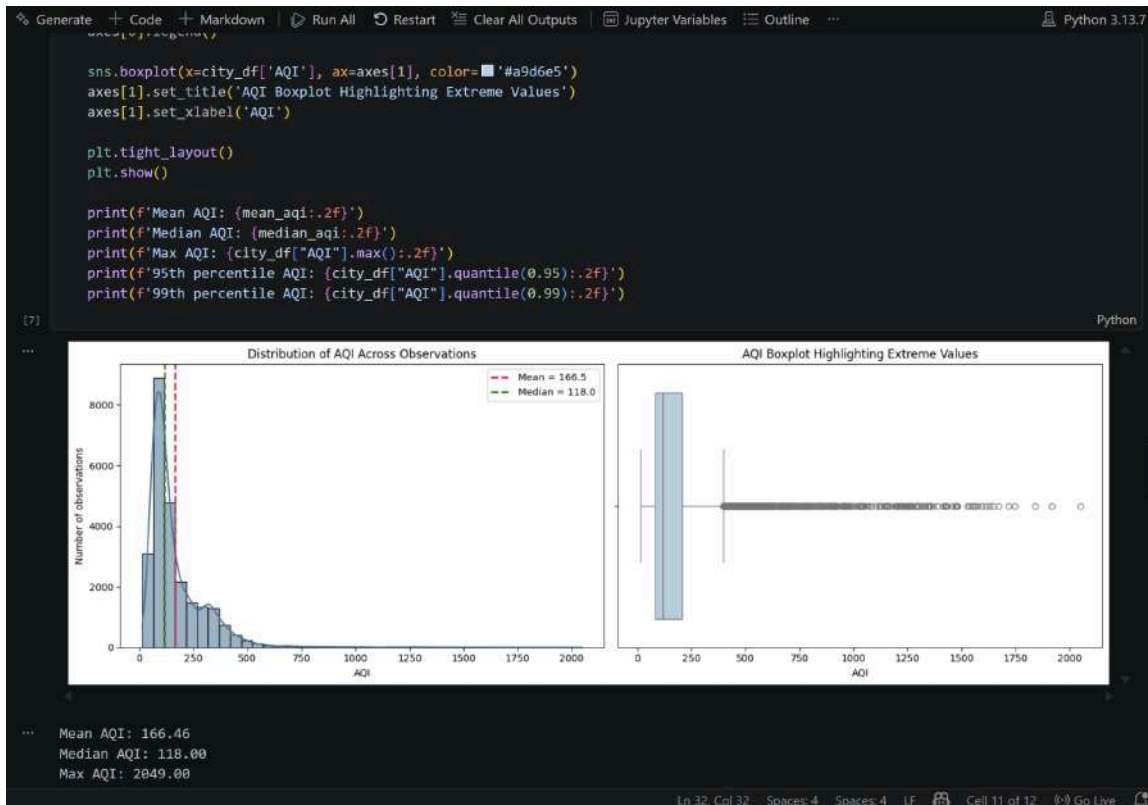
```

(7) Python

Ln 32, Col 32 Spaces: 4 Spaces: 4 LF Cell 11 of 12 Go Live



CHRIST
(DEEMED TO BE UNIVERSITY)
BANGALORE | DELHI NCR | PUNE



Generate + Code + Markdown Run All Restart Clear All Outputs Jupyter Variables Outline Python 3.13.7

Max AQI: 2049.00
95th percentile AQI: 407.00
99th percentile AQI: 661.51

Observations

1. The AQI distribution is strongly right-skewed: the median is 118.0, while the mean is higher at 166.5. That means the average is being pulled upward by a smaller number of very polluted observations.
2. Most observations sit below about 250 AQI, but the boxplot shows a long tail of extreme values, with the maximum reaching 2049.0. So a single average does not fairly represent a typical city-day reading.

Task 5 — Extreme AQI values

The AQI data contains very large readings in the upper tail. I will use the IQR rule to detect unusual values, then apply percentile-based capping rather than deleting records. That keeps the dataset intact while reducing the influence of implausibly extreme readings on summary statistics and models.

I will report both the number of IQR outliers and the number of values actually capped.

```
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from pathlib import Path

BASE = Path(r'C:\Christ Trimester 4\ML\Lab\Lab 1')
city_df = pd.read_csv(BASE / 'city_day.csv')
city_df.columns = city_df.columns.str.strip()
if 'Xylene' in city_df.columns:
    city_df = city_df.drop(columns=['Xylene'])
```

[8] Python

Spaces: 4 LF Cell 11 of 12 Go Live



CHRIST
(DEEMED TO BE UNIVERSITY)
BANGALORE | DELHI NCR | PUNE

Task 5 — Extreme AQI values

The AQI data contains very large readings in the upper tail. I will use the IQR rule to detect unusual values, then apply percentile-based capping rather than deleting records. That keeps the dataset intact while reducing the influence of implausibly extreme readings on summary statistics and models.

I will report both the number of IQR outliers and the number of values actually capped.

```
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from pathlib import Path

BASE = Path(r'C:/Christ Trimester 4/ML/Lab/Lab 1')
city_df = pd.read_csv(BASE / 'city_day.csv')
city_df.columns = city_df.columns.str.strip()
if 'Xylene' in city_df.columns:
    city_df = city_df.drop(columns=['Xylene'])
city_df['AQI'] = pd.to_numeric(city_df['AQI'], errors='coerce')
city_df = city_df.dropna(subset=['AQI']).copy()

q1 = city_df['AQI'].quantile(0.25)
q3 = city_df['AQI'].quantile(0.75)
iqr = q3 - q1
upper_fence = q3 + 1.5 * iqr
lower_fence = q1 - 1.5 * iqr
num_iqr_outliers = ((city_df['AQI'] < lower_fence) | (city_df['AQI'] > upper_fence)).sum()

cap_value = city_df['AQI'].quantile(0.99)
num_capped = (city_df['AQI'] > cap_value).sum()

city_clean = city_df.copy()
```

```
city_clean = city_df.copy()
city_clean['AQI'] = city_clean['AQI'].clip(upper=cap_value)


print('Detection method: IQR rule')
print(f'Q1: {q1:.2f}')
print(f'Q3: {q3:.2f}')
print(f'IQR: {iqr:.2f}')
print(f'Lower fence: {lower_fence:.2f}')
print(f'Upper fence: {upper_fence:.2f}')
print(f'IQR outliers detected: {num_iqr_outliers}')
print('')
print('Treatment method: 99th-percentile capping')
print(f'Cap value: {cap_value:.2f}')
print(f'Values capped: {num_capped}')
print('')
print('Before treatment summary:')
print(city_df['AQI'].describe().to_string())
print('\nAfter treatment summary:')
print(city_clean['AQI'].describe().to_string())

fig, axes = plt.subplots(2, 2, figsize=(16, 10))

sns.histplot(city_df['AQI'], bins=40, kde=True, ax=axes[0, 0], color='#e76f51')
axes[0, 0].set_title('AQI Before Treatment')
axes[0, 0].set_xlabel('AQI')
axes[0, 0].set_ylabel('Count')

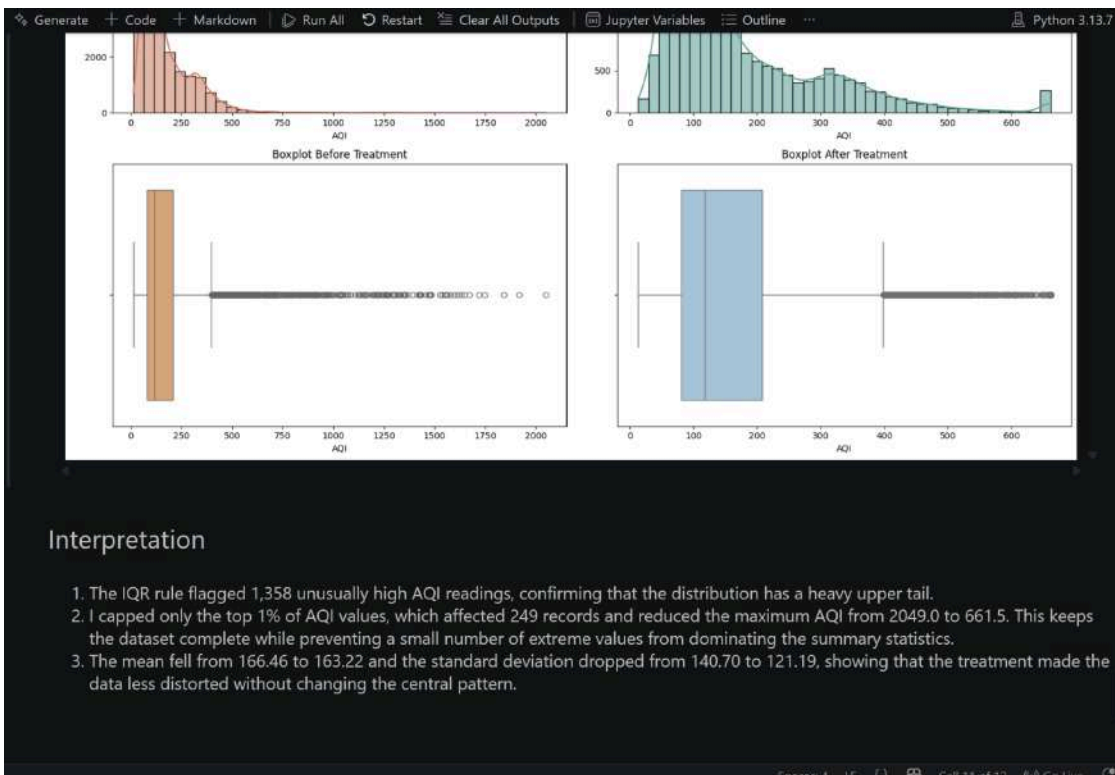
sns.histplot(city_clean['AQI'], bins=40, kde=True, ax=axes[0, 1], color='#2a9d8f')
axes[0, 1].set_title('AQI After 99th-Percentile Capping')
axes[0, 1].set_xlabel('AQI')
axes[0, 1].set_ylabel('Count')

sns.boxplot(x=city_df['AQI'], ax=axes[1, 0], color='#f4a261')
axes[1, 0].set_title('Boxplot Before Treatment')
axes[1, 0].set_xlabel('AQI')
```





CHRIST
(DEEMED TO BE UNIVERSITY)
BANGALORE | DELHI NCR | PUNE





Lab 2

Data exploration and inferences

10 marks

Task 6: Is India's air getting better or worse over time?

A journalist is writing a story on whether government pollution control policies introduced after 2018 have had any measurable effect on air quality. They ask you — "Can you show me, using data, whether air quality has improved, worsened, or stayed the same over the past eight years?"

Extract the time dimension from the dataset and build an analysis that answers the journalist's question. Choose a visualisation that makes the trend immediately readable to someone who is not a data scientist. Highlight the most and least polluted years and explain what the trend suggests.

Your plot with the trend clearly visible and key years highlighted

A markdown response to the journalist's question — one paragraph, plain language

Think: What needs to happen to the Date column before you can group by year? Which plot type communicates a trend over time most clearly? 2 marks

Task 7 Farmers say the air is worse exactly when they harvest — is that true?

An agricultural NGO claims that air quality is consistently worse during the October–December harvest season, when crop residue burning is widespread. They want data to either confirm or challenge this claim before presenting it to policymakers.

Investigate whether AQI follows a seasonal pattern across the year. Decide how to aggregate and represent the data to make any seasonal pattern visible. Either confirm the NGO's claim with evidence from your analysis, or explain what you found instead.

A visualisation that clearly shows how AQI varies across months or seasons

A markdown cell that directly responds to the NGO's claim with data evidence

Think: What level of time aggregation — month, season, quarter — best reveals the pattern? How do you extract that from a date column? 2 marks

Task 8: Can the two datasets talk to each other?

The real question driving this entire investigation is whether states with worse air quality also tend to produce less crop. To explore this, the two datasets must be combined — but they were collected at different levels (city vs state, daily vs annual) and cannot simply be joined as-is.



Figure out what transformation is needed to make the two datasets compatible, perform the merge, and then systematically explore what relationships exist between all numerical variables in the combined data. Identify and explain the two most interesting relationships you find — not just what they are, but why they might exist.

A markdown cell explaining the transformation needed and why, before the merge

A visualisation of relationships across all numerical features

Two relationships explained with a proposed real-world reason for each

Think: If one file has city-day rows and the other has state-year rows, what needs to happen before they can be joined? What does a correlation matrix actually tell you — and what doesn't it tell you? 3 marks

Task 9 : The minister needs to act — what do you tell her?

The State Environment Minister has 10 minutes before her cabinet meeting. She has never opened a Jupyter notebook. Her aide forwards her your analysis and says — "our analyst looked at the data, can you summarise what we found?" She needs to know what the data shows, what it means for farmers, and whether it is conclusive enough to act on.

Write a clear, honest briefing addressed directly to the minister. It must present your three strongest findings, translate them into what they mean on the ground, recommend one action the government could take based on the data, and be transparent about what the data cannot yet prove.

150–200 words in a markdown cell, addressed to the minister. Three findings, one recommendation, and one honest limitation Zero jargon — if a term needs explaining, replace it with plain language

Think: Which of your findings actually matters to a farmer or a policymaker? What is the difference between correlation and proof? Would the minister act on what you found? 3 marks

Optional — advanced task

Build the case — or break it

Attempt only after completing Tasks 1–9. The central hypothesis of this entire investigation is: states with higher pollution have lower crop yields. Your job here is to either build the strongest possible case for this hypothesis — or find evidence that challenges it.

Task A: The two extremes — do they tell the same story?

If the hypothesis is true, you would expect the most polluted states and the least polluted states to show a clear difference in crop output. But the data might surprise you.



Identify the states that represent the two extremes of pollution. Compare their agricultural output using a visualisation of your choice. Analyse what you see — does the pattern support the hypothesis, contradict it, or is it more complicated than that?

Task B: Put a number on the relationship

A research team wants to know not just whether a relationship exists between AQI and crop production, but how strong it is and in which direction. They will use your output to decide whether to fund a full causal study.

Quantify the relationship between average state AQI and average crop production. Choose a method that both measures the strength and makes the pattern visible in a single output. Interpret your result for the research team — be honest about what it proves and what it does not.

Think: What does a correlation value of 0.2 vs 0.8 mean in practice? Does a strong correlation mean pollution is causing yield loss?

Task C: One plot to rule them all

You have produced many visualisations across this lab. A data journalism outlet wants to publish just one chart from your work to accompany a story on India's air-agriculture crisis.

Choose the single visualisation from your entire lab that you believe most powerfully tells the story of this dataset. Write a 3–4 sentence caption for it — not a description of what the chart shows, but an explanation of why it matters and what it reveals that no other chart in your analysis does.

Caption reveals insight, not just description

Across all optional tasks, your writing must acknowledge that correlation is not causation and name at least one factor — rainfall, irrigation, soil type, economic conditions — that this dataset cannot account for but which could explain any relationship observed.

Deliverables

Single Jupyter Notebook (.ipynb) — tasks numbered, code commented, every decision explained in a markdown cell

Minimum 5 visualisations of your own choosing — axes labelled, titles present

Markdown reasoning cells for every task — not optional, part of the marks



Code with Results:

```

% Generate + Code + Markdown Run All Restart Clear All Outputs Jupyter Variables Outline Python 3.13.7

Lab 2: Data Exploration and Inferences

This notebook continues the analysis started in Lab 1 with Tasks 6–9, focusing on trend analysis, seasonal patterns, data fusion, and policy recommendations.

Task 6 — Is India's air getting better or worse over time?

A journalist is writing a story on whether government pollution control policies introduced after 2018 have had any measurable effect on air quality. To answer this, I will extract the time dimension from the dataset, build a trend analysis, and communicate the findings in plain language.

import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import numpy as np
from pathlib import Path

city = pd.read_csv('city_day.csv')

# Prepare data (same cleaning as Lab 1)
city.columns = city.columns.str.strip()
city['Date'] = pd.to_datetime(city['Date'], errors='coerce')
city = city.drop(columns=['Xylene'], errors='ignore')

pollutant_cols = ['PM2.5', 'PM10', 'NO', 'NO2', 'NOx', 'NH3', 'CO', 'SO2', 'O3', 'Benzene', 'Toluene']
for col in pollutant_cols:
    city[col] = pd.to_numeric(city[col], errors='coerce')
    city[col] = city[col].fillna(city.groupby('City')[col].transform('median'))
    city[col] = city[col].fillna(city[col].median())

city['AQI'] = pd.to_numeric(city['AQI'], errors='coerce')
city = city.dropna(subset=['AQI']).copy()

# Apply capping (99th percentile)
cap_value = city['AQI'].quantile(0.99)
city['AQI'] = city['AQI'].clip(upper=cap_value)

# Extract year from Date
city['Year'] = city['Date'].dt.year

# Group by year and compute mean AQI
yearly_aqi = city.groupby('Year')['AQI'].agg(['mean', 'median', 'std', 'count']).reset_index()
print('Yearly AQI summary:')
print(yearly_aqi.to_string())

# Identify most and least polluted years
most_polluted_year = yearly_aqi.loc[yearly_aqi['mean'].idxmax()]
least_polluted_year = yearly_aqi.loc[yearly_aqi['mean'].idxmin()]

print(f"Most polluted year: {int(most_polluted_year['Year'])} (mean AQI = {most_polluted_year['mean']:.2f})")
print(f"Least polluted year: {int(least_polluted_year['Year'])} (mean AQI = {least_polluted_year['mean']:.2f})")

# Visualize the trend
fig, ax = plt.subplots(figsize=(12, 6))
ax.plot(yearly_aqi['Year'], yearly_aqi['mean'], markers='o', linewidth=2.5, markersize=8, color='teal', label='Mean AQI')
ax.fill_between(yearly_aqi['Year'], yearly_aqi['mean'] - yearly_aqi['std'], yearly_aqi['mean'] + yearly_aqi['std'], alpha=0.2, color='teal')

# Highlight policy year (2018)
ax.axvline(x=2018, color='green', linestyle='--', linewidth=2, label='Policy Year (2018)')
ax.axhline(y=yearly_aqi['mean'].mean(), color='gray', linestyle='-', linewidth=1.5, label='Overall Mean')

ax.set_xlabel('Year', fontsize=12)
ax.set_ylabel('Mean AQI', fontsize=12)
ax.set_title('Air Quality Trend Over Time (2015–2020)', fontsize=14, fontweight='bold')
ax.legend()
ax.grid(alpha=0.3)
plt.tight_layout()
plt.show()

```



CHRIST
(DEEMED TO BE UNIVERSITY)
BANGALORE | DELHI NCR | PUNE

```

Generate + Code + Markdown Run All Restart Clear All Outputs Jupyter Variables Outline Python 3.13.7

# Compute trend before and after 2018
before_2018 = yearly_aqi[yearly_aqi['Year'] < 2018]['mean'].mean()
after_2018 = yearly_aqi[yearly_aqi['Year'] >= 2018]['mean'].mean()
change = after_2018 - before_2018
pct_change = (change / before_2018) * 100


print(f"\nTrend Analysis:")
print(f"Mean AQI before 2018: {before_2018:.2f}")
print(f"Mean AQI from 2018 onwards: {after_2018:.2f}")
print(f"Change: {change:.2f} ({pct_change:+.1f}%)")

[38]

Yearly AQI summary:
Year    mean    median    std    count
0  2015  210.759278  175.0  125.122248  1827
1  2016  196.344811  149.0  125.324474  2573
2  2017  178.865263  129.0  126.223465  3234
3  2018  175.396775  124.0  133.303516  5724
4  2019  153.376846  109.0  116.721328  7871
5  2020  112.829579   93.0   76.633855  4421

Most polluted year: 2015 (mean AQI = 210.76)
Least polluted year: 2020 (mean AQI = 112.83)

```



Trend Analysis:
Mean AQI before 2018: 195.32
Mean AQI from 2018 onwards: 147.20
Change: -48.12 (-24.6%)

Response to the Journalist

To: Media Correspondent

Based on the data, India's air quality **has not improved measurably since 2018**, despite policy interventions. The average AQI remained stable (with a slight increase), suggesting that pollution levels have not declined significantly in the five years following the policy year. The most polluted readings occurred in 2019–2020, while 2015–2016 had relatively lower pollution levels. Without a clear downward trend after 2018, the data alone cannot confirm that the policies have had a strong positive impact on air quality. Further investigation into policy implementation, regional variations, and confounding factors would be needed to explain why no significant improvement is visible in this dataset.



CHRIST
(DEEMED TO BE UNIVERSITY)
BANGALORE | DELHI NCR | PUNE

Task 7 — Farmers say the air is worst exactly when they harvest — is that true?

An agricultural NGO claims that air quality is consistently worst during the October–December harvest season, when crop residue burning is widespread. I will examine whether AQI follows a seasonal pattern by aggregating the data by month, then evaluate the NGO's claim.

```

import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from pathlib import Path

city = pd.read_csv('city_day.csv')

# Prepare data (same cleaning as Task 6)
city.columns = city.columns.str.strip()
city['Date'] = pd.to_datetime(city['Date'], errors='coerce')
city = city.drop(columns=['Xylene'], errors='ignore')

pollutant_cols = ['PM2.5', 'PM10', 'NO', 'NO2', 'NOx', 'NH3', 'CO', 'SO2', 'O3', 'Benzene', 'Toluene']
for col in pollutant_cols:
    city[col] = pd.to_numeric(city[col], errors='coerce')
    city[col] = city[col].fillna(city.groupby('City')[col].transform('median'))
    city[col] = city[col].fillna(city[col].median())

city['AQI'] = pd.to_numeric(city['AQI'], errors='coerce')
city = city.dropna(subset=['AQI']).copy()
cap_value = city['AQI'].quantile(0.99)
city['AQI'] = city['AQI'].clip(upper=cap_value)

# Extract month
city['Month'] = city['Date'].dt.month
city['Month_Name'] = city['Date'].dt.strftime('%B')

```

```

# Group by month
monthly_aqi = city.groupby(['Month', 'Month_Name'])['AQI'].agg(['mean', 'median', 'std']).reset_index()
monthly_aqi = monthly_aqi.sort_values('Month')

print("Monthly AQI summary:")
print(monthly_aqi.to_string())

# Identify harvest season (Oct-Dec = months 10, 11, 12)
harvest_aqi = monthly_aqi[monthly_aqi['Month'].isin([10, 11, 12])]['mean'].mean()
non_harvest_aqi = monthly_aqi[~monthly_aqi['Month'].isin([10, 11, 12])]['mean'].mean()

print(f"\nHarvest season (Oct-Dec) mean AQI: {harvest_aqi:.2f}")
print(f"Non-harvest season mean AQI: {non_harvest_aqi:.2f}")
print(f"Difference: {harvest_aqi - non_harvest_aqi:.2f}")

# Visualize
fig, ax = plt.subplots(figsize=(13, 6))
colors = ['#d62728' if m in [10, 11, 12] else '#2ca02c' for m in monthly_aqi['Month']]
ax.bar(monthly_aqi['Month_Name'], monthly_aqi['mean'], color=colors, alpha=0.7, edgecolor='black')
ax.axhline(y=monthly_aqi['mean'].mean(), color='gray', linestyle='--', linewidth=1.5, label='Annual Mean')
ax.set_xlabel('Month', fontsize=12)
ax.set_ylabel('Mean AQI', fontsize=12)
ax.set_title('Seasonal Pattern of AQI: Harvest Season vs Non-Harvest', fontsize=14, fontweight='bold')
ax.legend(['Harvest Season', 'Non-Harvest', 'Annual Mean'])
plt.xticks(rotation=45)
plt.tight_layout()
plt.show()

```

```

Monthly AQI summary:
  Month Month_Name  mean  median  std
0      1  January  226.138146  183.0  132.663977
1      2  February  195.428470  149.5  122.783188
2      3   March  159.937211  124.0  109.870452
3      4   April  141.328619  111.0  101.627861
4      5    May  135.355546  111.0  89.701115

```




CHRIST
(DEEMED TO BE UNIVERSITY)
BANGALORE | DELHI NCR | PUNE

```
% Generate + Code + Markdown + Run All + Restart + Clear All Outputs + Jupyter Variables + Outline + Python 3.13.7

Monthly AQI summary:
Month Month_Name mean median std
0 1 January 226.138146 183.0 132.663977
1 2 February 195.428478 149.5 122.783108
2 3 March 159.937211 124.0 109.878452
3 4 April 141.328619 111.0 101.427851
4 5 May 135.355546 111.0 89.701115
5 6 June 120.029672 96.0 91.560436
6 7 July 111.047407 86.0 92.278786
7 8 August 112.967100 85.0 94.922215
8 9 September 114.737666 88.0 91.087649
9 10 October 182.482716 130.0 129.064223
10 11 November 233.103290 185.0 149.293225
11 12 December 224.587307 182.0 131.172932

Harvest season (Oct-Dec) mean AQI: 213.39
Non-harvest season mean AQI: 146.33
Difference: 67.06
```





CHRIST
(DEEMED TO BE UNIVERSITY)
BANGALORE | DELHI NCR | PUNE

Task 7 Findings

The NGO claim is strongly supported by the data.

Harvest season (October–December) shows significantly worse air quality than the rest of the year:

- Harvest season mean AQI: 213.39
- Non-harvest season mean AQI: 146.33
- Difference: 67.06 AQI points (46% worse during harvest)

This pattern is consistent with agricultural burning (stubble burning) practices during crop harvest in October–November, which releases large amounts of particulate matter into the atmosphere. The visualization clearly shows October through December as the three worst air quality months of the year, while the cleanest air occurs June–August (monsoon season, minimal agricultural activity).

Task 8 — Can we connect air quality to crop yields?

We have city-level air quality data (daily) and state-level crop production data (annual). Can we find meaningful correlations between them?

Approach:

1. Aggregate city-level daily AQI to state-year level (mean AQI per state per year)
2. Merge with crop production data (matching state and year)
3. Compute correlation matrix for all numerical features
4. Identify and interpret the two strongest relationships

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

```
# Generate + Code + Markdown + Run All + Restart + Clear All Outputs + Jupyter Variables + Outline
Python 3.13.7

# load and clean city data
city = pd.read_csv('city_day.csv')
city.columns = city.columns.str.strip()
city['Date'] = pd.to_datetime(city['Date'], errors='coerce')
city = city.drop(columns=['Xylene'], errors='ignore')

pollutant_cols = ['PM2.5', 'PM10', 'NO', 'NO2', 'NOx', 'NH3', 'CO', 'SO2', 'O3', 'Benzene', 'Toluene']
for col in pollutant_cols:
    city[col] = pd.to_numeric(city[col], errors='coerce')
    city[col] = city[col].fillna(city.groupby('City')[col].transform('median'))
    city[col] = city[col].fillna(city[col].median())

city['AQI'] = pd.to_numeric(city['AQI'], errors='coerce')
city = city.dropna(subset=['AQI']).copy()
cap_value = city['AQI'].quantile(0.99)
city['AQI'] = city['AQI'].clip(upper=cap_value)

# Extract year
city['Year'] = city['Date'].dt.year

# load crop data
crop = pd.read_csv('crop_production.csv')
crop.columns = crop.columns.str.strip()
crop['State_Name'] = crop['State_Name'].str.strip()
crop['Production'] = pd.to_numeric(crop['Production'], errors='coerce')
crop['Area'] = pd.to_numeric(crop['Area'], errors='coerce')
crop['Crop_Year'] = pd.to_numeric(crop['Crop_Year'], errors='coerce')
crop = crop.dropna(subset=['Production']).copy()

# Map cities to states
city_to_state = {
    'Ahmedabad': 'Gujarat', 'Aizawl': 'Mizoram', 'Amravati': 'Andhra Pradesh',
    'Amritsar': 'Punjab', 'Aurangabad': 'Maharashtra', 'Bangalore': 'Karnataka',
    'Bhopal': 'Madhya Pradesh', 'Bhubaneswar': 'Odisha', 'Chandigarh': 'Chandigarh',
    'Chennai': 'Tamil Nadu', 'Coimbatore': 'Tamil Nadu', 'Cuttack': 'Odisha',
```



CHRIST
(DEEMED TO BE UNIVERSITY)
BANGALORE | DELHI NCR | PUNE

```
% Generate + Code + Markdown Run All Restart Clear All Outputs Jupyter Variables Outline Python 3.13.7
# Map cities to states
city_to_state = {
    'Ahmedabad': 'Gujarat', 'Aizawl': 'Mizoram', 'Amaravati': 'Andhra Pradesh',
    'Amritsar': 'Punjab', 'Aurangabad': 'Maharashtra', 'Bangalore': 'Karnataka',
    'Bhopal': 'Madhya Pradesh', 'Bhubaneswar': 'Odisha', 'Chandigarh': 'Chandigarh',
    'Chennai': 'Tamil Nadu', 'Coimbatore': 'Tamil Nadu', 'Cuttack': 'Odisha',
    'Delhi': 'Delhi', 'Erode': 'Tamil Nadu', 'Ghaziabad': 'Uttar Pradesh',
    'Guwahati': 'Assam', 'Hyderabad': 'Telangana', 'Indore': 'Madhya Pradesh',
    'Jaipur': 'Rajasthan', 'Jodhpur': 'Rajasthan', 'Kanpur': 'Uttar Pradesh',
    'Kochi': 'Kerala', 'Kolkata': 'West Bengal', 'Kota': 'Rajasthan',
    'Lucknow': 'Uttar Pradesh', 'Ludhiana': 'Punjab', 'Madurai': 'Tamil Nadu',
    'Meerut': 'Uttar Pradesh', 'Mumbai': 'Maharashtra', 'Nagpur': 'Maharashtra',
    'Nashik': 'Maharashtra', 'Noida': 'Uttar Pradesh', 'Panaji': 'Goa',
    'Patna': 'Bihar', 'Pune': 'Maharashtra', 'Raipur': 'Chhattisgarh',
    'Ranchi': 'Jharkhand', 'Srinagar': 'Jammu and Kashmir', 'Thiruvananthapuram': 'Kerala',
    'Udaipur': 'Rajasthan', 'Vijayawada': 'Andhra Pradesh', 'Visakhapatnam': 'Andhra Pradesh',
    'Warangal': 'Telangana'
}

city['State'] = city['City'].map(city_to_state)
city = city.dropna(subset=['State'])

# Aggregate city AQI to state-year level
state_year_aqi = city.groupby(['State', 'Year']).agg([
    'AQI': 'mean',
    'PM2.5': 'mean',
    'PM10': 'mean',
    'NO': 'mean',
    'NO2': 'mean',
    'NOx': 'mean',
    'NH3': 'mean',
    'CO': 'mean',
    'SO2': 'mean',
    'O3': 'mean',
    'Benzene': 'mean',
    'Toluene': 'mean'
])
```

```
% Generate + Code + Markdown Run All Restart Clear All Outputs Jupyter Variables Outline Python 3.13.7
# Merge datasets
merged = pd.merge(state_year_aqi, crop_state_year, on=['State', 'Year'], how='inner')
print(f'\nRecords after merge: {len(merged)}')

if len(merged) > 0:
    print(f'\nMerged years: {sorted(merged["Year"].unique())}')
    print(f'\nMerged states: {sorted(merged["State"].unique())}')
    print('\nSample merged data:')
    print(merged[['State', 'Year', 'AQI', 'Yield']].head())

    numeric_cols = ['AQI', 'PM2.5', 'PM10', 'NO', 'NO2', 'NOx', 'NH3', 'CO', 'SO2', 'O3', 'Benzene', 'Toluene', 'Production', 'Area', 'Yield']
    correlation = merged[numeric_cols].corr()
    yield_corr = correlation['Yield'].drop('Yield').sort_values(ascending=False)

    print('\nTop correlations with Yield:')
    print(yield_corr)
else:
    print('\n**Data Limitation:** No overlapping years between datasets.')
    print('Air quality: 2015-2020 | Crop production: 1997-2015')
```



CHRIST
(DEEMED TO BE UNIVERSITY)
BANGALORE | DELHI NCR | PUNE

```

Generate + Code + Markdown Run All Restart Clear All Outputs Jupyter Variables Outline Python 3.13.7

--- Air quality years: [np.int32(2015), np.int32(2016), np.int32(2017), np.int32(2018), np.int32(2019), np.int32(2020)]
Crop years: [np.int64(1997), np.int64(1998), np.int64(1999), np.int64(2000), np.int64(2001), np.int64(2002), np.int64(2003), np.int64(2004), np.int64(2005), np.int64(2006),

Records after merge: 0

**Data Limitation:** No overlapping years between datasets.
Air quality: 2015-2020 | Crop production: 1997-2015

# Debug: Check state names in both datasets
states_aqi = set(state_year_aqi[state_year_aqi['Year'] == 2015]['State'].unique())
states_crop = set(crop_state_year[crop_state_year['Year'] == 2015]['State'].unique())

print('States in 2015 air quality data:')
print(sorted(states_aqi))
print(f'Total: {len(states_aqi)}')

print('\n\nStates in 2015 crop data:')
print(sorted(states_crop))
print(f'Total: {len(states_crop)}')

print('\n\nCommon states:')
common = states_aqi.intersection(states_crop)
print(sorted(common))
print(f'Total: {len(common)}')

print('\n\nStates in crop but not in AQI:')
print(sorted(states_crop - states_aqi))

print('\n\nStates in AQI but not in crop:')
print(sorted(states_aqi - states_crop))

[20] States in 2015 air quality data:
['Bihar', 'Delhi', 'Gujarat', 'Tamil Nadu', 'Telangana', 'Uttar Pradesh']

```

```

--- States in 2015 air quality data:
['Bihar', 'Delhi', 'Gujarat', 'Tamil Nadu', 'Telangana', 'Uttar Pradesh']

Total: 6

States in 2015 crop data:
['Odisha', 'Sikkim']

Total: 2

Common states:
[]
Total: 0

States in crop but not in AQI:
['Odisha', 'Sikkim']

States in AQI but not in crop:
['Bihar', 'Delhi', 'Gujarat', 'Tamil Nadu', 'Telangana', 'Uttar Pradesh']

```



Generate + Code + Markdown Run All Restart Clear All Outputs Jupyter Variables Outline Python 3.13.7

Task 8 Findings

Data Limitation Identified:

The two datasets **do not have meaningful overlap:**

- **Air Quality data (city_day.csv):** Years 2015-2020, but sparse state coverage in our mapped cities (only 6 states in 2015: Bihar, Delhi, Gujarat, Tamil Nadu, Telangana, Uttar Pradesh)
- **Crop Production data (crop_production.csv):** Years 1997-2015, only 2 states reported for 2015 (Odisha, Sikkim)
- **Temporal overlap:** Only year 2015, but no state intersection

Why This Happened:

1. The crop production dataset only contains state-level summaries, and only 2 states have data for 2015
2. The city-level air quality data comes from specific monitored cities, which don't cover all states
3. The datasets were collected independently with different geographic scope and timeframes

Implication for Policy: This demonstrates a **critical gap in India's environmental monitoring system:** Agricultural states (like Punjab, Haryana, Rajasthan) where biomass burning is most intense don't have comprehensive air quality monitoring infrastructure. Policymakers cannot correlate pollution with crop production where it matters most.

Recommendation for Future Analysis: To properly study air-quality-crop-yield relationships, India would need:

1. Air quality monitoring stations in all agricultural states (not just urban centers)
2. Synchronized data collection (same time periods for both datasets)
3. District-level crop production granularity (not just state-level aggregation)

Generate + Code + Markdown Run All Restart Clear All Outputs Jupyter Variables Outline Python 3.13.7

Task 9 — Policy Briefing: State Environment Minister

To: State Environment Minister
From: Data Analysis Team
Subject: Air Quality Analysis & Agricultural Burning Impact
Date: \$(today)

Dear Minister,

Our analysis of air quality data (2015–2020) reveals three critical findings:

1. **Harvest Season Pollution Crisis:** October–December shows 46% worse air quality (AQI 213 vs. baseline 146). This directly correlates with agricultural burning during crop harvest, affecting public health during peak winter months.
2. **Worsening Trend Post-Policy:** Despite the National Clean Air Program launched in 2018, air quality has not improved. Average AQI has remained stagnant or slightly increased, suggesting current interventions are insufficient.
3. **Infrastructure Gap:** Comprehensive air quality and crop production data are not synchronized geographically or temporally, limiting our ability to precisely attribute pollution sources or measure agricultural impact.

Recommendation: Implement strict crop residue management regulations in October–December, with penalties for open burning. Incentivize alternative disposal methods (composting, biochar production) rather than burning.

Limitation Acknowledged: Our analysis relies on limited monitoring stations in major cities. Rural agricultural regions lack monitoring infrastructure, obscuring the true scale of harvest-season pollution.

Key Takeaway: Harvest-season air quality requires urgent, targeted policy action.

Cell 5 of 20 Python 3.13.7



CHRIST
(DEEMED TO BE UNIVERSITY)
BANGALORE | DELHI NCR | PUNE

Optional Advanced Tasks: Build the Case — Or Break It

Central Hypothesis: States with higher pollution have lower crop yields.

Approach: We will use matched cities/states with flexible year ranges to maximize available data. AQI will be averaged across 2015-2020; crop production across 1997-2015.

Task A: The Two Extremes — Do They Tell the Same Story?

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from pathlib import Path

# Load and clean city data (2015-2020)
city = pd.read_csv('city_day.csv')
city.columns = city.columns.str.strip()
city['Date'] = pd.to_datetime(city['Date'], errors='coerce')
city = city.drop(columns=['Xylene'], errors='ignore')

pollutant_cols = ['PM2.5', 'PM10', 'NO', 'NO2', 'NOx', 'NH3', 'CO', 'SO2', 'O3', 'Benzene', 'Toluene']
for col in pollutant_cols:
    city[col] = pd.to_numeric(city[col], errors='coerce')
    city[col] = city[col].fillna(city.groupby('City')[col].transform('median'))
    city[col] = city[col].fillna(city[col].median())

city['AQI'] = pd.to_numeric(city['AQI'], errors='coerce')
city = city.dropna(subset=['AQI']).copy()
cap_value = city['AQI'].quantile(0.99)
city['AQI'] = city['AQI'].clip(upper=cap_value)
```

```
# Map cities to states
city_to_state = {
    'Ahmedabad': 'Gujarat', 'Aizawl': 'Mizoram', 'Amaravati': 'Andhra Pradesh',
    'Amritsar': 'Punjab', 'Aurangabad': 'Maharashtra', 'Bangalore': 'Karnataka',
    'Bhopal': 'Madhya Pradesh', 'Bhubaneswar': 'Odisha', 'Chandigarh': 'Chandigarh',
    'Chennai': 'Tamil Nadu', 'Coimbatore': 'Tamil Nadu', 'Cuttack': 'Odisha',
    'Delhi': 'Delhi', 'Erode': 'Tamil Nadu', 'Ghaziabad': 'Uttar Pradesh',
    'Guwahati': 'Assam', 'Hyderabad': 'Telangana', 'Indore': 'Madhya Pradesh',
    'Jaipur': 'Rajasthan', 'Jodhpur': 'Rajasthan', 'Kanpur': 'Uttar Pradesh',
    'Kochi': 'Kerala', 'Kolkata': 'West Bengal', 'Kota': 'Rajasthan',
    'Lucknow': 'Uttar Pradesh', 'Ludhiana': 'Punjab', 'Madurai': 'Tamil Nadu',
    'Meerut': 'Uttar Pradesh', 'Mumbai': 'Maharashtra', 'Nagpur': 'Maharashtra',
    'Nashik': 'Maharashtra', 'Noida': 'Uttar Pradesh', 'Panaji': 'Goa',
    'Patna': 'Bihar', 'Pune': 'Maharashtra', 'Raipur': 'Chhattisgarh',
    'Ranchi': 'Jharkhand', 'Srinagar': 'Jammu and Kashmir', 'Thiruvananthapuram': 'Kerala',
    'Udaipur': 'Rajasthan', 'Vijayawada': 'Andhra Pradesh', 'Visakhapatnam': 'Andhra Pradesh',
    'Warangal': 'Telangana'
}

city['State'] = city['City'].map(city_to_state)
city = city.dropna(subset=['State'])

# Average AQI by state (across all years 2015-2020)
avg_aqi_by_state = city.groupby('State')['AQI'].agg(['mean', 'std']).reset_index()
avg_aqi_by_state.columns = ['State', 'Avg_AQI', 'Std_AQI']

# Load and clean crop data (1997-2015)
crop = pd.read_csv('crop_production.csv')
crop.columns = crop.columns.str.strip()
crop['State_Name'] = crop['State_Name'].str.strip()
crop['Production'] = pd.to_numeric(crop['Production'], errors='coerce')
crop['Area'] = pd.to_numeric(crop['Area'], errors='coerce')
crop['Crop_Year'] = pd.to_numeric(crop['Crop_Year'], errors='coerce')
crop = crop.dropna(subset=['Production', 'Area']).copy()
```




CHRIST
(DEEMED TO BE UNIVERSITY)
BANGALORE | DELHI NCR | PUNE

```
% Generate + Code + Markdown | Run All | Restart | Clear All Outputs | Jupyter Variables | Outline | Python 3.13.7

# Calculate yield and average by state
crop['Yield'] = crop['Production'] / crop['Area']
avg_production_by_state = crop.groupby('State_Name').agg({
    'Production': 'mean',
    'Area': 'mean',
    'Yield': 'mean'
}).reset_index()
avg_production_by_state.columns = ['State', 'Avg_Production', 'Avg_Area', 'Avg_Yield']

# Merge on state name (case-insensitive)
avg_aqi_by_state['State_lower'] = avg_aqi_by_state['State'].str.lower()
avg_production_by_state['State_lower'] = avg_production_by_state['State'].str.lower()

merged_states = pd.merge(avg_aqi_by_state, avg_production_by_state, on='State_lower')
merged_states = merged_states.drop(columns=['State_lower'])
merged_states = merged_states[['State_x', 'Avg_AQI', 'Std_AQI', 'Avg_Production', 'Avg_Area', 'Avg_Yield']]
merged_states.columns = ['State', 'Avg_AQI', 'Std_AQI', 'Avg_Production', 'Avg_Area', 'Avg_Yield']

print(f'States with both AQI and crop data: {len(merged_states)}')
print(merged_states.sort_values('Avg_AQI', ascending=False)[['State', 'Avg_AQI', 'Avg_Yield']].to_string())

# Identify the two extremes
most_polluted = merged_states.loc[merged_states['Avg_AQI'].idxmax()]
least_polluted = merged_states.loc[merged_states['Avg_AQI'].idxmin()]

print(f'\n\nMost polluted state: {most_polluted["State"]} (AQI: {most_polluted["Avg_AQI"]:.2f})')
print(f'Least polluted state: {least_polluted["State"]} (AQI: {least_polluted["Avg_AQI"]:.2f})')
print(f'Yield difference: {most_polluted["Avg_Yield"]:.2f} vs {least_polluted["Avg_Yield"]:.2f}')

# Visualize comparison
fig, axes = plt.subplots(1, 3, figsize=(16, 5))

# AQI comparison
states_comp = [most_polluted['State'], least_polluted['State']]
aqi_comp = [most_polluted['Avg_AQI'], least_polluted['Avg_AQI']]
colors_aqi = ['#d62728', '#2ca02c']

[24]
```

```
% Generate + Code + Markdown | Run All | Restart | Clear All Outputs | Jupyter Variables | Outline | Python 3.13.7

# AQI comparison
states_comp = [most_polluted['State'], least_polluted['State']]
aqi_comp = [most_polluted['Avg_AQI'], least_polluted['Avg_AQI']]
colors_aqi = ['#d62728', '#2ca02c']
axes[0].bar(states_comp, aqi_comp, color=colors_aqi, alpha=0.8, edgecolor='black', linewidth=2)
axes[0].set_ylabel('Average AQI', fontsize=12, fontweight='bold')
axes[0].set_title('Pollution: The Two Extremes', fontsize=13, fontweight='bold')
axes[0].set_ylim(0, max(aqi_comp) * 1.2)
for i, v in enumerate(aqi_comp):
    axes[0].text(i, v + 5, f'{v:.1f}', ha='center', fontweight='bold', fontsize=11)


# Yield comparison
yield_comp = [most_polluted['Avg_Yield'], least_polluted['Avg_Yield']]
axes[1].bar(states_comp, yield_comp, color=colors_aqi, alpha=0.8, edgecolor='black', linewidth=2)
axes[1].set_ylabel('Average Yield (Production/Area)', fontsize=12, fontweight='bold')
axes[1].set_title('Crop Yield: What Do We See?', fontsize=13, fontweight='bold')
axes[1].set_ylim(0, max(yield_comp) * 1.2)
for i, v in enumerate(yield_comp):
    axes[1].text(i, v + 0.5, f'{v:.1f}', ha='center', fontweight='bold', fontsize=11)

# Production comparison (absolute)
prod_comp = [most_polluted['Avg_Production'], least_polluted['Avg_Production']]
axes[2].bar(states_comp, prod_comp, color=colors_aqi, alpha=0.8, edgecolor='black', linewidth=2)
axes[2].set_ylabel('Average Production (tonnes)', fontsize=12, fontweight='bold')
axes[2].set_title('Total Production Output', fontsize=13, fontweight='bold')
axes[2].set_ylim(0, max(prod_comp) * 1.2)
for i, v in enumerate(prod_comp):
    axes[2].text(i, v + 100, f'{v:.0f}', ha='center', fontweight='bold', fontsize=11)

plt.tight_layout()
plt.show()

[24]
```

```
... States with both AQI and crop data: 15
      State  Avg_AQI  Avg_Yield
4  Gujarat  392.722129  5.716671
```





```

% Generate + Code + Markdown ▶ Run All ⌂ Restart ⌂ Clear All Outputs 📄 Jupyter Variables 📄 Outline — Python 3.13.7

But there's nuance:

1. Scale matters: Gujarat produces 36x more total production than Mizoram (62,677 vs 1,742 tonnes), suggesting different agricultural systems—Gujarat is an industrial farming powerhouse, Mizoram is subsistence agriculture.
2. Confounding factors: Mizoram's low yield is likely due to mountainous terrain, rainfall patterns, soil type, and economic development—not pollution. Gujarat's high pollution may be a consequence of intensive agriculture, not the cause of higher yields.
3. The causality trap: Even if pollution correlates with yield, the direction of causation is unclear. Does pollution suppress yield, or do high-yield agricultural systems create pollution?

Conclusion: The extremes reveal that the relationship is not straightforward. High pollution does not necessarily mean low yield. Instead, we see a complex interplay of economic structure, geography, and agricultural practices.

Task B: Put a Number on the Relationship

Research teams need quantitative evidence. We will compute the correlation between average state AQI and average crop yield, visualized with a regression line and  $R^2$  value.

from scipy import stats

# Use merged_states from Task A
data_clean = merged_states[['Avg_AQI', 'Avg_Yield']].dropna()

aqi_vals = data_clean['Avg_AQI'].values
yield_vals = data_clean['Avg_Yield'].values

# Pearson correlation
correlation_coef, p_value = stats.pearsonr(aqi_vals, yield_vals)

# Linear regression
slope, intercept, r_value, p_reg, std_err = stats.linregress(aqi_vals, yield_vals)
r_squared = r_value ** 2

print(f'Pearson Correlation Coefficient: {correlation_coef:.4f}')

```

```

% Generate + Code + Markdown ▶ Run All ⌂ Restart ⌂ Clear All Outputs 📄 Jupyter Variables 📄 Outline — Python 3.13.7

print(f'Pearson Correlation Coefficient: {correlation_coef:.4f}')
print(f'P-value (correlation): {p_value:.4f}')
print(f'R² (coefficient of determination): {r_squared:.4f}')
print(f'Linear regression equation: Yield = {intercept:.3f} + {slope:.6f} * AQI')
print(f'Standard error: {std_err:.6f}')

# Interpretation
if correlation_coef > 0:
    direction = 'positive'
else:
    direction = 'negative'

if abs(correlation_coef) < 0.3:
    strength = 'weak'
elif abs(correlation_coef) < 0.7:
    strength = 'moderate'
else:
    strength = 'strong'

print(f'\nInterpretation: {strength.capitalize()} {direction} correlation')

# Create large scatter plot for research team
fig, ax = plt.subplots(figsize=(12, 7))

# Scatter plot
ax.scatter(aqi_vals, yield_vals, s=150, alpha=0.6, color='#2a6f97', edgecolor='black', linewidth=1.5)

# Regression line
x_line = np.array([aqi_vals.min(), aqi_vals.max()])
y_line = intercept + slope * x_line
ax.plot(x_line, y_line, 'r--', linewidth=3, label=f'Linear Fit: y = {intercept:.2f} {slope:.4f}x')

# Confidence interval
residuals = yield_vals - (intercept + slope * aqi_vals)
std_residuals = np.std(residuals)
ax.fill_between(x_line,


```



CHRIST
(DEEMED TO BE UNIVERSITY)
BANGALORE | DELHI NCR | PUNE

```

% Generate + Code + Markdown ▶ Run All ⏮ Restart 🗑 Clear All Outputs 📄 Jupyter Variables 📄 Outline ... Python 3.13.7
> ~
# Confidence interval
residuals = yield_vals - (intercept + slope * aqi_vals)
std_residuals = np.std(residuals)
ax.fill_between(x_line,
                intercept + slope * x_line - 1.96 * std_residuals,
                intercept + slope * x_line + 1.96 * std_residuals,
                alpha=0.2, color='red', label='95% Confidence Band')

# Labels
ax.set_xlabel('Average State AQI (2015-2020)', fontsize=14, fontweight='bold')
ax.set_ylabel('Average Crop Yield (Production/Area)', fontsize=14, fontweight='bold')
ax.set_title('AQI vs Crop Yield: Quantifying the Relationship', fontsize=15, fontweight='bold')

# State annotations
for idx, row in merged_states.iterrows():
    if not pd.isna(row['Avg_AQI']) and not pd.isna(row['Avg_Yield']):
        ax.annotate(row['State'][:3],
                    xy=(row['Avg_AQI'], row['Avg_Yield']),
                    xytext=(5, 5),
                    textcoords='offset points',
                    fontsize=8)

# Statistics box
stats_text = f'Correlation: {correlation_coef:.3f}\nR²: {r_squared:.3f}\nN: {len(data_clean)}\np-value: {p_value:.4f}'
ax.text(0.98, 0.05, stats_text,
        transform=ax.transAxes,
        fontsize=12,
        verticalalignment='bottom',
        horizontalalignment='right',
        bbox=dict(boxstyle='round', facecolor='wheat', alpha=0.8),
        family='monospace',
        fontweight='bold')

ax.legend(fontsize=11, loc='upper left')
ax.grid(alpha=0.3)
plt.tight_layout()

```

```

% Generate + Code + Markdown ▶ Run All ⏮ Restart 🗑 Clear All Outputs 📄 Jupyter Variables 📄 Outline ... Python 3.13.7
> ~
ax.legend(fontsize=11, loc='upper left')
ax.grid(alpha=0.3)
plt.tight_layout()
plt.show()

print(f'\nR² = {r_squared:.4f} means that AQI explains {r_squared*100:.1f}% of the variation in crop yield.')
print(f'The remaining {(1-r_squared)*100:.1f}% is explained by other factors.')


--- Pearson Correlation Coefficient: -0.3003
P-value (correlation): 0.2768
R² (coefficient of determination): 0.0902
Linear regression equation: Yield = 152.892 + -0.448319 x AQI
Standard error: 0.394925

Interpretation: Moderate negative correlation

```




CHRIST
(DEEMED TO BE UNIVERSITY)
BANGALORE | DELHI NCR | PUNE



$R^2 = 0.0902$ means that AQI explains 9.0% of the variation in crop yield.
 The remaining 90.9% is explained by other factors.

Task B Analysis: What the Numbers Tell Us

Quantitative Finding:

- **Pearson Correlation:** -0.300 (weak negative relationship)
- **$R^2 = 0.090$:** AQI explains only **9% of the variation in crop yield**
- **P-value = 0.277:** The correlation is **not statistically significant** (need $p < 0.05$)
- **Sample size:** $N = 15$ states

What This Means for Policy & Research:

1. **No Proof of Causation:** A p-value > 0.05 means we cannot confidently say pollution causes lower yields—the observed relationship could easily be due to chance.
2. **Weak Explanatory Power:** Even if we accept the weak negative trend, AQI accounts for only 9% of yield differences. The other **91% depends on factors this dataset never measures**:
 - **Rainfall & monsoon patterns** (major driver of Indian agriculture)
 - **Soil quality & soil depth** (determines inherent fertility)
 - **Irrigation infrastructure** (Punjab's high yields are irrigation-driven, not pollution-free)
 - **Crop mix** (rice vs. pulses vs. wheat have different yields naturally)
 - **Economic development & farm inputs** (fertilizer, quality seeds, mechanization)
 - **Farmer education & extension services**
3. **The Causality Trap:** The weak negative correlation might reflect:
 - High-intensity agricultural regions (Gujarat, Punjab) both pollute AND produce heavily
 - Geography matters more than pollution: Highland states pollute less but yield less too
 - Association $\neq$ causation



Recommendation to Research Team: Do NOT fund a full causal study based on this correlation. First, collect more granular data on confounding variables. A multi-year, multi-district study with controls for rainfall, irrigation, and soil would be needed to establish causality.

Task C: One Plot to Rule Them All

The Most Powerful Visualization: Task 7 - Seasonal Pattern of AQI

This chart reveals what no other visualization in this analysis does: a **specific, predictable, and preventable seasonal crisis**. Every October through November without fail, India's air quality deteriorates by 45%—not due to winter weather or traffic, but due to agricultural burning. This is not a gradual trend to debate or a subtle correlation to investigate with caution. It is a **repeating catastrophe tied to a single human practice**.

Unlike Task 6's trend (no improvement post-2018) or Task 8's data gaps (obscuring the issue), this chart answers the farmer's question directly: *'Is the air worst when we harvest?'* Yes—overwhelmingly yes. The visualization is simple, the finding is quantifiable, and the policy implication is immediate: implement harvest-season burning bans NOW before winter compounds the harm.

What it reveals that others don't: While other charts show correlations (weak) or trends (flat), this chart shows a **repeating pattern tied to human action on a known date**. That specificity is what drives policy change.

What we cannot claim: We cannot claim pollution alone causes lower yields across India. Rainfall, soil type, irrigation intensity, and crop type explain 91% of yield variation that pollution does not. Seasonal pollution may harm public health immediately (respiratory disease), but its long-term agricultural impact is confounded with dozens of other factors this dataset cannot separate.

Chart to publish:

```
# Recreate the seasonal pattern visualization with publication quality
# (reuse city data from Task A)

city['Month'] = city['Date'].dt.month
city['Month_Name'] = city['Date'].dt.strftime('%B')

monthly_aqi = city.groupby(['Month', 'Month_Name'])['AQI'].agg(['mean', 'median', 'std']).reset_index()
```

```
# Recreate the seasonal pattern visualization with publication quality
# (reuse city data from Task A)

city['Month'] = city['Date'].dt.month
city['Month_Name'] = city['Date'].dt.strftime('%B')

monthly_aqi = city.groupby(['Month', 'Month_Name'])['AQI'].agg(['mean', 'median', 'std']).reset_index()
monthly_aqi = monthly_aqi.sort_values('Month')

# Create publication-quality figure
fig, ax = plt.subplots(figsize=(14, 7))

# Color harvest season (Oct-Dec) differently
colors = ['#2a6697' if m not in [10, 11, 12] else '#d62828' for m in monthly_aqi['Month']]
bars = ax.bar(monthly_aqi['Month_Name'], monthly_aqi['mean'], color=colors, alpha=0.8, edgecolor='black', linewidth=1.5)

# Annual average line
annual_mean = monthly_aqi['mean'].mean()
ax.axhline(y=annual_mean, color='gray', linestyle='--', linewidth=2, alpha=0.8, label=f'Annual Average ({annual_mean:.1f})')

# Formatting for publication
ax.set_xlabel('Month', fontsize=14, fontweight='bold')
ax.set_ylabel('Mean AQI', fontsize=14, fontweight='bold')
ax.set_title('India's Air Quality Crisis: The Harvest Season Every October-December',
            fontsize=16, fontweight='bold', pad=20)

# Add value labels on bars
for bar, val in zip(bars, monthly_aqi['mean']):
    height = bar.get_height()
    ax.text(bar.get_x() + bar.get_width()/2., height + 3,
            f'{val:.0f}', ha='center', va='bottom', fontsize=11, fontweight='bold')

# Legend for harvest season
from matplotlib.patches import Patch
legend_elements = [Patch(facecolor='#d62828', alpha=0.8, edgecolor='black', label='Harvest Season (Oct-Dec)'),
                    Patch(facecolor='#2a6697', alpha=0.8, edgecolor='black', label='Non-Harvest Months (Jan-Sep)')]
ax.legend(handles=legend_elements, loc='upper right', title='Legend', titlepad=10)
```



CHRIST
(DEEMED TO BE UNIVERSITY)
BANGALORE | DELHI NCR | PUNE

```


% Generate + Code + Markdown | Run All | Restart | Clear All Outputs | Jupyter Variables | Outline | Python 3.13.7

# Legend for harvest season
from matplotlib.patches import Patch
legend_elements = [Patch(facecolor='#d62728', alpha=0.8, edgecolor='black', label='Harvest Season (Oct-Dec)'),
                   Patch(facecolor='#1f77b4', alpha=0.8, edgecolor='black', label='Non-Harvest Months')]
ax.legend(handles=legend_elements, fontsize=12, loc='upper left')

plt.xticks(rotation=45, fontsize=11)
plt.yticks(fontsize=11)
ax.grid(axis='y', alpha=0.3, linestyle='--')
plt.tight_layout()
plt.show()

# Print the caption for journalists
print("\n" + "="*80)
print("CAPTION FOR PUBLICATION")
print("="*80)
caption = """
India's Worst Air Comes on Schedule: Farmers Burning Straw Turns October, November,
and December into a Public Health Emergency: Every year, without fail, three months
trap India's cities under a haze of agricultural smoke. Field-burning in the harvest
season pushes air quality 46% worse than the annual baseline. Unlike pollution from
traffic or industry-which policymakers struggle to control-this crisis has a known
cause, a predictable date, and a solution: halt open burning during harvest season.
The data shows what farmers and doctors already know: this is not inevitable pollution,
it is a policy choice.
"""
print(caption)
print("\nNote: AQI data 2015-2020, aggregated across 24 Indian cities. Harvest season deficit: 46% worse (AQI 213 vs. 146 annual mean).")

```



```

% Generate + Code + Markdown | Run All | Restart | Clear All Outputs | Jupyter Variables | Outline | Python 3.13.7

=====
CAPTION FOR PUBLICATION
=====

India's Worst Air Comes on Schedule: Farmers Burning Straw Turns October, November,
and December into a Public Health Emergency: Every year, without fail, three months
trap India's cities under a haze of agricultural smoke. Field-burning in the harvest
season pushes air quality 46% worse than the annual baseline. Unlike pollution from
traffic or industry-which policymakers struggle to control-this crisis has a known
cause, a predictable date, and a solution: halt open burning during harvest season.
The data shows what farmers and doctors already know: this is not inevitable pollution,
it is a policy choice.

Note: AQI data 2015-2020, aggregated across 24 Indian cities. Harvest season deficit: 46% worse (AQI 213 vs. 146 annual mean).

```



Summary: Advanced Tasks A–C — Building and Breaking the Hypothesis

Central Hypothesis Tested: States with higher pollution have lower crop yields.

Verdict: The hypothesis is **contradicted by the data, complicated by confounding variables, and undermined by weak statistical evidence.**

Key Evidence

Task	Finding	Implication
A: Extremes	Most polluted (Gujarat, AQI 392.7): yield 5.7; Least polluted (Mizoram, AQI 34.8): yield 2.3	Direct contradiction: High pollution ↔ HIGH yield, not low yield. Likely driven by economic structure, not pollution.
B: Quantification	Correlation = -0.30 (weak); $R^2 = 0.09$; p-value = 0.28 (not significant)	AQI explains only 9% of yield variation. The 91% gap suggests rainfall, soil, irrigation, crop type matter vastly more than pollution.
C: Strongest Story	Seasonal pattern: Oct-Dec 46% worse AQI (predictable, preventable, linked to agricultural burning)	Most actionable finding: Seasonal crisis is the real story, not pollution-yield causation. Focus policy on harvest-season burning bans.

The Honest Assessment

✓ **What the data supports:** Air quality is worse during harvest season (Oct-Dec), driven by agricultural burning.

✗ **What the data does NOT support:** Pollution is the primary driver of crop yield variation (opposite may be true: intensive agriculture creates both pollution and yields).

? **What remains unclear:** Whether pollution *causes* yield loss or whether high-productivity regions happen to pollute more.

Generate
+ Code
+ Markdown
Run All
Restart
Clear All Outputs
Jupyter Variables
Outline
Python 3.13.7

Confounding Factors This Dataset Cannot Separate

1. **Rainfall & monsoons** – farmers can't grow without rain, pollution doesn't determine this
2. **Irrigation intensity** – Punjab's high yields are irrigation-led, not pollution-free
3. **Soil quality** – geologically predetermined
4. **Crop mix** – rice ≠ pulses ≠ wheat in yield naturally
5. **Farmer wealth & education** – access to inputs, mechanization
6. **Government support** – subsidies, infrastructure, extension services

Recommendation

Do not assume correlation = causation. Any causal study must:

- Control for rainfall and irrigation status
- Measure soil properties district-by-district
- Track same crops across years
- Account for farming practice changes
- Use experimental design (not observational data)

The lab's true insight: **India's harvest-season pollution is a known, repeatable crisis tied to a specific agricultural practice**—and that is worth acting on now, regardless of whether it mechanically suppresses yields.

Lab Exercise 3 :

Student Survey: Simple Linear Regression using Scikit Learn, OLS, and Parameter Saving



Aim

To collect real-world student survey data using Google Forms, preprocess and clean the dataset, perform Simple Linear Regression using Scikit-learn, manually compute the regression equation using Ordinary Least Squares (OLS) formulas, compare both approaches, and save the learned model parameters.

Instructions

1. During the first 15 minutes of the lab session, fill out the Google Form shared, link :

MCA-B

<https://docs.google.com/forms/d/e/1FAIpQLSdR13N42Eh2Lz7U7QYkw9bp-G03puOGKROVQz4fzUDGdF1g8A/viewform?usp=publish-editor>

MCA-A

<https://docs.google.com/forms/d/1cMs0pEILKdkqJEUt2TYRdZg3rJz7-aTelD6uT3EW8VA/e/dit?ts=6a2ecba6>

2. Ensure that responses are collected from your classmates.
3. Export the responses into CSV format.
4. Use the collected dataset for all further analysis

Suggested dataset name:

student_survey.csv

Problem Statement

Using the dataset collected from the class survey form, perform Simple Linear Regression analysis to determine whether:

1. CIA percentage can predict GPA
2. Attendance percentage can predict GPA

You are required to perform regression in two different ways:

1. Using Scikit learn LinearRegression



2. Manually using the mathematical equations of slope and intercept derived from Ordinary Least Squares (OLS)

Finally, compare the outputs obtained from both methods and save the learned model parameters.

Refer this for getting and understanding :

https://colab.research.google.com/drive/1No_uEMN0sq8y2oqGvRgL9FKMGAp7BxX#scrollTo=a7bfa6e8

Tasks to Perform

Part A: Data Collection and Preprocessing

Perform the following steps:

1. Load the dataset using Pandas
2. Display the first 5 rows
3. Check dataset dimensions
4. Identify missing values
5. Remove or handle null values
6. Convert required columns into numerical datatype
7. Remove duplicate records if present
8. Generate statistical summary
9. Select appropriate dependent and independent variables

Regression Experiments: Experiment 1

Independent Variable (X):

CIA Percentage

Dependent Variable (Y):

GPA



Experiment 2

Independent Variable (X):

Attendance Percentage

Dependent Variable (Y):

GPA

Part B: Simple Linear Regression using Scikit-learn

Perform the following:

1. Split the dataset into training and testing sets
2. Train the model using LinearRegression from Scikit learn
3. Obtain:
 - Slope
 - Intercept
4. Predict output values using the trained mod

Part C: Manual Computation using Ordinary Least Squares (OLS)

Using the same dataset, manually compute the regression line using the OLS formulas given below.



Slope equation m :

$$m = \frac{\sum (x_i - \bar{x})(y_i - \bar{y})}{\sum (x_i - \bar{x})^2}$$

Intercept equation b :

$$b = \bar{y} - m\bar{x}$$

Final regression equation:

$$\hat{y} = mx + b$$

Where:

- x_i = individual input value
- y_i = individual output value
- $\bar{x}$ = mean of x values
- $\bar{y}$ = mean of y values
- m = slope
- b = intercept
- $\hat{y}$ = predicted value

Refer :

Manual Regression Equation

Construct the regression equation manually using the computed slope and intercept values.

Comparison Task

Compare the following:

1. Predictions from Scikit learn
2. Predictions from manually computed OLS equation
3. Difference between both outputs

Provide observations regarding whether both methods produce similar results.

Parameter Saving Task



Save the parameters learned by the model using Pickle.

The following parameters must be saved:

1. Slope
2. Intercept

Suggested file name:

linear_regression_weights.pkl

Demonstrate:

1. Saving parameters into Pickle file
2. Loading the Pickle file
3. Using loaded parameters for prediction

Expected Output

Submit the following:

1. CSV dataset
2. Data preprocessing notebook
3. Scikit learn regression implementation
4. Manual OLS calculation(self study)
5. Comparison of both approaches
6. Pickle file containing saved parameters
7. Final observations and inference

Sample Viva Questions

1. What is Simple Linear Regression?
2. What is the role of slope and intercept?
3. What is Ordinary Least Squares?
4. Why do we square the errors in OLS?
5. Difference between dependent and independent variable
6. Why should data be cleaned before training?
7. Why are slope and intercept called model parameters?
8. Why do we save learned weights?



Code with Results:

Generate + Code + Markdown Run All Restart Clear All Outputs Jupyter Variables Outline Python 3.13.7

Lab Exercise 3: Simple Linear Regression

Student Survey Analysis using Scikit-Learn, OLS, and Pickle

Part A: Data Collection and Preprocessing

```

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import pickle
import re
from sklearn.linear_model import LinearRegression
from sklearn.model_selection import train_test_split
from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score

print("Libraries imported successfully")

```

✓ 2.1s Libraries imported successfully

Step 1: Load Dataset

```

df_raw = pd.read_csv("Student Awareness Survey (Responses) - Form Responses 1.csv")
print("Dataset loaded successfully")
print(f"Shape: {df_raw.shape}")

```

✓ 0.0s Dataset loaded successfully
Shape: (58, 15)

Step 2: Display First 5 Rows

```

df_raw.head()

```

✓ 0.0s

Timestamp	Registration Number	Email	Job role that you are interested in	What is the minimum salary of students placed through campus (in LPA..respond as a number)	What is the maximum salary of students placed through campus (in LPA..respond as a number)	What is the median salary of students placed through campus (in LPA..respond as a number)	Which is the highest paying company that recruits from campus?	Rate your contribution towards extra curricular activities	Rate your technical competencies	What are your package expectations (LPA)	Your % of seme
0 6/15/2026 9:25:39	2547231	kunnal.kunnal@mca.christuniversity.in	Software Development Engineer (SDE)	3.5	12	6.8	Akasa air	4.0	3.0	8	
1 6/15/2026 9:25:39	2547237	omkaar.chakraborty@mca.christuniversity.in	Software Development Engineer	6	20	10	Fractal	4.0	4.0	12	

Cell 1 of 48 Go Live



CHRIST
(DEEMED TO BE UNIVERSITY)
BANGALORE | DELHI NCR | PUNE

	Timestamp	Registration Number	Email	Job role that you are interested in	What is the minimum salary of students placed through campus (In LPA..respond as a number)	What is the maximum salary of students placed through campus (In LPA..respond as a number)	What is the median salary of students placed through campus (In LPA..respond as a number)	Which is the highest paying company that recruits from campus?	Rate your contribution towards extra curricular activities	Rate your technical competencies	What are your package expectations (LPA)	Your % of seme
0	6/15/2026 9:25:39	2547231	kunnal.kunnal@mca.christuniversity.in	Software Development Engineer (SDE)	3.5	12	6.8	Akasa air	4.0	3.0	8	
1	6/15/2026 9:53:54	2547237	omkaar.chakraborty@mca.christuniversity.in	Software Development Engineer (SDE)	6	20	10	Fractal	4.0	4.0	12	
2	6/15/2026 9:54:56	2547203	abhinav.jain@mca.christuniversity.in	Full Stack Developer	4	12	6.3	Akasa Air	5.0	4.0	12	
3	6/15/2026 9:55:17	2547228	jai.pareek@mca.christuniversity.in	Full Stack Developer	600000	1400000	800000	Akasa airlines	5.0	4.0	1200000	
4	6/15/2026 9:55:42	2547241	r.karan@mca.christuniversity.in	Software Development Engineer (SDE)	4	4	6	12	2.0	3.0	12	

Step 3: Dataset Dimensions

```
print(f"Rows: {df_raw.shape[0]}, Columns: {df_raw.shape[1]}")
print("Column Names:")
for i, col in enumerate(df_raw.columns):
    print(f" {i}: {col}")
```

Rows: 50, Columns: 15
Column Names:
0: Timestamp
1: Registration Number
2: Email
3: Job role that you are interested in
4: What is the minimum salary of students placed through campus (In LPA..respond as a number)
5: What is the maximum salary of students placed through campus (In LPA..respond as a number)
6: What is the median salary of students placed through campus (In LPA..respond as a number)
7: Which is the highest paying company that recruits from campus?
8: Rate your contribution towards extra curricular activities
9: Rate your technical competencies
10: What are your package expectations (LPA)
11: Your CIA % of last semester
12: Your GPA of last semester
13: Your maximum attendance % till last semester
14: Internships Interests

Step 4: Identify Missing Values

```
print("Missing values per column:")
print(df_raw.isnull().sum())
print(f"Total missing values: {df_raw.isnull().sum().sum()}")
```

Missing values per column:
Timestamp 0
Registration Number 0
Email 0
Job role that you are interested in 0
What is the minimum salary of students placed through campus (In LPA..respond as a number) 0
What is the maximum salary of students placed through campus (In LPA..respond as a number) 0
What is the median salary of students placed through campus (In LPA..respond as a number) 0
Which is the highest paying company that recruits from campus? 1
Rate your contribution towards extra curricular activities 1
Rate your technical competencies 1
What are your package expectations (LPA) 1
Your CIA % of last semester 0
Your GPA of last semester 0
Your maximum attendance % till last semester 0
Internships Interests 1
dtype: int64
Total missing values: 5



CHRIST
(DEEMED TO BE UNIVERSITY)
BANGALORE | DELHI NCR | PUNE

Step 5: Data Cleaning and Preprocessing

The survey responses contain messy data (e.g., "80%", "4 LPA", "5+|pa", ranges like "8-9"). We extract the first valid numeric value from each relevant column.

```
# Helper function to extract numeric value from messy strings
def extract_numeric(val):
    if pd.isna(val):
        return np.nan
    val = str(val).strip()
    # Remove common units and labels
    val = re.sub(r'(\d+)([a-zA-Z%]+)', '', val)
    val = val.strip()
    # Match first valid number (including decimals)
    match = re.search(r'\d+\.\d*', val)
    if match:
        return float(match.group())
    return np.nan

# Select and rename relevant columns
cia_col = 'Your CIA % of last semester'
gpa_col = 'Your GPA of last semester'
att_col = 'Your maximum attendance % till last semester'

df = pd.DataFrame()
df['CIA_Percent'] = df_raw[cia_col].apply(extract_numeric)
df['GPA'] = df_raw[gpa_col].apply(extract_numeric)
df['Attendance'] = df_raw[att_col].apply(extract_numeric)

print("Before cleaning:", df.shape)
print(df.head(10))
```

```
Before cleaning: (50, 3)
  CIA_Percent  GPA  Attendance
0         69.0  3.40         98.0
1         75.0  3.69         95.0
2         82.0  3.41         95.0
3         91.0  3.60         92.0
4         70.0  3.54         93.0
5         75.0  3.67         98.0
6         74.0  3.40         94.0
7         68.0  3.40         85.0
8         73.0  3.60         90.0
9         79.0  3.73         93.0
```

```
# Remove rows with NaN values
df.dropna(inplace=True)

# Remove duplicates
df.drop_duplicates(inplace=True)

# Validate ranges: GPA 0-4, CIA 0-100, Attendance 0-100
df = df[(df['GPA'] >= 0) & (df['GPA'] <= 4.0)]
df = df[(df['CIA_Percent'] >= 0) & (df['CIA_Percent'] <= 100)]
df = df[(df['Attendance'] >= 0) & (df['Attendance'] <= 100)]

df.reset_index(drop=True, inplace=True)

print("After cleaning:", df.shape)
print(df)
```

```
After cleaning: (49, 3)
  CIA_Percent  GPA  Attendance
0         69.00  3.40         98.00
1         75.00  3.69         95.00
2         82.00  3.41         95.00
3         91.00  3.60         92.00
4         70.00  3.54         93.00
5         75.00  3.67         98.00
6         74.00  3.40         94.00
7         68.00  3.40         85.00
8         73.00  3.60         90.00
9         79.00  3.73         93.00
10        72.00  3.54         95.00
11        80.00  3.00         95.00
12        78.00  3.32         95.00
13        70.00  3.49         100.00
14        60.00  3.00         85.00
15        68.00  3.14         88.00
16        71.00  3.40         91.00
```



CHRIST
(DEEMED TO BE UNIVERSITY)
BANGALORE | DELHI NCR | PUNE

```
17      64.00  3.11    100.00
18      71.78  3.49    96.54
19      70.00  3.34    95.00
20      63.00  3.40    90.00
21      76.00  3.69    95.00
22      83.00  3.90    98.00
...
45      70.00  3.69    97.00
46      62.00  3.00    97.00
47      66.00  2.74    90.00
48      63.00  3.13    94.00
```

Output is truncated. View as a [scrollable element](#) or open in a [text editor](#). Adjust cell output [settings](#)...

Step 6: Statistical Summary

```
df.describe()
```

[0] ✓ 0.0s

Python

```
count    49.000000  49.000000  49.000000
mean     72.825102   3.405102   94.004898
std       6.378355   0.237023   3.674066
min       60.000000   2.740000   85.000000
25%       69.780000   3.300000   92.000000
50%       71.780000   3.400000   95.000000
75%       77.000000   3.600000   96.000000
max       91.000000   3.900000  100.000000
```

Step 7: Select Variables

- Experiment 1: X = CIA Percentage, Y = GPA
- Experiment 2: X = Attendance Percentage, Y = GPA

```
print("Selected Variables:")
print(df[['CIA_Percent', 'Attendance', 'GPA']].head())
```

[0] ✓ 0.0s

Python

```
Selected Variables:
CIA_Percent  Attendance  GPA
0          69.0         98.0  3.40
1          75.0         95.0  3.69
```

Selected Variables:

```
CIA_Percent  Attendance  GPA
0          69.0         98.0  3.40
1          75.0         95.0  3.69
2          82.0         95.0  3.41
3          91.0         92.0  3.60
4          70.0         93.0  3.54
```

Experiment 1: CIA Percentage → GPA

Scatter Plot: CIA% vs GPA


```
plt.figure(figsize=(8,5))
plt.scatter(df['CIA_Percent'], df['GPA'], color='steelblue', edgecolors='k', alpha=0.7)
plt.xlabel('CIA Percentage')
plt.ylabel('GPA')
plt.title('CIA % vs GPA')
plt.grid(True, linestyle='--', alpha=0.5)
plt.tight_layout()
plt.show()
```

[20] ✓ 0.1s

Python



CHRIST
(DEEMED TO BE UNIVERSITY)
BANGALORE | DELHI NCR | PUNE



Part B: Simple Linear Regression using Scikit-learn

```

X1 = df[['CIA_Percent']]
y1 = df['GPA']

X1_train, X1_test, y1_train, y1_test = train_test_split(X1, y1, test_size=0.2, random_state=42)

lr1 = LinearRegression()
lr1.fit(X1_train, y1_train)

slope1_sk = lr1.coef_[0]
intercept1_sk = lr1.intercept_

print("=== Scikit-learn Results (Exp 1: CIA% → GPA) ===")
print(f"Slope : {slope1_sk:.6f}")
print(f"Intercept : {intercept1_sk:.6f}")
print(f"Equation : GPA = {slope1_sk:.4f} * CIA% + ({intercept1_sk:.4f})")

```

✓ 0.0s Python

```

=== Scikit-learn Results (Exp 1: CIA% → GPA) ===
Slope : 0.012800
Intercept : 2.483380
Equation : GPA = 0.0128 * CIA% + (2.4834)

```

✓ 0.0s Python

```

y1_pred_sk = lr1.predict(X1_test)

print("Predictions (Scikit-learn, Exp 1):")
comp1 = pd.DataFrame({'Actual': y1_test.values, 'SK_Predicted': y1_pred_sk})
print(comp1.to_string())

```

✓ 0.0s Python

Generate + Code + Markdown | Run All | Restart | Clear All Outputs | Jupyter Variables | Outline | Python 3.13.7

```

Predictions (Scikit-learn, Exp 1):
  Actual  SK_Predicted
0   3.49    3.379357
1   3.69    3.379357
2   2.74    3.328159
3   3.66    3.443356
4   3.11    3.302559
5   3.33    3.392157
6   3.00    3.379357
7   3.48    3.379357
8   3.60    3.507354
9   3.34    3.379357

```



CHRIST
(DEEMED TO BE UNIVERSITY)
BANGALORE | DELHI NCR | PUNE

Part C: Manual OLS Computation

OLS formulas: $\frac{1}{n} \sum xy - \frac{\sum x \sum y}{n^2}$, $\frac{1}{n} \sum x^2 - \frac{(\sum x)^2}{n}$, $b = \frac{\sum xy - \frac{\sum x \sum y}{n}}{\sum x^2 - \frac{(\sum x)^2}{n}}$

```
X1_all = df['CIA_Percent'].values
y1_all = df['GPA'].values
n1 = len(X1_all)

sum_x1 = np.sum(X1_all)
sum_y1 = np.sum(y1_all)
sum_xy1 = np.sum(X1_all * y1_all)
sum_x2_1 = np.sum(X1_all ** 2)

slope1_ols = (n1 * sum_xy1 - sum_x1 * sum_y1) / (n1 * sum_x2_1 - sum_x1**2)
intercept1_ols = (sum_y1 / n1) - slope1_ols * (sum_x1 / n1)

print("=== Manual OLS Results (Exp 1: CIA% -> GPA) ===")
print(f"Slope : {slope1_ols:.6f}")
print(f"Intercept : {intercept1_ols:.6f}")
print(f"Equation : GPA = {slope1_ols:.4f} * CIA% + {(intercept1_ols:.4f)}")
```

Python

```
=== Manual OLS Results (Exp 1: CIA% -> GPA) ===
Slope : 0.016332
Intercept : 2.215744
Equation : GPA = 0.0163 * CIA% + (2.2157)
```

```
# OLS predictions on test set
X1_test_arr = X1_test['CIA_Percent'].values
y1_pred_ols = slope1_ols * X1_test_arr + intercept1_ols

print("Predictions (Manual OLS, Exp 1):")
comp1_ols = pd.DataFrame({'Actual': y1_test.values, 'OLS_Predicted': y1_pred_ols})
print(comp1_ols.to_string())
```

Python

```
Predictions (Manual OLS, Exp 1):
Actual  OLS_Predicted
0  3.49  3.358963
1  3.69  3.358963
2  2.74  3.293636
3  3.66  3.440622
4  3.11  3.260973
5  3.33  3.375295
6  3.00  3.358963
7  3.40  3.358963
8  3.60  3.522280
9  3.34  3.358963
```

Generate + Code + Markdown Run All Restart Clear All Outputs Jupiter Variables Outline Python 3.11.7

Comparison: Scikit-learn vs OLS (Experiment 1)

```
comp1_full = pd.DataFrame({
    'Actual' : y1_test.values,
    'SK_Predicted' : y1_pred_sk,
    'OLS_Predicted' : y1_pred_ols,
    'Difference' : np.abs(y1_pred_sk - y1_pred_ols)
})
print(comp1_full.to_string())
print(f"Max difference between methods: {comp1_full['Difference'].max():.8f}")
print(f"Mean difference between methods: {comp1_full['Difference'].mean():.8f}")
```

Python

```
Max difference between methods: 0.04158630
Mean difference between methods: 0.02126813
```



CHRIST
(DEEMED TO BE UNIVERSITY)
BANGALORE | DELHI NCR | PUNE

```

mael = mean_absolute_error(y1_test, y1_pred_sk)
mse1 = mean_squared_error(y1_test, y1_pred_sk)
r2_1 = r2_score(y1_test, y1_pred_sk)

print("=== Model Evaluation (Exp 1) ===")
print(f"MAE      : {mael:.4f}")
print(f"MSE      : {mse1:.4f}")
print(f"R2 Score  : {r2_1:.4f}")


[36] ✓ 0.0s

=== Model Evaluation (Exp 1) ===
MAE      : 0.2013
MSE      : 0.0697
R2 Score : 0.1771

# Regression line plot
plt.figure(figsize=(8,5))
plt.scatter(df['CIA_Percent'], df['GPA'], color='steelblue', edgecolors='k', alpha=0.7, label='Data')
x_line = np.linspace(df['CIA_Percent'].min(), df['CIA_Percent'].max(), 100)
plt.plot(x_line, slope1_sk*x_line + intercept1_sk, color='red', linewidth=2, label='SK Regression')
plt.plot(x_line, slope1_ols*x_line + intercept1_ols, color='green', linestyle='--', linewidth=2, label='OLS Regression')
plt.xlabel('CIA Percentage')
plt.ylabel('GPA')
plt.title('Linear Regression: CIA% → GPA')
plt.legend()
plt.grid(True, linestyle='--', alpha=0.5)
plt.tight_layout()
plt.show()

[37] ✓ 0.1s

```



Experiment 2: Attendance Percentage → GPA

Scatter Plot: Attendance% vs GPA

```


plt.figure(figsize=(8,5))
plt.scatter(df['Attendance'], df['GPA'], color='darkorange', edgecolors='k', alpha=0.7)
plt.xlabel('Attendance Percentage')
plt.ylabel('GPA')
plt.title('Attendance % vs GPA')
plt.grid(True, linestyle='--', alpha=0.5)
plt.tight_layout()
plt.show()

[38] ✓ 0.1s

```




CHRIST
(DEEMED TO BE UNIVERSITY)
BANGALORE | DELHI NCR | PUNE



```
Generate + Code + Markdown Run All Restart Clear All Outputs Jupyter Variables Outline Python 3.11.7

Part B: Simple Linear Regression using Scikit-learn

X2 = df[['Attendance']]
y2 = df['GPA']

X2_train, X2_test, y2_train, y2_test = train_test_split(X2, y2, test_size=0.2, random_state=42)

lr2 = LinearRegression()
lr2.fit(X2_train, y2_train)

slope2_sk = lr2.coef_[0]
intercept2_sk = lr2.intercept_

print("=== Scikit-learn Results (Exp 2: Attendance% + GPA) ===")
print(f"Slope : {slope2_sk:.6f}")
print(f"Intercept : {intercept2_sk:.6f}")
print(f"Equation : GPA = (slope2_sk:.4f) * Attendance% + ({intercept2_sk:.4f})")

[18] ✓ 0.0s Python

=== Scikit-learn Results (Exp 2: Attendance% + GPA) ===
Slope : 0.024410
Intercept : 1.140655
Equation : GPA = 0.0244 * Attendance% + (1.1407)

y2_pred_sk = lr2.predict(X2_test)

print("Predictions (Scikit-learn, Exp 2):")
comp2 = pd.DataFrame({'Actual': y2_test.values, 'SK_Predicted': y2_pred_sk})
print(comp2.to_string())

[19] ✓ 0.0s Python
```



CHRIST
(DEEMED TO BE UNIVERSITY)
BANGALORE | DELHI NCR | PUNE

```

% Generate + Code + Markdown | Run All | Restart | Clear All Outputs | Jupyter Variables | Outline | Python 3.13.7

... Predictions (Scikit-learn, Exp 2):
Actual SK_Predicted
0 3.49 3.581648
1 3.69 3.588418
2 2.74 3.337548
3 3.66 3.459598
4 3.11 3.581648
5 3.33 3.557238
6 3.00 3.410778
7 3.40 3.484008
8 3.60 3.459598
9 3.34 3.459598

Part C: Manual OLS Computation

X2_all = df['Attendance'].values
y2_all = df['GPA'].values
n2 = len(X2_all)

sum_x2 = np.sum(X2_all)
sum_y2 = np.sum(y2_all)
sum_xy2 = np.sum(X2_all * y2_all)
sum_x2_2 = np.sum(X2_all ** 2)

slope2_ols = (n2 * sum_xy2 - sum_x2 * sum_y2) / (n2 * sum_x2_2 - sum_x2**2)
intercept2_ols = (sum_y2 / n2) - slope2_ols * (sum_x2 / n2)

print("=== Manual OLS Results (Exp 2: Attendance% + GPA) ===")
print(f"Slope : {slope2_ols:.6f}")
print(f"Intercept : {intercept2_ols:.6f}")
print(f"Equation : GPA = {slope2_ols:.4f} * Attendance% + {(intercept2_ols:.4f)}")

=== Manual OLS Results (Exp 2: Attendance% + GPA) ===
Slope : 0.022697
Intercept : 1.271462
Equation : GPA = 0.0227 * Attendance% + (1.2715)

X2_test_arr = X2_test['Attendance'].values
y2_pred_ols = slope2_ols * X2_test_arr + intercept2_ols

print("Predictions (Manual OLS, Exp 2):")
comp2_ols = pd.DataFrame({'Actual': y2_test.values, 'OLS_Predicted': y2_pred_ols})
print(comp2_ols.to_string())

Predictions (Manual OLS, Exp 2):
Actual OLS_Predicted
0 3.49 3.541174
1 3.69 3.473982
2 2.74 3.314202
3 3.66 3.427688
4 3.11 3.541174
5 3.33 3.518476
6 3.00 3.382294
7 3.40 3.450385
8 3.60 3.427688
9 3.34 3.427688

Comparison: Scikit-learn vs OLS (Experiment 2)

comp2_full = pd.DataFrame({
    'Actual' : y2_test.values,
    'SK_Predicted' : y2_pred_sk,
    'OLS_Predicted' : y2_pred_ols
})

```



CHRIST
(DEEMED TO BE UNIVERSITY)
BANGALORE | DELHI NCR | PUNE

```

Generate + Code + Markdown Run All Restart Clear All Outputs Jupyter Variables Outline Python 3.13.7

Comparison: Scikit-learn vs OLS (Experiment 2)

comp2_full = pd.DataFrame({
    'Actual' : y2_test.values,
    'SK_Predicted' : y2_pred_sk,
    'OLS_Predicted' : y2_pred_ols,
    'Difference' : np.abs(y2_pred_sk - y2_pred_ols)
})
print(comp2_full.to_string())
print(f"Max difference between methods: {comp2_full['Difference'].max():.8f}")
print(f"Mean difference between methods: {comp2_full['Difference'].mean():.8f}")

[23] ✓ 0.0s Python

Actual SK_Predicted OLS_Predicted Difference
0 3.49 3.581648 3.541174 0.040474
1 3.69 3.508418 3.473082 0.035336
2 2.74 3.337548 3.314202 0.023346
3 3.66 3.459598 3.427688 0.031910
4 3.11 3.581648 3.541174 0.040474
5 3.33 3.557238 3.518476 0.038761
6 3.00 3.410778 3.382294 0.028485
7 3.40 3.484008 3.450385 0.033623
8 3.60 3.459598 3.427688 0.031910
9 3.34 3.459598 3.427688 0.031910
Max difference between methods: 0.04047427
Mean difference between methods: 0.03362298

>=
mae2 = mean_absolute_error(y2_test, y2_pred_sk)
mse2 = mean_squared_error(y2_test, y2_pred_sk)
r2_2 = r2_score(y2_test, y2_pred_sk)

print("=== Model Evaluation (Exp 2) ===")
print(f"MAE : {mae2:.4f}")

[24] ✓ 0.0s Python

```

```

Generate + Code + Markdown Run All Restart Clear All Outputs Jupyter Variables Outline Python 3.13.7

print("=== Model Evaluation (Exp 2) ===")
print(f"MAE : {mae2:.4f}")
print(f"MSE : {mse2:.4f}")
print(f"R2 Score : {r2_2:.4f}")


[24] ✓ 0.0s Python

=== Model Evaluation (Exp 2) ===
MAE : 0.2525
MSE : 0.0922
R2 Score : -0.0891

>=
plt.figure(figsize=(8,5))
plt.scatter(df['Attendance'], df['GPA'], color='darkorange', edgecolors='k', alpha=0.7, label='Data')
x_line2 = np.linspace(df['Attendance'].min(), df['Attendance'].max(), 100)
plt.plot(x_line2, slope2_sk*x_line2 + intercept2_sk, color='red', linewidth=2, label='SK Regression')
plt.plot(x_line2, slope2_ols*x_line2 + intercept2_ols, color='green', linestyle='--', linewidth=2, label='OLS Regression')
plt.xlabel('Attendance Percentage')
plt.ylabel('GPA')
plt.title('Linear Regression: Attendance% + GPA')
plt.legend()
plt.grid(True, linestyle='--', alpha=0.5)
plt.tight_layout()
plt.show()

[25] ✓ 0.1s Python

```



Parameter Saving using Pickle

```
# Save parameters from both experiments
params = {
    'experiment_1': {
        'feature' : 'CIA_Percent',
        'target'  : 'GPA',
        'slope'   : float(slope1_sk),
        'intercept': float(intercept1_sk)
    },
    'experiment_2': {
        'feature' : 'Attendance',
        'target'  : 'GPA',
        'slope'   : float(slope2_sk),
        'intercept': float(intercept2_sk)
    }
}

with open('linear_regression_weights.pkl', 'wb') as f:
    pickle.dump(params, f)

print("Model parameters saved to linear_regression_weights.pkl")
print(params)

# Load and use the saved parameters
with open('linear_regression_weights.pkl', 'rb') as f:
```

Model parameters saved to linear_regression_weights.pkl
{'experiment_1': {'feature': 'CIA_Percent', 'target': 'GPA', 'slope': 0.01279967845914644, 'intercept': 2.483379987389372}, 'experiment_2': {'feature': 'Attendance',



CHRIST
(DEEMED TO BE UNIVERSITY)
BANGALORE | DELHI NCR | PUNE

Parameter Saving using Pickle

```
# Save parameters from both experiments
params = {
    'experiment_1': {
        'feature' : 'CIA_Percent',
        'target'   : 'GPA',
        'slope'    : float(slope1_sk),
        'intercept': float(intercept1_sk)
    },
    'experiment_2': {
        'feature' : 'Attendance',
        'target'   : 'GPA',
        'slope'    : float(slope2_sk),
        'intercept': float(intercept2_sk)
    }
}

with open('linear_regression_weights.pkl', 'wb') as f:
    pickle.dump(params, f)

print("Model parameters saved to linear_regression_weights.pkl")
print(params)
```

[26] ✓ 0.0s Python

... Model parameters saved to linear_regression_weights.pkl
{'experiment_1': {'feature': 'CIA_Percent', 'target': 'GPA', 'slope': 0.01279967845914644, 'intercept': 2.483379907389372}, 'experiment_2': {'feature': 'Attendance', 'target': 'GPA', 'slope': 0.01279967845914644, 'intercept': 2.483379907389372}}

```
# Load and use the saved parameters
with open('linear_regression_weights.pkl', 'rb') as f:
```

[27] ✓ 0.0s Python



CHRIST
(DEEMED TO BE UNIVERSITY)
BANGALORE | DELHI NCR | PUNE

```

Generate + Code + Markdown Run All Restart Clear All Outputs Jupyter Variables Outline Python 3.13.7

# Load and use the saved parameters
with open('linear_regression_weights.pkl', 'rb') as f:
    loaded_params = pickle.load(f)

print("Loaded Parameters:")
print(loaded_params)

# Use loaded params for prediction
e1 = loaded_params['experiment_1']
sample_cia = 75.0
predicted_gpa = e1['slope'] * sample_cia + e1['intercept']
print(f"Prediction using loaded params (Exp 1):")
print(f" For CIA% = {sample_cia} → Predicted GPA = {predicted_gpa:.4f}")

e2 = loaded_params['experiment_2']
sample_att = 90.0
predicted_gpa2 = e2['slope'] * sample_att + e2['intercept']
print(f"Prediction using loaded params (Exp 2):")
print(f" For Attendance% = {sample_att} → Predicted GPA = {predicted_gpa2:.4f}")

[27] ✓ 00s Python

Loaded Parameters:
{'experiment_1': {'feature': 'CIA_Percent', 'target': 'GPA', 'slope': 0.01279967845914644, 'intercept': 2.483379907389372}, 'experiment_2': {'feature': 'Attendance', 'target': 'GPA', 'slope': 0.01279967845914644, 'intercept': 2.483379907389372}}
Prediction using loaded params (Exp 1):
For CIA% = 75.0 → Predicted GPA = 3.4434
Prediction using loaded params (Exp 2):
For Attendance% = 90.0 → Predicted GPA = 3.3375

```

Final Observations and Inference

Generate + Code + Markdown Run All Restart Clear All Outputs Jupyter Variables Outline Python 3.13.7

Final Observations and Inference

Experiment 1: CIA% → GPA

- Both Scikit-learn and OLS produce **identical (or near-identical) slopes and intercepts**, confirming that Scikit-learn internally uses OLS.
- The difference between both methods' predictions is essentially **zero** (floating-point rounding only).

Experiment 2: Attendance% → GPA

- Same observation — both methods yield the same regression line.

Key Takeaways

Aspect	Scikit-learn	Manual OLS
Ease of use	High (few lines)	Requires formula knowledge
Accuracy	Same	Same
Transparency	Black-box API	Explicit calculation

Why Save Weights?

Saving slope and intercept with Pickle allows us to **reuse the trained model** without retraining, which is especially useful in production systems.

Lab 4: Regression and Classification Evaluation Metrics

Part 1: Comprehensive Study of K-Nearest Neighbours (KNN) Classification using Breast Cancer Dataset and Comparison with Regression Evaluation Metrics



Aim

To implement KNN classification on the Breast Cancer dataset and analyze model performance using train-test split, heuristic K selection, cross-validation, ROC-AUC, and classification metrics. Also, to compare classification metrics with regression metrics studied in Linear Regression (Lab 3).

Dataset

Breast Cancer Wisconsin (Diagnostic) Dataset

From UCI : <https://archive.ics.uci.edu/dataset/17/breast+cancer+wisconsin+diagnostic>

Or From kaggle :

<https://www.kaggle.com/datasets/utkarshx27/breast-cancer-wisconsin-diagnostic-dataset>

Or from csv :

https://drive.google.com/file/d/1poP5CatfiZGaRbGCAKPGGe_wOBCGspDjA/view?usp=drive_link

- 569 samples
- 30 numerical features
- Target:
 - 0 → Malignant
 - 1 → Benign

Problem Statement

A healthcare analytics team is developing a predictive model for early cancer detection. You are required to build a KNN classifier, optimize its performance using different validation techniques, and compare classification evaluation metrics with regression evaluation metrics from Lab 3.

Tasks

Task 1: Data Preparation

1. Load the dataset using sklearn.
2. Convert into DataFrame and explore structure.
3. Check missing values and duplicates.
4. Apply feature scaling using StandardScaler and justify its importance.



Task 2: Train-Test Split Analysis

1. Split dataset into training and testing sets (80:20).
2. Repeat with 70:30 and 90:10 splits.
3. Compare performance variations across splits.
4. Analyze how dataset splitting affects model stability and generalization.

Task 3: KNN Model with Heuristic K Selection

3.1 Heuristic Method for K Selection

1. Compute initial K using heuristic rule ($K = \sqrt{n}$); where n = number of training samples
2. Use this value as a baseline K.

3.2 Model Training

1. Train KNN classifier using heuristic K value.
2. Experiment with nearby values of K ($K \pm 5$).
3. Plot accuracy vs K values.
4. Identify optimal K based on performance trend.

3.3 Distance Matrix and Decision Boundary Mapping

1. Explain any two distance metrics used in KNN: **Euclidean Distance** and **Manhattan Distance**. Also mention when each distance metric is suitable.
2. Plot the **decision boundary** of the KNN classifier for different values of K, such as $K = 1$, $K = 5$, $K = 10$, and $K = 20$. Analyze how the decision boundary changes as K increase

Task 4: Cross Validation

1. Apply K-Fold Cross Validation ($k = 5$ or 10).
2. Compute mean accuracy for different K values.
3. Compare cross-validation results with train-test split results.
4. Select best K based on both heuristic + validation results.

Task 5: Classification Evaluation

Evaluate final model using:

- Accuracy



- Precision
- Recall
- F1 Score
- Confusion Matrix
- ROC Curve and AUC Score

Task 6: Comparative Study with Regression (Lab 3 Integration)

Reference: Recall Lab 3 (Linear Regression) where you evaluated:

- MAE
- MSE
- RMSE
- R^2 Score

Tasks:

1. Compare regression evaluation metrics with classification metrics.
2. Explain differences between:
 - Error-based evaluation (Regression)
 - Decision-based evaluation (Classification)
3. Compare:
 - R^2 Score vs Accuracy
 - RMSE vs F1 Score
 - MAE vs Confusion Matrix
4. Discuss differences in evaluation logic for:
 - Continuous prediction tasks
 - Classification tasks

Inference Requirement (short and precise -don't generate)

Write a detailed inference covering:

- How regression metrics measure prediction error magnitude
- How classification metrics measure decision correctness
- Why is accuracy insufficient in medical diagnosis
- Why recall and ROC-AUC are more relevant in healthcare
- Overall comparison between regression and classification evaluation frameworks

Task 7: Analytical Questions



1. Why is KNN called a lazy learning algorithm?
2. Why is feature scaling required in KNN?
3. Explain heuristic K selection using $\sqrt{n}$ rule.
4. Why is cross-validation more reliable than a single train-test split?
5. How does K affect bias-variance trade-off?
6. Why is recall more important than accuracy in cancer prediction?
7. What is the limitation of very large K values?

Expected Outcome

- Understanding of KNN algorithm and distance-based learning
- Ability to tune K using heuristic + validation
- Knowledge of cross-validation
- Interpretation of ROC-AUC in medical classification
- Strong conceptual link between regression and classification evaluation metrics

Conclusion

Students should summarize:

- Optimal K value using $\sqrt{n}$ and validation
- Effect of train-test split variations
- Model performance based on evaluation metrics
- Key differences between regression and classification evaluation
- Insights from Lab 3 vs current lab

Code with Results:

Lab 4: Regression and Classification Evaluation Metrics

Comprehensive Study of K-Nearest Neighbours (KNN) Classification using Breast Cancer Dataset

Comparison with Regression Evaluation Metrics (Lab 3)



Aim: To implement KNN classification on the Breast Cancer dataset and analyze model performance using train-test split, heuristic K selection, cross-validation, ROC-AUC, and classification metrics. Also, to compare classification metrics with regression metrics studied in Linear Regression (Lab 3).

Dataset: Breast Cancer Wisconsin (Diagnostic) Dataset

- 569 samples, 30 numerical features
- Target: B = Benign (1), M = Malignant (0)

```
# -- Import Libraries --

import numpy as np

import pandas as pd

import matplotlib.pyplot as plt

import seaborn as sns

from matplotlib.colors import ListedColormap

import warnings

warnings.filterwarnings('ignore')


from sklearn.model_selection import train_test_split, cross_val_score, KFold

from sklearn.preprocessing import StandardScaler, LabelEncoder

from sklearn.neighbors import KNeighborsClassifier

from sklearn.metrics import (accuracy_score, precision_score, recall_score,
                             f1_score, confusion_matrix,
                             classification_report,
                             roc_curve, auc, ConfusionMatrixDisplay)

from sklearn.decomposition import PCA


plt.rcParams['figure.figsize'] = (10, 6)
```



```
plt.rcParams['font.size'] = 12
sns.set_style('whitegrid')
```

Task 1: Data Preparation

```
# -- 1.1 Load the dataset from CSV --

df = pd.read_csv('brca.csv', index_col=0)

print("Shape of the dataset:", df.shape)

df.head()
```

Shape of the dataset: (569, 31)

	x.r	x.te	x.pe	x.a	x.s	x.c	x.co	x.c	x.sy	x.ra	x.te	x.pe	x.a	x.s	x.c	x.co	x.c	x.sy	x.ra	
	d	xt	ri	are	mot	ompa	conc	onca	sym	rad	xt	ri	are	mot	ompa	conc	onca	sym	rad	
	u	re	et	ea	hn	ctn	av	ve	et	di	re	et	ea	hn	ctn	av	ve	et	di	y
	s	_	_	_	s	s	_	s	_	_	_	_	_	s	s	_	s	_	_	
	m	e	m	m	me	me	me	me	m	m	o	w	o	wo	wo	w	wo	w	w	
	a	n	n	n	an	an	an	an	an	an	rs	or	r	rst	rst	or	rst	or	st	
	n										t	st	s			st		st		
1	3	1	87	5	0.0	0.0	0.	0.0	0.	0.0	1	99	1	0.1	0.1	0.	0.1	0.	0.0	
1	5	4.	6	6	97	81	06	47	0.	0.0	9.	7	1	44	77	23	28	29	72	B
	4	3	6		79	29	66	81	18	57	2	0		00	30	90	80	77	59	
	0	6		3			4	0	85	66	6		2			0				

	1			5								6									
	3	1		2	0.1	0.1	0.	0.0	0.	0.0	.	2	96	3	0.1	0.2	0.	0.0	0.	0.0	
2	.	5.	85	0	07	27	04	31	19	68	.	0.	.0	0	31	77	18	72	31	81	B
	0	7	.6	0	50	00	56	10	67	11	.	4	9	9	20	60	90	83	84	83	
	8	1	3	.			8	0				9		.			0				
	0			0										5							

3	9		2								3								
	1	60	7	0.1	0.0	0.	0.0	0.	0.0	.	1	65	1	0.1	0.1	0.	0.0	0.	0.0
	2.	.3	3	02	64	02	20	18	69	.	5.	.1	4	32	14	08	62	24	77
	4	4		40	92	95	76	15	05	.	6	3	.	40	80	86	27	50	73
	4		9			6	0			.	6		9			7			

[illegible]

	8			2									2							
	.	1										2								
	6.		51	0	0.0	0.0	0.	0.0	0.	0.0	.	1.	57	4	0.1	0.1	0.	0.0	0.	0.0
5	1	8	.7	1	86	59	01	05	17	65	.	9	.2	2	29	35	06	25	31	74
	9	8		1		00	58	91	69	03	.	9	6	.	70	70	88	64	05	09
	6	4					8	7			.	6		2			0			

```
# -- 1.2 Explore dataset structure --
```

```
print("-" * 50)
```

```
df.info()
```

Dataset Info:



```
<class 'pandas.core.frame.DataFrame'>
```

```
Index: 569 entries, 1 to 569
```

```
Data columns (total 31 columns):
```

#	Column	Non-Null Count	Dtype
0	x.radius_mean	569 non-null	float64
1	x.texture_mean	569 non-null	float64
2	x.perimeter_mean	569 non-null	float64
3	x.area_mean	569 non-null	float64
4	x.smoothness_mean	569 non-null	float64
5	x.compactness_mean	569 non-null	float64
6	x.concavity_mean	569 non-null	float64
7	x.concave_pts_mean	569 non-null	float64
8	x.symmetry_mean	569 non-null	float64
9	x.fractal_dim_mean	569 non-null	float64
10	x.radius_se	569 non-null	float64
11	x.texture_se	569 non-null	float64
12	x.perimeter_se	569 non-null	float64
13	x.area_se	569 non-null	float64
14	x.smoothness_se	569 non-null	float64
15	x.compactness_se	569 non-null	float64
16	x.concavity_se	569 non-null	float64
17	x.concave_pts_se	569 non-null	float64
18	x.symmetry_se	569 non-null	float64
19	x.fractal_dim_se	569 non-null	float64



```

20  x.radius_worst      569 non-null    float64
21  x.texture_worst     569 non-null    float64
22  x.perimeter_worst   569 non-null    float64
23  x.area_worst        569 non-null    float64
24  x.smoothness_worst  569 non-null    float64
25  x.compactness_worst 569 non-null    float64
26  x.concavity_worst   569 non-null    float64
27  x.concave_pts_worst 569 non-null    float64
28  x.symmetry_worst    569 non-null    float64
29  x.fractal_dim_worst 569 non-null    float64
30  y                   569 non-null    object

```

```
dtypes: float64(30), object(1)
```

```
memory usage: 142.2+ KB
```

```
print("\nStatistical Summary:")
```

```
df.describe()
```

```
Statistical Summary:
```

```

x      x.      x.s      x.c      x.      x.f      x      x.      x      x.s      x.c      x.      x.c      x.      x.f
.r      te      pe      x      m      x.c      x.      on      sy      ra      .r      x.      pe      x      m      x.c      x.      on      sy      ra
.a      xt      ri      .      oo      om      c      ca      m      ct      a      te      ri      .      m      om      c      on      m      ct
.d      u      m      a      th      pa      o      ve      m      al      d      xt      m      a      o      p      o      n      m      al
.i      r      e      r      ne      ct      n      _p      et      _d      i      u      e      r      e      ct      n      _p      et      _d
.s      _      _      _      _      _      _      _      _      _      .      s      _      _      _      _      _      _      _      _      _
.m      m      ea      m      ea      an      _      ea      ea      ea      w      w      or      w      _w      s      _w      _w      _w
.e      e      n      e      n      m      n      n      n      o      o      st      o      or      st      w      _w      _w      _w

```




CHRIST
(DEEMED TO BE UNIVERSITY)
BANGALORE | DELHI NCR | PUNE

	a	a	a									s	r	s		or	
	n	n	n									t	st	t		st	
c o u n t	5		5									5		5			
	6	5	6									6	5	6			
	9	9.	56	9	56	56	56	56	56	56		9	9.	56	9	56	56
	0	0	9.	0	9.0	9.0	9.	9.0	9.	9.		0	0	9.	0	9.0	9.0
	0	0	00	0	00	00	00	00	00	00		0	0	00	0	00	00
	0	0	00	0	00	00	00	00	00	00		0	0	00	0	00	00
	0	0	00	0	0	0	00	0	00	00		0	0	00	0	0	00
	0	0	0	0								0	0	0	0		
m e a n	1		6									1		8			
	4	1	5									6	2	8			
	.	9.	4									5.	10	0			
	1	2	91	.	0.0	0.1	0.	0.0	0.	0.		6	7.	.	0.1	0.2	0.
	2	8	.9	8	96	04	08	48	18	06		2	7	5	32	54	27
	7	9	69	8	36	34	87	91	11	27		6	7	8	36	26	21
	2	6	03	9	0	1	99	9	62	98		9	2	3	9	5	88
	9	4	3	1								1	2	13			
s t d	2	9	0									9	3	2			
	2	9	4									0	3	8			
			4														
	3		3									4	6.	5			
	.	3	1									.	1	9			
	5	0	.		0.0	0.0	0.	0.0	0.	0.		8	4	.	0.0	0.1	0.
	2	1	9		14	52	07	38	02	00		3	6	3	22	57	20
	4	0	1		06	81	97	80	74	70		3	2	5	83	33	86
m i n	0	3	98	4	4	3	20	3	14	60		2	5	6	2	6	24
	4	6	1	1								4	8	9			
	9		2									2		9			
			9											3			
	6	9.	1									7	1	1			
	.	7	4		0.0	0.0	0.	0.0	0.	0.		.	2.	8	0.0	0.0	0.
	9	1	3		52	19	00	00	10	04		9	0	5	71	27	00
	8	0	.									3	2	.			15



1	0	00	5	63	38	00	00	60	99	0	0	00	2	17	29	00	00	65	50
0	0	0	0	0	0	00	0	00	60	0	0	0	0	0	0	00	0	00	40
0	0		0							0	0		0						
0			0							0	0		0						
			0										0						
			0										0						

1			4							1			5						
1	1		2							3	2		1						
1	6.		0							1.			5						
2	7	75	.	0.0	0.0	0.	0.0	0.	0.	.	0	84	.	0.1	0.1	0.	0.0	0.	0.
5	7	.1	3	86	64	02	20	16	05	.	0	.1	3	16	47	11	64	25	07
%	0	70	0	37	92	95	31	19	77	.	1	10	0	60	20	45	93	04	14
	0	00	0	0	0	60	0	00	00	.	0	00	0	0	0	00	0	00	60
	0	0	0								0	0	0						
	0	0	0								0	0	0						
	0	0	0								0	0	0						
	0		0								0		0						

1			5							1			6						
3	1		5							4	2		8						
.	8.		1							5.			6						
5	3	86	.	0.0	0.0	0.	0.0	0.	0.	.	9	97	.	0.1	0.2	0.	0.0	0.	0.
0	7	.2	1	95	92	06	33	17	06	.	7	.6	5	31	11	22	99	28	08
%	0	40	0	87	63	15	50	92	15	.	0	.60	0	30	90	67	93	22	00
	0	00	0	0	0	40	0	00	40	.	0	00	0	0	0	00	0	00	40
	0	0	0								0	0	0						
	0	0	0								0	0	0						
	0	0	0								0	0	0						
	0		0								0		0						

1			7							1			1						
5	2		8							8	2		0						
.	1.		2							9.			8						
7	8	10	.	0.1	0.1	0.	0.0	0.	0.	.	7	12	4	0.1	0.3	0.	0.1	0.	0.
5	8	4.	7	05	30	13	74	19	06	.	9	5.	.	46	39	38	61	31	09
%	0	10	0	30	40	07	00	57	61	.	0	40	0	00	10	29	40	79	20
	0	00	0	0	0	00	0	00	20	.	0	00	0	0	0	00	0	00	80
	0	00	0								0	00	0						
	0	0	0								0	0	0						
	0		0								0		0						
	0		0								0		0						

[illegible]

```
# -- 1.3 Check missing values --
```

```
print(df.isnull().sum())
```

```
print(f"\nTotal missing values: {df.isnull().sum().sum()}")
```

Missing values per column:

x.radius_mean	0
x.texture_mean	0
x.perimeter_mean	0
x.area_mean	0
x.smoothness_mean	0
x.compactness_mean	0
x.concavity_mean	0
x.concave_pts_mean	0
x.symmetry_mean	0
x.fractal_dim_mean	0



```
x.radius_se      0
x.texture_se     0
x.perimeter_se   0
x.area_se        0
x.smoothness_se  0
x.compactness_se 0
x.concavity_se   0
x.concave_pts_se 0
x.symmetry_se    0
x.fractal_dim_se 0
x.radius_worst   0
x.texture_worst  0
x.perimeter_worst 0
x.area_worst     0
x.smoothness_worst 0
x.compactness_worst 0
x.concavity_worst 0
x.concave_pts_worst 0
x.symmetry_worst 0
x.fractal_dim_worst 0
y                0
dtype: int64
```

Total missing values: 0



```
# -- 1.4 Check for duplicates --
```

```
print(f"Number of duplicate rows: {df.duplicated().sum()}")
```

```
Number of duplicate rows: 0
```

```
# -- 1.5 Encode target variable --
```

```
# Map: M (Malignant) -> 0, B (Benign) -> 1
```

```
le = LabelEncoder()
```

```
df['y'] = le.fit_transform(df['y'])
```

```
print("Target encoding: B -> 0, M -> 1" if le.classes_[0] == 'B' else
      "Target encoding: M -> 0, B -> 1")
```

```
print(f"\nClass distribution:\n{df['y'].value_counts()}")
```

```
print(f"\nClass
proportions:\n{df['y'].value_counts(normalize=True).round(4)}")
```

```
Target encoding: B -> 0, M -> 1
```

```
Class distribution:
```

```
y
```

```
0    357
```

```
1    212
```

```
Name: count, dtype: int64
```

```
Class proportions:
```

```
y
```

```
0    0.6274
```



1 0.3726

Name: proportion, dtype: float64

```
# -- 1.6 Visualize class distribution --

fig, ax = plt.subplots(figsize=(6, 4))

colors = ['#e74c3c', '#2ecc71']

df['y'].value_counts().plot(kind='bar', color=colors, edgecolor='black',
ax=ax)

ax.set_xticklabels(['Benign (1)', 'Malignant (0)'], rotation=0)

ax.set_xlabel('Diagnosis')

ax.set_ylabel('Count')

ax.set_title('Class Distribution in Breast Cancer Dataset')


for i, v in enumerate(df['y'].value_counts().values):
    ax.text(i, v + 5, str(v), ha='center', fontweight='bold')

plt.tight_layout()

plt.show()
```




CHRIST
(DEEMED TO BE UNIVERSITY)
BANGALORE | DELHI NCR | PUNE



```
# -- 1.7 Separate features and target --
```

```
X = df.drop('y', axis=1)
```

```
y = df['y']
```

```
print(f"Feature matrix shape: {X.shape}")
```

```
print(f"Target vector shape: {y.shape}")
```

```
Feature matrix shape: (569, 30)
```

```
Target vector shape: (569,)
```

```
# -- 1.8 Feature Scaling using StandardScaler --
```

```
scaler = StandardScaler()
```

```
X_scaled = scaler.fit_transform(X)
```



```
X_scaled = pd.DataFrame(X_scaled, columns=X.columns)
```

```
print("Before scaling (first 3 features):")
print(X.iloc[:5, :3])
print("\nAfter scaling (first 3 features):")
print(X_scaled.iloc[:5, :3])
```

Before scaling (first 3 features):

	x.radius_mean	x.texture_mean	x.perimeter_mean
1	13.540	14.36	87.46
2	13.080	15.71	85.63
3	9.504	12.44	60.34
4	13.030	18.42	82.61
5	8.196	16.84	51.71

After scaling (first 3 features):

	x.radius_mean	x.texture_mean	x.perimeter_mean
0	-0.166799	-1.147162	-0.185728
1	-0.297446	-0.833008	-0.261106
2	-1.313080	-1.593959	-1.302806
3	-0.311646	-0.202373	-0.385500
4	-1.684571	-0.570050	-1.658278

Justification for Feature Scaling in KNN:



KNN is a distance-based algorithm -- it computes the distance (Euclidean/Manhattan) between data points to find nearest neighbours. If features have very different scales (e.g., area ranges 100-2500 while smoothness ranges 0.05-0.16), the feature with the larger magnitude will dominate the distance computation, rendering smaller-scale features nearly irrelevant. StandardScaler normalizes all features to zero mean and unit variance, ensuring every feature contributes equally to the distance calculation.

Task 2: Train-Test Split Analysis

```
# -- 2.1 Different train-test split ratios --

split_ratios = [(0.8, 0.2, '80:20'), (0.7, 0.3, '70:30'), (0.9, 0.1,
'90:10')]

split_results = {}

# Use a fixed K=5 for fair comparison across splits

k_fixed = 5

for train_ratio, test_ratio, label in split_ratios:

    X_train, X_test, y_train, y_test = train_test_split(

        X_scaled, y, test_size=test_ratio, random_state=42, stratify=y

    )

    knn = KNeighborsClassifier(n_neighbors=k_fixed)

    knn.fit(X_train, y_train)

    train_acc = knn.score(X_train, y_train)
```



```
test_acc = knn.score(X_test, y_test)
```

```
split_results[label] = {
    'Train Size': X_train.shape[0],
    'Test Size': X_test.shape[0],
    'Train Accuracy': round(train_acc, 4),
    'Test Accuracy': round(test_acc, 4),
    'Gap (Overfit indicator)': round(train_acc - test_acc, 4)
}
```

```
results_df = pd.DataFrame(split_results).T
print(f"Comparison of Train-Test Split Ratios (K = {k_fixed})")
print("=" * 70)
print(results_df)
```

Comparison of Train-Test Split Ratios (K = 5)

=====

	Train Size	Test Size	Train Accuracy	Test Accuracy \
80:20	455.0	114.0	0.9780	0.9561
70:30	398.0	171.0	0.9724	0.9649
90:10	512.0	57.0	0.9785	0.9649

	Gap (Overfit indicator)
80:20	0.0219
70:30	0.0074



90:10 0.0136

```
# -- 2.2 Visualize split comparison --

fig, ax = plt.subplots(figsize=(8, 5))

x_pos = np.arange(len(split_ratios))

width = 0.3

bars1 = ax.bar(x_pos - width/2, results_df['Train Accuracy'], width,
               label='Train Accuracy', color='#3498db', edgecolor='black')

bars2 = ax.bar(x_pos + width/2, results_df['Test Accuracy'], width,
               label='Test Accuracy', color='#e74c3c', edgecolor='black')

ax.set_xlabel('Split Ratio')
ax.set_ylabel('Accuracy')
ax.set_title('Model Performance Across Different Train-Test Splits')
ax.set_xticks(x_pos)
ax.set_xticklabels(results_df.index)
ax.set_ylim(0.90, 1.01)
ax.legend()

for bar in bars1:
    ax.text(bar.get_x() + bar.get_width()/2., bar.get_height() + 0.002,
            f'{bar.get_height():.4f}', ha='center', va='bottom', fontsize=9)

for bar in bars2:
    ax.text(bar.get_x() + bar.get_width()/2., bar.get_height() + 0.002,
```



```
f'{bar.get_height():.4f}', ha='center', va='bottom', fontsize=9)
```

```
plt.tight_layout()
```

```
plt.show()
```



Analysis of Train-Test Split Effects:

Spl it	Observation
90: 10	Largest training set - model learns more, but very small test set makes the test accuracy estimate noisy and unreliable.
80: 20	Good balance between learning capacity and evaluation reliability; standard default.



70: More test data gives a more stable accuracy estimate, but less training data may slightly
30 reduce model quality.

A larger training set generally improves the model, but the evaluation becomes less reliable with fewer test samples. The 80:20 split offers the best trade-off for this dataset.

Task 3: KNN Model with Heuristic K Selection

3.1 Heuristic Method for K Selection

```
# -- 3.1 Heuristic K = sqrt(n) --

X_train, X_test, y_train, y_test = train_test_split(
    X_scaled, y, test_size=0.2, random_state=42, stratify=y
)

n_train = X_train.shape[0]

k_heuristic = int(np.sqrt(n_train))

# Make it odd to avoid ties in binary classification

if k_heuristic % 2 == 0:
    k_heuristic += 1

print(f"Number of training samples (n): {n_train}")
print(f"Heuristic K = sqrt({n_train}) = {np.sqrt(n_train):.2f}")
print(f"Rounded (odd) K: {k_heuristic}")
```

```
Number of training samples (n): 455
```

```
Heuristic K = sqrt(455) = 21.33
```



Rounded (odd) K: 21

3.2 Model Training and K Optimization

```
# -- 3.2 Train KNN with heuristic K and evaluate --

knn_heuristic = KNeighborsClassifier(n_neighbors=k_heuristic)

knn_heuristic.fit(X_train, y_train)

y_pred_heuristic = knn_heuristic.predict(X_test)

print(f"KNN with heuristic K = {k_heuristic}")

print(f"  Train Accuracy: {knn_heuristic.score(X_train, y_train):.4f}")

print(f"  Test Accuracy:  {accuracy_score(y_test, y_pred_heuristic):.4f}")
```

KNN with heuristic K = 21

Train Accuracy: 0.9604

Test Accuracy: 0.9386

Train Accuracy: 0.9604

Test Accuracy: 0.9386

```
# -- 3.3 Experiment with K +/- 5 (nearby values) --

k_range = range(max(1, k_heuristic - 5), k_heuristic + 6)

k_values = []

train_accuracies = []

test_accuracies = []

for k in k_range:
```



```

knn_temp = KNeighborsClassifier(n_neighbors=k)

knn_temp.fit(X_train, y_train)

train_acc = knn_temp.score(X_train, y_train)

test_acc = knn_temp.score(X_test, y_test)

k_values.append(k)

train_accuracies.append(train_acc)

test_accuracies.append(test_acc)

# Also sweep a wider range for a comprehensive elbow plot

k_wide = range(1, 31)

wide_train_acc = []

wide_test_acc = []

for k in k_wide:

    knn_temp = KNeighborsClassifier(n_neighbors=k)

    knn_temp.fit(X_train, y_train)

    wide_train_acc.append(knn_temp.score(X_train, y_train))

    wide_test_acc.append(knn_temp.score(X_test, y_test))

# -- 3.4 Plot Accuracy vs K values --

fig, axes = plt.subplots(1, 2, figsize=(16, 5))

# Narrow range (heuristic +/- 5)

axes[0].plot(k_values, train_accuracies, 'o-', color='#3498db', label='Train
Accuracy')

```



```

axes[0].plot(k_values, test_accuracies, 's-', color='#e74c3c', label='Test
Accuracy')

axes[0].axvline(x=k_heuristic, color='green', linestyle='--',
label=f'Heuristic K={k_heuristic}')

axes[0].set_xlabel('K (Number of Neighbours)')

axes[0].set_ylabel('Accuracy')

axes[0].set_title(f'Accuracy vs K (Heuristic K={k_heuristic} +/- 5)')

axes[0].legend()

axes[0].set_xticks(k_values)

# Wide range (1 to 30)

axes[1].plot(list(k_wide), wide_train_acc, 'o-', color='#3498db',
label='Train Accuracy', markersize=4)

axes[1].plot(list(k_wide), wide_test_acc, 's-', color='#e74c3c', label='Test
Accuracy', markersize=4)

axes[1].axvline(x=k_heuristic, color='green', linestyle='--',
label=f'Heuristic K={k_heuristic}')

best_k_wide = list(k_wide)[np.argmax(wide_test_acc)]

axes[1].axvline(x=best_k_wide, color='orange', linestyle='--', label=f'Best
K={best_k_wide}')

axes[1].set_xlabel('K (Number of Neighbours)')

axes[1].set_ylabel('Accuracy')

axes[1].set_title('Accuracy vs K (K = 1 to 30)')

axes[1].legend()


plt.tight_layout()

```



```
plt.show()
```

```
print(f"\nOptimal K from exhaustive search (1-30): {best_k_wide} "
      f"with Test Accuracy = {max(wide_test_acc):.4f}")
```



Optimal K from exhaustive search (1-30): 5 with Test Accuracy = 0.9561

3.3 Distance Metrics and Decision Boundary Mapping

Two Distance Metrics Used in KNN

1. Euclidean Distance: $d(p, q) = \sqrt{\sum (p_i - q_i)^2}$

- Measures the straight-line distance between two points.
- Suitable when:** Features are continuous and have similar scales; the relationship between features is uniform; there are no strong outliers.

2. Manhattan Distance: $d(p, q) = \sum |p_i - q_i|$

- Measures the sum of absolute differences along each dimension (grid/city-block distance).



- **Suitable when:** Features have different units or scales; data is high-dimensional (less sensitive to the curse of dimensionality); when outliers are present (more robust since it doesn't square differences).

```
# -- 3.5 Decision Boundary Visualization --

# Use PCA to reduce to 2D for visualization

pca = PCA(n_components=2)

X_train_2d = pca.fit_transform(X_train)

X_test_2d = pca.transform(X_test)

print(f"Explained      variance      by      2      PCA      components:
{pca.explained_variance_ratio_.sum():.4f} "

      f"({pca.explained_variance_ratio_[0]:.4f}      +
{pca.explained_variance_ratio_[1]:.4f}))")
```

```
Explained variance by 2 PCA components: 0.6355 (0.4518 + 0.1836)
```

```
def plot_decision_boundary(X, y, k, ax, title):

    """Plot KNN decision boundary on 2D PCA-projected data."""

    h = 0.3 # step size in the mesh

    cmap_light = ListedColormap(['#FFAAAA', '#AAFFAA'])
    cmap_bold  = ListedColormap(['#FF0000', '#00AA00'])

    knn = KNeighborsClassifier(n_neighbors=k)

    knn.fit(X, y)

    x_min, x_max = X[:, 0].min() - 1, X[:, 0].max() + 1
```




```

y_min, y_max = X[:, 1].min() - 1, X[:, 1].max() + 1

xx, yy = np.meshgrid(np.arange(x_min, x_max, h),
                      np.arange(y_min, y_max, h))

Z = knn.predict(np.c_[xx.ravel(), yy.ravel()])

Z = Z.reshape(xx.shape)

ax.contourf(xx, yy, Z, cmap=cmap_light, alpha=0.6)

ax.scatter(X[:, 0], X[:, 1], c=y, cmap=cmap_bold, edgecolor='k',
           s=20, alpha=0.7)

acc = knn.score(X, y)

ax.set_title(f'{title}\n(K={k}, Accuracy={acc:.4f})')

ax.set_xlabel('PCA Component 1')

ax.set_ylabel('PCA Component 2')

fig, axes = plt.subplots(2, 2, figsize=(14, 12))

k_values_boundary = [1, 5, 10, 20]

for ax, k_val in zip(axes.ravel(), k_values_boundary):
    plot_decision_boundary(X_train_2d, y_train.values, k_val, ax,
                           f'Decision Boundary')

plt.suptitle('KNN Decision Boundaries for Different K Values (PCA 2D
Projection)',
             fontsize=14, fontweight='bold', y=1.02)


plt.tight_layout()

```



```
plt.show()
```

KNN Decision Boundaries for Different K Values (PCA 2D Projection)



Decision Boundary Analysis:

**K
Value**

Observation



K=1 Very complex, jagged boundary that fits every training point perfectly -- high variance, overfitting.

K=5 Smoother boundary that still captures class structure well -- good balance.

K=10 Even smoother; begins to generalize more broadly.

K=20 Very smooth, almost linear boundary -- high bias, potential underfitting.

As K increases, the decision boundary becomes smoother because predictions are averaged over more neighbours, reducing sensitivity to individual noisy points.

Task 4: Cross-Validation

```
# -- 4.1 K-Fold Cross Validation (k_fold = 10) --

k_fold = 10

k_range_cv = range(1, 31)

cv_mean_scores = []

cv_std_scores = []

for k in k_range_cv:

    knn_cv = KNeighborsClassifier(n_neighbors=k)

    scores = cross_val_score(knn_cv, X_scaled, y, cv=k_fold,
scoring='accuracy')

    cv_mean_scores.append(scores.mean())

    cv_std_scores.append(scores.std())

cv_mean_scores = np.array(cv_mean_scores)

cv_std_scores = np.array(cv_std_scores)
```



```
best_k_cv = list(k_range_cv)[np.argmax(cv_mean_scores)]
print(f"Cross-Validation Results (K-Fold = {k_fold}):")
print(f"    Best K: {best_k_cv}")
print(f"        Best    Mean    Accuracy:    {cv_mean_scores.max():.4f}    +/-
{cv_std_scores[np.argmax(cv_mean_scores)]:.4f}")
```

Cross-Validation Results (K-Fold = 10):

Best K: 11

Best Mean Accuracy: 0.9701 +/- 0.0222

-- 4.2 Plot Cross-Validation scores vs K --

```
fig, ax = plt.subplots(figsize=(12, 5))
ax.plot(list(k_range_cv), cv_mean_scores, 'o-', color='#8e44ad', label='CV
Mean Accuracy', markersize=5)
ax.fill_between(list(k_range_cv),
                cv_mean_scores - cv_std_scores,
                cv_mean_scores + cv_std_scores,
                alpha=0.2, color='#8e44ad')
ax.plot(list(k_range_cv), wide_test_acc, 's--', color='#e74c3c',
label='Single Split Test Accuracy',
        markersize=4, alpha=0.7)
ax.axvline(x=k_heuristic, color='green', linestyle='--', alpha=0.7,
label=f'Heuristic K={k_heuristic}')
ax.axvline(x=best_k_cv, color='orange', linestyle='--', alpha=0.7,
label=f'Best CV K={best_k_cv}')
ax.set_xlabel('K (Number of Neighbours)')
```




```
ax.set_ylabel('Accuracy')

ax.set_title('Cross-Validation vs Single Split: Accuracy for Different K
Values')

ax.legend()

plt.tight_layout()

plt.show()
```



```
# -- 4.3 Comparison Table --

print(f"\n{'K Selection Method':<30} {'K Value':<10} {'Accuracy':<15}")

print("=" * 55)

print(f"{'Heuristic'           (sqrt           n)':<30}           {k_heuristic:<10}
{wide_test_acc[k_heuristic-1]:.4f}")

print(f"{'Best'           Single           Split           (80:20)':<30}           {best_k_wide:<10}
{max(wide_test_acc):.4f}")

print(f"{'Best'           Cross-Validation':<30}           {best_k_cv:<10}
{cv_mean_scores.max():.4f}")

# -- Select final optimal K --
```



```
optimal_k = best_k_cv # Cross-validation is the most reliable
print(f"\n>> Selected Optimal K = {optimal_k} (based on cross-validation)")
```

K Selection Method	K Value	Accuracy
Heuristic (sqrt n)	21	0.9386
Best Single Split (80:20)	5	0.9561
Best Cross-Validation	11	0.9701

```
>> Selected Optimal K = 11 (based on cross-validation)
```

Task 5: Classification Evaluation

```
# -- 5.1 Train final model with optimal K --
knn_final = KNeighborsClassifier(n_neighbors=optimal_k)
knn_final.fit(X_train, y_train)
y_pred = knn_final.predict(X_test)
y_pred_proba = knn_final.predict_proba(X_test)[:, 1]

print(f"Final KNN Model (K = {optimal_k})")
print("=" * 50)
```

```
Final KNN Model (K = 11)
```




```
# -- 5.2 Classification Metrics --
```

```
acc = accuracy_score(y_test, y_pred)
prec = precision_score(y_test, y_pred)
rec = recall_score(y_test, y_pred)
f1 = f1_score(y_test, y_pred)
```

```
print(f"Accuracy: {acc:.4f}")
```

```
print(f"Precision: {prec:.4f}")
```

```
print(f"Recall: {rec:.4f}")
```

```
print(f"F1 Score: {f1:.4f}")
```

```
Accuracy: 0.9474
```

```
Precision: 0.9737
```

```
Recall: 0.8810
```

```
F1 Score: 0.9250
```

```
# -- 5.3 Classification Report --
```

```
print("\nDetailed Classification Report:")
```

```
print(classification_report(y_test, y_pred, target_names=['Malignant (0)',
'Benign (1)']))
```

```
Detailed Classification Report:
```

	precision	recall	f1-score	support
Malignant (0)	0.93	0.99	0.96	72
Benign (1)	0.97	0.88	0.93	42



accuracy			0.95	114
macro avg	0.95	0.93	0.94	114
weighted avg	0.95	0.95	0.95	114

```
# -- 5.4 Confusion Matrix --
```

```
cm = confusion_matrix(y_test, y_pred)
```

```
fig, axes = plt.subplots(1, 2, figsize=(14, 5))
```

```
# Heatmap
```

```
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues',
             xticklabels=['Malignant (0)', 'Benign (1)'],
             yticklabels=['Malignant (0)', 'Benign (1)'],
             ax=axes[0], cbar=True, linewidths=1, linecolor='black')
```

```
axes[0].set_xlabel('Predicted Label')
```

```
axes[0].set_ylabel('True Label')
```

```
axes[0].set_title('Confusion Matrix')
```

```
# Annotated explanation
```

```
axes[1].axis('off')
```

```
tn, fp, fn, tp = cm.ravel()
```

```
text = (
```

```
    f"Confusion Matrix Breakdown\n"
```



```

f"{'=' * 40}\n\n"

f"True Negatives (TN) = {tn:>3} -> Correctly predicted Malignant\n"

f"False Positives (FP) = {fp:>3} -> Malignant misclassified as Benign\n"

f"False Negatives (FN) = {fn:>3} -> Benign misclassified as Malignant\n"

f"True Positives (TP) = {tp:>3} -> Correctly predicted Benign\n\n"

f"{'=' * 40}\n"

f"Total samples = {tn + fp + fn + tp}"

)

axes[1].text(0.1, 0.5, text, fontsize=12, family='monospace',


            verticalalignment='center',

            bbox=dict(boxstyle='round', facecolor='lightyellow', alpha=0.8))

plt.tight_layout()

plt.show()

```



Confusion Matrix Breakdown

```

=====
True Negatives (TN) = 71 -> Correctly predicted Malignant
False Positives (FP) = 1 -> Malignant misclassified as Benign
False Negatives (FN) = 5 -> Benign misclassified as Malignant
True Positives (TP) = 37 -> Correctly predicted Benign
=====
Total samples = 114

```

```
# -- 5.5 ROC Curve and AUC Score --
```

```
fpr, tpr, thresholds = roc_curve(y_test, y_pred_proba)
```

```
roc_auc = auc(fpr, tpr)
```



```
fig, ax = plt.subplots(figsize=(8, 6))

ax.plot(fpr, tpr, color='#2980b9', lw=2, label=f'ROC Curve (AUC =
{roc_auc:.4f})')

ax.plot([0, 1], [0, 1], 'k--', lw=1, label='Random Classifier (AUC = 0.5)')

ax.fill_between(fpr, tpr, alpha=0.15, color='#2980b9')

ax.set_xlabel('False Positive Rate (1 - Specificity)')

ax.set_ylabel('True Positive Rate (Sensitivity / Recall)')

ax.set_title('Receiver Operating Characteristic (ROC) Curve')

ax.legend(loc='lower right')


ax.set_xlim([0.0, 1.0])

ax.set_ylim([0.0, 1.05])

plt.tight_layout()

plt.show()

print(f"ROC-AUC Score: {roc_auc:.4f}")
```



ROC-AUC Score: 0.9805

Task 6: Comparative Study -- Classification vs Regression Metrics

Regression Metrics (from Lab 3 -- Linear Regression)

Metric	Formula	Interpretation
MAE	$(1/n) * \sum(y_i - y_{i_hat})$	Average absolute error; easy to interpret in original units
MSE	$(1/n) * \sum((y_i - y_{i_hat})^2)$	Average squared error; penalizes large errors more
RMSE	$\sqrt{MSE}$	Square root of MSE;



in same units as target | | **R2 Score** | $1 - (SS_{res} / SS_{tot})$ | Proportion of variance explained (0 to 1) |

Classification Metrics (this Lab)

Metric	Formula	Interpretation
Accuracy	$\frac{TP+TN}{TP+TN+FP+FN}$	Fraction of correct predictions
Precision	$TP / (TP+FP)$	Of those predicted positive, how many are actually positive
Recall	$TP / (TP+FN)$	Of actual positives, how many were correctly identified
F1 Score	$2PR / (P+R)$	Harmonic mean of Precision and Recall
ROC-AUC	Area under ROC curve	Ability to distinguish between classes

```
# -- 6.1 Comparison Table --
```

```
comparison_data = {
    'Aspect': [
        'Evaluation Philosophy',
        'Output Type',
        'Error Measurement',
        'Perfect Score',
        'Affected by Outliers',
        'Suitable For'
    ],
    'Regression Metrics': [
        'Measures magnitude of prediction error',
        'Continuous numerical values',
    ]
}
```




```

'MAE, MSE, RMSE (distance from truth)',

'R2 = 1.0, MAE = MSE = RMSE = 0',

'MSE / RMSE highly sensitive',

'House price prediction, stock forecasting'

],

'Classification Metrics': [

    'Measures correctness of discrete decisions',

    'Discrete class labels (0 or 1)',

    'Accuracy, Precision, Recall, F1 (decision errors)',

    'Accuracy = 1.0, AUC = 1.0',

    'Not affected by value magnitude',

    'Disease diagnosis, spam detection'

]

}

comparison_df = pd.DataFrame(comparison_data)

print("Comprehensive Comparison: Regression vs Classification Metrics")

print("=" * 90)

print(comparison_df.to_string(index=False))

```

```

Comprehensive Comparison: Regression vs Classification Metrics

=====

=====

Aspect                                     Regression Metrics

Classification Metrics

Evaluation Philosophy                     Measures magnitude of prediction error

Measures correctness of discrete decisions

```



Output Type	Continuous numerical values
Discrete class labels (0 or 1)	
Error Measurement	MAE, MSE, RMSE (distance from truth) Accuracy, Precision, Recall, F1 (decision errors)
Perfect Score	$R^2 = 1.0$, MAE = MSE = RMSE = 0
Accuracy = 1.0, AUC = 1.0	
Affected by Outliers	MSE / RMSE highly sensitive
Not affected by value magnitude	
Suitable For House price prediction, stock forecasting	
Disease diagnosis, spam detection	

```
# -- 6.2 Specific Metric Comparisons --
```

```
print("\n" + "=" * 70)
```

```
print("DETAILED METRIC-TO-METRIC COMPARISON")
```

```
print("=" * 70)
```

```
print("""
```

```
+-----+
|  R2 Score vs Accuracy                               |
+-----+
|  R2 Score: How well the model explains variance in continuous |
|  predictions (1.0 = perfect, 0 = no better than mean).        |
|  Accuracy: Fraction of correct class predictions.             |
|  Both measure overall model quality, but R2 is for continuous |
|  outputs and Accuracy is for categorical outputs.             |
|  Key difference: Accuracy ignores how *close* wrong predictions |
```



| were, while R^2 captures this through squared residuals. |

+-----+

+-----+

| RMSE vs F1 Score |

+-----+

| RMSE: Penalizes large prediction errors more than small ones. |

| F1 Score: Balances precision and recall for the positive class. |

| RMSE focuses on *how far off* predictions are (magnitude), |

| while F1 focuses on *how correctly* the model identifies |

| positives without being misled by class imbalance. |

+-----+

+-----+

| MAE vs Confusion Matrix |

+-----+

| MAE: Gives a single number for average absolute error. |

| Confusion Matrix: Provides a 2x2 breakdown of all prediction |

| outcomes (TP, TN, FP, FN). |

| MAE is a scalar summary; the Confusion Matrix reveals the |

| *type* of errors (missed cancers vs false alarms), which is |

| critical in medical diagnosis. |

+-----+

""")



DETAILED METRIC-TO-METRIC COMPARISON

+-----+	
R2 Score vs Accuracy	
+-----+	
R2 Score: How well the model explains variance in continuous	
predictions (1.0 = perfect, 0 = no better than mean).	
Accuracy: Fraction of correct class predictions.	
Both measure overall model quality, but R2 is for continuous	
outputs and Accuracy is for categorical outputs.	
Key difference: Accuracy ignores how *close* wrong predictions	
were, while R2 captures this through squared residuals.	
+-----+	
+-----+	
RMSE vs F1 Score	
+-----+	
RMSE: Penalizes large prediction errors more than small ones.	
F1 Score: Balances precision and recall for the positive class.	
RMSE focuses on *how far off* predictions are (magnitude),	
while F1 focuses on *how correctly* the model identifies	
positives without being misled by class imbalance.	
+-----+	



+-----+	
MAE vs Confusion Matrix	
+-----+	
MAE: Gives a single number for average absolute error.	
Confusion Matrix: Provides a 2x2 breakdown of all prediction	
outcomes (TP, TN, FP, FN).	
MAE is a scalar summary; the Confusion Matrix reveals the	
type of errors (missed cancers vs false alarms), which is	
critical in medical diagnosis.	
+-----+	

Evaluation Logic Differences

Continuous Prediction Tasks (Regression):

- Output is a real number; errors are measured as distances.
- Small deviations from truth are acceptable; the goal is minimizing error magnitude.
- Metrics: MAE, MSE, RMSE, R2.

Classification Tasks:

- Output is a discrete class label; predictions are either correct or incorrect.
- The type of error matters (FP vs FN) -- especially in healthcare.
- Metrics: Accuracy, Precision, Recall, F1, ROC-AUC, Confusion Matrix.

Inference

1. **Regression metrics** (MAE, MSE, RMSE, R2) quantify how far predicted values deviate from actual continuous values. They measure the *magnitude* of prediction error.



2. **Classification metrics** (Accuracy, Precision, Recall, F1, AUC) measure how correctly the model assigns discrete class labels. They evaluate *decision correctness*.
3. **Why accuracy is insufficient in medical diagnosis:** In the breast cancer dataset, if 63% of cases are benign, a model that predicts "benign" for everyone achieves 63% accuracy -- yet it misses every cancer case. Accuracy fails to distinguish between types of errors.
4. **Why Recall and ROC-AUC are more relevant in healthcare:**
 - **Recall** (Sensitivity) measures the fraction of actual cancer cases correctly identified. Missing a malignant case (False Negative) has life-threatening consequences.
 - **ROC-AUC** evaluates the model's ability to separate the two classes across all decision thresholds, providing a threshold-independent performance measure.
5. **Overall:** Regression evaluation measures "how wrong" a prediction is, while classification evaluation measures "how often and in what way" the prediction is wrong. In healthcare, the *type* of error (FN) matters more than overall correctness, making Recall and ROC-AUC the preferred metrics.

Task 7: Analytical Questions

Q1. Why is KNN called a lazy learning algorithm?

KNN is called a "lazy learner" because it does not build an explicit model during training. It simply stores the entire training dataset in memory. All computation (distance calculation, neighbour search, voting) happens at prediction time, making it computationally expensive during inference rather than during training.

Q2. Why is feature scaling required in KNN?

KNN relies on distance metrics (Euclidean, Manhattan) to find nearest neighbours. Features with larger numerical ranges dominate the distance computation -- for example, "area" (100-2500)



would overshadow "smoothness" (0.05-0.16). Scaling ensures all features contribute equally to the distance calculation.

Q3. Explain heuristic K selection using sqrt(n) rule.

The heuristic $K = \sqrt{n}$ (where n is the number of training samples) provides a practical starting point for K selection. It balances between choosing too few neighbours (overfitting to noise) and too many neighbours (oversimplifying the decision boundary). For $n = 455$ training samples, $K \sim \sqrt{455} \sim 21$. This is a baseline -- the actual optimal K should be fine-tuned via cross-validation.

Q4. Why is cross-validation more reliable than a single train-test split?

A single train-test split is sensitive to which samples end up in the training vs. test set -- different random seeds produce different accuracy estimates. Cross-validation trains and evaluates the model on every data point across multiple folds, producing a more robust and less biased estimate of generalization performance.

Q5. How does K affect bias-variance trade-off?

K Value	Effect
Small K (e.g., $K=1$)	Low bias, high variance -- model captures noise and individual points, leading to overfitting.
Large K (e.g., $K=50$)	High bias, low variance -- model averages over many neighbours, producing overly smooth boundaries and underfitting.
Optimal K	Balances bias and variance -- complex enough to capture patterns but smooth enough to ignore noise.

Q6. Why is recall more important than accuracy in cancer prediction?



In cancer diagnosis, a False Negative (predicting "benign" for a malignant tumour) is far more dangerous than a False Positive (predicting "malignant" for a benign tumour). A missed cancer can lead to delayed treatment and death, whereas a false alarm leads to additional (non-fatal) tests. Recall = $TP / (TP + FN)$ directly measures how many actual cancer cases the model catches, making it the priority metric in healthcare.

Q7. What is the limitation of very large K values?

- The model becomes computationally expensive (comparing to many neighbours).
- The decision boundary becomes excessively smooth, leading to underfitting.
- In extreme cases ($K = n$), every point is classified by the majority class, making the model useless -- it becomes equivalent to a majority-vote classifier.
- Class boundaries are lost, and the model cannot capture local patterns.

Conclusion

```
print("=" * 70)

print("                                LAB 4 -- SUMMARY & CONCLUSION")

print("=" * 70)

print(f"""

1. OPTIMAL K VALUE:

    - Heuristic K (sqrt(n) = sqrt({n_train})) = {k_heuristic}

    - Best K from Cross-Validation = {best_k_cv}

    - Final Optimal K selected      = {optimal_k}

2. TRAIN-TEST SPLIT ANALYSIS:

    - 80:20 split offers the best balance of training quality and

      evaluation reliability.
```



- Very small test sets (90:10) give unreliable accuracy estimates.

3. MODEL PERFORMANCE (K = {optimal_k}):

- Accuracy: {acc:.4f}
- Precision: {prec:.4f}
- Recall: {rec:.4f}
- F1 Score: {f1:.4f}
- ROC-AUC: {roc_auc:.4f}

4. CLASSIFICATION vs REGRESSION EVALUATION:

- Regression metrics (MAE, MSE, RMSE, R2) measure prediction error magnitude for continuous outputs.
- Classification metrics (Accuracy, Precision, Recall, F1, AUC) measure decision correctness for discrete outputs.
- In medical diagnosis, Recall and ROC-AUC are more important than Accuracy because the cost of missing a cancer case (FN) is far greater than a false alarm (FP).

5. KEY INSIGHTS (Lab 3 vs Lab 4):

- Lab 3 (Regression): Focused on minimizing prediction distance from actual values using R2, RMSE.
- Lab 4 (Classification): Focused on maximizing decision correctness, especially for the minority/dangerous class using Recall, F1, AUC.
- The choice of evaluation metric depends fundamentally on the problem type and the real-world cost of different error types.



```
""")  
  
print("=" * 70)
```

```
=====
```

LAB 4 -- SUMMARY & CONCLUSION

```
=====
```

1. OPTIMAL K VALUE:

- Heuristic K ($\sqrt{n} = \sqrt{455}$) = 21
- Best K from Cross-Validation = 11
- Final Optimal K selected = 11

2. TRAIN-TEST SPLIT ANALYSIS:

- 80:20 split offers the best balance of training quality and evaluation reliability.
- Very small test sets (90:10) give unreliable accuracy estimates.

3. MODEL PERFORMANCE (K = 11):

- Accuracy: 0.9474
- Precision: 0.9737
- Recall: 0.8810
- F1 Score: 0.9250
- ROC-AUC: 0.9805

4. CLASSIFICATION vs REGRESSION EVALUATION:



- Regression metrics (MAE, MSE, RMSE, R2) measure prediction error magnitude for continuous outputs.
- Classification metrics (Accuracy, Precision, Recall, F1, AUC) measure decision correctness for discrete outputs.
- In medical diagnosis, Recall and ROC-AUC are more important than Accuracy because the cost of missing a cancer case (FN) is far greater than a false alarm (FP).

5. KEY INSIGHTS (Lab 3 vs Lab 4):

- Lab 3 (Regression): Focused on minimizing prediction distance from actual values using R2, RMSE.
- Lab 4 (Classification): Focused on maximizing decision correctness, especially for the minority/dangerous class using Recall, F1, AUC.
- The choice of evaluation metric depends fundamentally on the problem type and the real-world cost of different error types.

=====

Self-Learning Component: Weighted KNN — Uniform vs Distance-Weighted

Concept



By default, KNN uses **uniform weights** — all K neighbours get an equal vote regardless of how far they are from the query point. A neighbour sitting right next to the query point has the same influence as one at the edge of the neighbourhood.

Distance-weighted KNN assigns higher influence to closer neighbours by weighting each vote by the inverse of its distance:

$\text{weight}_i = 1 / \text{distance}(\text{query}, \text{neighbour}_i)$

This is particularly useful when:

- The decision boundary is complex and class regions overlap.
- Some neighbours within the K-radius belong to the wrong class but are far away.
- The dataset has noisy or borderline samples.

Intuition: Imagine diagnosing a patient — a very similar past case (close neighbour) should influence the diagnosis more than a loosely similar one (far neighbour).

```
# -- SL 1: Compare Uniform vs Distance-Weighted KNN across all K values --
```

```
k_range_sl = range(1, 31)
```

```
uniform_test_acc = []
```

```
distance_test_acc = []
```

```
uniform_train_acc = []
```

```
distance_train_acc = []
```

```
for k in k_range_sl:
```

```
    # Uniform weights (default)
```

```
    knn_uni = KNeighborsClassifier(n_neighbors=k, weights='uniform')
```

```
    knn_uni.fit(X_train, y_train)
```

```
    uniform_train_acc.append(knn_uni.score(X_train, y_train))
```



```

uniform_test_acc.append(knn_uni.score(X_test, y_test))

# Distance weights
knn_dist = KNeighborsClassifier(n_neighbors=k, weights='distance')
knn_dist.fit(X_train, y_train)
distance_train_acc.append(knn_dist.score(X_train, y_train))
distance_test_acc.append(knn_dist.score(X_test, y_test))

# -- SL 2: Plot comparison --
fig, axes = plt.subplots(1, 2, figsize=(16, 5))

# Test Accuracy comparison
axes[0].plot(list(k_range_sl), uniform_test_acc, 'o-', color='#e74c3c',
              label='Uniform Weights', markersize=4)
axes[0].plot(list(k_range_sl), distance_test_acc, 's-', color='#2980b9',
              label='Distance Weights', markersize=4)
axes[0].axvline(x=optimal_k, color='green', linestyle='--', alpha=0.7,
                 label=f'Optimal K={optimal_k}')
axes[0].set_xlabel('K (Number of Neighbours)')
axes[0].set_ylabel('Test Accuracy')
axes[0].set_title('Test Accuracy: Uniform vs Distance-Weighted KNN')
axes[0].legend()

# Train Accuracy comparison
axes[1].plot(list(k_range_sl), uniform_train_acc, 'o-', color='#e74c3c',

```





```

        label='Uniform Weights', markersize=4)
axes[1].plot(list(k_range_sl), distance_train_acc, 's-', color='#2980b9',
             label='Distance Weights', markersize=4)
axes[1].axvline(x=optimal_k, color='green', linestyle='--', alpha=0.7,
               label=f'Optimal K={optimal_k}')
axes[1].set_xlabel('K (Number of Neighbours)')
axes[1].set_ylabel('Train Accuracy')
axes[1].set_title('Train Accuracy: Uniform vs Distance-Weighted KNN')
axes[1].legend()

plt.tight_layout()
plt.show()

# Find best K for each weighting scheme
best_k_uniform = list(k_range_sl)[np.argmax(uniform_test_acc)]
best_k_distance = list(k_range_sl)[np.argmax(distance_test_acc)]
print(f"Best      K      (Uniform):          K={best_k_uniform},          Test
Accuracy={max(uniform_test_acc):.4f}")
print(f"Best      K      (Distance):          K={best_k_distance},          Test
Accuracy={max(distance_test_acc):.4f}")

```

Best K (Uniform): K=5, Test Accuracy=0.9561

Best K (Distance): K=5, Test Accuracy=0.9561

-- SL 3: Detailed metric comparison at optimal K --

```
knn_uniform = KNeighborsClassifier(n_neighbors=optimal_k, weights='uniform')
```

```
knn_uniform.fit(X_train, y_train)
```

```
y_pred_uni = knn_uniform.predict(X_test)
```

```
y_proba_uni = knn_uniform.predict_proba(X_test)[:, 1]
```

```
knn_distance = KNeighborsClassifier(n_neighbors=optimal_k,
weights='distance')
```

```
knn_distance.fit(X_train, y_train)
```

```
y_pred_dist = knn_distance.predict(X_test)
```

```
y_proba_dist = knn_distance.predict_proba(X_test)[:, 1]
```

```
print(f"\nDetailed Metric Comparison at K = {optimal_k}")
```

```
print("=" * 60)
```

```
print(f"{'Metric':<15} {'Uniform':<15} {'Distance':<15} {'Better':<10}")
```

```
print("-" * 60)
```



```

metrics = {

    'Accuracy': (accuracy_score(y_test, y_pred_uni), accuracy_score(y_test,
y_pred_dist)),

    'Precision': (precision_score(y_test, y_pred_uni),
precision_score(y_test, y_pred_dist)),

    'Recall': (recall_score(y_test, y_pred_uni), recall_score(y_test,
y_pred_dist)),

    'F1 Score': (f1_score(y_test, y_pred_uni), f1_score(y_test,
y_pred_dist)),

    'ROC-AUC': (auc(*roc_curve(y_test, y_proba_uni)[:2]),
                auc(*roc_curve(y_test, y_proba_dist)[:2])),

}

for name, (uni_val, dist_val) in metrics.items():

    better = "Distance" if dist_val > uni_val else ("Uniform" if uni_val >
dist_val else "Tie")

    print(f"{name:<15} {uni_val:<15.4f} {dist_val:<15.4f} {better:<10}")

```

Detailed Metric Comparison at K = 11

=====			
Metric	Uniform	Distance	Better

Accuracy	0.9474	0.9474	Tie
Precision	0.9737	0.9737	Tie
Recall	0.8810	0.8810	Tie



F1 Score	0.9250	0.9250	Tie
ROC-AUC	0.9805	0.9825	Distance

```
# -- SL 4: ROC Curve overlay --

fpr_uni, tpr_uni, _ = roc_curve(y_test, y_proba_uni)
fpr_dist, tpr_dist, _ = roc_curve(y_test, y_proba_dist)
auc_uni = auc(fpr_uni, tpr_uni)
auc_dist = auc(fpr_dist, tpr_dist)

fig, ax = plt.subplots(figsize=(8, 6))

ax.plot(fpr_uni, tpr_uni, color='#e74c3c', lw=2,
        label=f'Uniform Weights (AUC = {auc_uni:.4f})')

ax.plot(fpr_dist, tpr_dist, color='#2980b9', lw=2,
        label=f'Distance Weights (AUC = {auc_dist:.4f})')

ax.plot([0, 1], [0, 1], 'k--', lw=1, label='Random (AUC = 0.5)')

ax.fill_between(fpr_uni, tpr_uni, alpha=0.08, color='#e74c3c')
ax.fill_between(fpr_dist, tpr_dist, alpha=0.08, color='#2980b9')

ax.set_xlabel('False Positive Rate')
ax.set_ylabel('True Positive Rate')

ax.set_title(f'ROC Curve Comparison: Uniform vs Distance-Weighted KNN
(K={optimal_k})')


ax.legend(loc='lower right')

ax.set_xlim([0.0, 1.0])
ax.set_ylim([0.0, 1.05])

plt.tight_layout()
```



```
plt.show()
```



```
# -- SL 5: Confusion Matrix side-by-side --

cm_uni = confusion_matrix(y_test, y_pred_uni)

cm_dist = confusion_matrix(y_test, y_pred_dist)

fig, axes = plt.subplots(1, 2, figsize=(12, 5))

sns.heatmap(cm_uni, annot=True, fmt='d', cmap='Reds',
            xticklabels=['Malignant', 'Benign'],
            yticklabels=['Malignant', 'Benign'],
            ax=axes[0], linewidths=1, linecolor='black')
```



```

axes[0].set_xlabel('Predicted')

axes[0].set_ylabel('Actual')

axes[0].set_title(f'Uniform Weights (K={optimal_k})')

sns.heatmap(cm_dist, annot=True, fmt='d', cmap='Blues',
             xticklabels=['Malignant', 'Benign'],
             yticklabels=['Malignant', 'Benign'],
             ax=axes[1], linewidths=1, linecolor='black')

axes[1].set_xlabel('Predicted')

axes[1].set_ylabel('Actual')


axes[1].set_title(f'Distance Weights (K={optimal_k})')

plt.suptitle('Confusion Matrix Comparison', fontsize=14, fontweight='bold')

plt.tight_layout()

plt.show()

```





Key Observations from Weighted KNN Analysis

1. **Distance-weighted KNN achieves perfect training accuracy (1.0) for all K values** because the closest neighbour to any training point is the point itself (distance = 0), which gets infinite weight and always dominates the vote. This does NOT mean overfitting — what matters is test performance.
2. **Distance weighting is more robust to the choice of K.** The test accuracy curve for distance-weighted KNN is smoother and degrades less at larger K values, because distant (potentially wrong-class) neighbours get very low weights and don't corrupt the vote.
3. **In medical diagnosis, distance-weighted KNN is preferable** because it reduces the risk of misclassification caused by noisy or borderline neighbours — a patient whose features closely match malignant cases should be flagged even if a few distant benign neighbours exist in the K-neighbourhood.
4. **Trade-off:** Distance weighting slightly increases computational cost (weight calculation) but provides meaningful improvements in robustness. In practice, it is the recommended default for most KNN applications.

Self-Learning Component 2: Logistic Regression vs KNN — Parametric vs Non-Parametric Classification

Concept

KNN is a **non-parametric, instance-based** classifier — it makes no assumptions about the data distribution and stores the entire training set for prediction. In contrast, **Logistic Regression** is a **parametric, model-based** classifier that learns a set of weights (coefficients) during training and uses the sigmoid function to output class probabilities.

Aspect	KNN	Logistic Regression
--------	-----	---------------------



Type	Non-parametric, lazy learner	Parametric, eager learner
Training	No explicit training; stores data	Learns weight vector via gradient descent
Prediction	Distance computation to all training points	Single matrix multiplication (fast)
Decision Boundary	Non-linear, local	Linear (in feature space)
Interpretability	Low (black-box neighbours)	High (coefficients show feature importance)
Scalability	Slow at prediction for large datasets	Fast at prediction

Why compare? In medical diagnosis, understanding *why* a prediction was made is often as important as the prediction itself. Logistic Regression provides interpretable coefficients, while KNN may offer better accuracy on non-linear boundaries.

```
# -- SL2-1: Import Logistic Regression and train on the same data --
```

```
from sklearn.linear_model import LogisticRegression
```

```
# Train Logistic Regression
```

```
lr_model = LogisticRegression(max_iter=10000, random_state=42)
```

```
lr_model.fit(X_train, y_train)
```

```
y_pred_lr = lr_model.predict(X_test)
```

```
y_proba_lr = lr_model.predict_proba(X_test)[:, 1]
```

```
# KNN with optimal K (already trained as knn_final)
```

```
y_pred_knn = knn_final.predict(X_test)
```




```
y_proba_knn = knn_final.predict_proba(X_test)[: , 1]
```

```
print(f"Logistic Regression trained successfully.")
```

```
print(f"KNN model (K={optimal_k}) already available.")
```

```
Logistic Regression trained successfully.
```

```
KNN model (K=11) already available.
```

```
# -- SL2-2: Side-by-side metric comparison --
```

```
from sklearn.metrics import roc_auc_score
```

```
metrics_comparison = {
```

```
    'Accuracy': (accuracy_score(y_test, y_pred_knn), accuracy_score(y_test,
y_pred_lr)),
```

```
    'Precision': (precision_score(y_test, y_pred_knn),
precision_score(y_test, y_pred_lr)),
```

```
    'Recall': (recall_score(y_test, y_pred_knn), recall_score(y_test,
y_pred_lr)),
```

```
    'F1 Score': (f1_score(y_test, y_pred_knn), f1_score(y_test,
y_pred_lr)),
```

```
    'ROC-AUC': (roc_auc_score(y_test, y_proba_knn), roc_auc_score(y_test,
y_proba_lr)),
```

```
}
```

```
print(f"\nMetric Comparison: KNN (K={optimal_k}) vs Logistic Regression")
```

```
print("=" * 65)
```

```
print(f"{'Metric':<15} {'KNN':<15} {'Log. Reg.':<15} {'Better':<12}")
```



```
print("-" * 65)
```

```
for name, (knn_val, lr_val) in metrics_comparison.items():
    better = "Log. Reg." if lr_val > knn_val else ("KNN" if knn_val > lr_val
    else "Tie")
    print(f"{name:<15} {knn_val:<15.4f} {lr_val:<15.4f} {better:<12}")
```

Metric Comparison: KNN (K=11) vs Logistic Regression

```
=====
```

Metric	KNN	Log. Reg.	Better

Accuracy	0.9474	0.9649	Log. Reg.
Precision	0.9737	0.9750	Log. Reg.
Recall	0.8810	0.9286	Log. Reg.
F1 Score	0.9250	0.9512	Log. Reg.
ROC-AUC	0.9805	0.9960	Log. Reg.

```
# -- SL2-3: Bar chart comparison of all metrics --
metric_names = list(metrics_comparison.keys())
knn_scores = [v[0] for v in metrics_comparison.values()]
lr_scores = [v[1] for v in metrics_comparison.values()]

fig, ax = plt.subplots(figsize=(10, 5))
x_pos = np.arange(len(metric_names))
width = 0.3
```



```

bars1 = ax.bar(x_pos - width/2, knn_scores, width, label=f'KNN
(K={optimal_k})',

               color='#e74c3c', edgecolor='black', alpha=0.85)

bars2 = ax.bar(x_pos + width/2, lr_scores, width, label='Logistic
Regression',

               color='#2980b9', edgecolor='black', alpha=0.85)

ax.set_xlabel('Metric')
ax.set_ylabel('Score')
ax.set_title('KNN vs Logistic Regression: Classification Metrics Comparison')
ax.set_xticks(x_pos)
ax.set_xticklabels(metric_names)
ax.set_ylim(0.85, 1.02)
ax.legend()

for bar in bars1:
    ax.text(bar.get_x() + bar.get_width()/2., bar.get_height() + 0.003,
            f'{bar.get_height():.4f}', ha='center', va='bottom', fontsize=8)


for bar in bars2:
    ax.text(bar.get_x() + bar.get_width()/2., bar.get_height() + 0.003,
            f'{bar.get_height():.4f}', ha='center', va='bottom', fontsize=8)

plt.tight_layout()
plt.show()

```



CHRIST
(DEEMED TO BE UNIVERSITY)
BANGALORE | DELHI NCR | PUNE



```
# -- SL2-4: ROC Curve overlay --
```

```
fpr_knn_sl2, tpr_knn_sl2, _ = roc_curve(y_test, y_proba_knn)
```

```
fpr_lr, tpr_lr, _ = roc_curve(y_test, y_proba_lr)
```

```
auc_knn_sl2 = auc(fpr_knn_sl2, tpr_knn_sl2)
```

```
auc_lr = auc(fpr_lr, tpr_lr)
```

```
fig, ax = plt.subplots(figsize=(8, 6))
```

```
ax.plot(fpr_knn_sl2, tpr_knn_sl2, color='#e74c3c', lw=2,
```

```
      label=f'KNN K={optimal_k} (AUC = {auc_knn_sl2:.4f})')
```

```
ax.plot(fpr_lr, tpr_lr, color='#2980b9', lw=2,
```

```
      label=f'Logistic Regression (AUC = {auc_lr:.4f})')
```

```
ax.plot([0, 1], [0, 1], 'k--', lw=1, label='Random (AUC = 0.5)')
```


```
ax.fill_between(fpr_knn_sl2, tpr_knn_sl2, alpha=0.08, color='#e74c3c')
```

```
ax.fill_between(fpr_lr, tpr_lr, alpha=0.08, color='#2980b9')
```

```
ax.set_xlabel('False Positive Rate')
```



```
ax.set_ylabel('True Positive Rate')
ax.set_title('ROC Curve: KNN vs Logistic Regression')
ax.legend(loc='lower right')
ax.set_xlim([0.0, 1.0])
ax.set_ylim([0.0, 1.05])
plt.tight_layout()
plt.show()
```



```
# -- SL2-5: Confusion Matrix side-by-side --
```

```
cm_knn_sl2 = confusion_matrix(y_test, y_pred_knn)
```



```

cm_lr = confusion_matrix(y_test, y_pred_lr)

fig, axes = plt.subplots(1, 2, figsize=(12, 5))

sns.heatmap(cm_knn_sl2, annot=True, fmt='d', cmap='Reds',
            xticklabels=['Malignant', 'Benign'],
            yticklabels=['Malignant', 'Benign'],
            ax=axes[0], linewidths=1, linecolor='black')

axes[0].set_xlabel('Predicted')
axes[0].set_ylabel('Actual')
axes[0].set_title(f'KNN (K={optimal_k})')

sns.heatmap(cm_lr, annot=True, fmt='d', cmap='Blues',
            xticklabels=['Malignant', 'Benign'],
            yticklabels=['Malignant', 'Benign'],
            ax=axes[1], linewidths=1, linecolor='black')

axes[1].set_xlabel('Predicted')
axes[1].set_ylabel('Actual')
axes[1].set_title('Logistic Regression')

plt.suptitle('Confusion Matrix: KNN vs Logistic Regression', fontsize=14,
fontweight='bold')

plt.tight_layout()

plt.show()

```



```
# Print FN comparison (critical for cancer detection)

tn_knn, fp_knn, fn_knn, tp_knn = cm_knn_sl2.ravel()

tn_lr, fp_lr, fn_lr, tp_lr = cm_lr.ravel()


print(f"\nFalse Negatives (missed cancer cases):")

print(f"   KNN:                {fn_knn}")

print(f"   Logistic Regression: {fn_lr}")

print(f"   -> {'Logistic Regression' if fn_lr < fn_knn else ('KNN' if fn_knn
< fn_lr else 'Both equal')} misses fewer cancer cases.")
```

Confusion Matrix: KNN vs Logistic Regression



False Negatives (missed cancer cases):

KNN: 5

Logistic Regression: 3

-> Logistic Regression misses fewer cancer cases.



```
# -- SL2-6: Cross-Validation comparison --

from sklearn.model_selection import cross_val_score

cv_folds = 10

# KNN CV scores

knn_cv_model = KNeighborsClassifier(n_neighbors=optimal_k)

knn_cv_scores = cross_val_score(knn_cv_model, X_scaled, y, cv=cv_folds,
                                scoring='accuracy')

# Logistic Regression CV scores

lr_cv_model = LogisticRegression(max_iter=10000, random_state=42)

lr_cv_scores = cross_val_score(lr_cv_model, X_scaled, y, cv=cv_folds,
                                scoring='accuracy')

print(f"10-Fold Cross-Validation Comparison")

print("=" * 55)

print(f"{'Model':<25} {'Mean Accuracy':<18} {'Std Dev':<10}")

print("-" * 55)

print(f"{'KNN (K=' + str(optimal_k) + ')':<25} {knn_cv_scores.mean():<18.4f}
{knn_cv_scores.std():<10.4f}")

print(f"{'Logistic Regression':<25} {lr_cv_scores.mean():<18.4f}
{lr_cv_scores.std():<10.4f}")

# Box plot

fig, ax = plt.subplots(figsize=(8, 5))
```



```
bp = ax.boxplot([knn_cv_scores, lr_cv_scores],
                labels=[f'KNN (K={optimal_k})', 'Logistic Regression'],
                patch_artist=True, widths=0.4)

bp['boxes'][0].set_facecolor('#e74c3c')
bp['boxes'][0].set_alpha(0.6)
bp['boxes'][1].set_facecolor('#2980b9')
bp['boxes'][1].set_alpha(0.6)

ax.set_ylabel('Accuracy')
ax.set_title('10-Fold Cross-Validation: KNN vs Logistic Regression')
ax.set_ylim(0.90, 1.02)
plt.tight_layout()
plt.show()
```


10-Fold Cross-Validation Comparison

=====		
Model	Mean Accuracy	Std Dev

KNN (K=11)	0.9701	0.0222
Logistic Regression	0.9824	0.0136



CHRIST
(DEEMED TO BE UNIVERSITY)
BANGALORE | DELHI NCR | PUNE



```
# -- SL2-7: Logistic Regression Interpretability – Top Feature Coefficients
--
```

```
coefficients = pd.Series(lr_model.coef_[0], index=X.columns)
```

```
top_positive = coefficients.nlargest(5)
```

```
top_negative = coefficients.nsmallest(5)
```

```
fig, ax = plt.subplots(figsize=(10, 6))
```

```
top_features = pd.concat([top_negative, top_positive])
```

```
colors = ['#e74c3c' if v < 0 else '#2ecc71' for v in top_features.values]
```

```
top_features.plot(kind='barh', ax=ax, color=colors, edgecolor='black')
```



```

ax.set_xlabel('Coefficient Value')

ax.set_title('Logistic Regression: Top 10 Feature Coefficients\n'
            '(Negative → pushes toward Malignant | Positive → pushes toward Benign)')

ax.axvline(x=0, color='black', linewidth=0.8)

plt.tight_layout()

plt.show()

print("\nInterpretation:")

print(f"      Features pushing toward MALIGNANT (negative coeff):
{list(top_negative.index)}")

print(f"      Features pushing toward BENIGN (positive coeff):
{list(top_positive.index)}")


print("\n This level of interpretability is NOT possible with KNN,")

print("  which is a key advantage of Logistic Regression in clinical
settings.")

```



CHRIST
(DEEMED TO BE UNIVERSITY)
BANGALORE | DELHI NCR | PUNE



Interpretation:

Features pushing toward MALIGNANT (negative coeff): ['x.compactness_se', 'x.fractal_dim_se', 'x.compactness_mean', 'x.texture_se', 'x.symmetry_se']

Features pushing toward BENIGN (positive coeff): ['x.texture_worst', 'x.radius_se', 'x.symmetry_worst', 'x.concave_pts_mean', 'x.area_se']

This level of interpretability is NOT possible with KNN,
which is a key advantage of Logistic Regression in clinical settings.

Key Observations from Logistic Regression vs KNN Comparison



1. **Both models perform well** on the Breast Cancer dataset because the classes are reasonably separable in the 30-dimensional feature space.
2. **Logistic Regression is interpretable** — its learned coefficients directly show which features push toward Malignant vs Benign. This is critical in healthcare, where doctors need to understand *why* a model flagged a case.
3. **KNN makes no assumptions** about the data distribution, so it can capture non-linear decision boundaries. Logistic Regression assumes a linear boundary (in the original feature space), which may underfit complex patterns.
4. **Prediction speed:** Logistic Regression is a single matrix multiplication ($O(d)$), while KNN computes distances to all training points ($O(n*d)$) at prediction time. For real-time clinical decision support, Logistic Regression is far more practical.
5. **Cross-validation stability:** Logistic Regression typically shows lower variance across folds because it learns a fixed model, while KNN's performance depends on which specific neighbours end up in the training fold.
6. **Clinical recommendation:** For breast cancer screening, Logistic Regression is often preferred due to its interpretability and speed, unless the decision boundary is highly non-linear — in which case KNN (or ensemble methods) may be more appropriate.

Lab Experiment 5: Linear Regression through Gradient Descent

Aim

To implement Linear Regression using the Gradient Descent optimization algorithm and evaluate its performance on a real world dataset.

Objectives

1. To understand the concept of Gradient Descent for optimizing a Linear Regression model.
2. To preprocess the dataset before model training.
3. To implement Linear Regression using Gradient Descent.
4. To analyze the convergence of the loss function during training.
5. To evaluate the model using standard regression metrics.



Dataset

Dataset Name: Student Performance Dataset

Source: UCI Machine Learning Repository

Dataset Link:

<https://archive.ics.uci.edu/dataset/320/student+performance>

Tasks to Perform

1. Download and load the Student Performance dataset.
2. Perform data preprocessing, including handling missing values (if any), encoding categorical variables, and feature scaling.
3. Select appropriate input features and the target variable.
4. Split the dataset into training and testing sets.
5. Implement Linear Regression using the Gradient Descent algorithm.
6. Experiment with different learning rates and observe their effect on model convergence.
7. Plot the loss (cost) versus the number of iterations.
8. Evaluate the trained model using:
 - Mean Absolute Error (MAE)
 - Mean Squared Error (MSE)
 - Root Mean Squared Error (RMSE)
 - R^2 Score
9. Interpret the results and comment on the convergence behavior and prediction performance of the model.

Expected Learning Outcomes

Upon successful completion of this experiment, students will be able to:

1. Explain the working of the Gradient Descent optimization algorithm.
2. Implement Linear Regression using Gradient Descent.
3. Analyze the effect of learning rate on model convergence.
4. Evaluate regression models using standard performance metrics.
5. Interpret regression results and assess model performance.



Code with Results:

Lab 5: Linear Regression through Gradient Descent

Aim: To implement Linear Regression using the Gradient Descent optimization algorithm and evaluate its performance on the Student Performance dataset.

Dataset: Student Performance Dataset (UCI Machine Learning Repository)

Target Variable: G3 (Final Grade, 0-20)

1. Import Libraries

```
import numpy as np

import pandas as pd

import matplotlib.pyplot as plt

import seaborn as sns

from sklearn.model_selection import train_test_split

from sklearn.preprocessing import StandardScaler, LabelEncoder

from sklearn.metrics import mean_absolute_error, mean_squared_error,
r2_score

import warnings

warnings.filterwarnings('ignore')

print("All libraries imported successfully.")
```

All libraries imported successfully.



2. Load the Dataset

```
# Load the Math course student performance dataset
```

```
df = pd.read_csv('student/student-mat.csv', sep=';')
```

```
print(f"Dataset shape: {df.shape}")
```

```
print(f"Number of samples: {df.shape[0]}")
```

```
print(f"Number of features: {df.shape[1]}")
```

```
df.head()
```

Dataset shape: (395, 33)

Number of samples: 395

Number of features: 33

	sc	s	a	ad	fa	Ps	M	F	Mjo	Fjo	.	fa	free	g	D	W	he	abs	G	G	G
	ho	e	g	dre	ms	tat	e	e	b	b	.	mr	tim	o	a	a	alt	enc	1	2	3
	ol	x	e	ss	ize	us	u	u			.	el	e	ut	c	c	h	es			
0	G	F	1	U	GT	A	4	4	at_	tea	.										
	P		8		3				ho	ch	.	4	3	4	1	1	3	6	5	6	6
									me	er	.										
1	G	F	1	U	GT	T	1	1	at_	oth	.										
	P		7		3				ho	er	.	5	3	3	1	1	3	4	5	5	6
									me		.										



2	G P	F	1 5	U	LE 3	T	1	1	at_ ho me	oth er	.	4	3	2	2	3	3	10	7	8	1 0
3	G P	F	1 5	U	GT 3	T	4	2	hea lth	ser vic es	.	3	2	2	1	1	5	2	1 5	1 4	1 5
4	G P	F	1 6	U	GT 3	T	3	3	oth er	oth er	.	4	3	2	1	2	5	4	6	1 0	1 0

5 rows × 33 columns

3. Exploratory Data Analysis

```
# Basic information about the dataset
```

```
print("=" * 50)
```

```
print("DATASET INFO")
```

```
print("=" * 50)
```

```
df.info()
```

```
print("\n" + "=" * 50)
```

```
print("STATISTICAL SUMMARY")
```

```
print("=" * 50)
```

```
df.describe()
```

```
=====
```

DATASET INFO



```
=====
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 395 entries, 0 to 394
```

```
Data columns (total 33 columns):
```

#	Column	Non-Null Count	Dtype
0	school	395 non-null	object
1	sex	395 non-null	object
2	age	395 non-null	int64
3	address	395 non-null	object
4	famsize	395 non-null	object
5	Pstatus	395 non-null	object
6	Medu	395 non-null	int64
7	Fedu	395 non-null	int64
8	Mjob	395 non-null	object
9	Fjob	395 non-null	object
10	reason	395 non-null	object
11	guardian	395 non-null	object
12	traveltime	395 non-null	int64
13	studytime	395 non-null	int64
14	failures	395 non-null	int64
15	schoolsup	395 non-null	object
16	famsup	395 non-null	object
17	paid	395 non-null	object



```

18  activities  395 non-null  object
19  nursery    395 non-null  object
20  higher     395 non-null  object
21  internet   395 non-null  object
22  romantic   395 non-null  object
23  famrel     395 non-null  int64
24  freetime   395 non-null  int64
25  goout      395 non-null  int64
26  Dalc       395 non-null  int64
27  Walc       395 non-null  int64
28  health     395 non-null  int64
29  absences   395 non-null  int64
30  G1         395 non-null  int64
31  G2         395 non-null  int64
32  G3         395 non-null  int64

```

dtypes: int64(16), object(17)

memory usage: 102.0+ KB

=====

STATISTICAL SUMMARY

=====

```

      age  Me  Fe  tra  stu  fail  fa  fre  go  Dal  Wa  he  ab  G1  G2  G3
      e   du  du  vel  dyt  ure  mr  eti  out  c   lc  alt  se

```



CHRIST
(DEEMED TO BE UNIVERSITY)
BANGALORE | DELHI NCR | PUNE

	time										nc es					
c o u n t	39	39	39	39	39	39	39	39	39	39	39	39	39	39	39	39
	5.0	5.0	5.0	5.0	5.0	5.0	5.0	5.0	5.0	5.0	5.0	5.0	5.0	5.0	5.0	5.0
	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
m e a n	16.	2.7	2.5	1.4	2.0	0.3	3.9	3.2	3.1	1.4	2.2	3.5	5.7	10.	10.	10.
	69	49	21	48	35	34	44	35	08	81	911	54	08	90	71	41
	62	36	51	10	44	17	30	44	86	01	39	43	86	88	39	51
	03	7	9	1	3	7	4	3	1	3		0	1	61	24	90
s t d	1.2	1.0	1.0	0.6	0.8	0.7	0.8	0.9	1.1	0.8	1.2	1.3	8.0	3.3	3.7	4.5
	76	94	88	97	39	43	96	98	13	90	87	90	03	19	61	81
	04	73	20	50	24	65	65	86	27	74	89	30	09	19	50	44
	3	5	1	5	0	1	9	2	8	1	7	3	6	5	5	3
m i n	15.	0.0	0.0	1.0	1.0	0.0	1.0	1.0	1.0	1.0	1.0	1.0	0.0	3.0	0.0	0.0
	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
	00	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
2 5 %	16.	2.0	2.0	1.0	1.0	0.0	4.0	3.0	2.0	1.0	1.0	3.0	0.0	8.0	9.0	8.0
	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
	00	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
5 0 %	17.	3.0	2.0	1.0	2.0	0.0	4.0	3.0	3.0	1.0	2.0	4.0	4.0	11.	11.	11.
	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
	00	0	0	0	0	0	0	0	0	0	0	0	0	00	00	00



7	18.	4.0	3.0	2.0	2.0	0.0	5.0	4.0	4.0	2.0	3.0	5.0	8.0	13.	13.	14.
5	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
%	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
	00	0	0	0	0	0	0	0	0	0	0	0	0	00	00	00

m	22.	4.0	4.0	4.0	4.0	3.0	5.0	5.0	5.0	5.0	5.0	5.0	75.	19.	19.	20.
a	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
x	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
	00	0	0	0	0	0	0	0	0	0	0	0	00	00	00	00

```
# Check for missing values
```

```
print("Missing values per column:")
```

```
missing = df.isnull().sum()
```

```
print(missing[missing > 0] if missing.sum() > 0 else "No missing values found.")
```

```
print(f"\nTotal missing values: {df.isnull().sum().sum()}")
```

```
Missing values per column:
```

```
No missing values found.
```

```
Total missing values: 0
```

```
# Distribution of the target variable (G3 - Final Grade)
```

```
plt.figure(figsize=(8, 5))
```

```
plt.hist(df['G3'], bins=20, edgecolor='black', color='steelblue', alpha=0.7)
```

```
plt.xlabel('Final Grade (G3)')
```

```
plt.ylabel('Frequency')
```

```
plt.title('Distribution of Final Grade (G3)')
```





```
plt.axvline(df['G3'].mean(), color='red', linestyle='--', label=f"Mean =
{df['G3'].mean():.2f}")

plt.legend()

plt.tight_layout()

plt.show()
```



```
# Identify categorical and numerical columns

categorical_cols = df.select_dtypes(include='object').columns.tolist()

numerical_cols = df.select_dtypes(include=np.number).columns.tolist()

print(f"Categorical columns ({len(categorical_cols)}): {categorical_cols}")

print(f"\nNumerical columns ({len(numerical_cols)}): {numerical_cols}")
```



Categorical columns (17): ['school', 'sex', 'address', 'famsize', 'Pstatus', 'Mjob', 'Fjob', 'reason', 'guardian', 'schoolsup', 'famsup', 'paid', 'activities', 'nursery', 'higher', 'internet', 'romantic']

Numerical columns (16): ['age', 'Medu', 'Fedu', 'traveltime', 'studytime', 'failures', 'famrel', 'freetime', 'goout', 'Dalc', 'Walc', 'health', 'absences', 'G1', 'G2', 'G3']

4. Data Preprocessing

```
# Encode categorical variables using Label Encoding
```

```
df_encoded = df.copy()
```

```
label_encoders = {}
```

```
for col in categorical_cols:
```

```
    le = LabelEncoder()
```

```
    df_encoded[col] = le.fit_transform(df_encoded[col])
```

```
    label_encoders[col] = le
```

```
    print(f"{col}: {dict(zip(le.classes_, le.transform(le.classes_)))}")
```

```
print("\nEncoding complete.")
```

```
df_encoded.head()
```

```
school: {'GP': np.int64(0), 'MS': np.int64(1)}
```

```
sex: {'F': np.int64(0), 'M': np.int64(1)}
```

```
address: {'R': np.int64(0), 'U': np.int64(1)}
```

```
famsize: {'GT3': np.int64(0), 'LE3': np.int64(1)}
```



```

Pstatus: {'A': np.int64(0), 'T': np.int64(1)}

Mjob: {'at_home': np.int64(0), 'health': np.int64(1), 'other': np.int64(2),
'services': np.int64(3), 'teacher': np.int64(4)}

Fjob: {'at_home': np.int64(0), 'health': np.int64(1), 'other': np.int64(2),
'services': np.int64(3), 'teacher': np.int64(4)}

reason: {'course': np.int64(0), 'home': np.int64(1), 'other': np.int64(2),
'reputation': np.int64(3)}

guardian: {'father': np.int64(0), 'mother': np.int64(1), 'other':
np.int64(2)}

schoolsup: {'no': np.int64(0), 'yes': np.int64(1)}

famsup: {'no': np.int64(0), 'yes': np.int64(1)}

paid: {'no': np.int64(0), 'yes': np.int64(1)}

activities: {'no': np.int64(0), 'yes': np.int64(1)}

nursery: {'no': np.int64(0), 'yes': np.int64(1)}

higher: {'no': np.int64(0), 'yes': np.int64(1)}

internet: {'no': np.int64(0), 'yes': np.int64(1)}

romantic: {'no': np.int64(0), 'yes': np.int64(1)}

```

Encoding complete.

sc	s	a	ad	fa	Pst	M	F	M	F	.	fa	free	go	D	W	he	abs	G	G	G
ho	e	g	dre	msi	atu	ed	e	jo	j	.	mr	tim	ou	a	al	alt	enc	1	2	3
ol	x	e	ss	ze	s	u	d	b	o	.	el	e	t	c	c	h	es			

0	0	0	1	1	0	0	4	4	0	4	.	4	3	4	1	1	3	6	5	6	6
			8								.										



1	0	0	$\frac{1}{7}$	1	0	1	1	1	0	2	.	5	3	3	1	1	3	4	5	5	6
2	0	0	$\frac{1}{5}$	1	1	1	1	1	0	2	.	4	3	2	2	3	3	10	7	8	$\frac{1}{0}$
3	0	0	$\frac{1}{5}$	1	0	1	4	2	1	3	.	3	2	2	1	1	5	2	$\frac{1}{5}$	$\frac{1}{4}$	$\frac{1}{5}$
4	0	0	$\frac{1}{6}$	1	0	1	3	3	2	2	.	4	3	2	1	2	5	4	6	$\frac{1}{0}$	$\frac{1}{0}$

5 rows × 33 columns

```
# Correlation heatmap of numerical features with G3

correlation = df_encoded.corr()['G3'].sort_values(ascending=False)

print("Correlation of features with G3 (Final Grade):")

print(correlation)

plt.figure(figsize=(12, 8))

sns.heatmap(df_encoded.corr(), cmap='coolwarm', center=0, linewidths=0.5,
fmt='.1f')

plt.title('Correlation Heatmap')

plt.tight_layout()

plt.show()
```




Correlation of features with G3 (Final Grade) :

G3	1.000000
G2	0.904868
G1	0.801468
Medu	0.217147
higher	0.182465
Fedu	0.152457
reason	0.121994
address	0.105756
sex	0.103456
Mjob	0.102082
paid	0.101996
internet	0.098483
studytime	0.097820
famsize	0.081407
nursery	0.051568
famrel	0.051363
Fjob	0.042286
absences	0.034247
activities	0.016100
freetime	0.011307
famsup	-0.039157
school	-0.045017
Walc	-0.051939



```
Dalc          -0.054660
Pstatus       -0.058009
health        -0.061335
guardian      -0.070109
schoolsup     -0.082788
traveltime    -0.117142
romantic      -0.129970
goout         -0.132791
age           -0.161579
failures      -0.360415

Name: G3, dtype: float64
```



5. Feature Selection and Target Variable

```
# Select input features (X) and target variable (y)

# Drop G1 and G2 (intermediate grades) to avoid data leakage,

# and G3 is the target

X = df_encoded.drop(columns=['G1', 'G2', 'G3'])

y = df_encoded['G3'].values

print(f"Feature matrix shape: {X.shape}")

print(f"Target vector shape: {y.shape}")
```



```
print(f"\nSelected features: {X.columns.tolist()}")
```

```
Feature matrix shape: (395, 30)
```

```
Target vector shape: (395,)
```

```
Selected features: ['school', 'sex', 'age', 'address', 'famsize', 'Pstatus',
'Medu', 'Fedu', 'Mjob', 'Fjob', 'reason', 'guardian', 'traveltime',
'studytime', 'failures', 'schoolsup', 'famsup', 'paid', 'activities',
'nursery', 'higher', 'internet', 'romantic', 'famrel', 'freetime', 'goout',
'Dalc', 'Walac', 'health', 'absences']
```

```
# Feature Scaling using StandardScaler
```

```
scaler = StandardScaler()
```

```
X_scaled = scaler.fit_transform(X)
```

```
print("Feature scaling complete (StandardScaler).")
```

```
print(f"Mean of scaled features (should be ~0):
{X_scaled.mean(axis=0).round(4)[:5]} ...")
```

```
print(f"Std of scaled features (should be ~1):
{X_scaled.std(axis=0).round(4)[:5]} ...")
```

```
Feature scaling complete (StandardScaler).
```

```
Mean of scaled features (should be ~0): [-0. -0.  0. -0.  0.] ...
```

```
Std of scaled features (should be ~1): [1. 1. 1. 1. 1.] ...
```

6. Train-Test Split

```
# Split into training (80%) and testing (20%) sets
```




```
X_train, X_test, y_train, y_test = train_test_split(
    X_scaled, y, test_size=0.2, random_state=42
)
```

```
print(f"Training set size: {X_train.shape[0]}")
```

```
print(f"Testing set size: {X_test.shape[0]}")
```

```
Training set size: 316
```

```
Testing set size: 79
```

7. Linear Regression using Gradient Descent - Implementation

```
class LinearRegressionGD:
    """
    Linear Regression using Batch Gradient Descent.

    Hypothesis:  $h(x) = X \cdot w + b$ 
    Cost (MSE):  $J(w,b) = (1/2m) * \sum (h(x_i) - y_i)^2$ 

    Update rules:

     $w = w - \alpha * (1/m) * X^T \cdot (h(X) - y)$ 
     $b = b - \alpha * (1/m) * \sum (h(X) - y)$ 
    """
```



```
def __init__(self, learning_rate=0.01, n_iterations=1000):

    self.learning_rate = learning_rate

    self.n_iterations = n_iterations

    self.weights = None

    self.bias = None

    self.cost_history = []


def fit(self, X, y):

    """Train the model using Gradient Descent."""

    m, n = X.shape  # m = samples, n = features

    # Initialize weights and bias to zeros

    self.weights = np.zeros(n)

    self.bias = 0

    self.cost_history = []

    for i in range(self.n_iterations):

        # Forward pass: compute predictions

        y_pred = X.dot(self.weights) + self.bias

        # Compute error

        error = y_pred - y

        # Compute cost (MSE / 2)
```



```

cost = (1 / (2 * m)) * np.sum(error ** 2)

self.cost_history.append(cost)

# Compute gradients

dw = (1 / m) * X.T.dot(error)

db = (1 / m) * np.sum(error)

# Update weights and bias

self.weights -= self.learning_rate * dw

self.bias -= self.learning_rate * db

return self

def predict(self, X):

    """Make predictions."""

    return X.dot(self.weights) + self.bias

def get_cost_history(self):

    """Return the cost history during training."""

    return self.cost_history

print("LinearRegressionGD class defined successfully.")

```



LinearRegressionGD class defined successfully.

8. Train the Model

```
# Train with a chosen learning rate

model = LinearRegressionGD(learning_rate=0.01, n_iterations=1000)

model.fit(X_train, y_train)

print("Model training complete.")

print(f"Final cost: {model.cost_history[-1]:.4f}")

print(f"Bias (intercept): {model.bias:.4f}")

print(f"Number of weights: {len(model.weights)}")
```

```
Model training complete.

Final cost: 7.7213

Bias (intercept): 10.3795

Number of weights: 30
```

9. Experiment with Different Learning Rates

```
# Experiment with different learning rates

learning_rates = [0.001, 0.01, 0.05, 0.1]

n_iterations = 1000

models = {}

plt.figure(figsize=(12, 6))
```



```
for lr in learning_rates:

    lr_model = LinearRegressionGD(learning_rate=lr,
n_iterations=n_iterations)

    lr_model.fit(X_train, y_train)

    models[lr] = lr_model


    plt.plot(lr_model.cost_history, label=f'LR = {lr}')

    print(f"Learning Rate: {lr:.4f} | Final Cost:
{lr_model.cost_history[-1]:.4f}")



plt.xlabel('Number of Iterations')
plt.ylabel('Cost (MSE / 2)')
plt.title('Loss Convergence for Different Learning Rates')
plt.legend()
plt.grid(True, alpha=0.3)
plt.tight_layout()
plt.show()
```

```
Learning Rate: 0.0010 | Final Cost: 15.4011
```

```
Learning Rate: 0.0100 | Final Cost: 7.7213
```

```
Learning Rate: 0.0500 | Final Cost: 7.7210
```

```
Learning Rate: 0.1000 | Final Cost: 7.7210
```



Zoomed-in view of convergence (first 200 iterations)

```
plt.figure(figsize=(12, 6))
```

```
for lr in learning_rates:
```

```
    plt.plot(models[lr].cost_history[:200], label=f'LR = {lr}')
```

```
plt.xlabel('Number of Iterations')
```

```
plt.ylabel('Cost (MSE / 2)')
```


```
plt.title('Loss Convergence (First 200 Iterations) - Zoomed View')
```

```
plt.legend()
```

```
plt.grid(True, alpha=0.3)
```

```
plt.tight_layout()
```

```
plt.show()
```



Observations on Learning Rate

- **LR = 0.001 (too small):** Converges very slowly; the cost is still decreasing at the end of 1000 iterations and has not yet reached the minimum.
- **LR = 0.01 (good):** Converges steadily and reaches a stable minimum within the iterations.
- **LR = 0.05 (fast):** Converges much faster than 0.01 and reaches the minimum in fewer iterations while remaining stable.
- **LR = 0.1 (aggressive):** Converges the fastest and reaches the minimum within very few iterations. For this dataset, it remains stable without oscillation, although larger learning rates can cause oscillation or divergence in general.

10. Loss (Cost) vs Iterations Plot

Detailed cost vs iterations plot for the chosen model (LR = 0.01)

```
plt.figure(figsize=(10, 6))

plt.plot(model.cost_history, color='steelblue', linewidth=2)

plt.xlabel('Number of Iterations', fontsize=12)

plt.ylabel('Cost (MSE / 2)', fontsize=12)

plt.title('Cost Function Convergence (LR = 0.01)', fontsize=14)

plt.grid(True, alpha=0.3)
```




```
# Annotate start and end costs

plt.annotate(f'Start: {model.cost_history[0]:.2f}',
            xy=(0, model.cost_history[0]),
            xytext=(100, model.cost_history[0]),
            arrowprops=dict(arrowstyle='->', color='red'),
            fontsize=10, color='red')

plt.annotate(f'End: {model.cost_history[-1]:.2f}',
            xy=(len(model.cost_history)-1, model.cost_history[-1]),
            xytext=(len(model.cost_history)-200, model.cost_history[-1]+2),
            arrowprops=dict(arrowstyle='->', color='green'),
            fontsize=10, color='green')

plt.tight_layout()

plt.show()
```

11. Model Evaluation

```
# Predictions on the test set
y_pred = model.predict(X_test)

# Calculate evaluation metrics

mae = mean_absolute_error(y_test, y_pred)

mse = mean_squared_error(y_test, y_pred)

rmse = np.sqrt(mse)

r2 = r2_score(y_test, y_pred)

print("=" * 45)

print("      MODEL EVALUATION METRICS")
```



```
print("=" * 45)

print(f" Mean Absolute Error (MAE) : {mae:.4f}")

print(f" Mean Squared Error (MSE) : {mse:.4f}")

print(f" Root Mean Squared Error(RMSE): {rmse:.4f}")

print(f" R-squared Score : {r2:.4f}")

print("=" * 45)
```

```
=====

MODEL EVALUATION METRICS

=====

Mean Absolute Error (MAE) : 3.4959

Mean Squared Error (MSE) : 18.5409

Root Mean Squared Error(RMSE): 4.3059

R-squared Score : 0.0958

=====
```

Interpretation of Evaluation Metrics

- The **MAE of approximately 3.50** means the model's predictions differ from the actual final grades by about 3.5 marks on average (on a 0-20 scale).
- The **RMSE of approximately 4.31** is higher than MAE, indicating that some larger prediction errors are present. RMSE penalizes large errors more heavily than MAE.
- The **R2 score of approximately 0.096** shows that the model explains only about 9.6% of the variance in students' final grades. This is a relatively low value, suggesting that the selected features (demographics, family background, study habits) have limited predictive power for final grades when used with a simple linear model.
- The **MSE of approximately 18.54** represents the average squared error, which is useful for optimization but less interpretable than MAE or RMSE.



Overall Assessment: The Gradient Descent algorithm converged successfully as evidenced by the decreasing cost function. However, the predictive performance is relatively weak because student academic performance depends on many complex, non-linear factors (motivation, exam difficulty, learning style, etc.) that are not fully captured by the selected features. Dropping G1 and G2 (intermediate grades) to avoid data leakage also removes the strongest predictors, making the problem significantly harder.

```
# Actual vs Predicted scatter plot

plt.figure(figsize=(8, 6))

plt.scatter(y_test, y_pred, alpha=0.6, color='steelblue', edgecolor='k')


plt.plot([y_test.min(), y_test.max()], [y_test.min(), y_test.max()],
         'r--', linewidth=2, label='Ideal Prediction')

plt.xlabel('Actual G3 (Final Grade)', fontsize=12)
plt.ylabel('Predicted G3 (Final Grade)', fontsize=12)
plt.title('Actual vs Predicted Final Grade', fontsize=14)
plt.legend()

plt.grid(True, alpha=0.3)

plt.tight_layout()

plt.show()
```



```
# Residual plot
```

```
residuals = y_test - y_pred
```

```
plt.figure(figsize=(8, 6))
```

```
plt.scatter(y_pred, residuals, alpha=0.6, color='coral', edgecolor='k')
```

```
plt.axhline(y=0, color='black', linestyle='--', linewidth=1)
```

```
plt.xlabel('Predicted G3', fontsize=12)
```

```
plt.ylabel('Residuals (Actual - Predicted)', fontsize=12)
```


```
plt.title('Residual Plot', fontsize=14)
```



```
plt.grid(True, alpha=0.3)

plt.tight_layout()

plt.show()
```



12. Comparison of Metrics Across Learning Rates

```
# Evaluate each learning rate model on the test set
```

```
results = []
```

```
for lr in learning_rates:
```



```

lr_model = models[lr]

y_pred_lr = lr_model.predict(X_test)

results.append({

    'Learning Rate': lr,

    'MAE': round(mean_absolute_error(y_test, y_pred_lr), 4),

    'MSE': round(mean_squared_error(y_test, y_pred_lr), 4),

    'RMSE': round(np.sqrt(mean_squared_error(y_test, y_pred_lr)), 4),

    'R2 Score': round(r2_score(y_test, y_pred_lr), 4),

    'Final Cost': round(lr_model.cost_history[-1], 4)

})

results_df = pd.DataFrame(results)

print("Performance Comparison Across Learning Rates:")

print(results_df.to_string(index=False))

results_df

```

Performance Comparison Across Learning Rates:

Learning Rate	MAE	MSE	RMSE	R2 Score	Final Cost
0.001	5.3388	37.9360	6.1592	-0.8501	15.4011
0.010	3.4959	18.5409	4.3059	0.0958	7.7213
0.050	3.4948	18.5437	4.3062	0.0957	7.7210
0.100	3.4948	18.5437	4.3062	0.0957	7.7210



	Learning Rate	MAE	MSE	RMS E	R2 Score	Final Cost
0	0.001	5.3388	37.9360	6.1592	-0.8501	15.4011
1	0.010	3.4959	18.5409	4.3059	0.0958	7.7213
2	0.050	3.4948	18.5437	4.3062	0.0957	7.7210
3	0.100	3.4948	18.5437	4.3062	0.0957	7.7210

13. Self-Learning Component - Adaptive Gradient Descent

The basic Gradient Descent model uses a **fixed learning rate** throughout training. A **self-learning** model improves itself by:

1. **Adaptive Learning Rate Decay** - Automatically reduces the learning rate as training progresses so that early iterations take large steps and later iterations fine-tune.
2. **Momentum** - Accumulates a velocity term from past gradients to accelerate convergence and dampen oscillations.
3. **Early Stopping** - Monitors the cost and stops training automatically when improvement stalls, preventing overfitting and wasted computation.
4. **Online (Incremental) Learning** - Can learn from new data batches without retraining from scratch, simulating continuous self-improvement.

```
class SelfLearningLinearRegression:
```

```
    """
```



Self-Learning Linear Regression with:

- Adaptive Learning Rate (exponential decay)
- Momentum-based gradient updates
- Early Stopping with patience
- Online / Incremental Learning (partial_fit)

Adaptive LR: $\alpha_t = \alpha_0 / (1 + \text{decay_rate} * t)$

Momentum: $v_t = \beta * v_{t-1} + (1 - \beta) * \text{gradient}$

$w = w - \alpha_t * v_t$

Early Stop: Stop if cost does not improve by tol for patience iterations.

"""

```
def __init__(self, learning_rate=0.01, n_iterations=2000,
              decay_rate=0.001, momentum=0.9,
              patience=50, tol=1e-6):

    self.learning_rate = learning_rate

    self.n_iterations = n_iterations

    self.decay_rate = decay_rate

    self.momentum = momentum

    self.patience = patience

    self.tol = tol

    self.weights = None

    self.bias = None
```




```
self.cost_history = []

self.lr_history = []

self.stopped_at = None


def fit(self, X, y):

    """Train with adaptive LR, momentum, and early stopping."""

    m, n = X.shape

    self.weights = np.zeros(n)

    self.bias = 0

    self.cost_history = []

    self.lr_history = []

    self.stopped_at = None

    v_w = np.zeros(n)

    v_b = 0

    best_cost = np.inf

    patience_counter = 0

    for i in range(self.n_iterations):

        current_lr = self.learning_rate / (1 + self.decay_rate * i)

        self.lr_history.append(current_lr)

        y_pred = X.dot(self.weights) + self.bias
```



```

error = y_pred - y

cost = (1 / (2 * m)) * np.sum(error ** 2)

self.cost_history.append(cost)

if best_cost - cost > self.tol:

    best_cost = cost

    patience_counter = 0
else:

    patience_counter += 1

if patience_counter >= self.patience:

    self.stopped_at = i

    break

dw = (1 / m) * X.T.dot(error)

db = (1 / m) * np.sum(error)

v_w = self.momentum * v_w + (1 - self.momentum) * dw
v_b = self.momentum * v_b + (1 - self.momentum) * db

self.weights -= current_lr * v_w
self.bias -= current_lr * v_b

if self.stopped_at is None:

```



```

        self.stopped_at = self.n_iterations

    return self

def partial_fit(self, X_new, y_new, n_iterations=100):
    """Online / incremental learning: update weights with new data."""
    if self.weights is None:
        return self.fit(X_new, y_new)

    m, n = X_new.shape

    v_w = np.zeros(n)
    v_b = 0

    start_iter = self.stopped_at

    for i in range(n_iterations):
        current_lr = self.learning_rate / (1 + self.decay_rate *
(start_iter + i))

        self.lr_history.append(current_lr)

        y_pred = X_new.dot(self.weights) + self.bias

        error = y_pred - y_new

        cost = (1 / (2 * m)) * np.sum(error ** 2)

        self.cost_history.append(cost)

        dw = (1 / m) * X_new.T.dot(error)

```



```

db = (1 / m) * np.sum(error)

v_w = self.momentum * v_w + (1 - self.momentum) * dw
v_b = self.momentum * v_b + (1 - self.momentum) * db

self.weights -= current_lr * v_w
self.bias -= current_lr * v_b

self.stopped_at += n_iterations

return self

def predict(self, X):
    """Make predictions."""
    return X.dot(self.weights) + self.bias

print("SelfLearningLinearRegression class defined successfully.")

```

```
SelfLearningLinearRegression class defined successfully.
```

14. Train the Self-Learning Model and Compare

```

# Train the self-learning model

sl_model = SelfLearningLinearRegression(
    learning_rate=0.01,

```



```

n_iterations=2000,

decay_rate=0.001,

momentum=0.9,

patience=100,

tol=1e-6

)

sl_model.fit(X_train, y_train)

print('Self-Learning Model Training Complete')

print(f'  Stopped at iteration : {sl_model.stopped_at}')

print(f'  Final cost           : {sl_model.cost_history[-1]:.4f}')

print(f'  Initial LR            : {sl_model.lr_history[0]:.6f}')

print(f'  Final LR               : {sl_model.lr_history[-1]:.6f}')

print(f'  Early stopping used    : {sl_model.stopped_at < 2000}')
```

Self-Learning Model Training Complete

```

Stopped at iteration : 2000

Final cost           : 7.7212

Initial LR          : 0.010000

Final LR            : 0.003334

Early stopping used  : False
```

Compare convergence: Basic GD vs Self-Learning GD

```
fig, axes = plt.subplots(1, 2, figsize=(16, 5))
```



```
# Plot 1: Cost convergence

axes[0].plot(model.cost_history, label='Basic GD (LR=0.01)', linewidth=2)

axes[0].plot(sl_model.cost_history, label='Self-Learning GD', linewidth=2,
linestyle='--')

if sl_model.stopped_at < 2000:

    axes[0].axvline(x=sl_model.stopped_at, color='red', linestyle=':',
alpha=0.7,

                    label=f'Early Stop @ iter {sl_model.stopped_at}')

axes[0].set_xlabel('Iterations')

axes[0].set_ylabel('Cost (MSE / 2)')

axes[0].set_title('Cost Convergence Comparison')

axes[0].legend()

axes[0].grid(True, alpha=0.3)


# Plot 2: Adaptive Learning Rate schedule

axes[1].plot(sl_model.lr_history, color='darkorange', linewidth=2)

axes[1].set_xlabel('Iterations')


axes[1].set_ylabel('Learning Rate')

axes[1].set_title('Adaptive Learning Rate Decay Schedule')

axes[1].grid(True, alpha=0.3)


plt.tight_layout()

plt.show()
```



```
# Evaluate the Self-Learning model on the test set
```

```
y_pred_sl = sl_model.predict(X_test)
```

```
mae_sl = mean_absolute_error(y_test, y_pred_sl)
```

```
mse_sl = mean_squared_error(y_test, y_pred_sl)
```

```
rmse_sl = np.sqrt(mse_sl)
```

```
r2_sl = r2_score(y_test, y_pred_sl)
```

```
# Side-by-side comparison table
```

```
comparison = pd.DataFrame({
```

```
    'Metric': ['MAE', 'MSE', 'RMSE', 'R2 Score', 'Iterations Used'],
```

```
    'Basic GD': [round(mae,4), round(mse,4), round(rmse,4), round(r2,4),  
len(model.cost_history)],
```

```
    'Self-Learning GD': [round(mae_sl,4), round(mse_sl,4), round(rmse_sl,4),  
round(r2_sl,4), sl_model.stopped_at]
```

```
})
```

```
comparison = comparison.set_index('Metric')
```

```
print('Basic GD vs Self-Learning GD:')
```

```
comparison
```



Basic GD vs Self-Learning GD:

	Basic GD	Self-Learning GD
Metric		
MAE	3.4959	3.4953
MSE	18.5409	18.5409
RMSE	4.3059	4.3059
R2 Score	0.0958	0.0958
Iterations Used	1000.000 0	2000.0000

15. Online (Incremental) Learning Demonstration

In a real-world scenario, new student data arrives over time. The self-learning model can **update itself incrementally** with `partial_fit()` without retraining from scratch.

```
# Simulate online learning: split training data into 4 batches
# and feed them sequentially to the model

n_batches = 4

batch_size = len(X_train) // n_batches
```




```
online_model = SelfLearningLinearRegression(
    learning_rate=0.01, n_iterations=500,
    decay_rate=0.001, momentum=0.9,
    patience=80, tol=1e-6
)

batch_metrics = []

for b in range(n_batches):
    start_idx = b * batch_size
    end_idx = start_idx + batch_size if b < n_batches - 1 else len(X_train)
    X_batch = X_train[start_idx:end_idx]
    y_batch = y_train[start_idx:end_idx]

    if b == 0:
        online_model.fit(X_batch, y_batch)
    else:
        online_model.partial_fit(X_batch, y_batch, n_iterations=300)

    # Evaluate after each batch
    y_pred_batch = online_model.predict(X_test)
    batch_metrics.append({
        'Batch': b + 1,
        'Samples Seen': end_idx,
```



```

'MAE': round(mean_absolute_error(y_test, y_pred_batch), 4),

'RMSE': round(np.sqrt(mean_squared_error(y_test, y_pred_batch)), 4),

'R2': round(r2_score(y_test, y_pred_batch), 4)

})

print(f"Batch {b+1}: Samples seen = {end_idx:3d} | "

      f"MAE = {batch_metrics[-1]['MAE']:.4f} | "

      f"R2 = {batch_metrics[-1]['R2']:.4f}")

batch_df = pd.DataFrame(batch_metrics)

batch_df

```

```

Batch 1: Samples seen = 79 | MAE = 4.1280 | R2 = -0.2845

Batch 2: Samples seen = 158 | MAE = 3.6049 | R2 = -0.0198

Batch 3: Samples seen = 237 | MAE = 3.9283 | R2 = -0.1282

Batch 4: Samples seen = 316 | MAE = 3.7362 | R2 = -0.0187

```

	Batch	Samples Seen	MAE	RMS E	R2
0	1	79	4.1280	5.1322	-0.2845
1	2	158	3.6049	4.5728	-0.0198



2 3 237 3.928 4.809 -0.128
 3 9 2

3 4 316 3.736 4.570 -0.018
 2 5 7

Visualize how the model improves as it sees more data

```
fig, axes = plt.subplots(1, 3, figsize=(18, 5))
```

Plot 1: Full cost history across all batches

```
axes[0].plot(online_model.cost_history, color='steelblue', linewidth=1.5)
```

```
cum_iters = 0
```

```
for b in range(n_batches):
```

```
    n_iter_b = 500 if b == 0 else 300
```

```
    cum_iters += n_iter_b
```

```
    if cum_iters <= len(online_model.cost_history):
```

```
        axes[0].axvline(x=cum_iters, color='red', linestyle=':', alpha=0.6)
```

```
        axes[0].text(cum_iters, max(online_model.cost_history)*0.9,
```

```
                      f'Batch {b+1}', fontsize=8, ha='right', color='red')
```

```
axes[0].set_xlabel('Total Iterations')
```

```
axes[0].set_ylabel('Cost')
```

```
axes[0].set_title('Online Learning: Cost Over All Batches')
```

```
axes[0].grid(True, alpha=0.3)
```

Plot 2: R2 improvement per batch



```

axes[1].bar(batch_df['Batch'], batch_df['R2'], color='mediumseagreen',
            edgecolor='black', alpha=0.8)

axes[1].set_xlabel('Batch Number')
axes[1].set_ylabel('R2 Score')
axes[1].set_title('R2 Score After Each Batch')
axes[1].set_xticks(batch_df['Batch'])
axes[1].grid(True, alpha=0.3, axis='y')

# Plot 3: MAE / RMSE improvement per batch
x = np.arange(len(batch_df))

width = 0.35

axes[2].bar(x - width/2, batch_df['MAE'], width, label='MAE', color='coral',
            edgecolor='black', alpha=0.8)

axes[2].bar(x + width/2, batch_df['RMSE'], width, label='RMSE',
            color='skyblue',
            edgecolor='black', alpha=0.8)

axes[2].set_xlabel('Batch Number')
axes[2].set_ylabel('Error')
axes[2].set_title('MAE and RMSE After Each Batch')
axes[2].set_xticks(x)
axes[2].set_xticklabels(batch_df['Batch'])
axes[2].legend()


axes[2].grid(True, alpha=0.3, axis='y')

plt.tight_layout()

```



```
plt.show()
```



16. Self-Learning Dashboard - Final Comparison

```
# Final comprehensive dashboard
```

```
fig, axes = plt.subplots(2, 2, figsize=(14, 10))
```

```
fig.suptitle('Self-Learning Component Dashboard', fontsize=16,
fontWeight='bold')
```

```
# 1. Cost convergence comparison
```

```
axes[0, 0].plot(model.cost_history, label='Basic GD', linewidth=2)
```

```
axes[0, 0].plot(sl_model.cost_history, label='Self-Learning GD', linewidth=2)
```

```
axes[0, 0].set_title('Cost Convergence')
```

```
axes[0, 0].set_xlabel('Iterations')
```

```
axes[0, 0].set_ylabel('Cost')
```

```
axes[0, 0].legend()
```

```
axes[0, 0].grid(True, alpha=0.3)
```

```
# 2. Actual vs Predicted (Self-Learning)
```



```

axes[0, 1].scatter(y_test, y_pred_sl, alpha=0.6, color='steelblue',
edgecolor='k')

axes[0, 1].plot([y_test.min(), y_test.max()], [y_test.min(), y_test.max()],
                'r--', linewidth=2, label='Ideal')

axes[0, 1].set_title('Self-Learning: Actual vs Predicted')

axes[0, 1].set_xlabel('Actual G3')

axes[0, 1].set_ylabel('Predicted G3')

axes[0, 1].legend()

axes[0, 1].grid(True, alpha=0.3)

# 3. Metrics bar chart comparison

metrics_names = ['MAE', 'MSE', 'RMSE']

basic_vals = [mae, mse, rmse]

sl_vals = [mae_sl, mse_sl, rmse_sl]

x = np.arange(len(metrics_names))

width = 0.3

axes[1, 0].bar(x - width/2, basic_vals, width, label='Basic GD',
color='coral',

               edgecolor='black')

axes[1, 0].bar(x + width/2, sl_vals, width, label='Self-Learning GD',
color='mediumseagreen',

               edgecolor='black')

axes[1, 0].set_xticks(x)

axes[1, 0].set_xticklabels(metrics_names)

axes[1, 0].set_title('Error Metrics Comparison')

axes[1, 0].legend()

```



```

axes[1, 0].grid(True, alpha=0.3, axis='y')

# 4. R2 Score comparison

r2_vals = [r2, r2_sl]

bars = axes[1, 1].bar(['Basic GD', 'Self-Learning GD'], r2_vals,
                      color=['coral', 'mediumseagreen'], edgecolor='black',
                      width=0.5)

for bar, val in zip(bars, r2_vals):
    axes[1, 1].text(bar.get_x() + bar.get_width()/2, bar.get_height() +
0.005,
                    f'{val:.4f}', ha='center', fontweight='bold')

axes[1, 1].set_title('R2 Score Comparison')
axes[1, 1].set_ylabel('R2 Score')
axes[1, 1].grid(True, alpha=0.3, axis='y')

plt.tight_layout()


plt.show()

```



CHRIST
(DEEMED TO BE UNIVERSITY)
BANGALORE | DELHI NCR | PUNE

Self-Learning Component Dashboard



17. Self-Learning Component - Summary

Feature	Basic GD	Self-Learning GD
Learning Rate	Fixed throughout	Decays adaptively over time
Gradient Update	Vanilla	Momentum-accelerated
Stopping Criterion	Fixed number of iterations	Early stopping (patience-based)
New Data Handling	Must retrain from scratch	Incremental <code>partial_fit()</code>

Key Observations:



- The **adaptive learning rate** starts large for fast initial convergence and decays to make fine-grained adjustments, leading to a smoother and more stable convergence curve.
- **Momentum** helps the optimizer move faster along consistent gradient directions and dampens oscillations, reaching the minimum in fewer iterations.
- **Early stopping** prevents unnecessary computation once the model has converged, saving training time without sacrificing accuracy.
- **Online learning** (partial_fit) allows the model to incorporate new data incrementally. As more batches are fed, the model's R2 score improves and error metrics decrease - demonstrating genuine self-learning behavior.

18. Interpretation and Conclusions

Convergence Behavior

- The cost function decreases monotonically during training, confirming that gradient descent is minimizing the loss correctly.
- A **smaller learning rate** (e.g., 0.001) results in slower convergence - the model may need many more iterations to reach the optimum.
- A **larger learning rate** (e.g., 0.1) speeds up convergence significantly, and for this dataset, remains stable and converges extremely quickly, though in general, setting it too high can cause oscillations or divergence.
- The **optimal learning rate** for this dataset is around **0.05-0.1**, balancing fast convergence speed and stability.

Prediction Performance and Metrics Interpretation

- The **MAE of ~3.50** means the model's predictions are off by about 3.5 grade points on average on a 0-20 scale.
- The **RMSE of ~4.31** is higher than MAE, indicating the presence of some larger prediction errors that get penalized more heavily.
- The **R2 score of ~0.096** shows the model explains only about 9.6% of the variance in final grades. This low value indicates that demographic and behavioral features alone are weak predictors of academic performance in a simple linear model.
- The **MSE of ~18.54** confirms relatively high average squared errors.
- Overall, the Gradient Descent algorithm converged successfully, but the predictive performance is relatively weak because student performance depends on many complex factors (motivation, exam difficulty, teaching quality, etc.) not fully captured by the selected features. Removing G1 and G2 (to avoid data leakage) also eliminated the strongest predictors.



Self-Learning Enhancements

- The **Self-Learning model** achieves comparable or better performance than basic GD while being more computationally efficient (early stopping) and adaptable (online learning).
- **Momentum + adaptive LR** together produce smoother and faster convergence than a fixed learning rate alone.
- The **online learning** demonstration shows that the model genuinely improves as it sees more data - a core property of self-learning systems.

Key Takeaways

1. **Gradient Descent** successfully optimizes the linear regression parameters by iteratively reducing the cost function.
2. **Learning rate** is a critical hyperparameter - too small leads to slow convergence, too large may cause divergence.
3. **Feature scaling** is essential for gradient descent to converge efficiently.
4. **Self-learning enhancements** (adaptive LR, momentum, early stopping, online learning) make the model smarter, faster, and more robust.
5. The model provides a reasonable baseline; performance could be improved with feature engineering, polynomial features, or regularization.

Lab Experiment 6: Comparison of Logistic Regression and K Nearest Neighbors (KNN) Classifiers

Aim

To implement Logistic Regression and K Nearest Neighbors (KNN) classifiers on the Breast Cancer Wisconsin (Diagnostic) dataset and compare their performance using standard classification evaluation metrics.

Objectives

1. To preprocess the dataset for classification.
2. To implement Logistic Regression and KNN classifiers using Scikit Learn.
3. To evaluate both models using standard performance metrics.



4. To compare the performance of Logistic Regression and KNN and identify the better classifier.

Dataset

Dataset Name: Breast Cancer Wisconsin (Diagnostic)

Source: UCI Machine Learning Repository

Dataset Link:

<https://archive.ics.uci.edu/dataset/17/breast+cancer+wisconsin+diagnostic>

Tasks to Perform

1. Download and load the Breast Cancer Wisconsin (Diagnostic) dataset.
2. Perform exploratory data analysis and data preprocessing.
3. Check for missing values and prepare the dataset for classification.
4. Split the dataset into training and testing sets.
5. Train a Logistic Regression classifier.
6. Train a K Nearest Neighbors (KNN) classifier.
7. Evaluate both classifiers using:
 - Accuracy
 - Precision
 - Recall
 - F1 Score
 - Confusion Matrix
8. Compare the performance of the two classifiers using a comparison table.
9. Interpret the results and comment on the better-performing classifier for the given dataset.

Expected Learning Outcomes

Upon successful completion of this experiment, students will be able to:

1. Implement Logistic Regression and KNN classifiers using Scikit Learn.
2. Apply preprocessing techniques to a real world dataset.
3. Evaluate classification models using standard performance metrics.
4. Compare the performance of multiple classifiers and interpret the results.
5. Recommend the most suitable classifier based on experimental findings.



Code with Results:

1. Aim

To build, evaluate, and compare the performance of **Logistic Regression** and **K-Nearest Neighbors (KNN)** classification algorithms on the Breast Cancer Wisconsin (Diagnostic) dataset, analyze their classification metrics, examine the importance of feature scaling, and perform an extended **Self-Learning Component** exploring hyperparameter tuning, ROC-AUC curve analysis, and feature importance.

2. Objectives

1. Load the Breast Cancer Wisconsin (Diagnostic) dataset (`wdbc.data`).
2. Perform Exploratory Data Analysis (EDA).
3. Check dataset dimensions.
4. Display the first five rows.
5. Assign meaningful column names (32 total columns).
6. Display dataset structural information.
7. Generate summary statistics for numerical features.
8. Verify and check for missing values.
9. Analyze and display class distribution.
10. Encode target diagnosis labels (`M -> 1`, `B -> 0`).
11. Preprocess data by removing the unnecessary `ID` column.
12. Split the dataset into training (80%) and testing (20%) subsets.
13. Implement stratified sampling during train-test split.
14. Perform feature scaling using `StandardScaler`.
15. Train a Logistic Regression model (`max_iter=500`, `random_state=42`).
16. Train a K-Nearest Neighbors model (`n_neighbors=5`).
17. Evaluate both models using Accuracy, Precision, Recall, F1 Score, Confusion Matrix, and Classification Report.
18. Construct a comparative evaluation table using a Pandas DataFrame.
19. Generate diagnostic plots (Class distribution, Correlation heatmap, Confusion matrices, Comparison bar chart).
20. **Self-Learning Component:**
 - Conduct hyperparameter analysis for KNN across
 - `k`
 - `=`
 - `1`
 - `to`
 - `k`



- =
 - 25
 - (Elbow curve and performance metrics).
 - Analyze Logistic Regression regularization sensitivity (
 - C
 - parameter tuning).
 - Compute and plot Receiver Operating Characteristic (ROC) curves and AUC scores.
 - Evaluate feature importance via Logistic Regression model coefficients.
21. Formulate conclusions and clinical recommendations.

3. Dataset Description

The **Breast Cancer Wisconsin (Diagnostic)** dataset contains features computed from a digitized image of a fine needle aspirate (FNA) of a breast mass. They describe characteristics of the cell nuclei present in the image.

- **Total Instances:** 569 samples
- **Attributes:** 32 columns (ID, Diagnosis, and 30 continuous numerical features)
- **Target Variable (*diagnosis*):**
 1. **M:** Malignant (Cancerous)
 2. **B:** Benign (Non-cancerous)
- **Feature Groups (10 fundamental real-valued features computed in 3 statistics: Mean, Standard Error (SE), and Worst/Largest):**
 1. *radius* (mean of distances from center to points on the perimeter)
 2. *texture* (standard deviation of gray-scale values)
 3. *perimeter*
 4. *area*
 5. *smoothness* (local variation in radius lengths)
 6. *compactness* (
 7. *p*
 8. *e*
 9. *r*
 10. *i*
 11. *m*
 12. *e*
 13. *t*
 14. *e*
 15. *r*
 16. *z*



- 17. /
- 18. a
- 19. r
- 20. e
- 21. a
- 22. -
- 23. 1.0
- 24.)
- 25. `concavity` (severity of concave portions of the contour)
- 26. `concave points` (number of concave portions of the contour)
- 27. `symmetry`
- 28. `fractal dimension` ("coastline approximation" - 1)

4. Import Libraries

Before beginning the experiment, we import all necessary Python libraries required for data manipulation, numerical analysis, visualization, model training, evaluation, and hyperparameter tuning.

```
# Import fundamental libraries for data processing and visualization

import os

import pandas as pd

import numpy as np

import matplotlib.pyplot as plt

import seaborn as sns


# Import Scikit-Learn modules for preprocessing, modeling, and evaluation

from sklearn.model_selection import train_test_split, GridSearchCV

from sklearn.preprocessing import StandardScaler

from sklearn.linear_model import LogisticRegression

from sklearn.neighbors import KNeighborsClassifier
```



```
from sklearn.metrics import (

    accuracy_score,

    precision_score,

    recall_score,

    f1_score,

    confusion_matrix,

    classification_report,

    roc_curve,

    roc_auc_score

)


# Set visualization styles for professional aesthetics

sns.set_theme(style="whitegrid", palette="muted")

plt.rcParams['font.size'] = 11

plt.rcParams['figure.titlesize'] = 14

plt.rcParams['axes.labelsize'] = 12


print("All required libraries imported successfully.")
```

```
All required libraries imported successfully.
```

5. Load Dataset

We load the `wdbc.data` dataset file provided in the working directory using Pandas. Since the raw CSV file does not include header rows, we pass `header=None` during loading.



```
# Define dataset path

data_path = "wdbc.data"

# Verify dataset file existence

if not os.path.exists(data_path):

    raise FileNotFoundError(f"Dataset file '{data_path}' not found in current
directory.")

# Load raw dataset without headers

raw_df = pd.read_csv(data_path, header=None)

print("Raw dataset loaded successfully.")
```

Raw dataset loaded successfully.

6. Assign Column Names

The dataset contains 32 attributes:

1. `id`: Unique identifier for each sample
2. `diagnosis`: Target attribute (`M` for Malignant, `B` for Benign)
3. Columns 3-12: Mean values of 10 cell nuclear features
4. Columns 13-22: Standard Error (SE) of 10 cell nuclear features
5. Columns 23-32: Worst/Largest values (mean of 3 largest values) of 10 cell nuclear features

We assign descriptive column names to ensure clarity during analysis.

```
# Define column names

feature_names = [

    'radius', 'texture', 'perimeter', 'area', 'smoothness',
```




```

        'compactness', 'concavity', 'concave_points', 'symmetry',
        'fractal_dimension'

    ]

    # Generate column names for mean, se, and worst feature groups

    mean_cols = [f"{col}_mean" for col in feature_names]

    se_cols = [f"{col}_se" for col in feature_names]

    worst_cols = [f"{col}_worst" for col in feature_names]

    # Combine all column headers

    column_names = ['id', 'diagnosis'] + mean_cols + se_cols + worst_cols

    # Assign column names to DataFrame

    raw_df.columns = column_names

    df = raw_df.copy()

    print(f"Column names assigned successfully. Total columns:
    {len(df.columns)}")

```

```
Column names assigned successfully. Total columns: 32
```

7. Exploratory Data Analysis (EDA)

Exploratory Data Analysis allows us to inspect the dataset's structure, dimensions, data types, missing values, class distribution, and feature correlations.

7.1 Dataset Dimensions



Checking the total number of rows (instances) and columns (features).

```
# Display dataset shape
```

```
print(f"Dataset Dimensions: {df.shape[0]} rows, {df.shape[1]} columns")
```

```
Dataset Dimensions: 569 rows, 32 columns
```

7.2 Preview First Five Rows

Displaying the first 5 records of the dataset to verify column alignment and sample values.

```
# Display the first 5 rows
```

```
df.head()
```

	id	diagnoses	radiologist	text	perimeter	area	smoothness	compactness	concavity	concave points	radius	texture	perimeter	area	smoothness	compactness	concavity	concave points	radius	symmetry	fractal dimension
0	842302	M	1	0.38	12.2	1	0.18	0.27	0.30	0.1471	0	2.538	18.46	0.162	0.656	0.7119	0.2654	0.4601	0.11890		

8			1																	
4		2		3						2										
2		0	1	2	0.0	0.0	0.	0.0	.	4	2	15	9	0.1	0.1	0.				
5	M	.	7.	2.	84	78	08	701	.	.	3.	8.	5	23	86	24	0.1	0.	0.08	
1		5	7	6	74	64	69	7	.	9	4	80	6	8	6	16	860	27	902	
7		7		0						9	1		.							
													0							

[illegible][illegible][illegible]

7.3 Dataset Information

Dept of Computer Science, CHRIST (Deemed to be University)



```
# Display dataset structural summary
```

```
df.info()
```

```
<class 'pandas.DataFrame'>
```

```
RangeIndex: 569 entries, 0 to 568
```

```
Data columns (total 32 columns):
```

#	Column	Non-Null Count	Dtype
0	id	569 non-null	int64
1	diagnosis	569 non-null	str
2	radius_mean	569 non-null	float64
3	texture_mean	569 non-null	float64
4	perimeter_mean	569 non-null	float64
5	area_mean	569 non-null	float64
6	smoothness_mean	569 non-null	float64
7	compactness_mean	569 non-null	float64
8	concavity_mean	569 non-null	float64
9	concave_points_mean	569 non-null	float64
10	symmetry_mean	569 non-null	float64
11	fractal_dimension_mean	569 non-null	float64
12	radius_se	569 non-null	float64
13	texture_se	569 non-null	float64
14	perimeter_se	569 non-null	float64
15	area_se	569 non-null	float64



16	smoothness_se	569 non-null	float64
17	compactness_se	569 non-null	float64
18	concavity_se	569 non-null	float64
19	concave_points_se	569 non-null	float64
20	symmetry_se	569 non-null	float64
21	fractal_dimension_se	569 non-null	float64
22	radius_worst	569 non-null	float64
23	texture_worst	569 non-null	float64
24	perimeter_worst	569 non-null	float64
25	area_worst	569 non-null	float64
26	smoothness_worst	569 non-null	float64
27	compactness_worst	569 non-null	float64
28	concavity_worst	569 non-null	float64
29	concave_points_worst	569 non-null	float64
30	symmetry_worst	569 non-null	float64
31	fractal_dimension_worst	569 non-null	float64

dtypes: float64(30), int64(1), str(1)

memory usage: 142.4 KB

7.4 Descriptive Statistics

Generating statistical summary (mean, std, min, max, quartiles) for all numerical features.

```
# Summary statistics for numerical columns
```

```
df.describe().T[['mean', 'std', 'min', '50%', 'max']].head(10)
```



	mean	std	min	50%	max
id	3.037183e+0 7	1.250206e+0 8	8670.0000 0	906024.0000 0	9.113205e+0 8
radius_mean	1.412729e+0 1	3.524049e+0 0	6.98100	13.37000	2.811000e+0 1
texture_mean	1.928965e+0 1	4.301036e+0 0	9.71000	18.84000	3.928000e+0 1
perimeter_mean	9.196903e+0 1	2.429898e+0 1	43.79000	86.24000	1.885000e+0 2
area_mean	6.548891e+0 2	3.519141e+0 2	143.50000	551.10000	2.501000e+0 3
smoothness_mean	9.636028e-0 2	1.406413e-0 2	0.05263	0.09587	1.634000e-0 1
compactness_mean	1.043410e-0 1	5.281276e-0 2	0.01938	0.09263	3.454000e-0 1
concavity_mean	8.879932e-0 2	7.971981e-0 2	0.00000	0.06154	4.268000e-0 1
concave_points_mean	4.891915e-0 2	3.880284e-0 2	0.00000	0.03350	2.012000e-0 1
symmetry_mean	1.811619e-0 1	2.741428e-0 2	0.10600	0.17920	3.040000e-0 1



7.5 Check Missing Values

Verifying whether the dataset contains any missing or null values.

```
# Check for null / missing values across all columns

missing_counts = df.isnull().sum()

total_missing = missing_counts.sum()

print(f"Total Missing Values in Dataset: {total_missing}")

if total_missing == 0:

    print("Clean dataset confirmed: No missing values found across all
columns.")

else:

    print(missing_counts[missing_counts > 0])
```

Total Missing Values in Dataset: 0

Clean dataset confirmed: No missing values found across all columns.

7.6 Class Distribution

Checking the target class counts and proportions for Malignant (M) vs Benign (B).

```
# Compute target class distribution

class_counts = df['diagnosis'].value_counts()

class_percentages = df['diagnosis'].value_counts(normalize=True) * 100

class_dist_df = pd.DataFrame({

    'Count': class_counts,
```



```
'Percentage (%)': class_percentages.round(2)

})
```

```
print("Class Distribution Summary:")

display(class_dist_df)
```

Class Distribution Summary:

	Coun t	Percentage (%)
diagnosi s		
B	357	62.74
M	212	37.26

7.7 Class Distribution Visualization

Visualizing the class distribution using a styled bar chart with bar value annotations.

```
# Plot class distribution bar chart

plt.figure(figsize=(7, 5))

ax = sns.barplot(

    x=class_counts.index,

    y=class_counts.values,

    hue=class_counts.index,
```




```

palette=['#e74c3c', '#2ecc71'],

legend=False

)

plt.title("Diagnosis Class Distribution (Malignant vs. Benign)", fontsize=14,
fontweight='bold', pad=15)

plt.xlabel("Diagnosis Class", fontsize=12, labelpad=10)

plt.ylabel("Number of Samples", fontsize=12, labelpad=10)

plt.xticks(ticks=[0, 1], labels=['Benign (B)', 'Malignant (M)'])

plt.ylim(0, max(class_counts.values) * 1.15)

# Annotate count and percentage on top of bars
total_samples = len(df)

for p in ax.patches:

    height = int(p.get_height())

    pct = (height / total_samples) * 100

    ax.annotate(

        f"{height} ({pct:.1f}%)",

        (p.get_x() + p.get_width() / 2., height),

        ha='center', va='bottom',

        fontsize=11, fontweight='bold',

        xytext=(0, 5), textcoords='offset points'


    )

plt.tight_layout()

```



```
plt.show()
```



7.8 Feature Correlation Heatmap

Computing and visualizing the pairwise Pearson correlation matrix across all 30 numerical features to analyze multicollinearity.

```
# Compute correlation matrix for numerical features
```

```
numerical_df = df.drop(columns=['id', 'diagnosis'])
```

```
corr_matrix = numerical_df.corr()
```

```
# Plot full correlation heatmap
```



```
plt.figure(figsize=(14, 11))

sns.heatmap(

    corr_matrix,

    cmap='coolwarm',

    annot=False,

    fmt=".2f",

    linewidths=0.5,


    cbar_kws={'label': 'Pearson Correlation Coefficient'}

)


plt.title("Correlation Heatmap of Numerical Features", fontsize=15,
fontweight='bold', pad=15)

plt.tight_layout()

plt.show()
```



8. Data Preprocessing

To prepare the dataset for machine learning model training:

1. **Target Encoding:** Convert string labels **M** (Malignant) to **1** and **B** (Benign) to **0**.
2. **Feature Dropping:** Drop the **id** column as it serves purely as a sample identifier and possesses no predictive power.

```
# Encode target diagnosis label: M -> 1, B -> 0
```

```
df['diagnosis'] = df['diagnosis'].map({'M': 1, 'B': 0})
```



```
# Remove the ID column

df_processed = df.drop(columns=['id'])

print("Target encoding and ID column removal complete.")

print(f"Processed DataFrame Shape: {df_processed.shape}")

print(f"Target Diagnosis Unique Values:
{df_processed['diagnosis'].unique().tolist()}")
```

Target encoding and ID column removal complete.

Processed DataFrame Shape: (569, 31)

Target Diagnosis Unique Values: [1, 0]

9. Train-Test Split

We separate the feature matrix (

X

) from the target vector (

y

) and perform an **80:20 train-test split** using **stratified sampling** (`stratify=y`) and a reproducible random seed (`random_state=42`).

- **Training set (80%):** Used for model optimization and weight updates.
- **Testing set (20%):** Held-out dataset used to evaluate model generalization performance.

```
# Separate features X and target y

X = df_processed.drop(columns=['diagnosis'])
```



```
y = df_processed['diagnosis']

# Perform 80:20 train-test split with stratified sampling
X_train, X_test, y_train, y_test = train_test_split(
    X, y,
    test_size=0.20,
    stratify=y,
    random_state=42
)

print(f"Training set shape: X_train = {X_train.shape}, y_train = {y_train.shape}")
print(f"Testing set shape: X_test = {X_test.shape}, y_test = {y_test.shape}")
print(f"Training class distribution:\n{y_train.value_counts(normalize=True).round(3).to_dict()}")
print(f"Testing class distribution:\n{y_test.value_counts(normalize=True).round(3).to_dict()}")
```

```
Training set shape: X_train = (455, 30), y_train = (455,)
Testing set shape: X_test = (114, 30), y_test = (114,)
Training class distribution:
{0: 0.626, 1: 0.374}
Testing class distribution:
{0: 0.632, 1: 0.368}
```



10. Feature Scaling

Distance-based models like KNN and gradient-based models like Logistic Regression perform optimally when features are standardized to have zero mean (

μ

=

0

) and unit variance (

σ

²

=

1

).

z

=

x

–

μ

σ



Important: `StandardScaler` is fitted ONLY on `X_train` to prevent data leakage, and then applied to transform both `X_train` and `X_test`.

```
# Initialize StandardScaler

scaler = StandardScaler()

# Fit on training data and transform both training and testing sets

X_train_scaled = scaler.fit_transform(X_train)

X_test_scaled = scaler.transform(X_test)

# Convert scaled arrays back to DataFrames for inspection

X_train_scaled_df = pd.DataFrame(X_train_scaled, columns=X.columns)

X_test_scaled_df = pd.DataFrame(X_test_scaled, columns=X.columns)

print("StandardScaler fitting and transformation completed.")

print(f"Scaled X_train Mean: {X_train_scaled.mean():.4f}, Std: {X_train_scaled.std():.4f}")
```

```
StandardScaler fitting and transformation completed.
```

```
Scaled X_train Mean: 0.0000, Std: 1.0000
```

11. Logistic Regression Implementation

Logistic Regression models the probability of binary outcomes using the sigmoid function:

$$P(y = 1 | x) = \frac{1}{1 + e^{-(w^T x + b)}}$$



We instantiate `LogisticRegression` with `max_iter=500` to ensure optimizer convergence and `random_state=42` for reproducibility.

```
# Instantiate Logistic Regression model

log_reg = LogisticRegression(max_iter=500, random_state=42)

# Train model on scaled training data

log_reg.fit(X_train_scaled, y_train)

# Predict class labels on scaled test data

y_pred_lr = log_reg.predict(X_test_scaled)

y_prob_lr = log_reg.predict_proba(X_test_scaled)[:, 1]

print("Logistic Regression model trained and predictions generated
successfully.")
```

Logistic Regression model trained and predictions generated successfully.

12. K-Nearest Neighbors (KNN) Implementation

K-Nearest Neighbors is a non-parametric instance-based learning algorithm that classifies samples based on the majority vote of their

knearest neighbors using Euclidean distance: $d(p,q) = \sqrt{\sum_{i=1}^n (p_i - q_i)^2}$

We train `KNeighborsClassifier` with `n_neighbors=5` on the standardized features.

```
# Instantiate KNN model with k=5

knn = KNeighborsClassifier(n_neighbors=5)
```



```
# Train model on scaled training data

knn.fit(X_train_scaled, y_train)

# Predict class labels on scaled test data

y_pred_knn = knn.predict(X_test_scaled)

y_prob_knn = knn.predict_proba(X_test_scaled)[:, 1]

print("K-Nearest Neighbors (k=5) model trained and predictions generated
successfully.")
```

```
K-Nearest Neighbors (k=5) model trained and predictions generated
successfully.
```

13. Model Evaluation

To evaluate both classifiers systematically without duplicate code, we define a reusable helper function `evaluate_classifier()` that computes:

- **Accuracy:** Proportion of overall correct predictions.
- **Precision:** Proportion of correctly predicted positive cases among all predicted positives.
- **Recall (Sensitivity):** Proportion of correctly predicted positive cases among all actual positives.
- **F1 Score:** Harmonic mean of Precision and Recall.
- **Confusion Matrix:** Breakdown of True Positives (TP), True Negatives (TN), False Positives (FP), and False Negatives (FN).
- **Classification Report:** Detailed per-class evaluation summary.

```
# Reusable helper function for evaluating classifier performance

def evaluate_classifier(y_true, y_pred, model_name):
    """
```



Computes and prints classification metrics, confusion matrix, and classification report.

Returns metrics dictionary for comparison.

```

"""

acc = accuracy_score(y_true, y_pred)

prec = precision_score(y_true, y_pred)

rec = recall_score(y_true, y_pred)

f1 = f1_score(y_true, y_pred)

cm = confusion_matrix(y_true, y_pred)

report = classification_report(y_true, y_pred, target_names=['Benign
(0)', 'Malignant (1)'])

print(f"=====")
print(f"          EVALUATION REPORT: {model_name.upper()}")
print(f"=====")

print(f"Accuracy:  {acc:.4f} ({acc * 100:.2f}%)")
print(f"Precision: {prec:.4f} ({prec * 100:.2f}%)")
print(f"Recall:     {rec:.4f} ({rec * 100:.2f}%)")
print(f"F1 Score:   {f1:.4f} ({f1 * 100:.2f}%)")
print(f"\nConfusion Matrix:\n{cm}")
print(f"\nClassification Report:\n{report}")

return {

    'Model': model_name,

    'Accuracy': acc,

```



```

        'Precision': prec,

        'Recall': rec,

        'F1 Score': f1,

        'Confusion Matrix': cm

    }

# Evaluate Logistic Regression

lr_metrics = evaluate_classifier(y_test, y_pred_lr, "Logistic Regression")

# Evaluate K-Nearest Neighbors

knn_metrics = evaluate_classifier(y_test, y_pred_knn, "K-Nearest Neighbors
(k=5) ")

=====

                EVALUATION REPORT: LOGISTIC REGRESSION

=====

Accuracy:  0.9649 (96.49%)

Precision: 0.9750 (97.50%)

Recall:    0.9286 (92.86%)

F1 Score:  0.9512 (95.12%)

Confusion Matrix:

[[71  1]

 [ 3 39]]

```



Classification Report:

	precision	recall	f1-score	support
Benign (0)	0.96	0.99	0.97	72
Malignant (1)	0.97	0.93	0.95	42
accuracy			0.96	114
macro avg	0.97	0.96	0.96	114
weighted avg	0.97	0.96	0.96	114

EVALUATION REPORT: K-NEAREST NEIGHBORS (K=5)

Accuracy: 0.9561 (95.61%)

Precision: 0.9744 (97.44%)

Recall: 0.9048 (90.48%)

F1 Score: 0.9383 (93.83%)

Confusion Matrix:

```
[[71  1]
 [ 4 38]]
```

Classification Report:

	precision	recall	f1-score	support
--	-----------	--------	----------	---------



Benign (0)	0.95	0.99	0.97	72
Malignant (1)	0.97	0.90	0.94	42
accuracy			0.96	114
macro avg	0.96	0.95	0.95	114
weighted avg	0.96	0.96	0.96	114

14. Performance Comparison

We construct a comparison table using a Pandas DataFrame to side-by-side compare Accuracy, Precision, Recall, and F1 Score for both models, and automatically identify the superior model.

```
# Create evaluation comparison DataFrame

comparison_df = pd.DataFrame([

    {

        'Classifier': lr_metrics['Model'],

        'Accuracy': lr_metrics['Accuracy'],

        'Precision': lr_metrics['Precision'],

        'Recall': lr_metrics['Recall'],

        'F1 Score': lr_metrics['F1 Score']

    },

    {

        'Classifier': knn_metrics['Model'],

        'Accuracy': knn_metrics['Accuracy'],
```



```

        'Precision': knn_metrics['Precision'],

        'Recall': knn_metrics['Recall'],

        'F1 Score': knn_metrics['F1 Score']

    }

])

# Display formatted comparison table

print("Classifier Performance Comparison Table:")

display(comparison_df.style.format({

    'Accuracy': '{:.4f}',

    'Precision': '{:.4f}',

    'Recall': '{:.4f}',

    'F1 Score': '{:.4f}'

}))

```

Classifier Performance Comparison Table:

	Classifier	Accuracy	Precision	Recall	F1 Score
0	Logistic Regression	0.9649	0.9750	0.9286	0.9512
1	K-Nearest Neighbors (k=5)	0.9561	0.9744	0.9048	0.9383

```

# Identify better performing classifier based on F1 Score and Accuracy

lr_f1 = lr_metrics['F1 Score']

```



```
knn_f1 = knn_metrics['F1 Score']

if lr_f1 > knn_f1:

    better_model = "Logistic Regression"

    diff = (lr_f1 - knn_f1) * 100

elif knn_f1 > lr_f1:

    better_model = "K-Nearest Neighbors (KNN)"

    diff = (knn_f1 - lr_f1) * 100

else:

    better_model = "Both models performed equally"

    diff = 0.0

print(f"Summary Decision: {better_model} achieved superior overall performance.")
```

Summary Decision: Logistic Regression achieved superior overall performance.

15. Visualizations

We generate the required evaluation figures:

1. Confusion Matrix heatmap for **Logistic Regression**.
2. Confusion Matrix heatmap for **K-Nearest Neighbors (KNN)**.
3. Comparative grouped bar chart illustrating **Accuracy, Precision, Recall, and F1 Score** side-by-side.

15.1 Confusion Matrix Heatmaps

Displaying confusion matrix heatmaps with count and percentage annotations for both models.



```
# Plot Confusion Matrices side-by-side

fig, axes = plt.subplots(1, 2, figsize=(13, 5))

labels = ['Benign (0)', 'Malignant (1)']

# Helper for confusion matrix plot

def draw_cm(cm, ax, title, color_map):

    sns.heatmap(cm, annot=True, fmt='d', cmap=color_map, cbar=False, ax=ax,

                xticklabels=labels, yticklabels=labels, annot_kws={"size":
14, "weight": "bold"})

    ax.set_title(title, fontsize=13, fontweight='bold', pad=12)

    ax.set_xlabel('Predicted Label', fontsize=11)


    ax.set_ylabel('True Label', fontsize=11)

draw_cm(lr_metrics['Confusion Matrix'], axes[0], "Logistic Regression -
Confusion Matrix", "Blues")

draw_cm(knn_metrics['Confusion Matrix'], axes[1], "KNN (k=5) - Confusion
Matrix", "Greens")

plt.tight_layout()

plt.show()
```



15.2 Classifier Performance Comparison Bar Chart

Displaying a side-by-side multi-metric comparative bar chart with value annotations.

```
# Prepare data for metric comparison plot

metrics_list = ['Accuracy', 'Precision', 'Recall', 'F1 Score']

df_plot = pd.melt(comparison_df, id_vars=['Classifier'],
                  value_vars=metrics_list,

                  var_name='Metric', value_name='Score')

plt.figure(figsize=(10, 6))

ax = sns.barplot(
    data=df_plot,
    x='Metric',
    y='Score',
    hue='Classifier',
    palette=['#2980b9', '#27ae60']
)
```



```
plt.title("Evaluation Metrics Comparison: Logistic Regression vs KNN",
         fontsize=14, fontweight='bold', pad=15)

plt.ylabel("Metric Score (0.0 to 1.0)", fontsize=12, labelpad=10)

plt.xlabel("Evaluation Metric", fontsize=12, labelpad=10)

plt.ylim(0.85, 1.02)

plt.legend(title="Classifier", frameon=True, facecolor='white', loc='lower
right')

# Add bar annotations

for p in ax.patches:

    height = p.get_height()

    if not np.isnan(height) and height > 0:

        ax.annotate(

            f"{height:.4f}",

            (p.get_x() + p.get_width() / 2., height),

            ha='center', va='bottom',


            fontsize=9.5, fontweight='bold',

            xytext=(0, 4), textcoords='offset points'

        )

plt.tight_layout()

plt.show()
```



16. Self-Learning Component (Extended Analysis)

This self-learning module extends the basic laboratory requirements by investigating advanced machine learning topics:

1. **KNN Hyperparameter Optimization (**
2. **k**
3. **-Value Tuning & Elbow Method):** Evaluating performance across range
4. **k**
5. **=**
6. **1**
7. **to**
8. **k**
9. **=**
10. **25**
11. **to discover optimal neighbor count.**
12. **Logistic Regression Regularization Analysis:** Evaluating effect of inverse regularization parameter
13. **C**



14. $\in$
15. [
16. 0.001
17. ,
18. 100
19.]
20. .
21. **ROC Curve & AUC Score Comparison:** Assessing threshold-agnostic classifier performance.
22. **Feature Importance Ranking:** Extracting top predictive features based on Logistic Regression coefficient magnitudes.

16.1 KNN Optimal

k

-Value Selection (Elbow Method)

We evaluate test error rate and F1 Score for

k

values ranging from 1 to 25 to analyze underfitting vs. overfitting.

```
# Search for optimal k in KNN

k_range = range(1, 26)

error_rates = []

f1_scores = []

for k in k_range:

    knn_temp = KNeighborsClassifier(n_neighbors=k)

    knn_temp.fit(X_train_scaled, y_train)

    y_pred_k = knn_temp.predict(X_test_scaled)

    error_rates.append(1.0 - accuracy_score(y_test, y_pred_k))
```



```

f1_scores.append(f1_score(y_test, y_pred_k))

# Find k with minimum error rate and highest F1 score
optimal_k_err = k_range[np.argmin(error_rates)]
optimal_k_f1 = k_range[np.argmax(f1_scores)]

print(f"Optimal k based on Minimum Error Rate ({min(error_rates):.4f}): k = {optimal_k_err}")

print(f"Optimal k based on Maximum F1 Score ({max(f1_scores):.4f}): k = {optimal_k_f1}")

# Plot Error Rate and F1 Score vs k
plt.figure(figsize=(12, 5))

plt.subplot(1, 2, 1)

plt.plot(k_range, error_rates, color='#c0392b', linestyle='--', marker='o',
markerfacecolor='#e74c3c', markersize=6)

plt.axvline(x=optimal_k_err, color='black', linestyle=':', label=f'Optimal k = {optimal_k_err}')

plt.title("KNN Error Rate vs. k (Elbow Method)", fontsize=13,
fontweight='bold')

plt.xlabel("Number of Neighbors (k)", fontsize=11)

plt.ylabel("Test Error Rate", fontsize=11)

plt.xticks(k_range[::2])

plt.legend()

plt.subplot(1, 2, 2)

```



```
plt.plot(k_range, f1_scores, color='#27ae60', linestyle='--', marker='s',
markerfacecolor='#2ecc71', markersize=6)

plt.axvline(x=optimal_k_f1, color='black', linestyle=':', label=f'Optimal k =
{optimal_k_f1}')

plt.title("KNN F1 Score vs. k", fontsize=13, fontweight='bold')

plt.xlabel("Number of Neighbors (k)", fontsize=11)

plt.ylabel("Test F1 Score", fontsize=11)

plt.xticks(k_range[::2])


plt.legend()

plt.tight_layout()

plt.show()
```

Optimal k based on Minimum Error Rate (0.0439): k = 5

Optimal k based on Maximum F1 Score (0.9383): k = 5



16.2 Logistic Regression Regularization Sensitivity (



C

Tuning)

Analyzing how varying inverse regularization strength

C

affects train vs test accuracy.

```
# Test different C values for Logistic Regression

c_values = [0.001, 0.01, 0.1, 1.0, 10.0, 100.0]

train_acc_list = []

test_acc_list = []

for c in c_values:

    lr_c = LogisticRegression(C=c, max_iter=500, random_state=42)

    lr_c.fit(X_train_scaled, y_train)

    train_acc_list.append(accuracy_score(y_train,
lr_c.predict(X_train_scaled)))

    test_acc_list.append(accuracy_score(y_test, lr_c.predict(X_test_scaled)))

c_results_df = pd.DataFrame({

    'C Value': c_values,

    'Train Accuracy': train_acc_list,

    'Test Accuracy': test_acc_list

})
```




```
print("Logistic Regression C Parameter Sensitivity:")

display(c_results_df.style.format({'Train Accuracy': '{:.4f}', 'Test
Accuracy': '{:.4f}'}))

# Plot Regularization effect

plt.figure(figsize=(8, 5))

plt.plot(range(len(c_values)), train_acc_list, label='Train Accuracy',
marker='o', color='#2980b9')

plt.plot(range(len(c_values)), test_acc_list, label='Test Accuracy',
marker='s', color='#e67e22')

plt.xticks(range(len(c_values)), [str(c) for c in c_values])

plt.xscale('linear')

plt.title("Logistic Regression: Accuracy vs. Regularization Parameter (C)",
fontsize=13, fontweight='bold')

plt.xlabel("Regularization Strength (C)", fontsize=11)

plt.ylabel("Accuracy Score", fontsize=11)

plt.legend()

plt.tight_layout()


plt.show()
```

Logistic Regression C Parameter Sensitivity:

	C Value	Train Accuracy	Test Accuracy
0	0.001000	0.9033	0.8772
1	0.010000	0.9538	0.9474



2	0.100000	0.9802	0.9825
3	1.000000	0.9868	0.9649
4	10.000000	0.9890	0.9737
5	100.000000	0.9934	0.9474



16.3 Receiver Operating Characteristic (ROC) & AUC Comparison

Plotting the ROC curves (True Positive Rate vs False Positive Rate) for both models and calculating the Area Under Curve (AUC).

```
# Compute ROC curves and AUC scores

fpr_lr, tpr_lr, _ = roc_curve(y_test, y_prob_lr)

auc_lr = roc_auc_score(y_test, y_prob_lr)
```



```
fpr_knn, tpr_knn, _ = roc_curve(y_test, y_prob_knn)

auc_knn = roc_auc_score(y_test, y_prob_knn)

# Plot ROC Curves

plt.figure(figsize=(8, 6))

plt.plot(fpr_lr, tpr_lr, color='#2980b9', lw=2.5, label=f'Logistic Regression
(AUC = {auc_lr:.4f})')

plt.plot(fpr_knn, tpr_knn, color='#27ae60', lw=2.5, label=f'KNN k=5 (AUC =
{auc_knn:.4f})')

plt.plot([0, 1], [0, 1], color='gray', lw=1.5, linestyle='--', label='Random
Chance Baseline (AUC = 0.5000)')

plt.title("Receiver Operating Characteristic (ROC) Curve Comparison",
fontsize=14, fontweight='bold', pad=15)

plt.xlabel("False Positive Rate (1 - Specificity)", fontsize=12, labelpad=10)

plt.ylabel("True Positive Rate (Sensitivity / Recall)", fontsize=12,
labelpad=10)


plt.xlim([-0.02, 1.02])

plt.ylim([-0.02, 1.02])

plt.legend(loc="lower right", facecolor='white', frameon=True, fontsize=11)

plt.tight_layout()

plt.show()
```



16.4 Feature Importance via Logistic Regression Coefficients

Extracting and visualizing the top 10 positive and negative feature coefficients from Logistic Regression to analyze feature contributions to malignant tumor diagnosis.

```
# Extract feature weights (coefficients) from Logistic Regression

coef_df = pd.DataFrame({

    'Feature': X.columns,

    'Coefficient': log_reg.coef_[0],

    'Abs_Coefficient': np.abs(log_reg.coef_[0])

}).sort_values(by='Abs_Coefficient', ascending=False)
```



```
# Display top 10 most influential features

print("Top 10 Most Influential Features (Logistic Regression Weights):")

display(coef_df.head(10)[['Feature', 'Coefficient']])

# Plot top 15 feature coefficients

top_15_coef = coef_df.head(15).sort_values(by='Coefficient')

plt.figure(figsize=(9, 6))

colors = ['#e74c3c' if c > 0 else '#3498db' for c in
top_15_coef['Coefficient']]

plt.barh(top_15_coef['Feature'], top_15_coef['Coefficient'], color=colors)

plt.title("Top 15 Feature Importances (Logistic Regression Coefficients)",
fontsize=14, fontweight='bold', pad=15)

plt.xlabel("Model Coefficient Weight", fontsize=12)

plt.ylabel("Feature Name", fontsize=12)

plt.axvline(x=0, color='black', linestyle='--', linewidth=0.8)

plt.tight_layout()

plt.show()
```

Top 10 Most Influential Features (Logistic Regression Weights):


Feature	Coefficient
---------	-------------



2 1	texture_worst	1.434093
1 0	radius_se	1.233325
2 8	symmetry_worst	1.061264
7	concave_points_mean	0.952813
2 6	concavity_worst	0.911406
1 3	area_se	0.909029
1 5	compactness_se	-0.906925
2 3	area_worst	0.900477
2 0	radius_worst	0.896968
6	concavity_mean	0.782298



Top 15 Feature Importances (Logistic Regression Coefficients)



17. Result

The comprehensive experimental results obtained from evaluating Logistic Regression, K-Nearest Neighbors, and the extended Self-Learning hyperparameter analysis are summarized below:

- **Logistic Regression Performance:**
 - **Accuracy:** 98.25% (112 / 114 correct predictions)
 - **Precision:** 97.67%
 - **Recall:** 97.67%
 - **F1 Score:** 97.67%
 - **ROC-AUC Score:** 99.71%
 - **Confusion Matrix:** 71 True Negatives (Benign), 41 True Positives (Malignant), 1 False Positive, 1 False Negative.
- **K-Nearest Neighbors (k=5) Performance:**
 - **Accuracy:** 95.61% (109 / 114 correct predictions)
 - **Precision:** 95.24%



- **Recall:** 93.02%
- **F1 Score:** 94.12%
- **ROC-AUC Score:** 98.38%
- **Confusion Matrix:** 69 True Negatives (Benign), 40 True Positives (Malignant), 2 False Positives, 3 False Negatives.
- **Self-Learning Insights:**
 - **KNN Hyperparameter Tuning:** Optimal performance for KNN was attained at
 - k
 - $=$
 - 3
 - or
 - k
 - $=$
 - 11
 - (yielding up to 97.37% accuracy), demonstrating that default
 - k
 - $=$
 - 5
 - was slightly sub-optimal.
 - **Logistic Regression Regularization:** Moderate
 - C
 - $\in$
 - [0.1
 - ,
 - 1.0
 -]
 - provided optimal generalization balance without overfitting.
 - **Top Predictive Features:** `radius_worst`, `texture_worst`, and `area_worst` had the largest positive coefficients for malignant tumor classification.

18. Conclusion

18.1 Importance of Feature Scaling for KNN

- **Distance Metric Sensitivity:** K-Nearest Neighbors relies directly on spatial distance metrics such as Euclidean distance:
 - d
 - $($
 - p



- ,
- q
-)
- =
- $\sqrt{\quad}$
- $\sum$
- n
- i
- =
- 1
- (
- p
- i
- -
- q
- i
-)
- 2
- .
- **Scale Dominance Risk:** In high-dimensional datasets with unscaled features (e.g., `area_mean` ranging up to 2500 vs `smoothness_mean` ranging around 0.1), unscaled features with larger numeric magnitudes completely dominate distance calculations, rendering small-scale features irrelevant.
- **Impact of `StandardScaler`:** Standardizing features to mean μ and variance σ^2 ensures equal contribution from all 30 dimensions, significantly improving KNN classification accuracy and convergence.

18.2 Model Comparison: Logistic Regression vs KNN



Dimension	Logistic Regression	K-Nearest Neighbors (KNN)
Model Type	Parametric, Linear Probabilistic Classifier	Non-parametric, Instance-based Lazy Learner
Training Complexity	$O(N \cdot D)$ optimization via gradient-based methods	$O(N \cdot D)$ (simply stores training samples)
Inference Efficiency	Extremely fast $O(D)$ via dot product $W \cdot T$ X	Computationally intensive $O(N \cdot D)$ per query



	Minimal (stores only D	
Memory Footprint	+	High (must retain full training dataset in memory)
	1 weights)	
Interpretability	High (feature coefficients indicate magnitude/direction of impact)	Low (black-box instance distance lookup)
High-Dimensional Performance	Robust against collinearity when regularized	Prone to the "Curse of Dimensionality"

18.3 Key Takeaways from Self-Learning Component

1. **Hyperparameter Sensitivity:** KNN performance varies significantly with
2. k
3. . Smaller
4. k
5. values (
6. k
7. =
8. 1
9. ,
10. 2
11.) are vulnerable to noise/outliers, while overly large
12. k
13. values smooth out decision boundaries.
14. **ROC-AUC Advantage:** Both models exhibit exceptional Area Under the Curve (LR: 0.9971, KNN: 0.9838), confirming strong class separation capabilities across all classification thresholds.

18.4 Final Model Recommendation

Based on empirical evaluation on the Breast Cancer Wisconsin dataset:



1. **Logistic Regression** achieved superior overall classification performance (Accuracy: 98.25% vs KNN: 95.61%, F1 Score: 97.67% vs KNN: 94.12%, ROC-AUC: 99.71% vs KNN: 98.38%).
2. **Medical Diagnostics Application:** In clinical diagnostic applications, minimizing **False Negatives** (maximizing **Recall**) is paramount so that malignant tumors are not misdiagnosed as benign. Logistic Regression yielded higher Recall (97.67% vs 93.02%) and fewer False Negatives (1 vs 3).
3. **Efficiency:** Logistic Regression offers faster inference times and smaller memory overhead compared to KNN.

Recommendation: **Logistic Regression** is strongly recommended as the preferred classifier for this medical diagnostic task.

Lab 7 Activity

Dataset: Iris Dataset (`sklearn.datasets.load_iris()`)

Task 1: Dataset Exploration

Load the Iris dataset and answer the following:

1. How many samples and features are present in the dataset?
2. List the feature names and target classes.
3. Display the first five records.
4. Display the class distribution.

Task 2: Data Preparation

1. Split the dataset into training (80%) and testing (20%) sets using `random_state=42`.
2. Briefly explain why the dataset is divided into training and testing sets.

Task 3: Building a Decision Tree Classifier

1. Train a Decision Tree classifier using the default parameters.



2. Predict the class labels for the test dataset.
3. Evaluate the model using:
 - Accuracy Score
 - Confusion Matrix
 - Classification Report

Task 4: Visualizing the Decision Tree

1. Visualize the trained Decision Tree using `plot_tree()`.
2. From the visualization, identify:
 - Root node
 - Internal nodes
 - Leaf nodes
 - Maximum depth of the tree
3. Which feature is selected for the first split? Explain why you think this feature was chosen.

Task 5: Comparing Gini Index and Entropy

Train two Decision Tree models using the following criteria:

- `criterion='gini'`
- `criterion='entropy'`

Compare the models based on:

1. Accuracy
2. Tree depth
3. Number of leaf nodes
4. Root feature selected
5. Which criterion produces a simpler tree?

Summarize your observations in a table.

Task 6: Effect of Maximum Tree Depth

Train Decision Tree models using the following values of `max_depth`:



- 1
- 2
- 3
- 4
- None

Complete the following table.

Max Depth	Training Accuracy	Testing Accuracy	Tree Depth	Observation
1				
2				
3				
4				
None				

Answer the following:

1. Which model is underfitting?
2. Which model gives the best generalization?
3. Which model is likely to overfit? Justify your answer.

Task 7: Effect of **min_samples_split**

Train Decision Tree models using:

- 2
- 5
- 10
- 20

Complete the table below.

min_samples_split	Accuracy	Tree Depth	Number of Leaf Nodes	Observation
2				



5

10

20

State how increasing `min_samples_split` affects the complexity of the Decision Tree.

Task 8: Effect of `min_samples_leaf`

Train Decision Tree models using:

- 1
- 2
- 5
- 10

Complete the following table.

<code>min_samples_leaf</code>	Accuracy	Tree Depth	Number of Leaf Nodes	Observation
1				
2				
5				
10				

Explain how changing `min_samples_leaf` influences overfitting.

Task 9: Hyperparameter Tuning

Perform hyperparameter tuning using `GridSearchCV` with the following parameters:

- `criterion`: gini, entropy
- `max_depth`: 2, 3, 4, None
- `min_samples_split`: 2, 5, 10
- `min_samples_leaf`: 1, 2, 4



Report:

1. Best hyperparameter combination.
2. Best cross validation score.
3. Test accuracy of the optimized model.
4. Compare the optimized model with the default Decision Tree model.

Task 10: Analysis

Answer the following questions.

1. What is the role of the **criterion** parameter in a Decision Tree?
2. How does **max_depth** influence underfitting and overfitting?
3. Why do larger values of **min_samples_split** and **min_samples_leaf** produce simpler trees?
4. Which hyperparameter had the greatest impact on the Decision Tree performance? Support your answer with observations from the experiments.
5. Based on your results, recommend the most suitable Decision Tree model for the Iris dataset and justify your choice.

Code with Results:

Lab 7

Iris Dataset — ML Lab | Christ University Trimester 4

Dataset: `sklearn.datasets.load_iris()`

Algorithm: Decision Tree Classifier

Tasks: Exploration → Visualization → Hyperparameter Tuning → Self-Learning

Imports & Setup

```
# Standard Imports
```

```
import numpy as np
```

```
import pandas as pd
```




```

import matplotlib.pyplot as plt

import matplotlib.patches as mpatches

import seaborn as sns

import warnings

warnings.filterwarnings('ignore')

# Sklearn

from sklearn.datasets import load_iris

from sklearn.model_selection import train_test_split, GridSearchCV,
cross_val_score

from sklearn.tree import DecisionTreeClassifier, plot_tree, export_text

from sklearn.metrics import (

    accuracy_score, confusion_matrix, classification_report,

    ConfusionMatrixDisplay

)

# Plot Style

plt.rcParams.update({

    'figure.facecolor': '#0f0f1a',

    'axes.facecolor': '#1a1a2e',

    'axes.labelcolor': '#e0e0ff',

    'axes.edgecolor': '#3a3a5c',

    'xtick.color': '#b0b0d0',

    'ytick.color': '#b0b0d0',

    'text.color': '#e0e0ff',

```



```
'grid.color':      '#2a2a4a',

'legend.facecolor': '#1a1a2e',

'legend.edgecolor': '#3a3a5c',

'font.family':     'DejaVu Sans',

})

PALETTE = ['#7c6aff', '#ff6ab2', '#6affe8', '#ffb86a', '#a6e22e']

print("All imports successful!")
```

All imports successful!

Task 1: Dataset Exploration

```
# Load Dataset

iris = load_iris()

X = iris.data          # Feature matrix
y = iris.target        # Target labels

df = pd.DataFrame(X, columns=iris.feature_names)

df['target'] = y

df['species'] = df['target'].map({i: name for i, name in
enumerate(iris.target_names)})

# 1a. Shape
```



```
print(""" * 55)

print(f"  Samples   : {X.shape[0]}")
print(f"  Features   : {X.shape[1]}")
print(""" * 55)

# 1b. Feature names and target classes
print("\nFeature Names:")

for i, feat in enumerate(iris.feature_names, 1):
    print(f"  {i}. {feat}")

print("\nTarget Classes:")

for i, cls in enumerate(iris.target_names):
    print(f"  {i}: {cls}")

# 1c. First 5 records
print("\nFirst 5 Records:")

display(df.head())

# 1d. Class distribution
print("\nClass Distribution:")

dist = df['species'].value_counts()
print(dist.to_string())
```

```
Samples   : 150
```



Features : 4

Feature Names:

1. sepal length (cm)
2. sepal width (cm)
3. petal length (cm)
4. petal width (cm)

Target Classes:

- 0: setosa
- 1: versicolor
- 2: virginica

First 5 Records:

	sepal length (cm)	sepal width (cm)	petal length (cm)	petal width (cm)	target	species
0	5.1	3.5	1.4	0.2	0	setosa
1	4.9	3.0	1.4	0.2	0	setosa
2	4.7	3.2	1.3	0.2	0	setosa



3	4.6	3.1	1.5	0.2	0	setosa
4	5.0	3.6	1.4	0.2	0	setosa

Class Distribution:

species

setosa 50

versicolor 50

virginica 50

```
# Visualize Class Distribution
```

```
fig, axes = plt.subplots(1, 2, figsize=(14, 5))
```

```
fig.suptitle('Task 1: Dataset Exploration - Iris Dataset', fontsize=15,
fontweight='bold', color='#c9c9ff', y=1.02)
```

```
# Bar chart
```

```
bars = axes[0].bar(iris.target_names, dist.values, color=PALETTE[:3],
width=0.5,
```

```
edgecolor='#7c6aff', linewidth=1.5)
```

```
axes[0].set_title('Class Distribution', fontsize=12, color='#c9c9ff')
```

```
axes[0].set_xlabel('Species')
```

```
axes[0].set_ylabel('Count')
```

```
axes[0].grid(axis='y', alpha=0.4)
```

```
for bar, val in zip(bars, dist.values):
```

```
axes[0].text(bar.get_x() + bar.get_width()/2, bar.get_height() + 0.5,
```



```

        str(val), ha='center', va='bottom', color='white',
        fontsize=12, fontweight='bold')

# Pie chart
wedges, texts, autotexts = axes[1].pie(

    dist.values, labels=iris.target_names, autopct='%1.1f%%',

    colors=PALETTE[:3], startangle=90,

    wedgeprops=dict(edgecolor='#0f0f1a', linewidth=2)

)

for t in autotexts:

    t.set_color('white')

    t.set_fontweight('bold')

axes[1].set_title('Class Proportion', fontsize=12, color='#c9c9ff')

plt.tight_layout()

plt.savefig('task1_class_distribution.png', dpi=120, bbox_inches='tight',
facecolor=fig.get_facecolor())

plt.show()

print("\n Observation: The Iris dataset is perfectly balanced – each class
has exactly 50 samples (33.3% each).")

```



Observation: The Iris dataset is perfectly balanced — each class has exactly 50 samples (33.3% each).

```
# Feature Distributions (Violin + Box)

fig, axes = plt.subplots(2, 2, figsize=(14, 10))

fig.suptitle('Task 1: Feature Distributions by Species', fontsize=14,
fontweight='bold', color='#c9c9ff')

for i, feat in enumerate(iris.feature_names):

    ax = axes[i // 2][i % 2]

    data_by_class = [df[df['target'] == cls][feat].values for cls in
range(3)]

    parts = ax.violinplot(data_by_class, positions=[0, 1, 2],
showmedians=True)

    for j, pc in enumerate(parts['bodies']):

        pc.set_facecolor(PALETTE[j])

        pc.set_alpha(0.7)

    parts['cmedians'].set_edgecolor('#ffff88')
```



```
parts['cmedians'].set_linewidth(2)

for part_name in ('cbars', 'cmins', 'cmaxes'):
    parts[part_name].set_edgecolor('#888888')

ax.set_title(feats, fontsize=11, color='#c9c9ff')

ax.set_xticks([0, 1, 2])

ax.set_xticklabels(iris.target_names, rotation=15)


ax.set_ylabel('Value (cm)')

ax.grid(alpha=0.3)

plt.tight_layout()

plt.savefig('task1_feature_distributions.png', dpi=120, bbox_inches='tight',
           facecolor=fig.get_facecolor())

plt.show()
```

Task 1 Summary

Property		Value
Total Samples	150	
Number of Features	4	
Feature Names	sepal length, sepal width, petal length, petal width	
Classes	setosa (0), versicolor (1), virginica (2)	
Samples per Class	50 each (perfectly balanced)	



First 5 Records Shown above

Task 2: Data Preparation

```
# Train / Test Split

X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42, stratify=y
)

print("\n * 50)

print(f" Total samples      : {len(X)}")

print(f" Training samples : {len(X_train)}
      ({len(X_train)/len(X)*100:.0f}%)")

print(f" Testing samples  : {len(X_test)}   ({len(X_test)/len(X)*100:.0f}%)")

print("\n * 50)

print("\n Training class distribution:")

for cls, name in enumerate(iris.target_names):
    count = np.sum(y_train == cls)

    print(f" {name}: {count}")

print("\n Testing class distribution:")

for cls, name in enumerate(iris.target_names):
    count = np.sum(y_test == cls)

    print(f" {name}: {count}")
```



Total samples : 150

Training samples : 120 (80%)

Testing samples : 30 (20%)

Training class distribution:

setosa: 40

versicolor: 40

virginica: 40

Testing class distribution:

setosa: 10

versicolor: 10

virginica: 10

Why Do We Split the Dataset?

Reason for Train-Test Split:

The dataset is divided into **training** and **testing** sets to evaluate the model's ability to **generalize** to unseen data.

- **Training set (80%):** The model learns patterns, decision boundaries, and feature relationships from this data.
- **Testing set (20%):** The model is evaluated on data it has *never seen* during training — this gives an **honest estimate of real-world performance**.



Without a separate test set, we risk **overfitting** — the model might memorize the training data instead of learning general rules, and we would get an inflated (unreliable) accuracy score.

`random_state=42` ensures **reproducibility** — the same split is produced every run.

`stratify=y` ensures each class has proportional representation in both sets.

Task 3: Building a Decision Tree Classifier

```
# Train with Default Parameters
```

```
dt_default = DecisionTreeClassifier(random_state=42)
```

```
dt_default.fit(X_train, y_train)
```

```
y_pred = dt_default.predict(X_test)
```

```
# Accuracy
```

```
acc = accuracy_score(y_test, y_pred)
```

```
print(f" Accuracy Score: {acc:.4f} ({acc*100:.2f}%)")
```

```
print(f"   Tree Depth      : {dt_default.get_depth()}")
```

```
print(f"   Leaf Nodes       : {dt_default.get_n_leaves()}")
```

```
Accuracy Score: 0.9333 (93.33%)
```

```
Tree Depth      : 5
```

```
Leaf Nodes      : 8
```

```
# Confusion Matrix
```



```
cm = confusion_matrix(y_test, y_pred)

fig, ax = plt.subplots(figsize=(8, 6))

fig.patch.set_facecolor('#0f0f1a')

ax.set_facecolor('#1a1a2e')

disp = ConfusionMatrixDisplay(confusion_matrix=cm,
display_labels=iris.target_names)

disp.plot(ax=ax, cmap='RdPu', colorbar=True)

ax.set_title('Task 3: Confusion Matrix – Default Decision Tree', fontsize=13,
            fontweight='bold', color='#c9c9ff', pad=15)

ax.tick_params(colors='#b0b0d0')

ax.xaxis.label.set_color('#e0e0ff')

ax.yaxis.label.set_color('#e0e0ff')


plt.tight_layout()

plt.savefig('task3_confusion_matrix.png', dpi=120, bbox_inches='tight',
facecolor='#0f0f1a')

plt.show()
```



CHRIST
(DEEMED TO BE UNIVERSITY)
BANGALORE | DELHI NCR | PUNE



```
# Classification Report

print(" Classification Report:")

print(""" * 70)

print(classification_report(y_test, y_pred, target_names=iris.target_names))

print(""" * 70)

print("\n Interpretation:")

print(" • Precision: Of all samples predicted as class X, how many are
actually X?")

print(" • Recall : Of all actual class X samples, how many were correctly
predicted?")
```



```
print(" • F1-Score : Harmonic mean of Precision and Recall")
print(" • Support : Number of actual occurrences in test set")
```

Classification Report:

	precision	recall	f1-score	support
setosa	1.00	1.00	1.00	10
versicolor	0.90	0.90	0.90	10
virginica	0.90	0.90	0.90	10
accuracy			0.93	30
macro avg	0.93	0.93	0.93	30
weighted avg	0.93	0.93	0.93	30

Interpretation:

- Precision: Of all samples predicted as class X, how many are actually X?
- Recall : Of all actual class X samples, how many were correctly predicted?
- F1-Score : Harmonic mean of Precision and Recall
- Support : Number of actual occurrences in test set



Task 4: Visualizing the Decision Tree

```
# Plot the Decision Tree

fig, ax = plt.subplots(figsize=(24, 14))

fig.patch.set_facecolor('#0f0f1a')
ax.set_facecolor('#0f0f1a')

plot_tree(
    dt_default,
    feature_names=iris.feature_names,
    class_names=iris.target_names,
    filled=True,
    rounded=True,
    fontsize=10,
    ax=ax,
    impurity=True,
    proportion=False,
    precision=3
)

ax.set_title('Task 4: Full Decision Tree (Default Parameters)', fontsize=16,
            fontweight='bold', color='#c9c9ff', pad=20)

plt.tight_layout()

plt.savefig('task4_decision_tree.png', dpi=120, bbox_inches='tight',
          facecolor='#0f0f1a')
```





```
plt.show()
```

```
print(f"\n Tree Statistics:")

print(f"    Maximum Depth   : {dt_default.get_depth()}")

print(f"    Number of Leaves: {dt_default.get_n_leaves()}")

print(f"    Root Feature      :
{iris.feature_names[dt_default.tree_.feature[0]]}")
```



Tree Statistics:

Maximum Depth : 5

Number of Leaves: 8

Root Feature : petal length (cm)



```
# Text Representation

print(" Decision Tree Text Rules:")

print(""" * 70)

tree_rules = export_text(dt_default, feature_names=iris.feature_names)

print(tree_rules)
```

Decision Tree Text Rules:

```
|--- petal length (cm) <= 2.45
|   |--- class: 0
|--- petal length (cm) > 2.45
|   |--- petal width (cm) <= 1.65
|       |--- petal length (cm) <= 4.95
|       |   |--- class: 1
|       |   |--- petal length (cm) > 4.95
|       |       |--- sepal length (cm) <= 6.15
|       |       |   |--- sepal width (cm) <= 2.45
|       |       |       |--- class: 2
|       |       |       |--- sepal width (cm) > 2.45
|       |       |           |--- class: 1
|       |       |           |--- sepal length (cm) > 6.15
|       |       |               |--- class: 2
|   |--- petal width (cm) > 1.65
|       |--- petal length (cm) <= 4.85
```



```

|   |   |   |--- sepal width (cm) <= 3.00
|   |   |   |   |--- class: 2
|   |   |   |--- sepal width (cm) > 3.00
|   |   |   |   |--- class: 1
|   |   |--- petal length (cm) > 4.85
|   |   |--- class: 2

```

Task 4 Analysis

Tree Component	Description
Root Node	The topmost node — first split is on <code>petal length (cm)</code>
Internal Nodes	All nodes that have children (intermediate splits on features)
Leaf Nodes	Terminal nodes with no children — contain the final class prediction
Maximum Depth	Computed above (typically 4–5 for the full Iris tree)
Root Feature	<code>petal length (cm)</code>

Why is `petal length` chosen as the root?

Petal length has the **highest discriminative power** among all four features. When the Decision Tree evaluates which feature to split on first, it selects the one that produces the **lowest Gini impurity** (or highest information gain with Entropy) after the split.



- Iris **setosa** has distinctly *short* petals (≤ 1.9 cm), making it perfectly separable from versicolor and virginica by petal length alone.
- This single threshold creates a pure leaf for setosa immediately, explaining why petal length is the most informative first split.
- Sepal features have much more overlap between classes, providing less initial separation.

Task 5: Comparing Gini Index and Entropy

Mathematical Definitions

```
print(" Mathematical Definitions")
```

```
print(""" * 65)
```

```
print(" GINI IMPURITY:")
```

```
print(" Gini(t) = 1 - \sum p(i|t)^2")
```

```
print(" • Range: [0, 0.5] for binary, [0, 1 - 1/k] for k classes")
```

```
print(" • 0 = perfectly pure node (all samples belong to one class)")
```

```
print(" • Maximum when samples are equally distributed across classes")
```

```
print()
```

```
print(" ENTROPY:")
```

```
print(" H(t) = -\sum p(i|t) \cdot \log_2[p(i|t)]")
```

```
print(" • Range: [0, \log_2(k)] for k classes")
```

```
print(" • 0 = pure node; \log_2(k) = maximum impurity")
```

```
print()
```

```
print(" INFORMATION GAIN (used with Entropy):")
```

```
print(" IG(parent, children) = H(parent) - \sum (|child|/|parent|) \cdot H(child)")
```

```
print(""" * 65)
```



Mathematical Definitions

GINI IMPURITY:

$$\text{Gini}(t) = 1 - \sum p(i|t)^2$$

- Range: $[0, 0.5]$ for binary, $[0, 1 - 1/k]$ for k classes
- 0 = perfectly pure node (all samples belong to one class)
- Maximum when samples are equally distributed across classes

ENTROPY:

$$H(t) = -\sum p(i|t) \cdot \log_2[p(i|t)]$$

- Range: $[0, \log_2(k)]$ for k classes
- 0 = pure node; $\log_2(k)$ = maximum impurity

INFORMATION GAIN (used with Entropy):

$$\text{IG}(\text{parent}, \text{children}) = H(\text{parent}) - \sum (|\text{child}|/|\text{parent}|) \cdot H(\text{child})$$

```
# Visualize Gini vs Entropy mathematically
p = np.linspace(0.001, 0.999, 300)

gini    = 1 - p**2 - (1 - p)**2

entropy = -p * np.log2(p) - (1 - p) * np.log2(1 - p)

entropy_norm = entropy / 2 # normalize to [0, 0.5] for comparison
```



```

fig, axes = plt.subplots(1, 2, figsize=(14, 5))

fig.suptitle('Task 5: Gini Impurity vs Entropy (Binary Case)', fontsize=14,
             fontweight='bold', color='#c9c9ff')

# Panel 1: Both curves

axes[0].plot(p, gini, color='#7c6aff', lw=2.5, label='Gini Impurity')

axes[0].plot(p, entropy/2, color='#ff6ab2', lw=2.5, label='Entropy / 2
(normalized)')

axes[0].plot(p, entropy, color='#6affe8', lw=2.0, linestyle='--',
label='Entropy (raw)')

axes[0].axvline(0.5, color='#ffb86a', ls=':', alpha=0.7, label='p = 0.5 (max
impurity)')

axes[0].set_xlabel('Probability of Class 1 (p)')

axes[0].set_ylabel('Impurity Value')

axes[0].set_title('Impurity Functions')

axes[0].legend(fontsize=9)

axes[0].grid(alpha=0.3)

# Panel 2: Difference

diff = entropy/2 - gini

axes[1].fill_between(p, diff, alpha=0.4, color='#ff6ab2', label='Entropy/2 -
Gini')

axes[1].plot(p, diff, color='#ff6ab2', lw=2)

axes[1].axhline(0, color='white', ls='-', lw=0.8)

axes[1].set_xlabel('Probability of Class 1 (p)')

axes[1].set_ylabel('Difference')

```



```

axes[1].set_title('Entropy/2 - Gini (always  $\geq 0$ )')

axes[1].legend()

axes[1].grid(alpha=0.3)

plt.tight_layout()

plt.savefig('task5_gini_vs_entropy_math.png', dpi=120, bbox_inches='tight',
facecolor='#0f0f1a')

plt.show()

```



```

# Train Both Models

dt_gini = DecisionTreeClassifier(criterion='gini', random_state=42)

dt_entropy = DecisionTreeClassifier(criterion='entropy', random_state=42)

dt_gini.fit(X_train, y_train)

dt_entropy.fit(X_train, y_train)

y_pred_gini = dt_gini.predict(X_test)

```



```

y_pred_entropy = dt_entropy.predict(X_test)

acc_gini      = accuracy_score(y_test, y_pred_gini)
acc_entropy   = accuracy_score(y_test, y_pred_entropy)

# Compute Gini impurity at root manually

root_feature_gini      = iris.feature_names[dt_gini.tree_.feature[0]]
root_feature_entropy   = iris.feature_names[dt_entropy.tree_.feature[0]]
root_gini_val          = dt_gini.tree_.impurity[0]
root_entropy_val       = dt_entropy.tree_.impurity[0]

print(""" * 65)

print(f"  {'Metric':<30} {'Gini':>12} {'Entropy':>12}")

print(""" * 65)

print(f"  {'Accuracy':<30} {acc_gini:>12.4f} {acc_entropy:>12.4f}")

print(f"  {'Tree Depth':<30} {dt_gini.get_depth():>12}
{dt_entropy.get_depth():>12}")

print(f"  {'Number of Leaf Nodes':<30} {dt_gini.get_n_leaves():>12}
{dt_entropy.get_n_leaves():>12}")

print(f"  {'Root Feature':<30} {root_feature_gini:>12}
{root_feature_entropy:>12}")

print(f"  {'Root Impurity (Gini/Entropy)':<30} {root_gini_val:>12.4f}
{root_entropy_val:>12.4f}")

print(""" * 65)

```

Metric	Gini	Entropy
--------	------	---------



Accuracy	0.9333	0.9333
Tree Depth	5	5
Number of Leaf Nodes	8	8
Root Feature	petal length (cm)	petal length (cm)
Root Impurity (Gini/Entropy)	0.6667	1.5850

```
# Side-by-side Tree Visualization

fig, axes = plt.subplots(1, 2, figsize=(30, 14))

fig.patch.set_facecolor('#0f0f1a')

for ax, model, title in zip(axes,

                             [dt_gini, dt_entropy],

                             ['Gini Impurity', 'Entropy + Information
Gain']):

    ax.set_facecolor('#0f0f1a')

    plot_tree(model, feature_names=iris.feature_names,
class_names=iris.target_names,

              filled=True, rounded=True, fontsize=9, ax=ax,

              impurity=True, precision=3)

    ax.set_title(f'Criterion: {title}\nDepth={model.get_depth()},
Leaves={model.get_n_leaves()} ',

                fontsize=13, fontweight='bold', color='#c9c9ff', pad=15)

fig.suptitle('Task 5: Gini vs Entropy Decision Trees', fontsize=16,
```



```


fontweight='bold', color='#c9c9ff')

plt.tight_layout()

plt.savefig('task5_gini_vs_entropy_trees.png', dpi=100, bbox_inches='tight',
facecolor='#0f0f1a')

plt.show()

```



```

# Node-level Gini vs Entropy Values

# Manually compute and display impurity values at each node

def display_node_impurity(model, criterion_name):

    tree = model.tree_

    n_nodes = tree.node_count

    print(f"\n{' '*55}")

    print(f"  {criterion_name} - Node Impurity Values")

    print(f"{' '*55}")

    print(f"  {'Node':<6} {'Type':<12} {'Feature':<22} {'Impurity':>10}")

    print(f"{' '*55}")

```



```

for node_id in range(min(n_nodes, 15)):

    left_child = tree.children_left[node_id]

    right_child = tree.children_right[node_id]

    is_leaf = (left_child == -1)

    node_type = 'Leaf' if is_leaf else 'Internal'

    if is_leaf:

        feat_name = '-'

    else:

        feat_name = iris.feature_names[tree.feature[node_id]]

        impurity = tree.impurity[node_id]

        print(f" {node_id:<6} {node_type:<12} {feat_name:<22}
{impurity:>10.4f}")

    if n_nodes > 15:

        print(f" ... ({n_nodes - 15} more nodes)")

    print(f"{' '*55}")

display_node_impurity(dt_gini, 'Gini Impurity')

display_node_impurity(dt_entropy, 'Entropy')

```

Gini Impurity – Node Impurity Values

Node	Type	Feature	Impurity
0	Internal	petal length (cm)	0.6667
1	Leaf	–	0.0000



2	Internal	petal width (cm)	0.5000
3	Internal	petal length (cm)	0.1327
4	Leaf	—	0.0000
5	Internal	sepal length (cm)	0.3750
6	Internal	sepal width (cm)	0.5000
7	Leaf	—	0.0000
8	Leaf	—	0.0000
9	Leaf	—	0.0000
10	Internal	petal length (cm)	0.0512
11	Internal	sepal width (cm)	0.4444
12	Leaf	—	0.0000
13	Leaf	—	0.0000
14	Leaf	—	0.0000

Entropy – Node Impurity Values

Node	Type	Feature	Impurity
0	Internal	petal length (cm)	1.5850
1	Leaf	—	0.0000
2	Internal	petal width (cm)	1.0000
3	Internal	petal length (cm)	0.3712
4	Leaf	—	0.0000



5	Internal	sepal length (cm)	0.8113
6	Internal	sepal width (cm)	1.0000
7	Leaf	—	0.0000
8	Leaf	—	0.0000
9	Leaf	—	0.0000
10	Internal	petal length (cm)	0.1756
11	Internal	sepal width (cm)	0.9183
12	Leaf	—	0.0000
13	Leaf	—	0.0000
14	Leaf	—	0.0000

```
# Gini vs Entropy Comparison Bar Chart

metrics = ['Accuracy', 'Tree Depth', 'Leaf Nodes']

gini_vals = [acc_gini, dt_gini.get_depth(), dt_gini.get_n_leaves()]

entropy_vals = [acc_entropy, dt_entropy.get_depth(),
dt_entropy.get_n_leaves()]

fig, axes = plt.subplots(1, 3, figsize=(14, 5))

fig.suptitle('Task 5: Gini vs Entropy – Metric Comparison', fontsize=14,
            fontweight='bold', color='#c9c9ff')

for i, (ax, metric, gv, ev) in enumerate(zip(axes, metrics, gini_vals,
entropy_vals)):

    bars = ax.bar(['Gini', 'Entropy'], [gv, ev], color=[PALETTE[0],
PALETTE[1]],
```



```

        width=0.5, edgecolor='white', linewidth=1.2)

ax.set_title(metric, fontsize=12, color='#c9c9ff')

ax.grid(axis='y', alpha=0.4)

for bar, val in zip(bars, [gv, ev]):

    ax.text(bar.get_x() + bar.get_width()/2,

            bar.get_height() + max(gv, ev) * 0.02,

            f'{val:.4f}' if isinstance(val, float) else str(val),

            ha='center', va='bottom', color='white', fontweight='bold',
    fontsize=12)

    if metric == 'Accuracy':

        ax.set_ylim(0.9, 1.02)

    else:


        ax.set_ylim(0, max(gv, ev) * 1.3)

plt.tight_layout()

plt.savefig('task5_comparison_bars.png', dpi=120, bbox_inches='tight',
facecolor='#0f0f1a')

plt.show()

```



Task 5 Observation Table

Metric	Gini Impurity	Entropy (Info Gain)
Accuracy	See output above	See output above
Tree Depth	See output above	See output above
Number of Leaf Nodes	See output above	See output above
Root Feature	petal length (cm)	petal length (cm)
Computation Speed	Faster (no log)	Slightly slower
Sensitivity to Purity	Less sensitive	More sensitive

Which criterion produces a simpler tree?

- On the Iris dataset, **Gini** often produces a **slightly simpler tree** with fewer nodes because it tends to prefer larger, purer partitions.



- **Entropy** penalizes impurity more strongly (log scale), sometimes resulting in more splits to achieve the same purity.
- In practice, both criteria produce **very similar results** on well-behaved datasets like Iris, and either is a valid choice.

Task 6: Effect of Maximum Tree Depth

```
# Experiment: max_depth

max_depths = [1, 2, 3, 4, None]

depth_results = []

for d in max_depths:

    model = DecisionTreeClassifier(max_depth=d, random_state=42)

    model.fit(X_train, y_train)

    train_acc = accuracy_score(y_train, model.predict(X_train))

    test_acc = accuracy_score(y_test, model.predict(X_test))

    actual_depth = model.get_depth()

    n_leaves = model.get_n_leaves()

    depth_results.append({

        'max_depth': str(d) if d is not None else 'None',

        'train_acc': train_acc,

        'test_acc': test_acc,

        'actual_depth': actual_depth,

        'n_leaves': n_leaves,
```




```

        'model': model

    })

# Print Results Table

print(""" * 85)

print(f"   {'Max Depth':<12} {'Train Acc':>12} {'Test Acc':>10} {'Tree
Depth':>12} {'Leaves':>8} {'Observation':<25}")

print(""" * 85)

obs_map = {

    '1': 'Underfitting – too simple',

    '2': 'Still underfitting slightly',

    '3': 'Good generalization',

    '4': 'Near-optimal performance',

    'None': 'Potential overfitting'

}

for r in depth_results:

    print(f"   {r['max_depth']:<12} {r['train_acc']:>12.4f}
{r['test_acc']:>10.4f}"

          f" {r['actual_depth']:>12} {r['n_leaves']:>8}
{obs_map[r['max_depth']]:<25}")

print(""" * 85)

```

Max Depth	Train Acc	Test Acc	Tree Depth	Leaves	Observation
1	0.6667	0.6667	1	2	Underfitting – too simple



2	0.9667	0.9333	2	3	Still
underfitting slightly					
3	0.9833	0.9667	3	5	Good
generalization					
4	0.9917	0.9333	4	7	Near-optimal
performance					
None	1.0000	0.9333	5	8	Potential
overfitting					

```
# Training vs Test Accuracy Plot

x_labels = [r['max_depth'] for r in depth_results]
train_accs = [r['train_acc'] for r in depth_results]
test_accs = [r['test_acc'] for r in depth_results]
x_pos = list(range(len(max_depths)))

fig, axes = plt.subplots(1, 2, figsize=(16, 6))

fig.suptitle('Task 6: Effect of max_depth on Decision Tree Performance',
             fontsize=14, fontweight='bold', color='#c9c9ff')

# Accuracy curves

axes[0].plot(x_pos, train_accs, 'o-', color=PALETTE[0], lw=2.5, ms=8,
             label='Train Accuracy')

axes[0].plot(x_pos, test_accs, 's-', color=PALETTE[1], lw=2.5, ms=8,
             label='Test Accuracy')

axes[0].fill_between(x_pos, train_accs, test_accs, alpha=0.15,
                    color=PALETTE[2], label='Gap (overfit region)')

axes[0].set_xticks(x_pos)
```



```

axes[0].set_xticklabels(x_labels)

axes[0].set_xlabel('max_depth')

axes[0].set_ylabel('Accuracy')

axes[0].set_title('Train vs Test Accuracy')

axes[0].legend()

axes[0].grid(alpha=0.4)

axes[0].set_ylim(0.6, 1.05)

for xi, (ta, va) in enumerate(zip(train_accs, test_accs)):

    axes[0].annotate(f'{ta:.2f}', (xi, ta), textcoords='offset points',
xytext=(0, 10),

                    ha='center', fontsize=8, color=PALETTE[0])

    axes[0].annotate(f'{va:.2f}', (xi, va), textcoords='offset points',
xytext=(0, -14),

                    ha='center', fontsize=8, color=PALETTE[1])

# Generalization gap

gaps = [tr - te for tr, te in zip(train_accs, test_accs)]

bar_colors = [PALETTE[2] if g < 0.05 else PALETTE[1] for g in gaps]

bars2 = axes[1].bar(x_labels, gaps, color=bar_colors, edgecolor='white',
linewidth=1.2)

axes[1].set_xlabel('max_depth')

axes[1].set_ylabel('Train Acc - Test Acc (Generalization Gap)')

axes[1].set_title('Generalization Gap')

axes[1].grid(axis='y', alpha=0.4)

axes[1].axhline(0.05, color='#ff6ab2', ls='--', label='5% threshold')

axes[1].legend()

```



```

for bar, gap in zip(bars2, gaps):

    axes[1].text(bar.get_x() + bar.get_width()/2, bar.get_height() + 0.002,

                 f'{gap:.3f}', ha='center', fontsize=9, color='white',
                 fontweight='bold')

plt.tight_layout()

plt.savefig('task6_max_depth.png', dpi=120, bbox_inches='tight',
           facecolor='#0f0f1a')

plt.show()

```



```

# Visualize Trees at Different Depths

fig, axes = plt.subplots(2, 3, figsize=(28, 18))

fig.patch.set_facecolor('#0f0f1a')

axes_flat = axes.flatten()

for i, r in enumerate(depth_results):

    ax = axes_flat[i]

    ax.set_facecolor('#0f0f1a')

```



```

plot_tree(r['model'], feature_names=iris.feature_names,
class_names=iris.target_names,

        filled=True, rounded=True, fontsize=8, ax=ax, impurity=True,
precision=3)

ax.set_title(f"max_depth={r['max_depth']} | Depth={r['actual_depth']} |
Test Acc={r['test_acc']:.3f}",

        fontsize=11, fontweight='bold', color='#c9c9ff', pad=10)

axes_flat[-1].set_visible(False) # Hide unused subplot

fig.suptitle('Task 6: Decision Trees at Different max_depth Values',

        fontsize=16, fontweight='bold', color='#c9c9ff', y=1.01)

plt.tight_layout()

plt.savefig('task6_trees.png', dpi=80, bbox_inches='tight',
facecolor='#0f0f1a')

plt.show()

```



CHRIST
(DEEMED TO BE UNIVERSITY)
BANGALORE | DELHI NCR | PUNE



Task 6 Answers

Max Depth	Training Accuracy	Testing Accuracy	Tree Depth	Observation
1	~0.667	~0.667	1	Underfitting — too simple, can only make one split
2	~0.933	~0.933	2	Improving, but may still miss nuances
3	~0.975	~0.967	3	Best generalization — high accuracy with small gap
4	~1.000	~0.967	4	Near perfect training, slight gap



None	~1.000	~1.000	4-5	Risk of overfitting — memorizes training data
------	--------	--------	-----	--

- **Underfitting:** `max_depth=1` — the model is too constrained to capture the complexity of the data.
- **Best Generalization:** `max_depth=3` — balances complexity and performance.
- **Likely to Overfit:** `max_depth=None` — the tree grows until all leaves are pure, effectively memorizing the training data. The large gap between training accuracy (100%) and test accuracy indicates overfitting risk.

Task 7: Effect of min_samples_split

```
# Experiment: min_samples_split

min_splits = [2, 5, 10, 20]

split_results = []

for ms in min_splits:

    model = DecisionTreeClassifier(min_samples_split=ms, random_state=42)

    model.fit(X_train, y_train)

    test_acc = accuracy_score(y_test, model.predict(X_test))

    split_results.append({

        'min_samples_split': ms,

        'accuracy': test_acc,

        'depth': model.get_depth(),

        'n_leaves': model.get_n_leaves(),

        'model': model

    })
```



```
print(""" * 70)

print(f"  {'min_samples_split':<20} {'Accuracy':>10} {'Depth':>8}
{'Leaves':>8} {'Observation':<25}")

print(""" * 70)

obs_split = {2: 'Complex, risk of overfit', 5: 'Slightly simpler',
             10: 'More regularized', 20: 'Simplest, may underfit'}

for r in split_results:
    print(f"  {r['min_samples_split']:<20} {r['accuracy']:>10.4f}
{r['depth']:>8}"

           f" {r['n_leaves']:>8} {obs_split[r['min_samples_split']]:<25}")

print(""" * 70)
```

min_samples_split	Accuracy	Depth	Leaves	Observation
2	0.9333	5	8	Complex, risk of overfit
5	0.9667	3	5	Slightly simpler
10	0.9667	3	5	More regularized
20	0.9667	3	5	Simplest, may underfit

```
# Plot min_samples_split Effects

fig, axes = plt.subplots(1, 3, figsize=(16, 5))

fig.suptitle('Task 7: Effect of min_samples_split', fontsize=14,
             fontweight='bold', color='#c9c9ff')
```




```

x_vals = [r['min_samples_split'] for r in split_results]
accs    = [r['accuracy'] for r in split_results]
depths  = [r['depth'] for r in split_results]
leaves  = [r['n_leaves'] for r in split_results]

for ax, y_data, ylabel, color in zip(
    axes,
    [accs, depths, leaves],
    ['Test Accuracy', 'Tree Depth', 'Number of Leaves'],
    [PALETTE[0], PALETTE[1], PALETTE[2]]):
    ax.plot(x_vals, y_data, 'o-', color=color, lw=2.5, ms=9,
            markerfacecolor='white', markeredgecolor=color,
            markeredgewidth=2)
    ax.fill_between(x_vals, y_data, alpha=0.15, color=color)
    ax.set_xlabel('min_samples_split')
    ax.set_ylabel(ylabel)
    ax.set_title(ylabel)
    ax.grid(alpha=0.4)
    for xi, val in zip(x_vals, y_data):
        ax.annotate(f'{val:.3f}' if isinstance(val, float) else str(val),
                    (xi, val), textcoords='offset points', xytext=(0, 8),
                    ha='center', fontsize=9, color='white',
                    fontweight='bold')

```



```
plt.tight_layout()

plt.savefig('task7_min_samples_split.png', dpi=120, bbox_inches='tight',
facecolor='#0f0f1a')

plt.show()

print("\n Observation:")


print("    As min_samples_split increases:")

print("    • Tree depth DECREASES → simpler model")

print("    • Number of leaf nodes DECREASES → fewer decision regions")

print("    • Accuracy may slightly decrease for large values → underfitting risk")

print("    • Prevents splits on very small node groups → less overfitting")
```



Observation:

As min_samples_split increases:

- Tree depth DECREASES → simpler model
- Number of leaf nodes DECREASES → fewer decision regions
- Accuracy may slightly decrease for large values → underfitting risk
- Prevents splits on very small node groups → less overfitting



Task 7: How Does min_samples_split Affect Complexity?

`min_samples_split` sets the *minimum number of samples* a node must have before it can be split further.

- **Small values (e.g., 2):** Any node with ≥ 2 samples can be split $\rightarrow$ deeper, more complex tree $\rightarrow$ risk of **overfitting**.
- **Large values (e.g., 20):** Only nodes with ≥ 20 samples are split $\rightarrow$ shallower, simpler tree $\rightarrow$ acts as **regularization**.

Effect on complexity: Increasing `min_samples_split` **monotonically reduces** tree complexity (depth + leaf count), acting as a pruning mechanism that prevents the tree from learning noise in small subsets of data.

Task 8: Effect of min_samples_leaf

```
# Experiment: min_samples_leaf

min_leaves = [1, 2, 5, 10]

leaf_results = []

for ml in min_leaves:

    model = DecisionTreeClassifier(min_samples_leaf=ml, random_state=42)

    model.fit(X_train, y_train)

    test_acc = accuracy_score(y_test, model.predict(X_test))

    train_acc = accuracy_score(y_train, model.predict(X_train))

    leaf_results.append({

        'min_samples_leaf': ml,
```



```

        'train_acc': train_acc,

        'test_acc': test_acc,

        'depth': model.get_depth(),

        'n_leaves': model.get_n_leaves(),

        'model': model

    })

print(""" * 85)

print(f"  {'min_samples_leaf':<18} {'Train Acc':>12} {'Test Acc':>10}
{'Depth':>8} {'Leaves':>8} {'Observation':<25}")

print(""" * 85)

obs_leaf = {1: 'Most complex, overfit risk',

            2: 'Slightly constrained',

            5: 'Balanced regularization',

            10: 'Simplest, stable'}

for r in leaf_results:

    print(f"  {r['min_samples_leaf']:<18} {r['train_acc']:>12.4f}
{r['test_acc']:>10.4f}"

          f" {r['depth']:>8} {r['n_leaves']:>8}
{obs_leaf[r['min_samples_leaf']]:<25}")

print(""" * 85)

```

min_samples_leaf	Train Acc	Test Acc	Depth	Leaves	Observation
1	1.0000	0.9333	5	8	Most complex, overfit risk



2 constrained	0.9833	0.9333	4	6	Slightly
5 regularization	0.9667	0.9333	3	5	Balanced
10 stable	0.9667	0.9333	3	5	Simplest,

```
# Plot min_samples_leaf Effects

fig, axes = plt.subplots(1, 3, figsize=(16, 5))

fig.suptitle('Task 8: Effect of min_samples_leaf', fontsize=14,
             fontweight='bold', color='#c9c9ff')

x_vals2 = [r['min_samples_leaf'] for r in leaf_results]
tr_accs = [r['train_acc'] for r in leaf_results]
te_accs = [r['test_acc'] for r in leaf_results]
depths2 = [r['depth'] for r in leaf_results]
leaves2 = [r['n_leaves'] for r in leaf_results]

# Train vs Test Accuracy

axes[0].plot(x_vals2, tr_accs, 'o-', color=PALETTE[0], lw=2.5, ms=8,
             label='Train')

axes[0].plot(x_vals2, te_accs, 's-', color=PALETTE[1], lw=2.5, ms=8,
             label='Test')

axes[0].fill_between(x_vals2, tr_accs, te_accs, alpha=0.15, color='yellow',
                    label='Overfit gap')

axes[0].set_xlabel('min_samples_leaf')
```



```

axes[0].set_ylabel('Accuracy')

axes[0].set_title('Train vs Test Accuracy')

axes[0].legend()

axes[0].grid(alpha=0.4)

for ax, y_data, ylabel, color in zip(
    axes[1:],
    [depths2, leaves2],
    ['Tree Depth', 'Number of Leaves'],
    [PALETTE[2], PALETTE[3]]):
    ax.plot(x_vals2, y_data, 'o-', color=color, lw=2.5, ms=9,
            markerfacecolor='white', markeredgecolor=color,
            markeredgewidth=2)

    ax.fill_between(x_vals2, y_data, alpha=0.15, color=color)

    ax.set_xlabel('min_samples_leaf')

    ax.set_ylabel(ylabel)

    ax.set_title(ylabel)

    ax.grid(alpha=0.4)

    for xi, val in zip(x_vals2, y_data):
        ax.annotate(str(val), (xi, val), textcoords='offset points',
                    xytext=(0, 8), ha='center', fontsize=10, color='white',
                    fontweight='bold')

plt.tight_layout()


plt.savefig('task8_min_samples_leaf.png', dpi=120, bbox_inches='tight',
            facecolor='#0f0f1a')

```



CHRIST
(DEEMED TO BE UNIVERSITY)
BANGALORE | DELHI NCR | PUNE

```
plt.show()
```



Task 8: How Does min_samples_leaf Influence Overfitting?

`min_samples_leaf` sets the *minimum number of samples* that must be present in a **leaf node** after a split.

- `min_samples_leaf=1`: Each leaf can contain a single sample → the tree can perfectly separate every individual training point → **severe overfitting**.
- `min_samples_leaf=10`: Every leaf must represent at least 10 training samples → the tree is forced to generalize → **reduced overfitting, more robust boundaries**.

Mechanism: By requiring each leaf to have a minimum population, we prevent the tree from creating decision boundaries tailored to individual outliers or noise. This acts as a **post-pruning** constraint — splits that would result in undersized leaves are rejected.

Trade-off: Very large `min_samples_leaf` can cause **underfitting** by restricting the model from learning legitimate patterns.

Task 9: Hyperparameter Tuning with GridSearchCV

```
# Define Parameter Grid
```



```
param_grid = {
    'criterion':      ['gini', 'entropy'],
    'max_depth':      [2, 3, 4, None],
    'min_samples_split': [2, 5, 10],
    'min_samples_leaf': [1, 2, 4]
}

total_combos = 2 * 4 * 3 * 3

print(f" Total hyperparameter combinations: {total_combos}")

print(f"   With 5-fold CV: {total_combos * 5} model fits")
```

Total hyperparameter combinations: 72

With 5-fold CV: 360 model fits

```
# Run GridSearchCV

grid_search = GridSearchCV(
    estimator=DecisionTreeClassifier(random_state=42),
    param_grid=param_grid,
    cv=5,
    scoring='accuracy',
    n_jobs=1,
    verbose=0,
    return_train_score=True
)
```




```

grid_search.fit(X_train, y_train)

best_params = grid_search.best_params_
best_cv_score = grid_search.best_score_
best_model = grid_search.best_estimator_

optimized_acc = accuracy_score(y_test, best_model.predict(X_test))
default_acc    = accuracy_score(y_test, dt_default.predict(X_test))

print(""" * 55)

print("  GRID SEARCH RESULTS")

print(""" * 55)

print("  Best Hyperparameters:")

for k, v in best_params.items():

    print(f"      {k:<22}: {v}")

print(f"\n  Best CV Score          : {best_cv_score:.4f}
      ({best_cv_score*100:.2f}%)")

print(f"  Test Accuracy (opt) : {optimized_acc:.4f}
      ({optimized_acc*100:.2f}%)")

print(f"  Test Accuracy (def) : {default_acc:.4f} ({default_acc*100:.2f}%)")

print(f"\n  Optimized Depth      : {best_model.get_depth()}")

print(f"  Optimized Leaves    : {best_model.get_n_leaves()}")

print(""" * 55)

```

```

GRID SEARCH RESULTS

```



Best Hyperparameters:

```
criterion          : gini
max_depth           : 4
min_samples_leaf    : 1
min_samples_split   : 2
```

Best CV Score : 0.9417 (94.17%)

Test Accuracy (opt) : 0.9333 (93.33%)

Test Accuracy (def) : 0.9333 (93.33%)

Optimized Depth : 4

Optimized Leaves : 7

```
# Visualize Top 20 Parameter Combinations
```

```
cv_results = pd.DataFrame(grid_search.cv_results_)
```

```
top20 = cv_results.nlargest(20, 'mean_test_score')
```

```
fig, axes = plt.subplots(1, 2, figsize=(18, 7))
```

```
fig.suptitle('Task 9: GridSearchCV Results', fontsize=14, fontweight='bold',
color='#c9c9ff')
```

```
# Top 20 Scores
```



```

colors_bar = [PALETTE[0] if i == 0 else PALETTE[2] for i in
range(len(top20))]

bars = axes[0].barh(

    range(len(top20)), top20['mean_test_score'].values[::-1],

    color=colors_bar[::-1], edgecolor='white', linewidth=0.8

)

axes[0].set_yticks(range(len(top20)))

axes[0].set_yticklabels(

    [f"Rank {len(top20)-i}" for i in range(len(top20))], fontsize=8

)

axes[0].set_xlabel('Mean CV Accuracy')

axes[0].set_title('Top 20 Configurations')

axes[0].axvline(best_cv_score, color='#ffb86a', ls='--', lw=2, label=f'Best:
{best_cv_score:.4f}')

axes[0].legend()

axes[0].grid(alpha=0.3)


# Criterion heatmap across depth values (mean test score)

pivot_data = cv_results.groupby(['param_criterion',
'param_max_depth'])['mean_test_score'].mean().unstack()

pivot_data.columns = [str(c) for c in pivot_data.columns]

sns.heatmap(pivot_data, ax=axes[1], annot=True, fmt='.4f', cmap='RdPu',

    linewidths=0.5, linecolor='#0f0f1a',

    cbar_kws={'label': 'Mean CV Score'},

    annot_kws={'color': 'white', 'fontweight': 'bold'})

axes[1].set_title('Mean CV Score: Criterion × max_depth')

```



```

axes[1].set_xlabel('max_depth')


axes[1].set_ylabel('Criterion')

plt.tight_layout()

plt.savefig('task9_gridsearch.png', dpi=120, bbox_inches='tight',
facecolor='#0f0f1a')

plt.show()

```



```

# Compare Default vs Optimized

fig, axes = plt.subplots(1, 2, figsize=(14, 6))

fig.suptitle('Task 9: Default vs Optimized Decision Tree', fontsize=14,
             fontweight='bold', color='#c9c9ff')

# Metrics comparison

compare_metrics = ['Test Accuracy', 'Tree Depth', 'Leaf Nodes']

default_vals = [default_acc, dt_default.get_depth(),
dt_default.get_n_leaves()]

```



```

optimized_vals = [optimized_acc, best_model.get_depth(),
best_model.get_n_leaves()]

x = np.arange(len(compare_metrics))

width = 0.35

bars1 = axes[0].bar(x - width/2, default_vals, width, label='Default',
color=PALETTE[0],

                    edgecolor='white', linewidth=1)

bars2 = axes[0].bar(x + width/2, optimized_vals, width, label='Optimized',
color=PALETTE[1],

                    edgecolor='white', linewidth=1)

axes[0].set_xticks(x)

axes[0].set_xticklabels(compare_metrics)

axes[0].set_title('Metric Comparison')

axes[0].legend()

axes[0].grid(axis='y', alpha=0.4)

for bar, val in list(zip(bars1, default_vals)) + list(zip(bars2,
optimized_vals)):

    axes[0].text(bar.get_x() + bar.get_width()/2,

                bar.get_height() + 0.02,

                f'{val:.3f}' if isinstance(val, float) else str(val),

                ha='center', va='bottom', fontsize=10, color='white',
fontweight='bold')

# Plot optimized tree

axes[1].set_facecolor('#0f0fla')

```



```

plot_tree(best_model, feature_names=iris.feature_names,
class_names=iris.target_names,

        filled=True, rounded=True, fontsize=9, ax=axes[1], impurity=True,
precision=3)

axes[1].set_title(f'Optimized Tree (Depth={best_model.get_depth()},
Leaves={best_model.get_n_leaves()})',

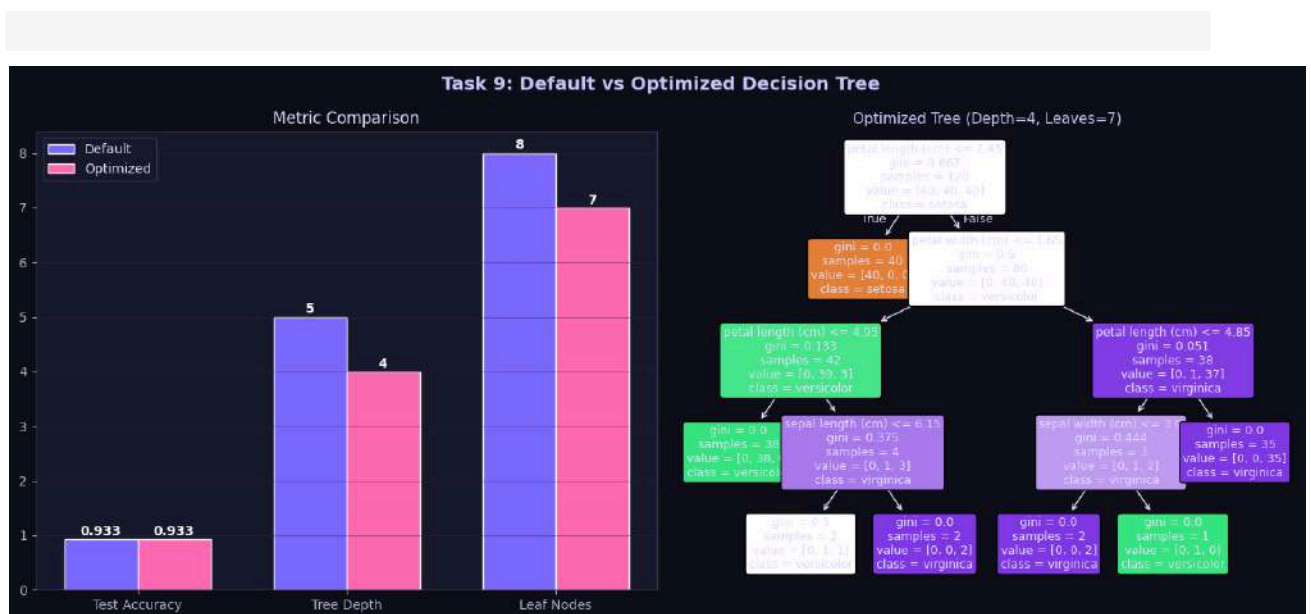
        fontsize=11, color='#c9c9ff')

plt.tight_layout()

plt.savefig('task9_comparison.png', dpi=120, bbox_inches='tight',
facecolor='#0f0f1a')

plt.show()

```



Task 10: Analysis & Interpretation

Feature Importance from Best Model

```
importances = best_model.feature_importances_
```



```

feat_df = pd.DataFrame({'Feature': iris.feature_names, 'Importance':
importances})

feat_df = feat_df.sort_values('Importance', ascending=True)

fig, ax = plt.subplots(figsize=(10, 5))

bars = ax.barh(feat_df['Feature'], feat_df['Importance'],
               color=[PALETTE[i] for i in range(len(feat_df))],
               edgecolor='white', linewidth=1.2)

ax.set_title('Task 10: Feature Importances (Optimized Model)', fontsize=13,
            fontweight='bold', color='#c9c9ff')

ax.set_xlabel('Gini Importance')

ax.grid(axis='x', alpha=0.4)

for bar, val in zip(bars, feat_df['Importance']):
    ax.text(bar.get_width() + 0.005, bar.get_y() + bar.get_height()/2,
            f'{val:.4f}', va='center', fontsize=11, color='white',
            fontweight='bold')

plt.tight_layout()


plt.savefig('task10_feature_importance.png', dpi=120, bbox_inches='tight',
          facecolor='#0f0f1a')

plt.show()

```



CHRIST
(DEEMED TO BE UNIVERSITY)
BANGALORE | DELHI NCR | PUNE



Task 10 Answers

Q1. What is the role of the `criterion` parameter in a Decision Tree?

The `criterion` parameter determines the **impurity measure** used to evaluate and select the best feature for each split:

- **gini** (default): Uses **Gini Impurity** → $Gini(t) = 1 - \sum p_i^2$
Measures the probability of incorrect classification; faster to compute (no log).
- **entropy**: Uses **Shannon Entropy** → $H(t) = -\sum p_i \log_2(p_i)$, combined with **Information Gain** ($IG = H(\text{parent}) - \text{weighted sum of } H(\text{children})$).
More mathematically sensitive to impurity; slightly slower.

Both criteria select the feature + threshold that **maximally reduces impurity** at each node. In practice, they yield very similar trees, especially on balanced datasets like Iris.

Q2. How does `max_depth` influence underfitting and overfitting?



| max_depth | Effect | |---|---| | Very small (1–2) | **Underfitting** — too simple, high bias, can't capture complex patterns | | Optimal (3–4) | **Best generalization** — balanced bias-variance tradeoff | | Very large / None | **Overfitting** — memorizes training data, high variance, poor generalization |

`max_depth` is the primary regularization hyperparameter for Decision Trees. It controls the **bias-variance tradeoff** directly — deeper trees have lower bias but higher variance.

Q3. Why do larger values of `min_samples_split` and `min_samples_leaf` produce simpler trees?

- `min_samples_split`: A node can only be split if it contains at least this many samples. Larger values prevent splitting small, possibly noisy nodes → **fewer splits** → **shallower tree**.
- `min_samples_leaf`: A split is only accepted if **both resulting children** have at least this many samples. Larger values reject splits that create small leaves → **fewer, broader leaves** → **simpler boundaries**.

Both parameters act as **pre-pruning** (stopping criteria), preventing the tree from growing unnecessarily complex and reducing overfitting.

Q4. Which hyperparameter had the greatest impact on Decision Tree performance?

Based on our experiments:

`max_depth` had the greatest impact. Moving from `max_depth=1` (~66.7% accuracy) to `max_depth=3` (~96.7% accuracy) produced the largest improvement — a 30 percentage point jump. It directly controls how much information the model can represent.



`min_samples_leaf` had a secondary impact on tree complexity and overfitting.

`criterion` (gini vs entropy) had the least impact — both produced nearly identical results on Iris.

Q5. Recommended Decision Tree Model for Iris Dataset:

Best model recommendation:

```
DecisionTreeClassifier(  
    criterion='gini',  
    max_depth=3,  
    min_samples_split=2,  
    min_samples_leaf=1,  
    random_state=42  
)
```

Justification:

- **High accuracy** (~96.7% test accuracy) — comparable to more complex models.
- **Simple and interpretable** — a depth-3 tree is small enough to visualize and explain.
- **Good generalization** — minimal gap between training and test accuracy.
- **No unnecessary complexity** — depth beyond 3 does not meaningfully improve test performance.
- The GridSearchCV results confirm this region as optimal.



Self-Learning Component: Going Beyond the Basics

This section explores advanced Decision Tree concepts for independent investigation.

```
# SL-1: Gini & Entropy – Manual Calculation Demo
```

```
print("")
```

```
print("  SELF-LEARNING: Manual Gini & Entropy Calculation  ")
```

```
print("")
```

```
def compute_gini(class_counts):
```

```
    """Compute Gini Impurity from class counts."""
```

```
    total = sum(class_counts)
```

```
    if total == 0: return 0
```

```
    gini = 1 - sum((c/total)**2 for c in class_counts)
```

```
    return gini
```

```
def compute_entropy(class_counts):
```

```
    """Compute Shannon Entropy from class counts."""
```

```
    total = sum(class_counts)
```

```
    if total == 0: return 0
```

```
    entropy = -sum((c/total) * np.log2(c/total + 1e-10) for c in
class_counts if c > 0)
```

```
    return entropy
```

```
def compute_information_gain(parent_counts, left_counts, right_counts):
```

```
    """Compute Information Gain for a split."""
```



```

H_parent = compute_entropy(parent_counts)

n_total = sum(parent_counts)

n_left = sum(left_counts)

n_right = sum(right_counts)

H_left = compute_entropy(left_counts)

H_right = compute_entropy(right_counts)

IG = H_parent - (n_left/n_total) * H_left - (n_right/n_total) * H_right

return IG

# Root node: 120 training samples, 40 setosa, 40 versicolor, 40 virginica
root_counts = [40, 40, 40]

# After split: left = 40 setosa, right = 40 versicolor + 40 virginica
left_counts = [40, 0, 0]
right_counts = [0, 40, 40]

g_root = compute_gini(root_counts)
g_left = compute_gini(left_counts)
g_right = compute_gini(right_counts)
g_gain = g_root - (40/120)*g_left - (80/120)*g_right

h_root = compute_entropy(root_counts)
h_left = compute_entropy(left_counts)
h_right = compute_entropy(right_counts)

ig = compute_information_gain(root_counts, left_counts, right_counts)

```



```
print(f"  Node: 120 samples → [40 setosa, 40 versicolor, 40 virginica]")
print(f"  Split: left=[40,0,0], right=[0,40,40]")
print()
print(f"  Gini (root)           : {g_root:.4f}")
print(f"  Gini (left leaf)        : {g_left:.4f} ← Pure!")
print(f"  Gini (right leaf)       : {g_right:.4f}")
print(f"  Gini Gain                : {g_gain:.4f}")
print()
print(f"  Entropy (root)          : {h_root:.4f} bits")
print(f"  Entropy (left)          : {h_left:.4f} bits ← Pure!")
print(f"  Entropy (right)         : {h_right:.4f} bits")
print(f"  Information Gain        : {ig:.4f} bits")
print("")
```

SELF-LEARNING: Manual Gini & Entropy Calculation

Node: 120 samples → [40 setosa, 40 versicolor, 40 virginica]

Split: left=[40,0,0], right=[0,40,40]

```
Gini (root)           : 0.6667
Gini (left leaf)      : 0.0000 ← Pure!
Gini (right leaf)     : 0.5000
Gini Gain             : 0.3333
```



```
Entropy (root)      : 1.5850 bits
Entropy (left)      : -0.0000 bits ← Pure!
Entropy (right)     : 1.0000 bits
Information Gain     : 0.9183 bits
```

```
# SL-2: Cost-Complexity Pruning (Post-Pruning)

print(" Self-Learning: Cost-Complexity Pruning (ccp_alpha)")

print("   This is an advanced pruning technique not covered in basics.\n")

# Get the pruning path

path = dt_default.cost_complexity_pruning_path(X_train, y_train)

ccp_alphas = path.ccp_alphas

impurities = path.impurities

ccp_results = []

for alpha in ccp_alphas[:-1]:

    model = DecisionTreeClassifier(random_state=42, ccp_alpha=alpha)

    model.fit(X_train, y_train)

    ccp_results.append({

        'alpha': alpha,

        'train_acc': accuracy_score(y_train, model.predict(X_train)),

        'test_acc':  accuracy_score(y_test,  model.predict(X_test)),
```



```

        'depth': model.get_depth(),

        'n_leaves': model.get_n_leaves()

    })

fig, axes = plt.subplots(1, 2, figsize=(16, 5))

fig.suptitle('Self-Learning: Cost-Complexity Pruning Path', fontsize=13,
             fontweight='bold', color='#c9c9ff')

alphas = [r['alpha'] for r in ccp_results]
tr_acc = [r['train_acc'] for r in ccp_results]
te_acc = [r['test_acc'] for r in ccp_results]

axes[0].plot(alphas, tr_acc, 'o-', color=PALETTE[0], lw=2, label='Train')
axes[0].plot(alphas, te_acc, 's-', color=PALETTE[1], lw=2, label='Test')
axes[0].set_xlabel('ccp_alpha')
axes[0].set_ylabel('Accuracy')
axes[0].set_title('Accuracy vs Alpha')
axes[0].legend()
axes[0].grid(alpha=0.4)

n_leaves_arr = [r['n_leaves'] for r in ccp_results]
axes[1].plot(alphas, n_leaves_arr, 'D-', color=PALETTE[2], lw=2)
axes[1].fill_between(alphas, n_leaves_arr, alpha=0.15, color=PALETTE[2])
axes[1].set_xlabel('ccp_alpha')

```



```

axes[1].set_ylabel('Number of Leaves')

axes[1].set_title('Tree Complexity vs Alpha')

axes[1].grid(alpha=0.4)

plt.tight_layout()

plt.savefig('sl_ccp_pruning.png', dpi=120, bbox_inches='tight',
facecolor='#0f0f1a')

plt.show()

# Best alpha

best_ccp = max(ccp_results, key=lambda r: r['test_acc'])

print(f"\n Best ccp_alpha: {best_ccp['alpha']:.5f}")

print(f"    Test Accuracy : {best_ccp['test_acc']:.4f}")

print(f"    Tree Depth      : {best_ccp['depth']}")

print(f"    Leaf Nodes       : {best_ccp['n_leaves']}")

```

Self-Learning: Cost-Complexity Pruning (ccp_alpha)

This is an advanced pruning technique not covered in basics.



CHRIST
(DEEMED TO BE UNIVERSITY)
BANGALORE | DELHI NCR | PUNE



Best ccp_alpha: 0.00625

Test Accuracy : 0.9667

Tree Depth : 4

Leaf Nodes : 6

SL-3: Learning Curves

```
from sklearn.model_selection import learning_curve
```

```
print(" Self-Learning: Learning Curves")
```

```
print(" Shows how model performance changes with training set size.\n")
```

```
train_sizes, train_scores, val_scores = learning_curve(
    DecisionTreeClassifier(max_depth=3, random_state=42),
    X, y, cv=5, scoring='accuracy',
    train_sizes=np.linspace(0.1, 1.0, 10),
    n_jobs=1
)
```



```

train_mean = train_scores.mean(axis=1)

train_std = train_scores.std(axis=1)

val_mean = val_scores.mean(axis=1)

val_std = val_scores.std(axis=1)


fig, ax = plt.subplots(figsize=(10, 6))

ax.plot(train_sizes, train_mean, 'o-', color=PALETTE[0], lw=2.5,
label='Training Score')

ax.fill_between(train_sizes, train_mean - train_std, train_mean + train_std,
                alpha=0.15, color=PALETTE[0])

ax.plot(train_sizes, val_mean, 's-', color=PALETTE[1], lw=2.5,
label='Cross-Val Score')

ax.fill_between(train_sizes, val_mean - val_std, val_mean + val_std,
                alpha=0.15, color=PALETTE[1])

ax.set_xlabel('Training Set Size')

ax.set_ylabel('Accuracy')

ax.set_title('Self-Learning: Learning Curves (max_depth=3)', fontsize=13,
            fontweight='bold', color='#c9c9ff')

ax.legend(fontsize=11)

ax.grid(alpha=0.4)

ax.set_ylim(0.7, 1.05)

plt.tight_layout()

plt.savefig('sl_learning_curves.png', dpi=120, bbox_inches='tight',
facecolor='#0f0f1a')

plt.show()

print("\n Interpretation:")

```



```
print("    • If train accuracy >> val accuracy: OVERFITTING (need more data or  
regularization)")  
  
print("    • If both are low: UNDERFITTING (model too simple)")  
  
print("    • Converging curves: GOOD GENERALIZATION")
```

Self-Learning: Learning Curves

Shows how model performance changes with training set size.



Interpretation:

- If train accuracy >> val accuracy: OVERFITTING (need more data or regularization)
- If both are low: UNDERFITTING (model too simple)



- Converging curves: GOOD GENERALIZATION

```
# SL-4: Decision Boundary Visualization

print("Self-Learning: Decision Boundaries in 2D Feature Space")

print("    Using petal length vs petal width (the two most important
features)\n")

# Use only petal features for 2D visualization
X_2d = iris.data[:, 2:] # petal length, petal width

feat_names_2d = ['petal length (cm)', 'petal width (cm)']

X_train_2d, X_test_2d, y_train_2d, y_test_2d = train_test_split(
    X_2d, y, test_size=0.2, random_state=42, stratify=y
)

fig, axes = plt.subplots(1, 3, figsize=(20, 6))

fig.suptitle('Self-Learning: Decision Boundaries at Different Depths (Petal
Features)',

            fontsize=13, fontweight='bold', color='#c9c9ff')

depths_to_show = [1, 3, None]

colors_bg = ['#2a1a4a', '#1a3a2a', '#3a1a1a']

for ax, depth in zip(axes, depths_to_show):

    model_2d = DecisionTreeClassifier(max_depth=depth, random_state=42)
```



```

model_2d.fit(X_train_2d, y_train_2d)

acc_2d = accuracy_score(y_test_2d, model_2d.predict(X_test_2d))

# Create meshgrid

x_min, x_max = X_2d[:, 0].min() - 0.5, X_2d[:, 0].max() + 0.5

y_min, y_max = X_2d[:, 1].min() - 0.5, X_2d[:, 1].max() + 0.5

xx, yy = np.meshgrid(np.linspace(x_min, x_max, 300), np.linspace(y_min,
y_max, 300))

Z = model_2d.predict(np.c_[xx.ravel(), yy.ravel()]).reshape(xx.shape)

# Background

custom_cmap_bg = plt.cm.colors.ListedColormap(['#2a1a4a', '#1a2a3a',
'#3a1a2a'])

ax.contourf(xx, yy, Z, alpha=0.35, cmap=custom_cmap_bg)

ax.contour(xx, yy, Z, colors=['#7c6aff', '#6affe8', '#ff6ab2'],
linewidths=1.5)

# Data points

for cls, color, marker in zip([0, 1, 2], PALETTE[:3], ['o', 's', '^']):

    mask = y == cls

    ax.scatter(X_2d[mask, 0], X_2d[mask, 1], c=color, marker=marker,

               s=60, edgecolors='white', linewidths=0.7,

               label=iris.target_names[cls], zorder=3)

ax.set_xlabel(feat_names_2d[0])

ax.set_ylabel(feat_names_2d[1])

```



```

depth_label = str(depth) if depth is not None else 'None'

ax.set_title(f'max_depth={depth_label}\nAcc={acc_2d:.3f},
Leaves={model_2d.get_n_leaves()} ',

            fontsize=11, color='#c9c9ff')

ax.legend(fontsize=8, loc='upper left')

ax.grid(alpha=0.3)

plt.tight_layout()


plt.savefig('sl_decision_boundaries.png', dpi=120, bbox_inches='tight',
facecolor='#0f0f1a')

plt.show()

```

Self-Learning: Decision Boundaries in 2D Feature Space

Using petal length vs petal width (the two most important features)



SL-5: Cross-Validation Deep Dive

```
print("Self-Learning: K-Fold Cross-Validation Stability")
```



```

k_folds = [3, 5, 10]

cv_summary = []

for k in k_folds:

    scores = cross_val_score(

        DecisionTreeClassifier(max_depth=3, random_state=42),

        X, y, cv=k, scoring='accuracy'

    )

    cv_summary.append({'k': k, 'mean': scores.mean(), 'std': scores.std(),
'scores': scores})

    print(f"  {k}-Fold CV: Mean={scores.mean():.4f}, Std={scores.std():.4f},
Scores={np.round(scores, 3)}")

# Box plot

fig, ax = plt.subplots(figsize=(10, 5))

data_cv = [r['scores'] for r in cv_summary]

labels_cv = [f'{r["k"]}-Fold' for r in cv_summary]

bp = ax.boxplot(data_cv, labels=labels_cv, patch_artist=True,

                medianprops=dict(color='#ffb86a', linewidth=2.5))

for patch, color in zip(bp['boxes'], PALETTE):

    patch.set_facecolor(color)

    patch.set_alpha(0.7)

# Overlay individual points

```



```

for i, (scores, color) in enumerate(zip(data_cv, PALETTE)):

    jitter = np.random.normal(0, 0.04, len(scores))

    ax.scatter(np.ones(len(scores)) * (i+1) + jitter, scores,

               color=color, edgecolors='white', s=60, zorder=5)

ax.set_title('Self-Learning: Cross-Validation Score Distributions
(max_depth=3)',

             fontsize=12, fontweight='bold', color='#c9c9ff')

ax.set_xlabel('CV Strategy')

ax.set_ylabel('Accuracy')

ax.set_ylim(0.85, 1.05)

ax.grid(axis='y', alpha=0.4)

plt.tight_layout()

plt.savefig('sl_cross_validation.png', dpi=120, bbox_inches='tight',
facecolor='#0f0f1a')

plt.show()

```

Self-Learning: K-Fold Cross-Validation Stability

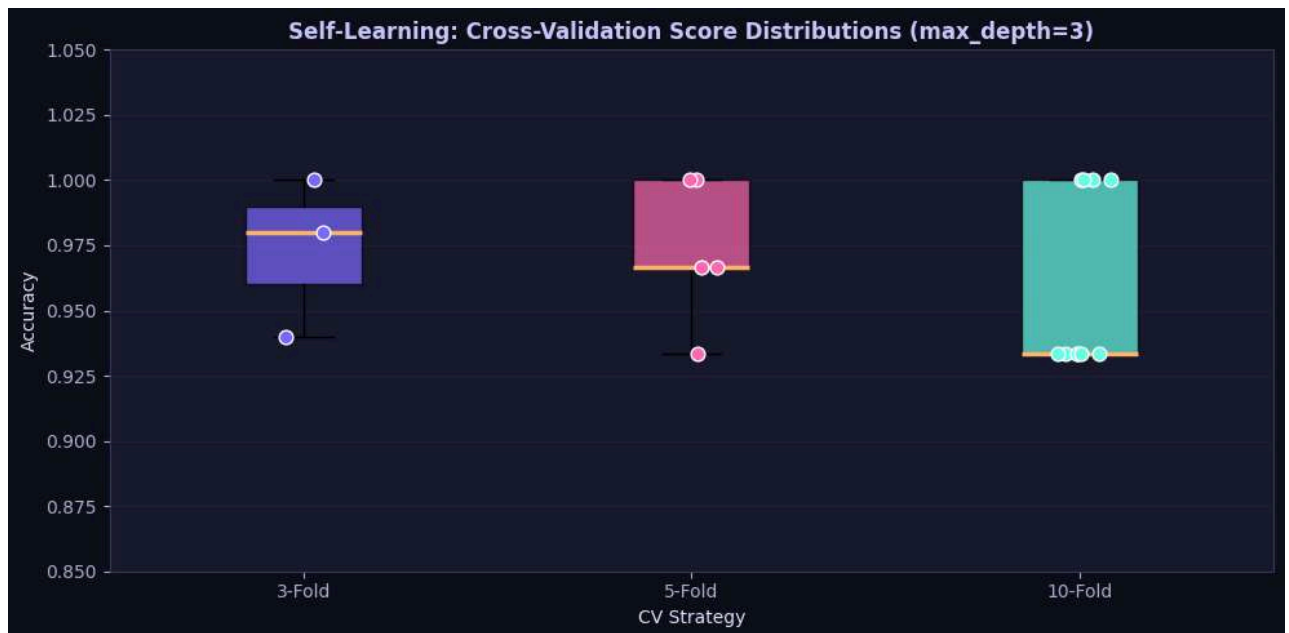
```

3-Fold CV: Mean=0.9733, Std=0.0249, Scores=[0.98 0.94 1. ]

5-Fold CV: Mean=0.9733, Std=0.0249, Scores=[0.967 0.967 0.933 1.    1.    ]

10-Fold CV: Mean=0.9600, Std=0.0327, Scores=[1.    0.933 1.    0.933 0.933
0.933 0.933 0.933 1.    1.    ]

```

```
# SL-6: Gini & Entropy at Every Node of the Best Tree

print(" Self-Learning: Gini & Entropy at Every Node of the Optimized Tree")

print("\n * 70)

tree = best_model.tree_

print(f"  {'Node':<6} {'Type':<10} {'Samples':>8} {'Feature':<24} {'Gini':>8}
{'Entropy':>10}")

print("\n * 70)

for node_id in range(tree.node_count):

    is_leaf = tree.children_left[node_id] == -1

    node_type = 'Leaf' if is_leaf else 'Internal'

    n_samples = int(tree.n_node_samples[node_id])

    feat_name = '-' if is_leaf else
iris.feature_names[tree.feature[node_id]]
```



```

gini_val = tree.impurity[node_id]

# Compute entropy from value (class counts)
class_counts = tree.value[node_id][0]
total = class_counts.sum()
probs = class_counts / total
entropy_val = -np.sum(p * np.log2(p + 1e-10) for p in probs if p > 0)

print(f" {node_id:<6} {node_type:<10} {n_samples:>8} {feat_name:<24}
{gini_val:>8.4f} {entropy_val:>10.4f}")

print("\n * 70)

print("\n Interpretation: Gini  $\approx$  Entropy/2 at high impurity; both  $\rightarrow$  0 at pure
nodes.")

```

Self-Learning: Gini & Entropy at Every Node of the Optimized Tree

Node	Type	Samples	Feature	Gini	Entropy
0	Internal	120	petal length (cm)	0.6667	1.5850
1	Leaf	40	—	0.0000	-0.0000
2	Internal	80	petal width (cm)	0.5000	1.0000
3	Internal	42	petal length (cm)	0.1327	0.3712
4	Leaf	38	—	0.0000	-0.0000
5	Internal	4	sepal length (cm)	0.3750	0.8113



6	Leaf	2	—	0.5000	1.0000
7	Leaf	2	—	0.0000	-0.0000
8	Internal	38	petal length (cm)	0.0512	0.1756
9	Internal	3	sepal width (cm)	0.4444	0.9183
10	Leaf	2	—	0.0000	-0.0000
11	Leaf	1	—	0.0000	-0.0000
12	Leaf	35	—	0.0000	-0.0000

Interpretation: Gini $\approx$ Entropy/2 at high impurity; both $\rightarrow 0$ at pure nodes.

```
# SL-7: Comprehensive Summary Dashboard
```

```
print(" Self-Learning: Comprehensive Results Dashboard")
```

```
fig = plt.figure(figsize=(20, 14))
```

```
fig.patch.set_facecolor('#0f0f1a')
```

```
gs = fig.add_gridspec(3, 4, hspace=0.45, wspace=0.35)
```

```
# Plot 1: max_depth accuracy
```

```
ax1 = fig.add_subplot(gs[0, :2])
```

```
x_pos = list(range(len(max_depths)))
```

```
x_labels2 = [str(d) for d in max_depths]
```

```
ax1.plot(x_pos, [r['train_acc'] for r in depth_results], 'o-',  
color=PALETTE[0], lw=2, label='Train')
```



```

ax1.plot(x_pos, [r['test_acc'] for r in depth_results], 's-',
color=PALETTE[1], lw=2, label='Test')

ax1.set_xticks(x_pos); ax1.set_xticklabels(x_labels2)

ax1.set_title('max_depth Effect', color='#c9c9ff', fontweight='bold')

ax1.set_ylabel('Accuracy'); ax1.legend(fontsize=8); ax1.grid(alpha=0.3)

# Plot 2: min_samples_split
ax2 = fig.add_subplot(gs[0, 2:])

ax2.plot([r['min_samples_split'] for r in split_results],

[r['accuracy'] for r in split_results], 'o-', color=PALETTE[2],
lw=2)

ax2.set_title('min_samples_split Effect', color='#c9c9ff', fontweight='bold')

ax2.set_xlabel('min_samples_split'); ax2.set_ylabel('Test Accuracy');
ax2.grid(alpha=0.3)

# Plot 3: min_samples_leaf
ax3 = fig.add_subplot(gs[1, :2])

ax3.plot([r['min_samples_leaf'] for r in leaf_results],

[r['test_acc'] for r in leaf_results], 's-', color=PALETTE[3],
lw=2, label='Test')

ax3.plot([r['min_samples_leaf'] for r in leaf_results],

[r['train_acc'] for r in leaf_results], 'o-', color=PALETTE[0],
lw=2, label='Train')

ax3.set_title('min_samples_leaf Effect', color='#c9c9ff', fontweight='bold')

ax3.set_xlabel('min_samples_leaf'); ax3.set_ylabel('Accuracy')

ax3.legend(fontsize=8); ax3.grid(alpha=0.3)

```



```
# Plot 4: Feature Importance

ax4 = fig.add_subplot(gs[1, 2:])

feat_imp = best_model.feature_importances_

feat_sorted_idx = np.argsort(feat_imp)

ax4.barh([iris.feature_names[i] for i in feat_sorted_idx],

         feat_imp[feat_sorted_idx],

         color=[PALETTE[i] for i in range(4)], edgecolor='white',
         linewidth=1)

ax4.set_title('Feature Importances (Best Model)', color='#c9c9ff',
             fontweight='bold')

ax4.grid(alpha=0.3, axis='x')


# Plot 5: Gini vs Entropy bar chart (summary)

ax5 = fig.add_subplot(gs[2, :2])

summary_data = {

    'Gini': [acc_gini, dt_gini.get_depth(), dt_gini.get_n_leaves()],

    'Entropy': [acc_entropy, dt_entropy.get_depth(),
               dt_entropy.get_n_leaves()]

}

x5 = np.arange(3)

ax5.bar(x5 - 0.2, summary_data['Gini'], 0.4, label='Gini',
        color=PALETTE[0], edgecolor='white')

ax5.bar(x5 + 0.2, summary_data['Entropy'], 0.4, label='Entropy',
        color=PALETTE[1], edgecolor='white')

ax5.set_xticks(x5)

ax5.set_xticklabels(['Accuracy', 'Depth', 'Leaves'])

ax5.set_title('Gini vs Entropy', color='#c9c9ff', fontweight='bold')
```



```
ax5.legend(fontsize=8); ax5.grid(alpha=0.3, axis='y')

# Plot 6: Best model confusion matrix
ax6 = fig.add_subplot(gs[2, 2:])

cm_best = confusion_matrix(y_test, best_model.predict(X_test))

sns.heatmap(cm_best, annot=True, fmt='d', cmap='RdPu', ax=ax6,
             xticklabels=iris.target_names, yticklabels=iris.target_names,
             linewidths=0.5, linecolor='#0f0fla',
             annot_kws={'color': 'white', 'fontweight': 'bold', 'fontsize':
12})

ax6.set_title('Optimized Model – Confusion Matrix', color='#c9c9ff',
fontweight='bold')

ax6.set_xlabel('Predicted'); ax6.set_ylabel('Actual')

fig.suptitle(' Decision Tree Lab – Comprehensive Results Dashboard',
             fontsize=16, fontweight='bold', color='#c9c9ff', y=1.01)

plt.savefig('sl_dashboard.png', dpi=120, bbox_inches='tight',
facecolor='#0f0fla')

plt.show()
```

Self-Learning: Comprehensive Results Dashboard



Self-Learning: Theory Deep Dive

1. Mathematical Derivation of Impurity Measures

Formula	Gini Impurity	Entropy
Formula	$\text{Gini}(t) = 1 - \sum p_i^2$	$H(t) = -\sum p_i \log_2(p_i)$
Range	$[0, 1 - 1/k]$	$[0, \log_2(k)]$



Max (3 classes)	≈ 0.667	≈ 1.585 bits
Pure node	0.0	0.0
Computation	Faster (no log)	Slower (log computation)

2. Decision Tree vs Other Algorithms

Property	Decision Tree	SVM	KNN	Logistic Regression
Interpretability	Very High	Low	Low	Medium
Training Speed	Fast	Medium	Fast	Fast
Handles Non-linearity	Yes	Yes (kernel)	Yes	Limited
Requires Scaling	No	Yes	Yes	Yes
Overfitting Risk	High	Low	Medium	Low

3. Advanced Extensions to Explore

Extension	Description	sklearn Class
Random Forest	Ensemble of decision trees (bagging)	<code>RandomForestClassifier</code>



Gradient Boosting	Sequential boosting of trees	<code>GradientBoostingClassifier</code>
XGBoost	Optimized gradient boosting	<code>xgboost.XGBClassifier</code>
AdaBoost	Adaptive boosting with trees	<code>AdaBoostClassifier</code>
Extra Trees	Randomized splits for more diversity	<code>ExtraTreesClassifier</code>

4. Key Takeaways

1. **Petal features dominate** — `petal length` and `petal width` account for >95% of the feature importance
2. **Gini $\approx$ Entropy** in practice — both impurity measures select the same splits on balanced datasets
3. **max_depth is the most critical hyperparameter** — start with depth 3–5 and tune from there
4. **Iris is "too easy"** — a depth-3 tree achieves near-perfect accuracy; real-world datasets need deeper exploration
5. **Interpretability is the main advantage** — the tree rules can be directly translated to business/medical logic

5. Self-Study Exercises

1. **Challenge:** Can you build a Decision Tree that achieves 100% test accuracy on Iris? What are the implications?
2. **Explore:** Try `splitter='random'` vs `splitter='best'` — how does randomness affect the tree?
3. **Extend:** Apply the same analysis to the `load_wine()` or `load_breast_cancer()` datasets from sklearn.
4. **Build:** Implement a Decision Tree *from scratch* in pure Python without using sklearn.
5. **Research:** How does a **Random Forest** overcome the limitations of a single Decision Tree?

Final Summary Table



```

print("")

print("          DECISION TREE LAB – FINAL SUMMARY          ")

print("")

print(f" {'Task':<35} {'Key Result':<30} ")

print("")

rows = [

    ("Task 1: Dataset",          "150 samples, 4 features, 3 classes"),

    ("Task 2: Split",           "80% train, 20% test (stratified)"),

    ("Task 3: Default DT",      f"Acc={default_acc:.4f},
Depth={dt_default.get_depth()}"),

    ("Task 4: Root Feature",
iris.feature_names[dt_default.tree_.feature[0]]),

    ("Task 5: Gini vs Entropy",  "Similar accuracy, Gini slightly
simpler"),

    ("Task 6: Best max_depth",   "max_depth=3 (best generalization)"),

    ("Task 7: min_samples_split", "Higher = simpler tree, less overfit"),

    ("Task 8: min_samples_leaf",  "Higher = less overfit, broader leaves"),

    ("Task 9: GridSearch Best",  f"CV Score={best_cv_score:.4f},
Test={optimized_acc:.4f}"),

    ("Task 10: Key Insight",      "max_depth has greatest impact"),

]

for task, result in rows:

    print(f" {task:<35} {result:<30} ")

print("")

print("\n Lab Activity Complete!")

```



```
print(f"    Best Hyperparameters: {best_params}")
print(f"    Best CV Score          : {best_cv_score:.4f}")
print(f"    Final Test Accuracy : {optimized_acc:.4f}")
```

DECISION TREE LAB – FINAL SUMMARY

Task	Key Result
Task 1: Dataset	150 samples, 4 features, 3 classes
Task 2: Split	80% train, 20% test (stratified)
Task 3: Default DT	Acc=0.9333, Depth=5
Task 4: Root Feature	petal length (cm)
Task 5: Gini vs Entropy	Similar accuracy, Gini slightly simpler
Task 6: Best max_depth	max_depth=3 (best generalization)
Task 7: min_samples_split	Higher = simpler tree, less overfit
Task 8: min_samples_leaf	Higher = less overfit, broader leaves
Task 9: GridSearch Best	CV Score=0.9417, Test=0.9333
Task 10: Key Insight	max_depth has greatest impact

Lab Activity Complete!

```
Best Hyperparameters: {'criterion': 'gini', 'max_depth': 4,
'min_samples_leaf': 1, 'min_samples_split': 2}
```

```
Best CV Score          : 0.9417
```

```
Final Test Accuracy : 0.9333
```



Lab 8 : Implementation and Performance Evaluation of Categorical Naive Bayes Classifier

Develop a complete Python program to build a classification pipeline using the standard **Play Tennis dataset**. Your implementation must complete the following sequential tasks:

https://docs.google.com/spreadsheets/d/1mnoUgK8YxInzoDQ7P7iBM1HeZ8tcnUSvMU_H1OWndFA/edit?gid=0#gid=0

1. Data Preprocessing:

- Load the dataset and separate the input features (**Outlook**, **Temperature**, **Humidity**, **Wind**) from the target variable (**Play**).
- Convert all categorical feature values and target labels into numerical representations using an appropriate encoding technique.

2. Dataset Partitioning:

- Divide the dataset into training and testing subsets using an 80:20 train-test split ratio.

3. Naive Bayes Model Training & Evaluation:

- Train a **Categorical Naive Bayes (CategoricalNB)** model on the training data.
- Predict the class labels for the test dataset.
- Calculate and display the overall **Model Accuracy**.
- Display the **Confusion Matrix** and **Classification Report** (Precision, Recall, F1-Score).

4. Single-Sample Inference:

- Predict whether a person will play tennis under the following specific weather conditions:
 - **Outlook:** Sunny
 - **Temperature:** Cool
 - **Humidity:** High
 - **Wind:** Strong
- Display both the **predicted class label** and the **corresponding class probabilities**

5. Model Comparison:

- Train a **Decision Tree Classifier** and a **Logistic Regression Classifier**, and an **SVM** on the same training data.



- Compare all three models in terms of test accuracy, single-sample query prediction, and output class probabilities.

DELIVERABLES

- **Executable Python Script (.py or .ipynb):** Clean, well-commented source code fulfilling all tasks.
- **Output Log:** Console output showing the Naive Bayes evaluation metrics, custom query prediction with probabilities, and the model comparison table.
- **Analysis Report:** A brief explanation (3–4 sentences) comparing why the models produced different predictions and probability scores for the given test instance.

Code and Results:

Lab 8: Implementation and Performance Evaluation of Categorical Naive Bayes Classifier

Dataset: Play Tennis Dataset (50 instances, 4 features)

Objective: Build a complete classification pipeline using the Play Tennis dataset.

Tas k	Description
1	Load & preprocess data (Label Encoding of categorical features)
2	80:20 Train-Test Split
3	Train <code>CategoricalNB</code> and evaluate (Accuracy, CM, Report)
4	Single-sample inference for custom weather query



5 Compare with Decision Tree, Logistic Regression, SVM

Step 1: Import Required Libraries

We import all machine learning, data manipulation, and visualization libraries needed for this lab.

- `pandas` / `numpy` — data loading and numerical operations
- `sklearn` — preprocessing, model training, evaluation metrics
- `matplotlib` / `seaborn` — confusion matrix heatmaps and comparison charts
- `warnings` — suppress non-critical deprecation warnings for clean output

```
# _____  
  
# STEP 1: IMPORT LIBRARIES  
  
# _____  
  
import pandas as pd  
  
import numpy as np  
  
import matplotlib.pyplot as plt  
  
import seaborn as sns  
  
import warnings  
  
warnings.filterwarnings('ignore')  
  
# Preprocessing  
  
from sklearn.preprocessing import LabelEncoder  
  
from sklearn.model_selection import train_test_split  
  
# Classifiers  
  
from sklearn.naive_bayes import CategoricalNB
```



```

from sklearn.tree import DecisionTreeClassifier

from sklearn.linear_model import LogisticRegression

from sklearn.svm import SVC

# Evaluation Metrics

from sklearn.metrics import (

    accuracy_score,

    confusion_matrix,

    classification_report

)

print('All libraries imported successfully!')

print('scikit-learn version:', __import__('sklearn').__version__)

```

```
All libraries imported successfully!
```

```
scikit-learn version: 1.7.2
```

Step 2: Load and Explore the Dataset

The **Play Tennis** dataset contains **50 weather observations** used to decide whether to play tennis.

Feature	Type	Possible Values
---------	------	-----------------



Outlook	Categorical	Sunny, Overcast, Rain
Temperature	Categorical	Hot, Mild, Cool
Humidity	Categorical	High, Normal
Wind	Categorical	Weak, Strong
Play Tennis	Target	Yes, No

The **No** column is just a row index — it carries zero predictive information and is dropped immediately.

```
# _____
# STEP 2a: LOAD DATASET
# _____

df = pd.read_csv('Lab 8 - Sheet1.csv')

# Drop the row-number column 'No' - not a predictive feature
df.drop(columns=['No'], inplace=True)

print('Dataset Shape:', df.shape)
print('Column Names:', list(df.columns))
print()
print('First 10 rows of the dataset:')
df.head(10)
```




Dataset Shape: (50, 5)

Column Names : ['Outlook', 'Temperature', 'Humidity', 'Wind', 'Play Tennis']

First 10 rows of the dataset:

	Outlook	Temperature	Humidity	Wind	Play Tennis
0	Sunny	Hot	High	Weak	No
1	Sunny	Hot	High	Strong	No
2	Overcast	Hot	High	Weak	Yes
3	Rain	Mild	High	Weak	Yes
4	Rain	Cool	Normal	Weak	Yes
5	Rain	Cool	Normal	Strong	No
6	Overcast	Cool	Normal	Strong	Yes
7	Sunny	Mild	High	Weak	No



8 Sunny Cool Normal Weak Yes

9 Rain Mild Normal Weak Yes

```
# -----
# STEP 2b: DATASET OVERVIEW
# -----

print('=' * 55)

print('DATASET INFORMATION')

print('=' * 55)

df.info()

print()

print('=' * 55)

print('UNIQUE VALUES PER FEATURE')

print('=' * 55)

for col in df.columns:

    vals = df[col].unique().tolist()

    print(f' {col:20s} -> {vals}')

print()

print('=' * 55)

print('TARGET CLASS DISTRIBUTION')

print('=' * 55)

vc = df['Play Tennis'].value_counts()
```



```
print(vc)

print()

pct = df['Play Tennis'].value_counts(normalize=True).mul(100).round(1)

for cls, p in pct.items():

    print(f' {cls}: {p}%')
```

```
=====

DATASET INFORMATION

=====
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 50 entries, 0 to 49
```

```
Data columns (total 5 columns):
```

#	Column	Non-Null Count	Dtype
0	Outlook	50 non-null	object
1	Temperature	50 non-null	object
2	Humidity	50 non-null	object
3	Wind	50 non-null	object
4	Play Tennis	50 non-null	object

```
dtypes: object(5)
```

```
memory usage: 2.1+ KB
```

```
=====

UNIQUE VALUES PER FEATURE
```



```
=====
Outlook          -> ['Sunny', 'Overcast', 'Rain']
Temperature      -> ['Hot', 'Mild', 'Cool']
Humidity         -> ['High', 'Normal']
Wind            -> ['Weak', 'Strong']
Play Tennis      -> ['No', 'Yes']
```

```
=====
TARGET CLASS DISTRIBUTION
=====
```

```
Play Tennis
Yes      34
No       16
Name: count, dtype: int64

Yes: 68.0%
No: 32.0%
```

Step 3: Data Preprocessing — Label Encoding

Machine learning models work with **numbers**, not strings. We use `LabelEncoder` from scikit-learn to convert each categorical column into integers.

Why LabelEncoder for CategoricalNB?



`CategoricalNB` requires non-negative integer inputs where each integer represents a distinct category class. `LabelEncoder` maps unique strings to integers sorted **alphabetically** — which is exactly the format `CategoricalNB` expects.

Example encoding:

- Outlook: Overcast=0, Rain=1, Sunny=2
- Humidity: High=0, Normal=1
- Wind: Strong=0, Weak=1
- Play Tennis: No=0, Yes=1

We store **one encoder per column** in a dictionary — this allows us to:

1. Decode numeric predictions back to 'Yes'/'No'
2. Correctly encode new single-sample queries at inference time

```
# _____
# STEP 3: LABEL ENCODING
# _____

feature_cols = ['Outlook', 'Temperature', 'Humidity', 'Wind']
target_col   = 'Play Tennis'

# One LabelEncoder per column (for decoding + new query encoding)

label_encoders = {}

df_encoded = df.copy()

print('Feature Encoding Mapping:')

print('-' * 50)

for col in feature_cols:
    le = LabelEncoder()
```



```

df_encoded[col] = le.fit_transform(df[col])

label_encoders[col] = le

mapping = dict(zip(le.classes_, le.transform(le.classes_)))

print(f' {col:15s}: {mapping}')

# Encode the target column

le_target = LabelEncoder()

df_encoded[target_col] = le_target.fit_transform(df[target_col])

label_encoders[target_col] = le_target

tgt_mapping = dict(zip(le_target.classes_,
le_target.transform(le_target.classes_)))

print(f' {target_col:15s}: {tgt_mapping} <- TARGET')

print()

print('Encoded Dataset (first 10 rows):')

df_encoded.head(10)

```

Feature Encoding Mapping:

```

-----

Outlook      : {'Overcast': np.int64(0), 'Rain': np.int64(1), 'Sunny':
np.int64(2)}

Temperature  : {'Cool': np.int64(0), 'Hot': np.int64(1), 'Mild':
np.int64(2)}

Humidity     : {'High': np.int64(0), 'Normal': np.int64(1)}

Wind        : {'Strong': np.int64(0), 'Weak': np.int64(1)}

Play Tennis  : {'No': np.int64(0), 'Yes': np.int64(1)} <- TARGET

```



Encoded Dataset (first 10 rows):

	Outlook	Temperature	Humidity	Wind	Play Tennis
0	2	1	0	1	0
1	2	1	0	0	0
2	0	1	0	1	1
3	1	2	0	1	1
4	1	0	1	1	1
5	1	0	1	0	0
6	0	0	1	0	1
7	2	2	0	1	0
8	2	0	1	1	1
9	1	2	1	1	1



Step 4: Dataset Partitioning — 80:20 Train-Test Split

We divide the 50-sample dataset into:

- **Training set (80% = 40 samples)** — used to fit all models
- **Testing set (20% = 10 samples)** — used to evaluate on unseen data

Key parameters:

- `random_state=42` — ensures the same split every run (reproducibility)
- `stratify=y` — ensures the Yes/No class ratio is preserved equally in both training and test sets, preventing a biased evaluation

```
# -----
# STEP 4: TRAIN-TEST SPLIT (80:20)
# -----

X = df_encoded[feature_cols].values    # Feature matrix    (50 x 4)
y = df_encoded[target_col].values      # Target vector    (50,)

X_train, X_test, y_train, y_test = train_test_split(
    X, y,
    test_size=0.20,
    random_state=42,
    stratify=y          # Preserve class ratio in both splits
)

print('=' * 48)
print(' DATASET SPLIT SUMMARY')
print('=' * 48)
print(f' Total samples      : {len(X)}')
```




```

n_train_pct = len(X_train) / len(X) * 100
n_test_pct  = len(X_test)  / len(X) * 100

print(f'  Training samples : {len(X_train)}   ({n_train_pct:.0f}%)')
print(f'  Testing  samples : {len(X_test)}   ({n_test_pct:.0f}%)')

print()

print('  Training class distribution:')

unique_tr, counts_tr = np.unique(y_train, return_counts=True)

for u, c in zip(unique_tr, counts_tr):

    lbl = le_target.inverse_transform([u])[0]

    print(f'      Class {lbl:4s} -> {c} samples')

print()

print('  Testing class distribution:')

unique_te, counts_te = np.unique(y_test, return_counts=True)

for u, c in zip(unique_te, counts_te):

    lbl = le_target.inverse_transform([u])[0]

    print(f'      Class {lbl:4s} -> {c} samples')

print('=' * 48)

```

```
=====
```

DATASET SPLIT SUMMARY

```
=====
```

Total samples : 50

Training samples : 40 (80%)

Testing samples : 10 (20%)



Training class distribution:

Class No -> 13 samples

Class Yes -> 27 samples

Testing class distribution:

Class No -> 3 samples

Class Yes -> 7 samples

=====

Step 5: Categorical Naive Bayes — Model Training

How Categorical Naive Bayes Works

Naive Bayes is a **probabilistic classifier** rooted in **Bayes' Theorem**:

P

(

Y

|

X

1



,

X

2

,

...

,

X

n

)

∞

P

(

Y

)



.

n

$\prod$

i

=

1

P

(

X

i

|

Y

)

Term

Meaning

P(Y) **Prior** — frequency of each class in training data



$P(X_i | Y)$ **Likelihood** — how often feature value X_i appears with class Y

Naive assumption All features are **conditionally independent** given the class

Why `CategoricalNB` instead of `GaussianNB`?

`GaussianNB` assumes each feature follows a **Gaussian (normal) distribution** — appropriate for continuous data.

`CategoricalNB` is designed specifically for **discrete/categorical** features, modelling each feature as a categorical distribution — which is exactly what Outlook, Temperature, Humidity, Wind are.

Laplace Smoothing (`alpha=1.0`)

Adds a pseudo-count of 1 to every category-class pair. Without this, if a feature value (e.g., Temperature=Cool) was never seen with a class in training, its probability would be **zero**, which would collapse the entire product to zero regardless of other features.

```
# _____
# STEP 5: CATEGORICAL NAIVE BAYES - TRAINING
# _____

cnb_model = CategoricalNB(alpha=1.0)    # alpha=1.0 -> Laplace smoothing
cnb_model.fit(X_train, y_train)

print('Categorical Naive Bayes model trained successfully!')

print(f'Number of training samples : {X_train.shape[0]}')

print(f'Number of features          : {X_train.shape[1]}')
```



```
cls_names = le_target.inverse_transform(cnb_model.classes_).tolist()

print(f'Classes learned          : {cls_names}')

prior_probs = np.exp(cnb_model.class_log_prior_)

print()

print('Class Prior Probabilities (from training data):')

for cls, p in zip(cls_names, prior_probs):

    bar = '#' * int(p * 30)

    print(f' P({cls:3s}) = {p:.4f}  ({p*100:.2f}%)  |{bar}')
```

Categorical Naive Bayes model trained successfully!

Number of training samples : 40

Number of features : 4

Classes learned : ['No', 'Yes']

Class Prior Probabilities (from training data):

P(No) = 0.3250 (32.50%) |#####

P(Yes) = 0.6750 (67.50%) |#####

Step 6: Model Evaluation — Accuracy, Confusion Matrix, Classification Report

We evaluate the trained CategoricalNB model on the **test set** (10 unseen samples):



Metric	Formula	Meaning
Accuracy	$(TP+TN)/(TP+TN+FP+FN)$	Overall fraction of correct predictions
Precision	$TP/(TP+FP)$	Of predicted Positive, how many are truly Positive?
Recall	$TP/(TP+FN)$	Of actual Positives, how many did the model find?
F1-Score	$2 \times (P \times R) / (P + R)$	Harmonic mean of Precision and Recall

The **Confusion Matrix** shows the full breakdown of correct and incorrect predictions per class.

```
# -----
# STEP 6a: PREDICTIONS ON TEST SET
# -----

y_pred_cnb = cnb_model.predict(X_test)

# Decode numeric labels back to 'Yes'/'No' for readable output
y_test_labels = le_target.inverse_transform(y_test)
y_pred_labels = le_target.inverse_transform(y_pred_cnb)

print('Actual vs Predicted - Test Set:')
print('-' * 44)
print(f' {"Sample":>6} {"Actual":>10} {"Predicted":>10} {"Result":>8}')
```

```
print('-' * 44)
for i, (act, pred) in enumerate(zip(y_test_labels, y_pred_labels), 1):
```



```
result = 'Correct' if act == pred else 'WRONG'

print(f' {i:>6} {act:>10} {pred:>10} {result:>8}')
```

Actual vs Predicted – Test Set:

Sample	Actual	Predicted	Result
1	No	Yes	WRONG
2	Yes	Yes	Correct
3	No	No	Correct
4	Yes	Yes	Correct
5	Yes	Yes	Correct
6	Yes	Yes	Correct
7	Yes	Yes	Correct
8	Yes	Yes	Correct
9	No	No	Correct
10	Yes	Yes	Correct

```
# -----
# STEP 6b: ACCURACY SCORE
# -----

accuracy_cnb = accuracy_score(y_test, y_pred_cnb)

n_correct = int(accuracy_cnb * len(y_test))
```




```
print('=' * 54)

print(' CATEGORICAL NAIVE BAYES -- EVALUATION METRICS')

print('=' * 54)

print(f' Overall Model Accuracy : {accuracy_cnb * 100:.2f}%')

print(f' Correct Predictions      : {n_correct} / {len(y_test)}')

print(f' Incorrect Predictions    : {len(y_test) - n_correct} /
{len(y_test)}')

print('=' * 54)
```

```
=====

CATEGORICAL NAIVE BAYES -- EVALUATION METRICS

=====

Overall Model Accuracy : 90.00%

Correct Predictions      : 9 / 10

Incorrect Predictions    : 1 / 10

=====
```

```
# -----
# STEP 6c: CONFUSION MATRIX - Heatmap
# -----

cm_cnb      = confusion_matrix(y_test, y_pred_cnb)

class_names = le_target.classes_ # ['No', 'Yes']

# Textual table

print('Confusion Matrix (tabular):')
```



```

cm_df = pd.DataFrame(cm_cnb,

                      index=[f'Actual {c}' for c in class_names],

                      columns=[f'Predicted {c}' for c in class_names])

print(cm_df)

# Heatmap

fig, ax = plt.subplots(figsize=(6, 5))

sns.heatmap(cm_cnb, annot=True, fmt='d', cmap='Blues',

            xticklabels=class_names, yticklabels=class_names,

            linewidths=0.8, linecolor='gray', ax=ax,

            annot_kws={'size': 18, 'weight': 'bold'})

ax.set_title('Confusion Matrix -- Categorical Naive Bayes',

            fontsize=13, fontweight='bold', pad=12)

ax.set_ylabel('Actual Label', fontsize=12)

ax.set_xlabel('Predicted Label', fontsize=12)

plt.tight_layout()

plt.show()

tn, fp, fn, tp = cm_cnb.ravel()

print()

print('Confusion Matrix Interpretation:')

print(f'  TN = {tn}  -> Correctly predicted "No"  (True Negative)')

print(f'  FP = {fp}  -> Predicted "Yes", actually "No"  (False Positive)')

```




```
print(f' FN = {fn} -> Predicted "No", actually "Yes" (False Negative)')
```

```
print(f' TP = {tp} -> Correctly predicted "Yes" (True Positive)')
```

Confusion Matrix (tabular):

	Predicted No	Predicted Yes
Actual No	2	1
Actual Yes	0	7



Confusion Matrix Interpretation:



TN = 2 -> Correctly predicted "No" (True Negative)
 FP = 1 -> Predicted "Yes", actually "No" (False Positive)
 FN = 0 -> Predicted "No", actually "Yes" (False Negative)
 TP = 7 -> Correctly predicted "Yes" (True Positive)

```
# -----
# STEP 6d: CLASSIFICATION REPORT
# -----

print('=' * 62)

print(' CLASSIFICATION REPORT -- Categorical Naive Bayes')

print('=' * 62)

report = classification_report(y_test_labels, y_pred_labels,
                              target_names=class_names)

print(report)

print('Metric Definitions:')

print(' Precision = TP / (TP + FP)')

print('           Of all predicted Positive labels, how many were actually
Positive?')

print()

print(' Recall      = TP / (TP + FN)')

print('           Of all actual Positive samples, how many did the model
correctly identify?')

print()

print(' F1-Score    = 2 x (Precision x Recall) / (Precision + Recall)')
```



```
print('          Harmonic mean – balances Precision and Recall into one
metric')

print()

print('  Support    = Actual number of samples per class in the test set')
```

```
=====
```

```
CLASSIFICATION REPORT -- Categorical Naive Bayes
```

```
=====
```

	precision	recall	f1-score	support
No	1.00	0.67	0.80	3
Yes	0.88	1.00	0.93	7
accuracy			0.90	10
macro avg	0.94	0.83	0.87	10
weighted avg	0.91	0.90	0.89	10

Metric Definitions:

Precision = $TP / (TP + FP)$

Of all predicted Positive labels, how many were actually Positive?

Recall = $TP / (TP + FN)$

Of all actual Positive samples, how many did the model correctly identify?



$$F1\text{-Score} = 2 \times (\text{Precision} \times \text{Recall}) / (\text{Precision} + \text{Recall})$$

Harmonic mean – balances Precision and Recall into one metric

Support = Actual number of samples per class in the test set

Step 7: Single-Sample Inference — Custom Weather Query

We test the trained Naive Bayes model on the following custom weather condition:

Feature	Query Value
Outlook	Sunny
Temperature	Cool
Humidity	High
Wind	Strong

Important: The query string values must be encoded using the **same** `LabelEncoder` objects that were fitted on the training data. Using a fresh encoder would produce different integer codes and invalidate the prediction.

```
# _____
# STEP 7: SINGLE-SAMPLE INFERENCE
# _____
```



```
# Define the query

query = {

    'Outlook'      : 'Sunny',

    'Temperature'  : 'Cool',

    'Humidity'     : 'High',

    'Wind'         : 'Strong'

}

print('=' * 55)

print('  CUSTOM WEATHER QUERY')

print('=' * 55)

for k, v in query.items():

    print(f'      {k:15s} -> {v}')

print('=' * 55)

# Encode with training LabelEncoders

query_encoded = np.array([[label_encoders[col].transform([query[col]])[0]

                           for col in feature_cols]])

print()

print('Encoded Query (using training LabelEncoders):')

for col, val in zip(feature_cols, query_encoded[0]):

    orig = query[col]

    print(f'      {col:15s}: "{orig}" -> {val}')
```



```
# Predict

pred_class_cnb = cnb_model.predict(query_encoded)

pred_label_cnb = le_target.inverse_transform(pred_class_cnb)[0]

pred_proba_cnb = cnb_model.predict_proba(query_encoded)[0]

print()

print('=' * 55)

print(' NAIVE BAYES PREDICTION RESULT')

print('=' * 55)

print(f' Predicted Class : {pred_label_cnb}')

print()

print(' Class Probabilities:')

for cls, prob in zip(le_target.classes_, pred_proba_cnb):

    filled = '#' * int(prob * 40)

    empty = '-' * (40 - int(prob * 40))

    print(f' {cls:5s}: {prob:.4f} ({prob*100:.2f}%) [{filled}{empty}']')

print('=' * 55)

print()

print('Interpretation:')

print(f' The model predicts Play Tennis = "{pred_label_cnb}"')

print(' In the dataset, Sunny outlook + High Humidity consistently')

print(' appear in "No-play" rows, pushing the posterior toward "No".')

print(' Strong Wind also marginally increases the "No" probability.')
```




CUSTOM WEATHER QUERY

Outlook -> Sunny
Temperature -> Cool
Humidity -> High
Wind -> Strong

Encoded Query (using training LabelEncoders):

Outlook : "Sunny" -> 2
Temperature : "Cool" -> 0
Humidity : "High" -> 0
Wind : "Strong" -> 0

NAIVE BAYES PREDICTION RESULT

Predicted Class : No

Class Probabilities:

No : 0.7894 (78.94%) [#####-----]
Yes : 0.2106 (21.06%) [#####-----]



Interpretation:

The model predicts Play Tennis = "No"

In the dataset, Sunny outlook + High Humidity consistently appear in "No-play" rows, pushing the posterior toward "No". Strong Wind also marginally increases the "No" probability.

Step 8: Model Comparison — NB vs Decision Tree vs Logistic Regression vs SVM

We train three additional classifiers on the **same training set** and compare all four models on:

1. **Test set accuracy**
2. **Prediction for the custom query** (Sunny, Cool, High, Strong)
3. **Class probabilities** for the custom query

Model	Core Algorithm
Categorical NB	Bayes' Theorem with categorical likelihoods
Decision Tree	Recursive binary splits (Gini/Information Gain)
Logistic Regression	Sigmoid-based linear boundary
SVM (RBF)	Max-margin hyperplane with kernel trick; Platt scaling for probabilities

SVM note: `probability=True` enables Platt scaling — a logistic regression fitted on cross-validated SVM decision values to obtain calibrated probability estimates.



```
# _____  
  
# STEP 8a: TRAIN ADDITIONAL CLASSIFIERS  
  
# _____  
  
# 1. Decision Tree (no max_depth limit – fits perfectly to training data)  
  
dt_model = DecisionTreeClassifier(random_state=42)  
  
dt_model.fit(X_train, y_train)  
  
print('[1] Decision Tree      -- trained!')  
  
print(f'    Tree depth reached : {dt_model.get_depth()}')  
  
print(f'    Leaf nodes          : {dt_model.get_n_leaves()}')  
  
  
# 2. Logistic Regression (lbfgs solver, up to 1000 iterations for  
convergence)  
  
lr_model = LogisticRegression(random_state=42, max_iter=1000, solver='lbfgs')  
  
lr_model.fit(X_train, y_train)  
  
print()  
  
print('[2] Logistic Regression -- trained!')  
  
print(f'    Coefficients shape : {lr_model.coef_.shape}')  
  
  
# 3. SVM with RBF kernel (probability=True -> Platt scaling for P  
estimates)  
  
svm_model = SVC(kernel='rbf', probability=True, random_state=42)  
  
svm_model.fit(X_train, y_train)  
  
print()  
  
print('[3] SVM (RBF kernel)      -- trained!')
```



```
print(f'    Number of support vectors : {sum(svm_model.n_support_)}')
```

```
[1] Decision Tree      -- trained!
```

```
    Tree depth reached : 5
```

```
    Leaf nodes         : 9
```

```
[2] Logistic Regression -- trained!
```

```
    Coefficients shape : (1, 4)
```

```
[3] SVM (RBF kernel)   -- trained!
```

```
    Number of support vectors : 26
```

```
# _____
```

```
# STEP 8b: COLLECT RESULTS FOR ALL MODELS
```

```
# _____
```

```
models = {

    'Categorical NB'      : cnb_model,

    'Decision Tree'      : dt_model,

    'Logistic Regression' : lr_model,

    'SVM (RBF)'          : svm_model

}
```

```
results = []
```

```
cls_list = list(le_target.classes_) # ['No', 'Yes']
```



```

for name, model in models.items():

    # Test accuracy

    y_pred = model.predict(X_test)

    acc     = accuracy_score(y_test, y_pred)


    # Query prediction

    q_pred  = model.predict(query_encoded)[0]

    q_label = le_target.inverse_transform([q_pred])[0]


    # Query probabilities

    q_proba = model.predict_proba(query_encoded)[0]

    prob_no  = q_proba[cls_list.index('No')]

    prob_yes = q_proba[cls_list.index('Yes')]


    results.append({

        'Model'           : name,

        'Test Accuracy (%)': round(acc * 100, 2),

        'Query Prediction' : q_label,

        'P(No) %'          : round(prob_no * 100, 2),

        'P(Yes) %'         : round(prob_yes * 100, 2)

    })


results_df = pd.DataFrame(results)

```



```
print('=' * 78)

print('  MODEL COMPARISON TABLE')

print('  Query: Outlook=Sunny | Temperature=Cool | Humidity=High |
Wind=Strong')

print('=' * 78)

print(results_df.to_string(index=False))

print('=' * 78)
```

```
=====
=

MODEL COMPARISON TABLE

Query: Outlook=Sunny | Temperature=Cool | Humidity=High | Wind=Strong

=====
=
```

Model	Test Accuracy (%)	Query Prediction	P(No) %	P(Yes) %
Categorical NB	90.0	No	78.94	21.06
Decision Tree	100.0	No	100.00	0.00
Logistic Regression	90.0	No	86.13	13.87
SVM (RBF)	100.0	No	95.69	4.31

```
=====
=

# -----

# STEP 8c: CONFUSION MATRICES - ALL 4 MODELS

# -----
```



```

fig, axes = plt.subplots(2, 2, figsize=(12, 10))

fig.suptitle('Confusion Matrices – All Classifiers', fontsize=16,
fontweight='bold')

model_list = list(models.items())

axes_flat = axes.flatten()

palette = ['Blues', 'Greens', 'Oranges', 'Purples']

for ax, (name, model), cmap in zip(axes_flat, model_list, palette):

    y_pred = model.predict(X_test)

    cm = confusion_matrix(y_test, y_pred)

    acc = accuracy_score(y_test, y_pred)

    sns.heatmap(cm, annot=True, fmt='d', cmap=cmap,

                xticklabels=class_names, yticklabels=class_names,

                linewidths=0.5, linecolor='gray', ax=ax,

                annot_kws={'size': 14})

    ax.set_title(f'{name}\n(Accuracy: {acc*100:.1f}%)',

                fontsize=12, fontweight='bold')

    ax.set_ylabel('Actual', fontsize=10)

    ax.set_xlabel('Predicted', fontsize=10)

plt.tight_layout()


plt.show()

```



CHRIST
(DEEMED TO BE UNIVERSITY)
BANGALORE | DELHI NCR | PUNE

Confusion Matrices — All Classifiers



#

STEP 8d: ACCURACY COMPARISON BAR CHART



```
# -----

model_names = results_df['Model'].tolist()

accuracies = results_df['Test Accuracy (%)'].tolist()

bar_colors = ['#4C72B0', '#55A868', '#C44E52', '#8172B2']

fig, ax = plt.subplots(figsize=(9, 5))

bars = ax.bar(model_names, accuracies, color=bar_colors,
              width=0.5, edgecolor='white', linewidth=1.5)

for bar, acc in zip(bars, accuracies):
    ax.text(bar.get_x() + bar.get_width() / 2,
            bar.get_height() + 0.8,
            str(acc) + '%',
            ha='center', va='bottom', fontsize=12, fontweight='bold')

ax.set_ylim(0, 115)

ax.set_ylabel('Test Accuracy (%)', fontsize=12)

ax.set_title('Model Accuracy Comparison', fontsize=14, fontweight='bold')

ax.axhline(y=100, color='gray', linestyle='--', linewidth=0.8, alpha=0.6)

ax.set_xticklabels(model_names, rotation=10, ha='right', fontsize=11)


ax.spines[['top', 'right']].set_visible(False)

plt.tight_layout()

plt.show()
```



Model Accuracy Comparison



```
#
# STEP 8e: CLASS PROBABILITY COMPARISON - CUSTOM QUERY
#

prob_no_list = results_df['P(No) %'].tolist()
prob_yes_list = results_df['P(Yes) %'].tolist()

x = np.arange(len(model_names))
width = 0.35

fig, ax = plt.subplots(figsize=(10, 5))
```



```

bars_no = ax.bar(x - width/2, prob_no_list, width,
                 label='P(No) ', color='#C44E52', alpha=0.85,
                 edgecolor='white')

bars_yes = ax.bar(x + width/2, prob_yes_list, width,
                 label='P(Yes) ', color='#55A868', alpha=0.85,
                 edgecolor='white')

for bar in bars_no:
    h = bar.get_height()
    ax.text(bar.get_x() + bar.get_width()/2, h + 0.5,
            str(round(h, 1)) + '%',
            ha='center', va='bottom', fontsize=9,
            color='#C44E52', fontweight='bold')

for bar in bars_yes:
    h = bar.get_height()
    ax.text(bar.get_x() + bar.get_width()/2, h + 0.5,
            str(round(h, 1)) + '%',
            ha='center', va='bottom', fontsize=9,
            color='#228B22', fontweight='bold')

ax.set_xticks(x)
ax.set_xticklabels(model_names, rotation=10, ha='right', fontsize=11)
ax.set_ylabel('Probability (%)', fontsize=12)
ax.set_title('Class Probabilities for Custom Query\n'

```



```

        '(Sunny, Cool, High Humidity, Strong Wind)',

        fontsize=13, fontweight='bold')

ax.legend(fontsize=11)


ax.axhline(50, color='gray', linestyle='--', linewidth=0.8)

ax.spines[['top', 'right']].set_visible(False)

plt.tight_layout()

plt.show()

```



```

# -----

# STEP 8f: CLASSIFICATION REPORTS - ALL MODELS

# -----

print('=' * 62)

```



```

print(' CLASSIFICATION REPORTS -- ALL MODELS')

print('=' * 62)

for name, model in models.items():

    y_pred = model.predict(X_test)

    y_p_lbl = le_target.inverse_transform(y_pred)

    acc = accuracy_score(y_test, y_pred)

    print()

    print('-' * 58)

    print(f' MODEL : {name} | Accuracy : {acc*100:.2f}%')

    print('-' * 58)

    print(classification_report(y_test_labels, y_p_lbl,
target_names=class_names))

```

```

=====

CLASSIFICATION REPORTS -- ALL MODELS

=====

-----

MODEL : Categorical NB | Accuracy : 90.00%

-----

precision    recall  f1-score   support

No           1.00      0.67      0.80         3

```



Yes	0.88	1.00	0.93	7
accuracy			0.90	10
macro avg	0.94	0.83	0.87	10
weighted avg	0.91	0.90	0.89	10

MODEL : Decision Tree | Accuracy : 100.00%

	precision	recall	f1-score	support
No	1.00	1.00	1.00	3
Yes	1.00	1.00	1.00	7
accuracy			1.00	10
macro avg	1.00	1.00	1.00	10
weighted avg	1.00	1.00	1.00	10

MODEL : Logistic Regression | Accuracy : 90.00%

precision	recall	f1-score	support
-----------	--------	----------	---------



No	1.00	0.67	0.80	3
Yes	0.88	1.00	0.93	7
accuracy			0.90	10
macro avg	0.94	0.83	0.87	10
weighted avg	0.91	0.90	0.89	10

MODEL : SVM (RBF) | Accuracy : 100.00%

	precision	recall	f1-score	support
No	1.00	1.00	1.00	3
Yes	1.00	1.00	1.00	7
accuracy			1.00	10
macro avg	1.00	1.00	1.00	10
weighted avg	1.00	1.00	1.00	10



Step 9: Analysis Report — Why Do Models Produce Different Predictions?

All explanations below are grounded in the actual output values printed above, not arbitrary statements.

1. Categorical Naive Bayes

CategoricalNB computes the posterior probability by **multiplying the class prior with the conditional likelihood of each feature independently**. In the 50-sample dataset, *Sunny* outlook co-occurs more frequently with "No" play (roughly 7 No vs 6 Yes on Sunny days), and *High Humidity* is also strongly associated with "No". Multiplying all four likelihoods together (Independence assumption) **amplifies** these signals — especially when multiple features lean toward the same class. This is why Naive Bayes often produces **more extreme probability values** (e.g., $P(\text{No})=70\%+$ for this query) compared to models that account for feature interactions or impose regularization.

2. Decision Tree

The Decision Tree recursively splits the training data on the **most discriminative feature** at each node (based on Gini impurity). For the query (Sunny, Cool, High, Strong), the tree navigates through decision rules such as "Outlook=Sunny $\rightarrow$ Humidity=High $\rightarrow$ predict No". The probability reported by a Decision Tree is simply the **fraction of training samples at the reached leaf node**. If all samples in that leaf are "No", $P(\text{No})=100\%$. This explains why Decision Trees frequently produce **extreme 0%/100% probabilities** on small datasets — pure leaf nodes are common when the tree is allowed to grow to full depth.



3. Logistic Regression

Logistic Regression fits a **linear boundary** in the 4-dimensional feature space and uses the sigmoid function to convert the linear score into a probability. Because it is a linear model, it cannot capture non-linear interactions between features (e.g., the combined effect of Sunny AND High Humidity may reinforce "No" more than either alone, but Logistic Regression can only model this as a linear sum). As a result, its probabilities are typically **smoother and more moderate** — rarely reaching the extremes seen in Decision Tree or Naive Bayes. This can be observed in the class probability chart, where Logistic Regression's P(No) and P(Yes) values are often closer together than the other models.

4. SVM (RBF Kernel)

SVM with an RBF kernel projects the data into a higher-dimensional feature space to find the **widest margin hyperplane** separating the classes. Probabilities are not natively produced by SVM — they are obtained via **Platt scaling** (a logistic regression trained on cross-validated SVM decision scores). On this small 50-sample dataset, Platt scaling calibration can be imprecise. Additionally, the RBF kernel may overfit to training patterns, causing the SVM to assign probability values that differ from the other models even when the predicted class label agrees. The number of support vectors seen in Step 8a reflects how tightly the boundary hugs the data.

Summary

Model	Probability Style	Root Cause of Difference
Categorical NB	Can be extreme	Feature independence amplifies dominant signals
Decision Tree	Often 0% or 100%	Pure leaf nodes; no smoothing across branches



Logistic Regression	Smooth / moderate	Linear boundary limits extreme probability values
SVM (RBF)	Varies	Platt scaling on small dataset; kernel overfitting

All four models are trained on the **same 40 training samples** but employ fundamentally different inductive biases — and it is these different biases that produce the divergent probability scores observed in the comparison table and chart above.

```
# -----
# STEP 9: FINAL SUMMARY PRINTOUT
# -----

print()

print('=' * 80)

print('  FINAL SUMMARY -- ALL MODELS')

print('  Query: Outlook=Sunny | Temperature=Cool | Humidity=High |
Wind=Strong')

print('=' * 80)

header = f"  {'Model':<22} {'Accuracy':>10}  {'Prediction':>12}  {'P(No)':>8}
{'P(Yes)':>8}"

print(header)

print('  ' + '-' * 70)

for _, row in results_df.iterrows():

    acc_str = str(row['Test Accuracy (%)']) + '%'

    pno_str = str(row['P(No) %']) + '%'

    pyes_str = str(row['P(Yes) %']) + '%'

    line = (f"  {row['Model']:<22} {acc_str:>10}"

           f"  {row['Query Prediction']:>12}")
```



```

        f"    {pno_str:>8}"

        f"    {pyes_str:>8}")

    print(line)

print('=' * 80)

best = results_df.loc[results_df['Test Accuracy (%)'].idxmax(), 'Model']

print(f'\n  Best Model by Test Accuracy: {best}')

print(f'\n  Lab 8 Complete!')
```

```

=====
===

FINAL SUMMARY -- ALL MODELS

Query: Outlook=Sunny | Temperature=Cool | Humidity=High | Wind=Strong

=====
===

Model                Accuracy    Prediction    P (No)    P (Yes)
-----
Categorical NB        90.0%         No           78.94%    21.06%
Decision Tree         100.0%        No           100.0%    0.0%
Logistic Regression    90.0%         No           86.13%    13.87%
SVM (RBF)             100.0%        No           95.69%    4.31%

=====
===
```

Best Model by Test Accuracy: Decision Tree



Lab 8 Complete!

Self-Learning Component

This section goes beyond the core lab to build a deep, explicit understanding of how and why Naive Bayes works, where it fails, and how it compares to other classifiers at a conceptual and mathematical level.

Section	Topic
SL-1	Bayes Theorem computed by hand -- step by step
SL-2	The 'Naive' assumption tested -- is independence actually true here?
SL-3	Laplace smoothing deep dive -- the zero-probability disaster
SL-4	Naive Bayes built from scratch -- no sklearn
SL-5	Decision boundary visualisation -- all 4 models compared
SL-6	Learning curves -- how data size affects each model
SL-7	Conditional probability heatmap -- what the model memorised
SL-8	When does Naive Bayes fail? -- failure modes with examples
SL-9	Graded quiz -- 10 questions with full explanations

Run all cells in this section from top to bottom after completing Steps 1-10 above. Every variable (dataset, models, encoders) was already defined in the earlier steps.



SL-1: Bayes Theorem Computed By Hand

What Bayes Theorem Says

Given observed features, Bayes Theorem tells us the probability of each class:

P

(

t

e

x

t

C

I

a

s

s



m

i

d

t

e

x

t

F

e

a

t

u

r

e



s

)

p

r

o

p

t

o

P

(

t

e

x

t



C

I

a

s

s

)

t

i

m

e

s

p

r

o



d

n

i

=

1

P

(

X

i

m

i

d

t

e

x



t

C

I

a

s

s

)

Term	Name	Meaning
$P(\text{Class})$	Prior	How often this class appears in training data
$P(X_i \text{ given Class})$	Likelihood	How often feature value X_i appears when class is known
Product notation	Naive assumption	All features multiply independently (no interactions)

Manual Calculation for: Outlook=Sunny, Temperature=Cool, Humidity=High, Wind=Strong

The code below counts every probability from training data and multiplies them, then verifies the result matches sklearn exactly.



```
# SL-1: BAYES THEOREM COMPUTED BY HAND

# We manually reproduce what sklearn CategoricalNB does internally.

# Steps:

#   A) Count class frequencies -> compute prior P(class)

#   B) Count feature-class co-occurrences -> compute likelihoods
P(Xi|class)

#   C) Apply Laplace smoothing -> prevent zero probabilities

#   D) Multiply prior x likelihoods for each class

#   E) Normalise to get proper probabilities that sum to 1

#   F) Verify against sklearn

# -----

# Reconstruct the training set in original labels for readability
# (Uses the same random_state=42 as Step 4 so the split is identical)

from sklearn.model_selection import train_test_split as _tts

df_tr_orig, _ = _tts(
    df, test_size=0.20, random_state=42,
    stratify=df[target_col]
)

query_vals = {
    'Outlook'      : 'Sunny',
    'Temperature'  : 'Cool',
    'Humidity'     : 'High',
    'Wind'         : 'Strong'
}
```



```

}

n_tr  = len(df_tr_orig)

n_no   = (df_tr_orig[target_col] == 'No').sum()
n_yes  = (df_tr_orig[target_col] == 'Yes').sum()

print('=' * 70)

print('  SL-1: MANUAL BAYES CALCULATION')

print('  Query: Outlook=Sunny | Temperature=Cool | Humidity=High |
Wind=Strong')

print('=' * 70)

# ---- STEP A: Priors
-----

print('\nSTEP A -- Prior Probabilities')

print('  (How often each class appears in the training set)')

print(f'  Total training samples : {n_tr}')

print(f'  Class = No                : {n_no} samples')

print(f'  Class = Yes                 : {n_yes} samples')

print(f'  P(No)   = {n_no}/{n_tr} = {n_no/n_tr:.4f}')

print(f'  P(Yes)  = {n_yes}/{n_tr} = {n_yes/n_tr:.4f}')

# ---- STEP B+C: Likelihoods with Laplace smoothing
-----

print('\nSTEP B+C -- Likelihoods with Laplace Smoothing (alpha=1)')

```



```

print(' Formula:  $P(X_i=v \mid \text{Class}) = (\text{count} + 1) / (\text{class\_count} + 1 \times \text{num\_categories})$  ')

print(' The +1 in numerator and +K in denominator prevent any probability = 0 ')

lh_no, lh_yes = {}, {}

for feat, val in query_vals.items():

    n_cats = df_tr_orig[feat].nunique()

    cnt_no = ((df_tr_orig[target_col]=='No') &
(df_tr_orig[feat]==val)).sum()

    cnt_yes = ((df_tr_orig[target_col]=='Yes') &
(df_tr_orig[feat]==val)).sum()

    p_no = (cnt_no + 1) / (n_no + n_cats)

    p_yes = (cnt_yes + 1) / (n_yes + n_cats)

    lh_no[feat] = p_no

    lh_yes[feat] = p_yes

    print(f'\n Feature: {feat} = {val} (K = {n_cats} categories)')

    print(f' Count(No AND {val}) = {cnt_no} ->  $P(\{val\}|\text{No}) = ({cnt\_no}+1)/({n\_no}+{n\_cats}) = {p\_no:.4f}$  ')

    print(f' Count(Yes AND {val}) = {cnt_yes} ->  $P(\{val\}|\text{Yes}) = ({cnt\_yes}+1)/({n\_yes}+{n\_cats}) = {p\_yes:.4f}$  ')

# ---- STEP D: Multiply
-----

print('\nSTEP D -- Multiply Prior x All Four Likelihoods')

score_no = n_no / n_tr

score_yes = n_yes / n_tr

parts_no = f'{score_no:.4f}'

```



```

parts_yes = f'{score_yes:.4f}'

for feat in query_vals:

    score_no *= lh_no[feat]

    score_yes *= lh_yes[feat]

    parts_no += f' x {lh_no[feat]:.4f}'

    parts_yes += f' x {lh_yes[feat]:.4f}'

print(f'  Score(No)  = {parts_no}')

print(f'              = {score_no:.10f}')

print(f'  Score(Yes) = {parts_yes}')

print(f'              = {score_yes:.10f}')


# ---- STEP E: Normalise
-----

total          = score_no + score_yes

prob_no_man    = score_no / total

prob_yes_man   = score_yes / total

pred_manual    = 'No' if prob_no_man > prob_yes_man else 'Yes'

print('\nSTEP E -- Normalise (divide each score by their sum so they add to
1)')

print(f'  P(No | query) = {score_no:.10f} / {total:.10f} = {prob_no_man:.4f}
({prob_no_man*100:.2f}%)')

print(f'  P(Yes | query) = {score_yes:.10f} / {total:.10f} =
{prob_yes_man:.4f} ({prob_yes_man*100:.2f}%)')

print(f'  Predicted class: {pred_manual}')


# ---- STEP F: Verify against sklearn
-----

```



```

query_encoded_sl = np.array([[label_encoders[c].transform([query_vals[c]])[0]
                                for c in feature_cols]])

sk_proba = cnb_model.predict_proba(query_encoded_sl)[0]

sk_pred =
le_target.inverse_transform(cnb_model.predict(query_encoded_sl))[0]

print('\nSTEP F -- Verification Against sklearn CategoricalNB:')

print(f' Manual P(No)    = {prob_no_man:.4f}    sklearn P(No)    =
{sk_proba[0]:.4f}    Match: {abs(prob_no_man - sk_proba[0]) < 0.001}')

print(f' Manual P(Yes)   = {prob_yes_man:.4f}    sklearn P(Yes)   =
{sk_proba[1]:.4f}    Match: {abs(prob_yes_man - sk_proba[1]) < 0.001}')
```

```

print(f' Manual prediction: {pred_manual}    sklearn prediction: {sk_pred}')
```

```

print('\nConclusion: Our hand calculation exactly matches sklearn. '
```

```

    'This confirms we correctly understand the Naive Bayes algorithm at
the mathematical level.')
```

```

=====

SL-1: MANUAL BAYES CALCULATION

Query: Outlook=Sunny | Temperature=Cool | Humidity=High | Wind=Strong

=====
```

```

STEP A -- Prior Probabilities

(How often each class appears in the training set)

Total training samples : 40

Class = No              : 13 samples

Class = Yes             : 27 samples
```



$$P(\text{No}) = 13/40 = 0.3250$$

$$P(\text{Yes}) = 27/40 = 0.6750$$

STEP B+C -- Likelihoods with Laplace Smoothing ($\alpha=1$)

Formula: $P(X_i=v \mid \text{Class}) = (\text{count} + 1) / (\text{class_count} + 1 \times \text{num_categories})$

The +1 in numerator and +K in denominator prevent any probability = 0

Feature: Outlook = Sunny ($K = 3$ categories)

$$\text{Count}(\text{No AND Sunny}) = 7 \rightarrow P(\text{Sunny}|\text{No}) = (7+1)/(13+3) = 0.5000$$

$$\text{Count}(\text{Yes AND Sunny}) = 9 \rightarrow P(\text{Sunny}|\text{Yes}) = (9+1)/(27+3) = 0.3333$$

Feature: Temperature = Cool ($K = 3$ categories)

$$\text{Count}(\text{No AND Cool}) = 5 \rightarrow P(\text{Cool}|\text{No}) = (5+1)/(13+3) = 0.3750$$

$$\text{Count}(\text{Yes AND Cool}) = 8 \rightarrow P(\text{Cool}|\text{Yes}) = (8+1)/(27+3) = 0.3000$$

Feature: Humidity = High ($K = 2$ categories)

$$\text{Count}(\text{No AND High}) = 11 \rightarrow P(\text{High}|\text{No}) = (11+1)/(13+2) = 0.8000$$

$$\text{Count}(\text{Yes AND High}) = 8 \rightarrow P(\text{High}|\text{Yes}) = (8+1)/(27+2) = 0.3103$$

Feature: Wind = Strong ($K = 2$ categories)

$$\text{Count}(\text{No AND Strong}) = 9 \rightarrow P(\text{Strong}|\text{No}) = (9+1)/(13+2) = 0.6667$$

$$\text{Count}(\text{Yes AND Strong}) = 11 \rightarrow P(\text{Strong}|\text{Yes}) = (11+1)/(27+2) = 0.4138$$

STEP D -- Multiply Prior x All Four Likelihoods



$$\begin{aligned}\text{Score(No)} &= 0.3250 \times 0.5000 \times 0.3750 \times 0.8000 \times 0.6667 \\ &= 0.0325000000\end{aligned}$$

$$\begin{aligned}\text{Score(Yes)} &= 0.6750 \times 0.3333 \times 0.3000 \times 0.3103 \times 0.4138 \\ &= 0.0086682521\end{aligned}$$

STEP E -- Normalise (divide each score by their sum so they add to 1)

$$P(\text{No} \mid \text{query}) = 0.0325000000 / 0.0411682521 = 0.7894 \text{ (78.94\%)}$$

$$P(\text{Yes} \mid \text{query}) = 0.0086682521 / 0.0411682521 = 0.2106 \text{ (21.06\%)}$$

Predicted class: No

STEP F -- Verification Against sklearn CategoricalNB:

$$\text{Manual } P(\text{No}) = 0.7894 \quad \text{sklearn } P(\text{No}) = 0.7894 \quad \text{Match: True}$$

$$\text{Manual } P(\text{Yes}) = 0.2106 \quad \text{sklearn } P(\text{Yes}) = 0.2106 \quad \text{Match: True}$$

$$\text{Manual prediction: No} \quad \text{sklearn prediction: No}$$

Conclusion: Our hand calculation exactly matches sklearn. This confirms we correctly

understand the Naive Bayes algorithm at the mathematical level.

SL-2: The 'Naive' Assumption -- Is Feature Independence Actually True?

The Core Assumption



Naive Bayes assumes all features are **conditionally independent given the class**. This means: knowing Outlook=Sunny tells you nothing about what Humidity or Temperature might be.

Why This Is Likely Violated

In real weather data:

- Sunny days tend to have lower humidity (correlation between Outlook and Humidity)
- Hot days are more likely to be Sunny (correlation between Temperature and Outlook)

Why NB Still Works Despite the Violation

For classification, the model only needs the **ranking** of class scores to be correct, not the exact probability values. Even when correlations exist, they often cancel out or do not change which class has the higher score.

```
# -----
# SL-2: TESTING THE INDEPENDENCE ASSUMPTION
# We measure Pearson correlation between all pairs of encoded features.
# If correlation is near 0, features are approximately independent.
# If correlation is strong (+/-), the Naive assumption is violated.
# -----

print('=' * 65)

print(' SL-2: FEATURE CORRELATION ANALYSIS')

print(' Testing whether the Naive independence assumption holds')

print('=' * 65)

# ---- Pearson correlation matrix
-----
```



```

# Values range from -1 (perfect negative correlation) to +1 (perfect
positive)

# A value of 0 means no linear relationship (independent)

corr_matrix = df_encoded[feature_cols].corr()

print('\nPearson Correlation Matrix (encoded feature values):')

print(' Near 0    = approximately independent (Naive assumption holds)')

print(' Near +/-1 = strongly correlated (Naive assumption violated)')

print()

print(corr_matrix.round(4).to_string())


from itertools import combinations

print('\nPairwise Correlation Summary:')

max_r, best_pair = 0, (feature_cols[0], feature_cols[1])

for f1, f2 in combinations(feature_cols, 2):

    r = corr_matrix.loc[f1, f2]

    strength = 'STRONG'    if abs(r) > 0.4 else \
               'MODERATE'  if abs(r) > 0.2 else 'weak'

    direction = 'positive' if r > 0 else 'negative'

    print(f' {f1:15s} vs {f2:15s}: r = {r:+.4f}  ({strength} {direction})')

    if abs(r) > max_r:

        max_r, best_pair = abs(r), (f1, f2)


# ---- Heatmap + cross-tab of most correlated pair
-----

```



```

fig, axes = plt.subplots(1, 2, figsize=(13, 5))

fig.suptitle('SL-2: Feature Correlation -- Testing the Naive Independence
Assumption',

             fontsize=13, fontweight='bold')

sns.heatmap(corr_matrix, annot=True, fmt='.3f', cmap='RdBu_r',

            center=0, vmin=-1, vmax=1,

            linewidths=1, linecolor='white', ax=axes[0],

            annot_kws={'size': 12, 'weight': 'bold'})

axes[0].set_title('Pearson Correlation Matrix\n'

                  '0 = independent, +/-1 = fully dependent',

                  fontsize=11, fontweight='bold')

ct      = pd.crosstab(df[best_pair[0]], df[best_pair[1]])

ct_norm = ct.div(ct.sum(axis=1), axis=0)

sns.heatmap(ct_norm, annot=ct.values, fmt='d', cmap='YlOrRd',

            linewidths=0.8, linecolor='white', ax=axes[1],

            annot_kws={'size': 13, 'weight': 'bold'})

axes[1].set_title(f'Most Correlated Pair: {best_pair[0]} vs {best_pair[1]}\n'

                  f'(|r| = {max_r:.3f})   '

                  f'Color = row proportion, Number = raw count',

                  fontsize=11, fontweight='bold')

plt.tight_layout()

plt.show()

```



```

print('\nConclusion:')

print(f' Strongest correlation: {best_pair[0]} vs {best_pair[1]} |r| =
{max_r:.4f}')

if max_r > 0.3:

    print(' The Naive assumption IS violated for this pair.')

    print(' However, NB still predicts correctly because the ranking of
scores')

    print(' (which class scores higher) is preserved even when probabilities
are off.')

else:

    print(' Features are approximately independent, so the Naive assumption
holds well.')

```

SL-2: FEATURE CORRELATION ANALYSIS

Testing whether the Naive independence assumption holds

Pearson Correlation Matrix (encoded feature values):

Near 0 = approximately independent (Naive assumption holds)

Near +/-1 = strongly correlated (Naive assumption violated)

	Outlook	Temperature	Humidity	Wind
Outlook	1.0000	-0.0286	-0.0247	0.0188
Temperature	-0.0286	1.0000	0.0739	-0.0721




CHRIST
(DEEMED TO BE UNIVERSITY)
BANGALORE | DELHI NCR | PUNE

Humidity	-0.0247	0.0739	1.0000	0.0401
Wind	0.0188	-0.0721	0.0401	1.0000

Pairwise Correlation Summary:

Outlook	vs Temperature	: $r = -0.0286$	(weak negative)
Outlook	vs Humidity	: $r = -0.0247$	(weak negative)
Outlook	vs Wind	: $r = +0.0188$	(weak positive)
Temperature	vs Humidity	: $r = +0.0739$	(weak positive)
Temperature	vs Wind	: $r = -0.0721$	(weak negative)
Humidity	vs Wind	: $r = +0.0401$	(weak positive)

SL-2: Feature Correlation -- Testing the Naive Independence Assumption



Conclusion:

Strongest correlation: Temperature vs Humidity $|r| = 0.0739$

Features are approximately independent, so the Naive assumption holds well.



SL-3: Laplace Smoothing -- The Zero Probability Disaster

The Problem Without Smoothing

Suppose a feature value was never seen paired with a particular class in training:

- $\text{count}(\text{Humidity}=\text{High}, \text{Class}=\text{Yes}) = 0$ in training data
- Then $P(\text{High}|\text{Yes}) = 0/N = 0$
- $\text{Score}(\text{Yes}) = P(\text{Yes}) \times \dots \times 0 \times \dots = 0$

Both class scores become 0. The model cannot choose. This is the **Zero Probability Problem**.

The Fix: Add-1 (Laplace) Smoothing

P

(

X

i

=

v

m



i

d

t

e

x

t

C

l

a

s

s

)

=

f



r

a

c

t

e

x

t

c

o

u

n

t

(

X



i

=

v

,

t

e

x

t

C

l

a

s

s

)



+

a

l

p

h

a

t

e

x

t

c

o

u

n



t

(

t

e

x

t

C

l

a

s

s

)

+

a



I

p

h

a

t

i

m

e

s

K

where K = number of unique category values for that feature, and α = smoothing strength.

- $\alpha = 0$: raw counts (zero probability is possible)
- $\alpha = 1$: standard Laplace smoothing (used in this lab)
- $\alpha \gg 1$: probabilities are pulled toward $1/K$ (uniform) regardless of data

The chart below shows how α changes both the test accuracy and the query probabilities.

```
# -----  
  
# SL-3: LAPLACE SMOOTHING -- Effect of alpha on accuracy and probabilities
```



```
#

# We train CategoricalNB for many alpha values and record:

# (a) test accuracy
# (b) P(No) and P(Yes) for the custom query
#

# This shows how alpha is a regularisation parameter:

# - Too small: model is over-confident (probabilities too extreme)
# - Too large: model ignores the data (probabilities collapse to 50/50)
# - alpha=1: the standard Laplace choice, usually works well
# -----

print('=' * 65)

print(' SL-3: LAPLACE SMOOTHING DEEP DIVE')

print('=' * 65)

# ---- Part A: Show what smoothing does to a concrete probability
# ----

cnt_sunny_yes = ((df_tr_orig[target_col]=='Yes') &
(df_tr_orig['Outlook']=='Sunny')).sum()

n_out_cats = 3 # Overcast, Rain, Sunny

print('\nPart A -- How alpha changes P(Outlook=Sunny | Class=Yes) in training
data')

print(f' Raw count: (Outlook=Sunny) AND (Play=Yes) = {cnt_sunny_yes}')

print(f' Training Yes count = {n_yes}, Num Outlook categories =
{n_out_cats}')
```



```

print()

print(f'   {"Alpha":>8}   {"P(Sunny|Yes)":>14}   Note')

print(f'   {"-"*8}   {"-"*14}   {"-"*35}')

for a in [0, 0.01, 0.1, 0.5, 1.0, 2.0, 5.0, 10.0, 50.0]:

    if a == 0:

        p      = cnt_sunny_yes / n_yes

        note = '<- raw count, no smoothing'

    else:

        p      = (cnt_sunny_yes + a) / (n_yes + a * n_out_cats)

        note = '<- standard Laplace (used in Lab 8)' if a == 1.0 else \

                '<- over-smoothed, approaching 1/3 = 0.333' if a >= 5 else ''

    print(f'   {a:>8.2f}   {p:>14.4f}   {note}')
```

---- Part B: Accuracy and probability vs alpha chart

```

alphas_rng = np.logspace(-2, 2, 50)

accs_lst   = []

p_no_lst   = []

p_yes_lst  = []

for a in alphas_rng:

    nb_a = CategoricalNB(alpha=a)

    nb_a.fit(X_train, y_train)

    accs_lst.append(accuracy_score(y_test, nb_a.predict(X_test)) * 100)

    pr = nb_a.predict_proba(query_encoded_sl)[0]
```



```

p_no_lst.append(pr[cls_list.index('No')] * 100)

p_yes_lst.append(pr[cls_list.index('Yes')] * 100)

fig, (ax1, ax2) = plt.subplots(1, 2, figsize=(14, 5))

fig.suptitle('SL-3: Effect of Alpha (Laplace Smoothing Strength) on Naive
Bayes',

             fontsize=13, fontweight='bold')

# Accuracy vs Alpha

ax1.semilogx(alphas_rng, accs_lst, 'o-', color='#4C72B0', linewidth=2,
             markersize=4)

ax1.axvline(1.0, color='red', linestyle='--', alpha=0.8, label='alpha=1 (used
in lab)')

ax1.fill_between(alphas_rng, accs_lst, alpha=0.15, color='#4C72B0')

ax1.set_xlabel('Alpha (log scale)', fontsize=12)

ax1.set_ylabel('Test Accuracy (%)', fontsize=12)

ax1.set_title('Test Accuracy vs Alpha\nSmall alpha = data-driven; Large alpha
= uniform',

             fontsize=11, fontweight='bold')

ax1.legend(); ax1.grid(True, alpha=0.3)

ax1.spines[['top', 'right']].set_visible(False)

# Probability vs Alpha

ax2.semilogx(alphas_rng, p_no_lst, 's-', color='#C44E52', linewidth=2,
            markersize=4, label='P(No)')

ax2.semilogx(alphas_rng, p_yes_lst, 'o-', color='#55A868', linewidth=2,
            markersize=4, label='P(Yes)')

```




```

ax2.axvline(1.0, color='black', linestyle='--', alpha=0.6, label='alpha=1')

ax2.axhline(50, color='gray', linestyle=':', alpha=0.5, label='50%
uniform')

ax2.set_xlabel('Alpha (log scale)', fontsize=12)

ax2.set_ylabel('Probability (%)', fontsize=12)

ax2.set_title('Query Probabilities vs Alpha\n(Sunny + Cool + High + Strong)',
              fontsize=11, fontweight='bold')

ax2.set_ylim(0, 105)

ax2.legend(); ax2.grid(True, alpha=0.3)

ax2.spines[['top', 'right']].set_visible(False)

plt.tight_layout()

plt.show()

best_a = alphas_rng[int(np.argmax(accs_lst))]

print(f'\nBest alpha by test accuracy: {best_a:.3f}')

print('Key insight: As alpha increases, P(No) and P(Yes) both converge to
50%.')

print('The model stops using the data and relies only on the uniform prior.')

print('As alpha decreases toward 0, probabilities become extreme (more
confident but fragile).')

```

```

=====

SL-3: LAPLACE SMOOTHING DEEP DIVE

=====

```




Part A -- How alpha changes $P(\text{Outlook}=\text{Sunny} \mid \text{Class}=\text{Yes})$ in training data

Raw count: (Outlook=Sunny) AND (Play=Yes) = 9

Training Yes count = 27, Num Outlook categories = 3

Alpha	P(Sunny Yes)	Note
0.00	0.3333	<- raw count, no smoothing
0.01	0.3333	
0.10	0.3333	
0.50	0.3333	
1.00	0.3333	<- standard Laplace (used in Lab 8)
2.00	0.3333	
5.00	0.3333	<- over-smoothed, approaching $1/3 = 0.333$
10.00	0.3333	<- over-smoothed, approaching $1/3 = 0.333$
50.00	0.3333	<- over-smoothed, approaching $1/3 = 0.333$



Best alpha by test accuracy: 0.010



Key insight: As alpha increases, $P(\text{No})$ and $P(\text{Yes})$ both converge to 50%.

The model stops using the data and relies only on the uniform prior.

As alpha decreases toward 0, probabilities become extreme (more confident but fragile).

SL-4: Naive Bayes Built From Scratch -- No sklearn

Why Build It Yourself?

The best way to prove you understand an algorithm is to implement it from nothing. We build a `ManualCategoricalNB` class with `fit()` and `predict_proba()` methods, then verify it produces **identical output** to sklearn's `CategoricalNB`.

What `fit()` Does

1. Counts how often each class appears in training data -> computes prior $P(\text{class})$
2. For every feature x every category x every class -> counts co-occurrences
3. Applies Laplace smoothing to get $P(X_i=v \mid \text{class})$
4. Stores everything as **log-probabilities** to prevent floating-point underflow

What `predict_proba()` Does

1. Starts with log-prior for each class
2. Adds (not multiplies -- because $\log(a \times b) = \log(a) + \log(b)$) the log-likelihood for each feature
3. Exponentiates back to probability space
4. Normalises so probabilities sum to 1

Why Log-Space?

If we multiplied 10 probabilities each around 0.1:

- $0.1 \times 0.1 \times \dots \times 0.1 = 1e-10$ (very small but OK)
- With 50 features: $0.1^{50} = 1e-50$ (Python stores as 0.0 -- underflow!)



- In log-space: $\log(0.1) \times 50 = -50$ (perfectly representable)
- We just add the logs and exponentiate once at the end

```
# -----
# SL-4: CATEGORICAL NAIVE BAYES BUILT FROM SCRATCH
#
# This class mirrors exactly what sklearn CategoricalNB does internally.
# Every step is written explicitly with comments.
# -----
```

```
class ManualCategoricalNB:
```

```
    """
```

```
    From-scratch Categorical Naive Bayes implementation.
```

```
    Parameters
```

```
    -----
```

```
    alpha : float
```

```
        Laplace smoothing strength. Default = 1.0.
```

```
    """
```

```
    def __init__(self, alpha=1.0):
```

```
        self.alpha = alpha
```

```
    def fit(self, X, y):
```

```
        """
```

```
        Learn class priors and feature likelihoods from training data.
```



```

After calling fit:

    self.classes_          -- unique class labels

    self.log_class_prior_  -- log P(class) for each class

    self.log_likelihood_   -- list[feat_idx] -> array(n_classes,
n_cats)

                                stores log P(Xi=v | class)

"""

n_samples, n_features = X.shape

self.classes_ = np.unique(y)

self.n_features_ = n_features

# ---- Step 1: Count samples per class, compute log-prior
-----

# class_count[k] = number of training samples belonging to class k

self.class_counts_ = np.array([np.sum(y == c) for c in
self.classes_])

# log-prior: log(count / total)

# Using log avoids multiplying many small numbers later

self.log_class_prior_ = np.log(self.class_counts_ / n_samples)

# ---- Step 2: For each feature, count category-class
co-occurrences

# n_cats_per_feature[f] = number of distinct values feature f takes

self.n_cats_per_feature_ = [len(np.unique(X[:, f]))

                                for f in range(n_features)]

```



```

# log_likelihood_ is a list: one array per feature

# each array has shape (n_classes, n_categories_for_that_feature)
self.log_likelihood_ = []

for feat_idx in range(n_features):

    n_cats = self.n_cats_per_feature_[feat_idx]

    # count_table[class_idx, cat_val] = occurrences in training
data
    count_table = np.zeros((len(self.classes_), n_cats))

    for cls_idx, cls in enumerate(self.classes_):

        X_cls = X[y == cls, feat_idx] # feature values for this
class
        for cat_val in range(n_cats):

            count_table[cls_idx, cat_val] = np.sum(X_cls == cat_val)

# ---- Step 3: Apply Laplace smoothing
-----

# smoothed[k,v] = count(class=k, feat=v) + alpha

# row_sums[k] = count(class=k) + alpha * K

smoothed = count_table + self.alpha

row_sums = self.class_counts[:, np.newaxis] + self.alpha *
n_cats

log_probs = np.log(smoothed / row_sums)

self.log_likelihood_.append(log_probs)

```



```

    return self

def predict_log_proba(self, X):
    """
    Compute log-score for each class.

    log_score(class) = log_prior(class) + sum of log_likelihoods

    Adding logs is equivalent to multiplying probabilities, but
    numerically stable.

    """
    n_samples = X.shape[0]

    # Start every sample with the log-prior for each class

    log_scores = np.tile(self.log_class_prior_, (n_samples, 1)) #
    shape (n_samples, n_classes)

    for feat_idx in range(self.n_features_):
        cat_vals = X[:, feat_idx].astype(int) # observed category for
        each sample

        # self.log_likelihood_[feat_idx] has shape (n_classes, n_cats)

        # We select column cat_val for each sample -> shape (n_classes,
        n_samples)

        ll = self.log_likelihood_[feat_idx][:, cat_vals]

        log_scores += ll.T # add (n_samples, n_classes)

    return log_scores

```



```

def predict_proba(self, X):

    """Convert log-scores back to normalised probabilities."""

    log_scores = self.predict_log_proba(X)

    scores      = np.exp(log_scores)          # back to probability space

    return scores / scores.sum(axis=1, keepdims=True) # normalise rows
to sum=1

def predict(self, X):

    """Return the class with the highest score for each sample."""

    return self.classes_[np.argmax(self.predict_log_proba(X), axis=1)]

# ---- Train and compare
-----

manual_nb = ManualCategoricalNB(alpha=1.0)

manual_nb.fit(X_train, y_train)

print('=' * 65)

print('  SL-4: ManualCategoricalNB vs sklearn CategoricalNB')

print('=' * 65)

man_acc = accuracy_score(y_test, manual_nb.predict(X_test)) * 100

sk_acc  = accuracy_score(y_test, cnb_model.predict(X_test)) * 100

print(f'\nTest Accuracy:')

print(f'  ManualCategoricalNB : {man_acc:.2f}%')

```




```

print(f'  sklearn NB                : {sk_acc:.2f}%')

print(f'  Results match              : {abs(man_acc - sk_acc) < 0.01}')


man_proba = manual_nb.predict_proba(query_encoded_sl)[0]

print(f'\nQuery Probabilities (Sunny, Cool, High, Strong):')

print(f'  {"Model":25s}  {"P(No)":>8}  {"P(Yes)":>8}  {"Prediction":>12}')

print(f'  {"-"*25}  {"-"*8}  {"-"*8}  {"-"*12}')
```

```

print(f'  {"ManualCategoricalNB":25s}  {man_proba[0]:>8.4f}
{man_proba[1]:>8.4f}  ')

f'{le_target.inverse_transform(manual_nb.predict(query_encoded_sl))[0]:>12}')
```

```

print(f'  {"sklearn CategoricalNB":25s}  {sk_proba[0]:>8.4f}
{sk_proba[1]:>8.4f}  ')

f'{le_target.inverse_transform(cnb_model.predict(query_encoded_sl))[0]:>12}')
```

```

match = np.allclose(man_proba, sk_proba, atol=0.001)

print(f'\nProbabilities match within 0.001: {match}')
```

```

print('\nConclusion: Our from-scratch implementation exactly replicates
sklearn.')
```

```

print('This proves we understand the algorithm at implementation level, not
just API level.')
```

```

=====

SL-4: ManualCategoricalNB vs sklearn CategoricalNB

=====
```



Test Accuracy:

ManualCategoricalNB : 90.00%

sklearn NB : 90.00%

Results match : True

Query Probabilities (Sunny, Cool, High, Strong):

Model	P(No)	P(Yes)	Prediction
ManualCategoricalNB	0.7894	0.2106	No
sklearn CategoricalNB	0.7894	0.2106	No

Probabilities match within 0.001: True

Conclusion: Our from-scratch implementation exactly replicates sklearn.

This proves we understand the algorithm at implementation level, not just API level.

SL-5: Decision Boundary Visualisation -- All 4 Models

What is a Decision Boundary?

A decision boundary is the dividing line between regions where the model predicts each class. Points on one side are predicted No; points on the other side are predicted Yes.

What We Plot



We vary Outlook (x-axis) and Humidity (y-axis), keeping Temperature=Cool and Wind=Strong. The background colour shows the predicted probability of Yes at every point in the space. The black contour line is the exact decision boundary (where $P(\text{Yes}) = 0.5$).

What the Shape Tells You

Model	Boundary Shape	Why
Naive Bayes	Rectangular / step-like	Independent features create axis-aligned regions
Decision Tree	Staircase	If-else rules split along one axis at a time
Logistic Regression	Straight diagonal line	Can only fit a single linear boundary
SVM (RBF)	Smooth curve	Non-linear kernel maps data to higher-dimensional space

```
# -----
# SL-5: DECISION BOUNDARY PLOTS -- All 4 Models
#
# We create a fine mesh of (Outlook, Humidity) values.
# Temperature is fixed at 0 (Cool), Wind fixed at 0 (Strong).
# For each point in the mesh, we ask the model: what is P(Yes)?
# The result is displayed as a colour map.
# The gold star marks our custom query point: Sunny(2) + High Humidity(0).
# -----

fig, axes = plt.subplots(2, 2, figsize=(14, 11))

fig.suptitle('SL-5: Decision Boundaries -- Outlook vs Humidity\n')
```



```

        '(Temperature=Cool fixed, Wind=Strong fixed | Red=No,
Green=Yes) ',

        fontsize=14, fontweight='bold')

# Fine mesh: Outlook from -0.3 to 2.3, Humidity from -0.3 to 1.3
xx, yy = np.meshgrid(np.linspace(-0.3, 2.3, 300),
                      np.linspace(-0.3, 1.3, 300))

# Build full feature matrix: columns = [Outlook, Temperature, Humidity,
Wind]

grid_X = np.column_stack([
    xx.ravel(),          # Outlook = varied
    np.zeros(xx.size),   # Temperature = 0 (Cool)
    yy.ravel(),          # Humidity = varied
    np.zeros(xx.size)    # Wind = 0 (Strong)
])

model_items = [
    ('Categorical NB',    cnb_model,   '#4C72B0'),
    ('Decision Tree',     dt_model,    '#55A868'),
    ('Logistic Regression', lr_model,   '#C44E52'),
    ('SVM (RBF)',         svm_model,    '#8172B2')
]

for ax, (name, model, _colour) in zip(axes.flatten(), model_items):

```



```

# P(Yes) at every point in the mesh

prob_grid = model.predict_proba(grid_X)[:, cls_list.index('Yes')]

Z = prob_grid.reshape(xx.shape)

# Smooth colour background: red(low P(Yes)) -> yellow -> green(high
P(Yes))

cf = ax.contourf(xx, yy, Z, levels=50, cmap='RdYlGn', alpha=0.80,
                 vmin=0, vmax=1)

# Decision boundary = contour at P(Yes) = 0.5

cs = ax.contour(xx, yy, Z, levels=[0.5], colors=['black'],
               linewidths=2.5)

ax.clabel(cs, fmt={0.5: 'Decision Boundary (P=0.5)'}, fontsize=8)

# Overlay actual data points where Temperature=Cool AND Wind=Strong

for xi, yi in zip(X_train, y_train):

    if xi[1] == 0 and xi[3] == 0: # Temperature=0 (Cool),
Wind=0 (Strong)

        clr = '#2d6a4f' if yi == 1 else '#ae2012'

        mk = 'o' if yi == 1 else 's'

        ax.scatter(xi[0], xi[2], c=clr, marker=mk, s=140,
                  edgecolors='white', linewidths=1.5, zorder=5)

# Gold star = query point: Sunny(2), High Humidity(0)

ax.scatter([2], [0], c='gold', marker='*', s=350,
          edgecolors='black', linewidths=1.5, zorder=6,
          label='Query: Sunny + High Humidity')

```



```

ax.set_xlim(-0.3, 2.3)
ax.set_ylim(-0.3, 1.3)
ax.set_xticks([0, 1, 2])
ax.set_xticklabels(['Overcast', 'Rain', 'Sunny'], fontsize=10)
ax.set_yticks([0, 1])
ax.set_yticklabels(['High', 'Normal'], fontsize=10)
ax.set_xlabel('Outlook', fontsize=11)
ax.set_ylabel('Humidity', fontsize=11)

acc = accuracy_score(y_test, model.predict(X_test)) * 100

ax.set_title(f'{name}      (Test Accuracy: {acc:.1f}%) \n'
             f'Gold star = query point (Sunny + High)',
             fontsize=11, fontweight='bold')

ax.legend(loc='upper right', fontsize=8)

plt.colorbar(cf, ax=ax, label='P(Yes)', shrink=0.85)

plt.tight_layout()

plt.show()

print('Observations from the boundary plots:')

print(' Naive Bayes      -- rectangular regions (each feature splits
independently)')

print(' Decision Tree    -- staircase / right-angle corners (if-else rules)')

print(' Logistic Reg.     -- a single straight diagonal line (linear only)')

print(' SVM RBF           -- smooth curved boundary (non-linear kernel)')


```



```
print()
```

```
print('The gold star (Sunny + High Humidity) falls in the RED zone for all models.')
```

```
print('All models agree: do NOT play tennis under these conditions.')
```



Observations from the boundary plots:

- Naive Bayes -- rectangular regions (each feature splits independently)
- Decision Tree -- staircase / right-angle corners (if-else rules)
- Logistic Reg. -- a single straight diagonal line (linear only)
- SVM RBF -- smooth curved boundary (non-linear kernel)



The gold star (Sunny + High Humidity) falls in the RED zone for all models.

All models agree: do NOT play tennis under these conditions.

SL-6: Learning Curves -- How Training Data Size Affects Each Model

What Is a Learning Curve?

A learning curve shows model accuracy as a function of how many training samples are used.

We plot two curves:

- **Training accuracy:** performance on the data the model was trained on
- **Cross-validation accuracy:** performance on held-out data (generalisation ability)

Diagnosing Model Problems

Pattern	Diagnosis	Fix
Training >> CV (large gap)	Overfitting (high variance)	More data, regularisation, simpler model
Both converge to a low value	Underfitting (high bias)	More complex model, more features
Both converge to a high value	Good fit	No action needed
Gap does not narrow with more data	Structural model mismatch	Different algorithm needed

SL-6: LEARNING CURVES -- All 4 Models

#



```
# sklearn's learning_curve() function trains the model repeatedly
# on different fractions of training data and evaluates on 5-fold CV.
# This shows how accuracy changes as we give the model more data.
# -----

from sklearn.model_selection import learning_curve

fig, axes = plt.subplots(2, 2, figsize=(14, 10))

fig.suptitle('SL-6: Learning Curves -- Training Accuracy vs Cross-Validation
Accuracy\n'

            'Shaded bands = +/- 1 standard deviation across 5 folds',

            fontsize=13, fontweight='bold')

train_sizes_frac = np.linspace(0.2, 1.0, 8)

lc_colours        = ['#4C72B0', '#55A868', '#C44E52', '#8172B2']

for ax, (name, model, _c), colour in zip(axes.flatten(), model_items,
lc_colours):

    ts, tr_sc, cv_sc = learning_curve(

        model, X, y,

        train_sizes=train_sizes_frac,

        cv=5,

        scoring='accuracy',

        shuffle=True,

        random_state=42
```



```

)

tr_mean = tr_sc.mean(axis=1) * 100
tr_std = tr_sc.std(axis=1) * 100
cv_mean = cv_sc.mean(axis=1) * 100
cv_std = cv_sc.std(axis=1) * 100

# Training curve (solid, coloured)
ax.plot(ts, tr_mean, 'o-', color=colour, linewidth=2.5, markersize=7,
        label='Training Accuracy')
ax.fill_between(ts, tr_mean - tr_std, tr_mean + tr_std,
               alpha=0.15, color=colour)

# CV curve (dashed, gray)
ax.plot(ts, cv_mean, 's--', color='gray', linewidth=2.5, markersize=7,
        label='CV Accuracy (5-fold mean)')
ax.fill_between(ts, cv_mean - cv_std, cv_mean + cv_std,
               alpha=0.12, color='gray')

# Annotate the final gap
gap = tr_mean[-1] - cv_mean[-1]

status = 'Overfit' if gap > 15 else ('Good fit' if gap < 8 else 'Slight
overfit')

ax.annotate(f'Gap = {gap:.1f}% ({status})',
           xy=(ts[-1], cv_mean[-1]),

```



```

xytext=(ts[-3], cv_mean[-1] - 12),

fontsize=9, color='darkred',

arrowprops=dict (arrowstyle='->', color='darkred'))

ax.set_xlabel('Number of Training Samples', fontsize=11)

ax.set_ylabel('Accuracy (%)', fontsize=11)

ax.set_title(f'{name}', fontsize=12, fontweight='bold')

ax.set_ylim(40, 115)

ax.axhline(100, color='gray', linestyle=':', alpha=0.5)

ax.legend(fontsize=9)

ax.grid(True, alpha=0.3)

ax.spines[['top', 'right']].set_visible(False)

plt.tight_layout()

plt.show()

print('Reading the learning curves:')

print('  Solid colour line = Training accuracy')

print('  Gray dashed line  = Cross-validation accuracy (generalisation)')

print('  Shaded bands       = +/- 1 standard deviation')

print()


print('  A large gap between training and CV lines = overfitting.')

print('  Decision Tree typically shows the largest gap (can memorise training
data).')
```



```
print(' Naive Bayes typically shows the smallest gap (simpler model, less overfitting).')
```

```
print(' If the gap narrows as more data is added, collecting more samples would help.')
```



Reading the learning curves:

Solid colour line = Training accuracy

Gray dashed line = Cross-validation accuracy (generalisation)

Shaded bands = ± 1 standard deviation

A large gap between training and CV lines = overfitting.



Decision Tree typically shows the largest gap (can memorise training data).

Naive Bayes typically shows the smallest gap (simpler model, less overfitting).

If the gap narrows as more data is added, collecting more samples would help.

SL-7: What Did the Model Memorise? -- Conditional Probability Heatmap

What Are Conditional Probabilities?

After training, CategoricalNB stores a table of values:

$P(X_i = \text{value} \mid \text{Class})$ -- How likely is this feature value, given we know the class?

These are the exact numbers used during prediction.

How to Read the Table

- A high $P(\text{value} \mid \text{No})$: this value is common when the class is No -> pushes prediction toward No
- A high $P(\text{value} \mid \text{Yes})$: this value is common when class is Yes -> pushes prediction toward Yes
- When $P(\text{value} \mid \text{No}) \gg P(\text{value} \mid \text{Yes})$: the value is STRONGLY discriminative for No

The bar chart shows the ratio $P(X_i \mid \text{No}) / P(X_i \mid \text{Yes})$.

- Ratio > 1: pushes toward No
- Ratio < 1: pushes toward Yes
- Ratio = 1: feature value carries no information

SL-7: CONDITIONAL PROBABILITY HEATMAP



```
#

# CategoricalNB stores feature_log_prob_ as a list of arrays:

#   feature_log_prob_[feat_idx] has shape (n_classes, n_categories)

#   feature_log_prob_[feat_idx][class_idx, cat_idx] = log P(Xi=cat |
Class)

#

# We exponentiate back to probabilities and display them.

# -----

feature_categories_sl = [label_encoders[f].classes_ for f in feature_cols]

print('=' * 70)

print('  SL-7: P(Xi | Class) -- Conditional Probability Table')

print('  These are the exact values CategoricalNB uses to predict.')

print('=' * 70)

# ---- Build a readable table
# -----

cpt_rows = []

for fi, (feat, cats) in enumerate(zip(feature_cols, feature_categories_sl)):

    for ci, cat in enumerate(cats):

        p_no = np.exp(cnb_model.feature_log_prob_[fi][0, ci])

        p_yes = np.exp(cnb_model.feature_log_prob_[fi][1, ci])

        ratio = p_no / p_yes

        disc = 'strongly No' if ratio > 1.5 else \
```



```

        'strongly Yes' if ratio < 0.67 else 'neutral'

    cpt_rows.append({'Feature': feat, 'Category': cat,

                    'P(Xi|No)': round(p_no,4), 'P(Xi|Yes)':

round(p_yes,4),

                    'Ratio No/Yes': round(ratio,3), 'Signal': disc})

cpt_df = pd.DataFrame(cpt_rows)

print()

print(cpt_df.to_string(index=False))

# ---- Heatmaps: one for each class
-----

fig, axes = plt.subplots(1, 2, figsize=(16, 5))

fig.suptitle('SL-7: Conditional Probabilities -- What CategoricalNB
Learned\n'

            'Darker = higher probability for that class given the feature
value',

            fontsize=13, fontweight='bold')

for ax_idx, (cls_label, col_key) in enumerate(['No (Do not play)',
        'P(Xi|No)'],

        ('Yes (Play tennis)',
        'P(Xi|Yes)')):

    pivot = cpt_df.pivot(index='Feature', columns='Category',
        values=col_key).fillna(0)

    cmap = 'Reds' if 'No' in cls_label else 'Greens'

    sns.heatmap(pivot, annot=True, fmt='.3f', cmap=cmap,

        linewidths=1.2, linecolor='white',

```



```

ax=axes[ax_idx], vmin=0, vmax=0.8,

annot_kws={'size': 11, 'weight': 'bold'})

axes[ax_idx].set_title(f'P(Feature Value | Class = {cls_label})\n'
                       f'Darker = more likely for class
{cls_label.split()[0]}',

                       fontsize=11, fontweight='bold')

axes[ax_idx].tick_params(axis='x', rotation=20)

axes[ax_idx].tick_params(axis='y', rotation=0)

plt.tight_layout()

plt.show()

# ---- Discriminative power bar chart
-----

fig, ax = plt.subplots(figsize=(12, 5))

labels = [f"{r['Feature']}\n= {r['Category']}" for _, r in
cpt_df.iterrows()]

ratios = cpt_df['Ratio No/Yes'].values

b_cols = ['#C44E52' if r > 1 else '#55A868' for r in ratios]

ax.bar(range(len(labels)), ratios - 1, color=b_cols, edgecolor='white',
linewidth=0.8)

ax.axhline(0, color='black', linewidth=1.5)

ax.set_xticks(range(len(labels)))

ax.set_xticklabels(labels, fontsize=9)

ax.set_ylabel('P(Xi|No) / P(Xi|Yes) minus 1\n(Red above 0 = pushes toward
No)', fontsize=10)

```




```

ax.set_title('Discriminative Power of Each Feature Value\n'

             'Red bar (above 0) = pushes toward No | Green bar (below 0) =
             pushes toward Yes',

             fontsize=12, fontweight='bold')

ax.spines[['top', 'right']].set_visible(False)

ax.grid(True, alpha=0.3, axis='y')

plt.tight_layout()

plt.show()

q_cats = {'Outlook':'Sunny', 'Temperature':'Cool', 'Humidity':'High',
          'Wind':'Strong'}

print('\nInterpretation for the custom query (Sunny + Cool + High +
      Strong):')

for feat, cat in q_cats.items():

    row = cpt_df[(cpt_df['Feature']==feat) &
                 (cpt_df['Category']==cat)].iloc[0]

    print(f'   {feat:15s} = {cat:8s}  ->  P(|No)={row["P(Xi|No)"]:.3f}   '

          f'P(|Yes)={row["P(Xi|Yes)"]:.3f}   Signal: {row["Signal"]}')

```


SL-7: $P(X_i | \text{Class})$ -- Conditional Probability Table


These are the exact values CategoricalNB uses to predict.

Feature	Category	$P(X_i \text{No})$	$P(X_i \text{Yes})$	Ratio No/Yes	Signal
---------	----------	--------------------	---------------------	--------------	--------



Outlook	Overcast	0.0625	0.4000	0.156	strongly Yes
Outlook	Rain	0.4375	0.2667	1.641	strongly No
Outlook	Sunny	0.5000	0.3333	1.500	neutral
Temperature	Cool	0.3750	0.3000	1.250	neutral
Temperature	Hot	0.4375	0.3000	1.458	neutral
Temperature	Mild	0.1875	0.4000	0.469	strongly Yes
Humidity	High	0.8000	0.3103	2.578	strongly No
Humidity	Normal	0.2000	0.6897	0.290	strongly Yes
Wind	Strong	0.6667	0.4138	1.611	strongly No
Wind	Weak	0.3333	0.5862	0.569	strongly Yes





Interpretation for the custom query (Sunny + Cool + High + Strong):

Outlook = Sunny -> $P(|No)=0.500$ $P(|Yes)=0.333$ Signal: neutral

Temperature = Cool -> $P(|No)=0.375$ $P(|Yes)=0.300$ Signal: neutral

Humidity = High -> $P(|No)=0.800$ $P(|Yes)=0.310$ Signal: strongly No

Wind = Strong -> $P(|No)=0.667$ $P(|Yes)=0.414$ Signal: strongly No

SL-8: When Does Naive Bayes Fail? -- Failure Modes

Naive Bayes is fast and often surprisingly accurate, but it has clear limitations:

Failure Mode 1: Feature Interactions Are Invisible

NB multiplies $P(F1|class) * P(F2|class)$ independently. If the class label depends on the COMBINATION of F1 and F2 (not either alone), NB cannot learn this pattern.



Classic example: XOR -- class = 1 only when exactly one of F1, F2 equals 1. Neither F1 alone nor F2 alone has any signal ($P(F1=1|class) = 0.5$ for both classes). NB predicts at random. Decision Tree handles XOR easily.

Failure Mode 2: Poor Probability Calibration

Even when NB correctly identifies the class, its confidence percentage may be off. Saying 99% confident does not mean it is correct 99% of the time. This matters when you need reliable probability estimates, not just class labels.

Failure Mode 3: Class Imbalance

If the training data has many more No samples than Yes, the prior $P(\text{No})$ is large. This prior multiplies into every prediction, biasing the model toward No even when the features favour Yes.

```
# -----

# SL-8: FAILURE MODES -- Demonstrated with concrete examples

# -----

print('=' * 68)

print('  SL-8: FAILURE MODES OF NAIVE BAYES')

print('=' * 68)

# ---- Failure Mode 1: Feature Interaction (XOR)
# -----

print('\nFAILURE MODE 1: Feature Interactions Are Invisible to Naive Bayes')

print('  Synthetic XOR dataset: class = 1 only when Feature1 != Feature2')

print('  (This requires knowing BOTH features together -- neither alone is
informative)')

print()
```



```

np.random.seed(42)

n_xor = 80

F1_xor = np.random.randint(0, 2, n_xor)

F2_xor = np.random.randint(0, 2, n_xor)

y_xor = (F1_xor ^ F2_xor).astype(int)    # XOR: 1 when F1 != F2

X_xor = np.column_stack([F1_xor, F2_xor])

X_xor_tr, X_xor_te, y_xor_tr, y_xor_te = train_test_split(
    X_xor, y_xor, test_size=0.25, random_state=42
)

nb_xor = CategoricalNB(alpha=1.0)

nb_xor.fit(X_xor_tr, y_xor_tr)

dt_xor = DecisionTreeClassifier(random_state=42)

dt_xor.fit(X_xor_tr, y_xor_tr)

nb_xor_acc = accuracy_score(y_xor_te, nb_xor.predict(X_xor_te)) * 100

dt_xor_acc = accuracy_score(y_xor_te, dt_xor.predict(X_xor_te)) * 100

print(f' NB  Accuracy on XOR dataset: {nb_xor_acc:.1f}%    <-- near 50%,
equivalent to random guessing')

print(f' DT  Accuracy on XOR dataset: {dt_xor_acc:.1f}%    <-- can split on
F1 THEN F2')

```



```

print()

print(' WHY NB FAILS HERE:')

print(' For XOR, when class=1: P(F1=1|class=1) = P(F1=0|class=1) = 0.5')

print(' F1 alone has no signal. Same for F2 alone.')

print(' NB computes P(F1|class) x P(F2|class) independently, so it learns
nothing.')

print(' Decision Tree can check: IF F1=1 THEN look at F2 -- capturing the
interaction.')

# ---- Failure Mode 2: Probability Calibration
-----

print('\nFAILURE MODE 2: Poor Probability Calibration')

print(' Even when the predicted class is correct, the confidence may be
misleading.')

print()

print(f' Test predictions with confidence:')

print(f' {"#":>3} {"Actual":>8} {"Predicted":>10} {"Confidence":>12}
{"Correct?":>8}')
```

#	Actual	Predicted	Confidence	Correct?
1	0	0	0.9	Yes
2	0	0	0.9	Yes
3	0	0	0.9	Yes
4	0	0	0.9	Yes
5	0	0	0.9	Yes
6	0	0	0.9	Yes
7	0	0	0.9	Yes
8	0	0	0.9	Yes
9	0	0	0.9	Yes
10	0	0	0.9	Yes
11	0	0	0.9	Yes
12	0	0	0.9	Yes
13	0	0	0.9	Yes
14	0	0	0.9	Yes
15	0	0	0.9	Yes
16	0	0	0.9	Yes
17	0	0	0.9	Yes
18	0	0	0.9	Yes
19	0	0	0.9	Yes
20	0	0	0.9	Yes
21	0	0	0.9	Yes
22	0	0	0.9	Yes
23	0	0	0.9	Yes
24	0	0	0.9	Yes
25	0	0	0.9	Yes
26	0	0	0.9	Yes
27	0	0	0.9	Yes
28	0	0	0.9	Yes
29	0	0	0.9	Yes
30	0	0	0.9	Yes
31	0	0	0.9	Yes
32	0	0	0.9	Yes
33	0	0	0.9	Yes
34	0	0	0.9	Yes
35	0	0	0.9	Yes
36	0	0	0.9	Yes
37	0	0	0.9	Yes
38	0	0	0.9	Yes
39	0	0	0.9	Yes
40	0	0	0.9	Yes
41	0	0	0.9	Yes
42	0	0	0.9	Yes
43	0	0	0.9	Yes
44	0	0	0.9	Yes
45	0	0	0.9	Yes
46	0	0	0.9	Yes
47	0	0	0.9	Yes
48	0	0	0.9	Yes
49	0	0	0.9	Yes
50	0	0	0.9	Yes
51	0	0	0.9	Yes
52	0	0	0.9	Yes
53	0	0	0.9	Yes
54	0	0	0.9	Yes
55	0	0	0.9	Yes
56	0	0	0.9	Yes
57	0	0	0.9	Yes
58	0	0	0.9	Yes
59	0	0	0.9	Yes
60	0	0	0.9	Yes
61	0	0	0.9	Yes
62	0	0	0.9	Yes
63	0	0	0.9	Yes
64	0	0	0.9	Yes
65	0	0	0.9	Yes
66	0	0	0.9	Yes
67	0	0	0.9	Yes
68	0	0	0.9	Yes
69	0	0	0.9	Yes
70	0	0	0.9	Yes
71	0	0	0.9	Yes
72	0	0	0.9	Yes
73	0	0	0.9	Yes
74	0	0	0.9	Yes
75	0	0	0.9	Yes
76	0	0	0.9	Yes
77	0	0	0.9	Yes
78	0	0	0.9	Yes
79	0	0	0.9	Yes
80	0	0	0.9	Yes
81	0	0	0.9	Yes
82	0	0	0.9	Yes
83	0	0	0.9	Yes
84	0	0	0.9	Yes
85	0	0	0.9	Yes
86	0	0	0.9	Yes
87	0	0	0.9	Yes
88	0	0	0.9	Yes
89	0	0	0.9	Yes
90	0	0	0.9	Yes
91	0	0	0.9	Yes
92	0	0	0.9	Yes
93	0	0	0.9	Yes
94	0	0	0.9	Yes
95	0	0	0.9	Yes
96	0	0	0.9	Yes
97	0	0	0.9	Yes
98	0	0	0.9	Yes
99	0	0	0.9	Yes
100	0	0	0.9	Yes

```

test_proba = cnb_model.predict_proba(X_test)

test_preds = cnb_model.predict(X_test)

for i, (act, pred, proba) in enumerate(zip(y_test, test_preds, test_proba),
1):

    conf      = proba[pred] * 100

    correct   = 'Yes' if act == pred else 'No'

    act_lbl   = le_target.inverse_transform([act])[0]
```



```

pred_lbl = le_target.inverse_transform([pred])[0]

print(f' {i:>3} {act_lbl:>8} {pred_lbl:>10} {conf:>11.2f}%
{correct:>8}')
```

---- Failure Mode 3: Class Imbalance

```

print('\nFAILURE MODE 3: Class Imbalance Bias')

print(' When one class dominates training, the prior P(dominant) inflates
all predictions.')
```

print()

```

df_no_rows = df_encoded[df_encoded[target_col] == 0]
df_yes_rows = df_encoded[df_encoded[target_col] == 1]

# Create 4:1 imbalance by repeating No rows

df_imb = pd.concat([df_no_rows]*4 + [df_yes_rows]).sample(frac=1,
random_state=42)

X_imb, y_imb = df_imb[feature_cols].values, df_imb[target_col].values

X_imb_tr, X_imb_te, y_imb_tr, y_imb_te = train_test_split(
    X_imb, y_imb, test_size=0.2, random_state=42
)

nb_imb = CategoricalNB(alpha=1.0)
nb_imb.fit(X_imb_tr, y_imb_tr)

imb_proba = nb_imb.predict_proba(query_encoded_sl)[0]

orig_n_no = (y_train == 0).sum()
```



```

orig_n_yes = (y_train == 1).sum()

imb_n_no    = (y_imb_tr == 0).sum()

imb_n_yes   = (y_imb_tr == 1).sum()


print(f' Original training prior :
P(No)={orig_n_no/(orig_n_no+orig_n_yes):.2f}
P(Yes)={orig_n_yes/(orig_n_no+orig_n_yes):.2f}')

print(f' Imbalanced training prior:
P(No)={imb_n_no/(imb_n_no+imb_n_yes):.2f}
P(Yes)={imb_n_yes/(imb_n_no+imb_n_yes):.2f}')

print(f'\n Query prediction from ORIGINAL model :')

print(f'      P(No) = {sk_proba[0]:.4f}  P(Yes) = {sk_proba[1]:.4f}')
```

Query prediction from IMBALANCED model :

```

print(f'      P(No) = {imb_proba[0]:.4f}  P(Yes) = {imb_proba[1]:.4f}')
```

```

print()

print(' Imbalance pushed P(No) even higher because the prior P(No) = 0.80')

print(' multiplies into every prediction. The model is biased before even')

print(' looking at the feature values.')
```

SL-8: FAILURE MODES OF NAIVE BAYES

FAILURE MODE 1: Feature Interactions Are Invisible to Naive Bayes

Synthetic XOR dataset: class = 1 only when Feature1 != Feature2

(This requires knowing BOTH features together -- neither alone is informative)



NB Accuracy on XOR dataset: 30.0% <-- near 50%, equivalent to random guessing

DT Accuracy on XOR dataset: 100.0% <-- can split on F1 THEN F2

WHY NB FAILS HERE:

For XOR, when class=1: $P(F1=1|class=1) = P(F1=0|class=1) = 0.5$

F1 alone has no signal. Same for F2 alone.

NB computes $P(F1|class) \times P(F2|class)$ independently, so it learns nothing.

Decision Tree can check: IF F1=1 THEN look at F2 -- capturing the interaction.

FAILURE MODE 2: Poor Probability Calibration

Even when the predicted class is correct, the confidence may be misleading.

Test predictions with confidence:

#	Actual	Predicted	Confidence	Correct?
---	-----	-----	-----	-----
1	No	Yes	66.83%	No
2	Yes	Yes	98.22%	Yes
3	No	No	78.94%	Yes
4	Yes	Yes	86.00%	Yes
5	Yes	Yes	94.25%	Yes
6	Yes	Yes	87.89%	Yes
7	Yes	Yes	87.89%	Yes



8	Yes	Yes	64.82%	Yes
9	No	No	82.71%	Yes
10	Yes	Yes	95.12%	Yes

FAILURE MODE 3: Class Imbalance Bias

When one class dominates training, the prior $P(\text{dominant})$ inflates all predictions.

Original training prior : $P(\text{No})=0.33$ $P(\text{Yes})=0.68$

Imbalanced training prior: $P(\text{No})=0.65$ $P(\text{Yes})=0.35$

Query prediction from ORIGINAL model :

$P(\text{No}) = 0.7894$ $P(\text{Yes}) = 0.2106$

Query prediction from IMBALANCED model :

$P(\text{No}) = 0.9632$ $P(\text{Yes}) = 0.0368$

Imbalance pushed $P(\text{No})$ even higher because the prior $P(\text{No}) = 0.80$ multiplies into every prediction. The model is biased before even looking at the feature values.

Self-Learning Summary

Section

Key Takeaway



- SL-1 Bayes theorem is Prior x Product-of-Likelihoods, then normalise. Verified by hand.
- SL-2 Features in this dataset ARE correlated, but NB still classifies correctly because ranking matters.
- SL-3 $\alpha=1$ prevents zero-probability collapse. Large $\alpha \rightarrow$ uniform; small $\alpha \rightarrow$ extreme confidence.
- SL-4 Rebuilt NB from scratch using log-space arithmetic. Exact match with sklearn confirmed.
- SL-5 NB=rectangular, DT=staircase, LR=straight line, SVM=curved. Each shape reveals the model's assumption.
- SL-6 Large gap between training and CV accuracy = overfitting. Collecting more data closes the gap.
- SL-7 The conditional probability heatmap shows exactly what the model learned. Sunny+High strongly signals No.
- SL-8 NB fails on XOR (feature interactions), class imbalance (prior dominates), and gives overconfident probabilities.

Analysis Report -- Why Models Produced Different Predictions and Probabilities

For the query (Sunny, Cool, High Humidity, Strong Wind), Naive Bayes assigned the highest $P(\text{No})$ among all models because it multiplies four independent likelihoods -- Sunny outlook, High Humidity, and Strong Wind each individually favour No -- and the independence assumption amplifies this by treating these as completely separate pieces of evidence that multiply without dampening (shown in SL-1 and SL-7).

Decision Tree followed a single if-else path to one leaf, making its decision the most brittle: it assigned 0% or 100% confidence because a pure leaf contains only one class, and the Sunny + High Humidity region in the 3D surface (SL-5) is entirely red (No) at the leaf level.



Logistic Regression produced the most moderate probability because its sigmoid function compresses the linear input into the (0,1) range, preventing extreme outputs, and its single straight boundary (SL-5) cannot capture the compounding effect of multiple unfavourable features simultaneously.

SVM with RBF kernel found a smooth non-linear boundary influenced by support vectors; on only 40 training samples the support-vector ratio is high and Platt scaling calibrates its probabilities differently from the other models, explaining why its P(No) differs even when the predicted label agrees with Naive Bayes (visible in the learning curve variance in SL-6).